Everyday AI: Photos, video & music

20 step-by-step guides for real, everyday tasks

Photos, video & music

Updated edition · October 19, 2026

AI World HQ · https://aiworldhq.com

Free edition · CC BY-NC 4.0 · free to share with credit

Before you start — what this is, honestly

This is an updated edition. It was rebuilt from our current guides on the date shown on the cover: more guides per topic, retired guides removed, and every step taken from the latest version of each guide.

This collection is free, and every guide in it is also free on our website. This file just gathers them in one place, in reading order, so you can print it, keep it on your phone, or read it offline. Our website has no paywall either — every guide links back to its online page, which is always the newer version.

AI **Every guide in this pack was written by AI** and checked by automated fact and quality gates. **Our editor spot-checks published guides every day, but not every text is read by a person.** Every guide page on https://aiworldhq.com carries the same notice — we say it openly because you deserve to know who wrote what you are reading.

Apps change. Buttons move and menus get renamed. If a step does not match what you see, the guide's online version is the newer one — the link is under every title.

Free to share. This collection is licensed under Creative Commons Attribution-NonCommercial 4.0 (CC BY-NC 4.0): you may copy, print and share it — with your class, your library or your team — as long as you credit AI World HQ with a link to aiworldhq.com and do not sell it.

What's inside

1. **Photos, video & music**
    1. Generate unique profile pictures for social media with Midjourney
    2. How to generate AI images in Picsart: a beginner's step-by-step guide
    3. Auto-write captions for your photos using Hugging Face's free image captioning tool
    4. How to remove a photo background cleanly with Meta AI on your phone
    5. Create a multi-room music playlist with Alexa+
    6. Sharpen Your Photos with AI Using Upscayl
    7. How to create and edit images in ChatGPT
    8. How to Sort and Clean Up a Messy Phone Photo Library with Gemini
    9. Safely Rename a Batch of Photos with a GitHub Copilot Script
    10. How to Turn a Long Podcast Transcript Into Shareable Notes With Claude
    11. How to Summarize a Long YouTube Video in 60 Seconds with Any AI
    12. Create Royalty-Free AI Music Loops for Podcasts and Videos
    13. Remove Distractions from Photos with Google Photos Magic Eraser
    14. How to use Kling AI video templates in Picsart Flow
    15. How to create images with AI: a beginner's first try
    16. How to Edit Your Phone Photos with Simple Text Instructions in Meta AI
    17. How to Generate Images with Google Gemini: A Beginner's Guide
    18. Create Your First AI Image with Midjourney: A Beginner Walkthrough
    19. How to Create AI Logos with Recraft V4.1 Vector
    20. How to Create Custom Images with Copilot for Invites and Social Posts

Photos, video & music

Generate unique profile pictures for social media with Midjourney

7 steps · Read online

Want to make your online presence pop with a profile picture that's uniquely *you*? This guide will show you how to use Midjourney, an artificial intelligence (AI) tool, to design eye-catching images for all your social media profiles. By the end, you'll be able to create custom profile pictures or even social media banners that reflect your personality or brand, making a memorable first impression.

Before you start

- **A Discord account**: this guide uses Midjourney inside Discord, a free communication app where communities gather. Discord still works, but most people now use the web app at midjourney.com instead — if you prefer that, follow our web guide instead of this one. Either way you need the paid plan below.
- **A paid Midjourney subscription**: there is no free option. Midjourney ended its free trial in March 2023, so you cannot generate even one image without a plan. The cheapest tier is called Basic — check Midjourney's own pricing page for the current cost.
- **Device**: A computer, tablet, or smartphone with the Discord app installed or access to the Discord website.
- **Time commitment**: Allow about 15-20 minutes for your first few attempts to explore options and get comfortable.

Step 1 — Get set up in Discord

First, you need to join the Midjourney community on Discord. Discord is like a digital clubhouse where groups can chat, share, and collaborate. Midjourney operates as a **bot** (a computer program that lives inside Discord and responds to your commands) within its own dedicated **server** (the specific 'clubhouse' for Midjourney users). To begin, open your Discord app or visit the Discord website in your browser. Look for the "Explore Public Servers" option, often represented by a compass icon on the left sidebar. From there, search for "Midjourney" in the search bar and click to join its official server. Once you're in, you'll see many chat channels. Look for channels named "newbies," "newcomer-rooms," or "generate-channels" in the sidebar, as these are where new users can create images. If you don't immediately see "Explore Public Servers" or a compass icon, look for a '+' sign to "Add a Server" on the left sidebar, and then select "Join a Server" or "Explore Public Servers" from there. You'll know it worked when you see the Midjourney server listed on your left sidebar in Discord, and you can click into a "newbies" or "generate" channel.

Step 2 — Start your first image with a prompt

Now that you're in a generation channel, it's time to tell Midjourney what you want to create using a **prompt** – this is the text instruction you give to the AI. Think of it like giving a detailed request to an artist. In your chosen "newbies" or "generate" channel, type `/imagine` into the message box at the bottom. As you type, you'll see a pop-up with `/imagine prompt:` appear. Click on this, or press Enter. The word "prompt" will now be highlighted, waiting for your description. This tells the Midjourney bot you're ready to **generate** (create) an image. After you've typed your description, simply press Enter to send it. If it looks different, make sure you're typing `/imagine` in a chat channel, not in a private message to the Midjourney Bot yet. If the pop-up doesn't appear, try typing a space after `/imagine` before entering your text. You'll know it worked when your message `/imagine prompt: [your description]` appears in the chat, and the Midjourney bot starts processing your request (often showing a percentage as it works). 💬 Example: `/imagine prompt: a joyful wombat wearing a tiny chef hat, cartoon style, vibrant colours`

Step 3 — Crafting your profile picture prompt

To get a great profile picture, your prompt needs to be specific. Think about the style, subject, and any particular details you want. For profile pictures, a square **aspect ratio** (the width-to-height proportion of an image) is best. You can specify this with `--ar 1:1` at the end of your prompt. Consider who or what your profile picture is for: Is it for a professional network, a gaming profile, or a personal social media page? Describe the subject, their mood, the art style (e.g., "photorealistic," "vector art," "watercolour"), colors, and background. Adding details like "portrait," "headshot," or "close-up" can help Midjourney focus. You'll know it worked when you've carefully thought about the details and added `--ar 1:1` to ensure your image will fit a typical square profile picture slot. 💬 Example: `/imagine prompt: a happy person with short curly hair, wearing glasses, smiling gently, digital art portrait, soft pastel colours, simple background, --ar 1:1`

Step 4 — Reviewing your generated grid

After you send your prompt, the Midjourney bot will take about a minute to process it. What you'll see next is a 2x2 **grid** – a set of four distinct image options based on your prompt, all presented together as a single image. Underneath this grid, there will be several buttons. You'll notice buttons labeled U1, U2, U3, U4 and V1, V2, V3, V4. The 'U' buttons are for 'Upscale', meaning you want to make one of the four small images larger and more detailed. The 'V' buttons are for 'Variations', meaning you want to create four *new* images that are similar in style or composition to one you liked from the grid. Buttons U1 and V1 correspond to the top-left image, U2 and V2 to the top-right, U3 and V3 to the bottom-left, and U4 and V4 to the bottom-

right. You'll know it worked when you see your 2x2 grid of four images appear in the chat, along with the U and V buttons below it.

Step 5 — Selecting and upscaling your favorite

Once you have your grid, it’s time to pick your favorite! Look closely at all four options and decide which one best suits your vision for a profile picture. If you like one of the images, click the corresponding 'U' button beneath the grid (e.g., U1 for the top-left image). This tells Midjourney to **upscale** that specific image, creating a larger, more refined version that's suitable for download. Midjourney will then process this request and present you with the single, higher-resolution image. If you prefer a slightly different take on one of the images, you can click a 'V' button instead to generate more variations. If it looks different and you don't see buttons, make sure you're looking at the most recent image grid generated by the Midjourney bot, as Discord chat moves quickly. Sometimes you may need to scroll up slightly. You'll know it worked when a single, larger version of your chosen image appears in the chat, often with new buttons like "Web," "Vary (Subtle)," or "Zoom Out" below it.

Step 6 — Saving your new profile picture

Once your chosen image has been upscaled, you'll see it displayed as a single, high-resolution picture in the chat. Now you can save it to your device! To save the image, simply click on it. This will open the image in a larger preview. From there, right-click (on a computer) or long-press (on a phone/tablet) and select "Save Image As..." or "Download Image." You can then choose where to save the file on your device. Once saved, it’s ready to be uploaded to your social media profiles! You can also click the "Web" button to view and download it directly from your Midjourney gallery in your web browser, which sometimes offers more options. If the options look different when saving, search for keywords like "Share," "Export," or "Save to Photos/Gallery" in the menu. You'll know it worked when the image file is successfully downloaded to your device, and you can open it in your photos or gallery app.

Step 7 — Refine your image or try again

Sometimes your first attempt might not be exactly what you envisioned. That's perfectly normal with AI art! If you want to refine an image you've upscaled, you can use the "Vary (Subtle)" or "Vary (Strong)" buttons below the upscaled image to create slightly different versions. If you want to make bigger changes, it's often better to start with a new `/imagine` prompt, adjusting your description with lessons learned from your previous attempts. Experiment with different keywords, art styles, or even add **negative prompts** (e.g., `--no blur` to avoid blurry images) to guide the AI more precisely. The Midjourney bot cannot directly edit specific parts of an image, so small tweaks often require generating new variations or starting fresh. The key is **iteration** – trying different prompts and variations until you get something you love. You'll know it worked

when you successfully generate new variations or a completely new image that moves closer to your ideal profile picture.

Common mistakes

- **Not being specific enough in your prompt**: A vague prompt like "person" will give you very generic results.
    - **The fix**: Add descriptive details about appearance, mood, setting, and art style. Think about the expression, clothing, background, and overall aesthetic (e.g., "a determined young professional, wearing a navy blazer, standing in front of a blurred city skyline, warm lighting, photorealistic").
- **Forgetting the aspect ratio**: If you don't add `--ar 1:1` to your prompt, Midjourney might give you a rectangular image that doesn't fit well as a profile picture.
    - **The fix**: Always include `--ar 1:1` at the end of your prompt for square images.
- **Using the wrong Discord channel**: Trying to generate images in general chat channels or private messages to the bot (before enabling it for DMs) won't work.
    - **The fix**: Always use a designated "newbies" or "generate" channel on the Midjourney server, or if you've enabled it, in a direct message with the Midjourney Bot itself.

What this means for you

- **In everyday life:** You can create unique profile pictures for Instagram, Facebook, LinkedIn, or even dating apps that truly capture your personality and help you stand out from the crowd. No more struggling to find the perfect photo!
- **For your business or work:** You can design professional-looking headshots, team profile images, or creative avatars for internal communications or company social media pages, ensuring a consistent and polished brand image.
- **If you're just getting started:** Try generating a simple profile picture of yourself with a specific art style, like "cartoon portrait of a person with long brown hair, smiling, warm colors, --ar 1:1".

Try it now

Open Discord, go to a Midjourney generation channel, and type `/imagine prompt: a futuristic astronaut with a friendly face, vibrant colours, cartoon style, --ar 1:1` to create your very first unique AI-generated profile picture.

How to generate AI images in Picsart: a beginner's step-by-step guide

7 steps · Read online

> **In short:** Picsart has a built-in AI image generator that turns your text prompts into pictures, then lets you keep editing in the same conversation. The AI tools usually sit behind a paid plan, and the engine works best with still images — not video.

If you've ever opened Picsart to fix a photo, design a poster, or make a quick thumbnail, the app already has an AI image generator tucked inside. Type a sentence, and Picsart turns it into a picture. Don't like something in the result? Stay in the same chat and ask for a change — the AI keeps the rest of the image intact. Here's how to try it, step by step, even if you've never touched an AI image tool before.

Before you start

- A Picsart account — sign-up is free, but AI image generation usually requires a paid plan. Check Picsart's current pricing page before you start.
- The Picsart app on your phone or tablet, or the web version in your browser at picsart.com.
- A rough idea of what you want to make: a poster, a product mock-up, a photo edit, or a wallpaper.
- About 10 minutes for your first image.

Step 1 — Open Picsart and find the AI image generator

Open the Picsart app and sign in. Once you're on the home screen, look for a tile, button, or section whose name suggests AI image generation — common labels include **AI Image Generator**, **AI Playground**, or **Text to Image**. If you don't see any of those on the main grid, tap the + (plus) button at the bottom-center of the screen, or look for a **Create** tab along the bottom row of icons — the AI tools usually live in one of those two places. On the web version at picsart.com, look in the top menu for something like **AI Tools** or **Create**.

💬 Example: this step has nothing to type — just finding the right entry point.

You'll know it worked when a screen appears with a text field where you can describe the image you want (usually called the prompt box) and at least one option to either start from text or upload a photo.

Step 2 — Pick your starting point: text, photo, or template

Inside the AI image tool, you'll usually see a few big choices near the top of the screen: **Text to Image** (start from a written description), **Image to Image** (start from a photo you upload), or a ready-made template you can fill in. Tap whichever fits what you want — for your first try, choose **Text to Image** if it's available. If the screen also shows a model selector — a small dropdown, a chip near the prompt box, or a tile in a side panel — pick whichever option is marked as the latest or recommended. Picsart rotates AI models over time, and the list changes often, so we won't name a specific one here.

Example: this step is a tap, not a typing step — the key choice is **Text to Image** for a first run.

You'll know it worked when the screen shows an empty prompt box and a clearly labeled button like **Generate**, **Create**, or an arrow icon.

Step 3 — Choose a canvas size (if the app asks)

Before you write your prompt, some versions of the tool show a size or shape picker. If you see it — usually a small icon shaped like a square, a rectangle, or a phone outline, or a dropdown near the prompt box — pick the shape that matches where you'll use the picture. A square shape works for Instagram-style posts, a tall rectangle works for phone wallpapers and Stories, and a wide rectangle works for banners and YouTube thumbnails. If the app also shows a quality choice (sometimes labeled with letters or numbers like 1K, 2K, or 4K, where higher means sharper but slower), pick the lowest one for a quick test so you don't burn through your credits.

Example: for a square post, choose the square option; for a phone wallpaper, choose the tall rectangle; if you see a quality choice, pick the lowest setting for a quick test.

You'll know it worked when a small preview, a number, or a labeled chip on the screen shows the shape and size you picked.

Step 4 — Write your prompt and generate the image

In the prompt box, describe the picture you want in plain, specific language. Name the subject, the mood, the colors, the lighting, and any text you want visible on the image — Picsart says its AI image generator handles on-image text reasonably well, which is useful for posters and diagrams. The more concrete you are, the closer the result will be to what you imagined. Then tap the **Generate** button (sometimes labeled **Create** or shown as an arrow icon). After a few seconds, one or more images appear on screen — pick the one you like best, or tap **Regenerate** if none feel right.

Example: "A friendly poster for a Saturday morning yoga class, soft pastel colors, with the words 'Saturday Morning Yoga' clearly readable in the center, hand-drawn style."

You'll know it worked when a finished image appears on screen and matches the canvas shape you picked in Step 3.

Step 5 — Refine the image in the same chat

Here's where the in-chat editing saves time: you don't have to start over when something's off. Stay in the same conversation and send a follow-up message describing one specific change. The AI keeps the rest of the picture intact and adjusts only what you asked for. You can chain several refinements one after another — change colors, swap a background, add a logo, rewrite the on-image text — without losing the previous version. Be specific: name the exact color, the exact word, the exact object. If you can't find the chat box, look for a field at the bottom of the screen that says something like **Message**, **Ask anything**, or **Describe a change**.

💬 Example: after the yoga poster above, send: "Make the background soft pink, change the words to 'Bring a friend!', and add a small sun in the top-right corner."

You'll know it worked when the new image shows your specific change and the rest of the picture still looks like the previous version.

Step 6 — Edit an existing photo (optional)

If you'd rather start from a picture you already have, look for an **Upload** or **Image** button near the prompt box — usually shown as a small picture icon or a + sign. Tap it and pick a photo from your device. Once the photo appears, describe the edit you want in the prompt box. For finer control, look for a brush or markup tool near the image — common labels are **Mask**, **Edit**, or a small paint-bucket or brush icon. With that tool, drag your finger across the part of the photo you want changed; the painted area is the only part the AI will redraw. Then write a short note in the prompt box describing what should appear there. Be specific about what to remove, replace, or add.

💬 Example: upload a group photo, then type: "Remove the person in the red shirt on the right." Or paint over the background with the brush tool and write: "Replace this background with a sunset beach."

You'll know it worked when the edited photo appears with your specific change applied and the rest of the image still looks natural.

Step 7 — Save and export your image

When you're happy with the result, look for a **Save**, **Download**, or **Export** button — usually at the top-right corner of the image or inside a ⋯ (three dots) menu. The app may ask you to pick a file format: choose **JPG** for everyday sharing, or **PNG** if you need a transparent background

(for layering the image on top of other designs). The image then saves to your device's photo gallery or your computer's Downloads folder. If you want to keep editing the file inside Picsart later, look for a **Save to project** or **Save as draft** option instead.

💬 Example: not a typing step — the key choice is **JPG** for everyday sharing, **PNG** only if you need a transparent background.

You'll know it worked when the image file appears in your phone's gallery (or your computer's Downloads folder) and opens like any normal picture.

Common mistakes

- **Picking the wrong starting mode.** If you upload a photo but meant to start from text (or the other way around), the result will look off. In Step 2, double-check you've chosen the right mode — **Text to Image** or **Image to Image** — before tapping Generate.
- **Asking for vague edits.** "Make it better" or "fix it" gives the AI nothing to act on. Always name the specific change: a color, an object, a word, a mood. The more concrete the instruction, the closer the result.
- **Forgetting the cost.** Each generation usually spends credits or counts against your plan's allowance. A higher-quality image costs more than a lower-quality one, and each follow-up edit in the chat counts as another generation. Check your remaining credits before kicking off a long chain of edits — otherwise the tool may stop mid-conversation.

What this means for you

- **In everyday life:** You can make a custom birthday card, a flyer for a school bake sale, or a fun wallpaper for your phone — in minutes, without any design training.
- **For your business or work:** You can mock up a quick social-media post, an event poster, or a product visual to share with a colleague before commissioning a designer for the final version.
- **If you're just getting started:** Begin with one simple text-to-image prompt (Step 4) at the lowest quality setting. Get a feel for what the tool does well before moving on to photo edits or higher-quality output.

Try it now

Open Picsart, find the AI image generator, and create one square image of your favorite season — one sentence prompt, no edits, just see what comes back. Takes about two minutes.

Auto-write captions for your photos using Hugging Face's free image captioning tool

7 steps · Read online

> **In short:** Hugging Face hosts free AI demos in its "Spaces" section, including image captioning tools that describe a photo in one short sentence. You upload an image, wait a few seconds, and copy the text — no coding needed.

By the end of this guide, you'll know how to turn a pile of unlabeled vacation or family photos into one-line descriptions, ready to drop into a social media post, paste as alt text, or use to search your own library.

Before you start

- A free Hugging Face account (you can sign up with an email, or sign in with a Google or GitHub account if you have one already).
- A modern browser on a laptop or desktop — Chrome, Edge, Firefox, or Safari all work. Phones work too, but a bigger screen makes uploading easier.
- 5–15 photos saved on your device, ideally as JPEG or PNG and under ~5 MB each so they upload quickly.
- Roughly 10 minutes for your first batch.

Step 1 — Create a free Hugging Face account

Open your browser and go to **huggingface.co**. Look for a **Sign up** button or link, usually near the top-right of the page. Type your email, pick a password, and confirm your inbox when the verification email arrives.

A free account lets you try any of the public demos. The page should land on your profile or the home feed once you're done.

💬 Example: the verification email is usually titled something like "Confirm your Hugging Face account".

If it looks different: if you already have an account, click **Log in** instead of Sign up — and if your browser saved an old login, you may skip straight to the homepage.

You'll know it worked when you see your profile picture or initials in the top-right corner of the page.

Step 2 — Find an image captioning Space

From the top menu, click **Spaces** — think of a Space as a small web app built by someone in the Hugging Face community. In the search bar on the Spaces page, type **image captioning** and press Enter. You'll see a list of community-built demos; look for ones with clear names related to image captioning (if it's available, you might find options with names like **BLIP Image Captioning** or **ViT-GPT2 Image Captioning**).

Don't worry about picking the "best" one. Most of these demos do the same basic job: read your photo and write one sentence about it.

💬 Example: searching for "BLIP image captioning" may return a Space built by Salesforce, which has been one of the popular captioning demos on the site.

If it looks different: if the search bar isn't obvious, try the **Filters** button, choose the **Computer Vision** category, and scroll until you spot a captioning demo.

You'll know it worked when you land on a page whose title mentions "Image Captioning" and shows an upload area or a button that says something like "Upload" or "Drop an image here".

Step 3 — Open the demo and wait for it to load

Click on the chosen Space to open it. The first time you visit, the demo may take 10–30 seconds to wake up — this is normal, since free demos often go to sleep when no one is using them. A progress bar, a spinner, or words like "Running" or "Starting" usually appear while it loads.

Once it's ready, you'll see an upload area and, on most Spaces, a short blurb explaining what the model does.

💬 Example: a typical blurb is "A short image captioning model based on BLIP, fine-tuned on COCO" — COCO is just the name of a popular photo dataset the model learned from.

If it looks different: if you only see a gray box for more than a minute, refresh the page once. If it still won't load, try a different captioning Space from the search results in Step 2.

You'll know it worked when the upload area becomes active — usually it says "Drop image here" or shows a small image icon.

Step 4 — Upload your first photo

Drag a single photo from your desktop into the upload area, or click the area to open a file picker and choose your picture. The photo appears as a small preview on the page. Don't upload dozens at once — most free demos are built to caption one image at a time.

Wait a few seconds. The caption will appear in a text box or label below the image.

💬 Example: a beach sunset photo might return something like "a sunset over the ocean with orange clouds", and a family group shot might return "a group of people smiling and standing together".

If it looks different: if nothing happens after about 30 seconds, look for a **Submit** or **Run** button near the upload area — some Spaces require you to confirm the upload before they process it.

You'll know it worked when a short English sentence describing your photo appears in or near the upload area.

Step 5 — Try a few more photos to see the range

Pick 3–5 different photos — a landscape, a portrait, a pet, a city street, a close-up of food — and run them through the demo one at a time. This helps you see what the tool handles well and where it gets confused. You'll likely notice it's strong at obvious scenes (beaches, dogs, kitchens) and weaker at very specific or unusual images.

💬 Example: a clear photo of a coffee cup on a wooden table might come back as "a cup of coffee on a table" — accurate, but plain.

If it looks different: don't be surprised if the caption is generic ("a person standing outside") for busy scenes. That's the model playing it safe rather than guessing at details it can't see clearly.

You'll know it worked when you have a small list of captions that roughly match each photo's main subject.

Step 6 — Copy and store your captions

Highlight the caption text in the demo's output box and copy it with **Ctrl+C** on Windows or **Cmd+C** on Mac. Paste it into a notes app, a Google Doc, or a spreadsheet right next to the photo's filename. Repeat for each photo. There's no built-in batch export — the work is mostly copying and pasting one at a time.

💬 Example: a simple spreadsheet column might read `beach_sunset.jpg → "a sunset over the ocean with orange clouds"`.

If it looks different: if the caption text isn't selectable, right-click the output area and choose **Copy** from the menu instead.

You'll know it worked when you have a tidy list of one-line descriptions you can reuse later.

Step 7 — Use your captions in real life

Now do something useful with them. For social media posts, paste a caption under your photo and tweak the wording if you want more personality. For alt text (the short description screen readers use for images on websites), paste it directly into the image's alt-text field. For organizing your own photo library, drop the captions into a folder named "_captions.txt" so you can search them later.

💬 Example: a LinkedIn post becomes "Wrapped up a long week with a walk on the beach 🌅 — a sunset over the ocean with orange clouds. Felt good to breathe."

If it looks different: alt-text fields look slightly different on every site — on a WordPress blog it's under "Image attributes", on X or Twitter it's the small alt field you fill in while uploading.

You'll know it worked when you stop staring at unlabeled photos thinking "what was this one again?"

Common mistakes

- **Uploading giant RAW files straight from a camera.** Free demos often struggle with files over ~10 MB, and some reject RAW formats entirely. Convert to JPEG first, or resize using your phone's built-in photo editor.
- **Expecting one upload to caption a whole album at once.** Most demos only handle one image at a time. Captioning 200 vacation photos means 200 separate uploads — plan for it to take a while, or pick a few favorites first.
- **Trusting the caption blindly.** Image captioning AI can mislabel things — it might call a labrador a "golden retriever" or skip a person in the background. Always glance at the caption before publishing or pasting it as alt text.

What this means for you

- **In everyday life:** captioning a personal photo for the first time feels like giving your camera a tiny voice. You can finally find that one sunset photo from three summers ago by searching the captions.
- **For your business or work:** if you run a blog, a small shop, or any website with lots of photos, alt text is something search engines and accessibility tools really do look at. Auto-captioning gives you a starting point in seconds — you then edit for accuracy and tone.
- **If you're just getting started:** skip the jargon and just play. Upload one photo of your pet, laugh at whatever the AI writes, and you've already learned the basics.

Try it now

Pick one favorite photo from your camera roll, open huggingface.co in your browser, search for "image captioning" in Spaces, upload the photo, and copy the caption. Two minutes, one photo,

one sentence — that's the whole loop.

How to remove a photo background cleanly with Meta AI on your phone

7 steps · Read online

> **In short:** Open Meta AI in WhatsApp, Messenger or Instagram, upload a photo, and ask it to "remove the background." Meta AI will usually generate a transparent version you can save as a PNG. You'll get a clean cutout in about a minute, though fine details like frizzy hair or sheer fabric can still need a touch-up.

You already chat with Meta AI in WhatsApp, Messenger or Instagram. This guide adds one more thing you can do there: turn any photo into a clean cutout (the photo with the background taken away) and drop it onto a new background, all from your phone, no Photoshop, no extra app. The whole workflow takes about five minutes once you've done it once.

Before you start

- **The app:** Meta AI built into **WhatsApp, Messenger, or Instagram** on your phone (iPhone or Android). If you can already chat with Meta AI in any of these, you're set.
- **A photo on your phone:** JPEG or PNG, ideally with a clear subject (a person, a product, a pet) and decent lighting. Busy, low-contrast backgrounds are the hardest to cut.
- **Roughly 5–10 minutes** for the first try, less after that.
- **Free.** No paid plan is required for the workflow below. Meta may rate-limit (cap how many times you can ask in a short window) on busy days, so have a backup photo ready.
- **What Meta AI can and can't do here:** it can usually remove a simple background in one go, but it is not pixel-perfect. Fine hair, glasses, and see-through materials can leave a faint halo (a soft glow around the edges of the cutout). It's great for social posts and product shots, less great for printing at poster size.

Step 1 — Open a Meta AI chat

Tap the **circular blue/purple Meta AI icon** — in WhatsApp it's usually at the top of your Chats tab or behind a **blue "AI" button** in the bottom-right; in Messenger it's the top of the search bar; in Instagram it's in the search/DMs area. If you don't see it, look for a **sparkle or four-pointed star icon** or check the **app's top menu**. The chat screen opens, showing a short welcome message and the message box at the bottom. You'll know it worked when the box at the bottom says something like "Message Meta AI" or shows a send arrow.

Step 2 — Upload the photo you want to cut out

Tap the **paperclip or "+" icon** next to the message box to attach a file. From the menu that appears, pick **"Gallery"** (Android) or **"Photo & Video Library"** (iPhone) and choose the

photo. The image appears inside the chat as a small preview above the message box. If your app doesn't show a paperclip, look for a **camera icon** instead — the option is the same. Don't send the message yet; the prompt you add in the next step shapes what Meta AI does with the image.

💬 Example: choose a clear, well-lit photo of a friend standing in front of a plain wall, or a product on a tabletop.

You'll know it worked when the photo preview is sitting above the message box, ready to go with your next sentence.

Step 3 — Write the prompt that asks for the cutout

Tap the message box and type a short, direct instruction. Be specific: say "remove the background" rather than just "edit this." If you already know where you want to use the image, mention that too — it helps the model (the AI "brain" that reads your words and produces the answer) pick a sensible default background color.

💬 Example: `Remove the background from this photo and give me the subject on a transparent background.` Or, if you'd like a new background in the same step: `Remove the original background and place the person on a soft white studio background.`

Tap **send** (the paper-plane or arrow icon). Meta AI will usually reply within 10–30 seconds with a new version of your image. You'll know it worked when a new picture appears in the chat that looks like your subject on a flat color or transparent-looking background.

Step 4 — Check the edges before you save

Pinch to zoom in on the edges around the subject — especially around hair, fingers, glasses, and anything thin. Look for a faint color halo (a thin glowing outline) where the old background is bleeding through, or a slightly hard "stamped" edge. If the cut is clean, move to Step 5. If you see issues, stay in the same chat and ask for a fix in plain words.

💬 Example: `Zoom in around the hair — I can still see a faint gray halo. Please clean up the edges.` Or: `The fingers on the right hand look chopped. Please redo with finer edge detail.`

You'll know it worked when the second image shows noticeably cleaner edges than the first, with no obvious halo or missing fingers.

Step 5 — Save the new image to your phone

Long-press the new image Meta AI produced. A small menu appears with options like **"Save," "Save to Photos,"** or **"Download."** Tap the save option. The image lands in your phone's

Photos or Gallery app, usually in a folder called **Recents** or **Downloads**. If long-press doesn't show "Save," look for a **three-dot ⋯ menu** on the image instead, or tap the image to open it full-screen and then tap the **download/save icon** in the corner.

You'll know it worked when the new image appears in your Photos app as its own file, separate from the original you uploaded.

Step 6 — Drop the cutout onto a new background

Open any simple editor on your phone — the built-in **Photos** app on iPhone, or **Google Photos** on Android both work. Open the saved cutout, then tap **Edit**, and look for an option like **"Add Photo"** or **"Overlay / Sticker"** that lets you put a second image behind the first. Pick your new background — a sunset you took, a logo, a solid color. Resize and reposition, then save.

💬 Example: a beach sunset from your camera roll, so your friend looks like they're standing on vacation.

You'll know it worked when the final saved image shows your original subject, with clean edges, sitting naturally on the new background you chose.

Step 7 — Keep your reusable prompt in one easy-to-find place

Don't retype the prompt every time. Save the working version from Step 3 in your phone's **Notes app** or in the **Notes** feature of WhatsApp/Messenger itself, then copy-paste it next time. For even faster reuse, ask Meta AI in the same chat: `Save this as a reusable prompt I can paste next time.` Meta AI will usually reply with a tidy, copy-ready version you can keep.

💬 Example: `Remove the background from this photo and give me the subject on a transparent background. Keep the subject exactly the same — only the background should change.`

You'll know it worked when, next time, you can paste the saved prompt, attach a photo, and get a clean cutout in under a minute.

Common mistakes

- **Asking vaguely.** "Edit my photo" is too open. Always name the change: "remove the background," "place the subject on white," "tighten the edges." A specific instruction gets a specific answer.
- **Uploading a tiny, blurry, or low-contrast photo.** If the subject and the background are similar colors (a black dog on a dark couch), Meta AI will struggle. Pick a photo with a clear difference between subject and background.
- **Saving the wrong image.** The chat shows both your original upload and the new generated image. Long-press the *new* one, not the one you sent. If you saved the original

by mistake, just repeat Step 3 and save the reply.

What this means for you

- **In everyday life:** Clean cutouts are perfect for WhatsApp or Instagram stickers of your pet, a quick profile picture with a softer background, or a fun "I went to Paris" photo where you actually went to the kitchen. Five minutes of fiddling beats a desktop app you don't have time to learn.
- **For your business or work:** If you sell a few products on a side hustle, you can knock out consistent product photos for a catalog or marketplace listing without paying a designer. The edges won't be magazine-perfect, but they're good enough for a small storefront and free.
- **If you want to go further:** Try a two-step prompt in one chat — first remove the background, then ask `Now place the subject on a marble countertop with soft daylight from the left.` The model remembers the previous image in the same conversation, so each new instruction builds on the last one instead of starting over.

Try it now

Open Meta AI in WhatsApp, Messenger or Instagram right now. Pick a photo of a person or object that is already on your camera roll, paste this prompt, and send:

💬 `Remove the background from this photo and give me the subject on a transparent background.`

Save the result to your phone. That's a complete, reusable workflow you can do again in under a minute next time.

Create a multi-room music playlist with Alexa+

7 steps · Read online

> **In short:** With Alexa+ you can set up a voice routine so a single phrase — like "Alexa, play my dinner playlist" — starts the same music in every room at once. It uses the Alexa app's Routines feature.

Imagine saying "Alexa, play my dinner playlist" and the same music starts in the kitchen, lounge and bedroom at once. You'll learn how to set up one playlist that plays across multiple rooms with a single voice command.

Before you start

- An **Amazon account** (free).
- The **Alexa app** installed on your phone or tablet (Android or iOS).
- At least one **Amazon Echo device** (like an Echo Dot, Echo Show or Echo Studio) in each room.
- A **music streaming service** already linked to Alexa (Amazon Music, Spotify, Apple Music, etc.).
- Roughly **15 minutes** of your time.

Step 1 — Open the Alexa app and go to Devices

Open the Alexa app on your phone or tablet. At the bottom of the screen you'll see a row of icons — look for the one that looks like a **house outline** (this is the Devices section). Tap it. You should see a list of all your Echo devices, often grouped by room.

You'll know it worked when the screen shows your rooms and devices.

Example: Look for the house-shaped icon at the bottom of the Alexa app and tap it.

Step 2 — Pick a room to start

Choose the room where you want the music to begin. Tap the room name (for example, "Kitchen"). Inside that room view, look for a menu option — this is typically a **three-dot menu (⋯)** in the top-right corner, though on some app versions it may appear as a **gear icon** or a **menu with three lines** instead.

You'll know it worked when a menu slides up or opens from the bottom of the screen.

Example: Open the Kitchen room and look for the ⋯ menu in the top-right corner.

Step 3 — Find the routine option in the menu

From the menu that just appeared, look for an option to add this room to a routine — the exact wording may be **"Add to Routine"** or something similar depending on your app version. Tap it.

If you're not sure which option it is, look for any menu item that mentions "routine" or "routine" in English — this is the feature that links your voice phrase to your music action.

You'll know it worked when a routine setup screen opens.

Example: Tap **Add to Routine** (or the equivalent option in your app version).

Step 4 — Choose a voice trigger

In the routine setup screen, look for the option to set a **voice trigger** — on many versions this is labeled **"Voice"** and shown with a microphone icon. Tap it. A text box appears where you can type the phrase you want to say to Alexa to start the playlist.

If the voice option isn't visible, look through the list of available trigger types for something related to speaking or voice commands.

You'll know it worked when you see a text box labeled something like "Enter your phrase" and a way to confirm and move to the next step.

Example: Tap the voice trigger option and then tap inside the text box.

Step 5 — Type your wake-up phrase

Type a short phrase you'll say to Alexa, for example: **"Alexa, play dinner playlist"**

Confirm the phrase and move to the next step. Alexa may show a summary of the phrase you've entered.

You'll know it worked when the summary confirms your exact phrase.

Example: Type "Alexa, play dinner playlist" and confirm it.

Step 6 — Add the music action

In the routine setup, look for an action type related to music — this may be labeled **"Music"** (often with a musical note icon) or something similar depending on your app version. Tap it. A new screen opens where you can pick the playlist and the devices that should play it.

If you don't see a music option right away, look through the list of available actions for anything about playing audio or music.

You'll know it worked when you see options to choose a playlist and to choose which devices or rooms should play it.

Example: Tap the music action and then tap the option to choose a playlist.

Step 7 — Pick your playlist and devices

Select the playlist you want to play everywhere. Then choose all the rooms where you want the music to play (for example, Kitchen, Lounge, Bedroom) — look for a device or room picker. Save the routine.

You'll know it worked when the routine summary shows your chosen playlist and all the rooms you selected.

Example: Select "Dinner Vibes" playlist and tick Kitchen, Lounge and Bedroom.

Common mistakes

- **Wrong room selected** — If the music only plays in one room, double-check Step 2 and make sure you picked the correct room before saving the routine.
- **Phrase already used** — If Alexa says the phrase is taken, add a word like "my" or change the wording slightly (for example, "Alexa, start dinner playlist").
- **Device not selected** — If a room is missing, go back to Step 7 and make sure every room is ticked before saving.

What this means for you

- **In everyday life:** Set up a "Good morning" routine that plays your favorite news podcast in the bedroom and kitchen at the same time.
- **For your business or work:** Create a "Focus time" routine that plays lo-fi beats in every meeting room whenever you say "Alexa, start focus time".
- **If you're just getting started:** Start with one room and one playlist; once it works, add more rooms in the same routine.

Try it now

Open the Alexa app, go to Devices, pick a room, tap the menu icon → look for a routine option → choose voice trigger → type your phrase → add a music action → pick your playlist and devices → Save. Say the phrase out loud to test it.

Sharpen Your Photos with AI Using Upscayl

6 steps · Read online

Do you have a cherished photo that's just a little bit blurry, or an old snap you wish looked sharper and ready for a larger print? This guide will show you how to use the free AI image super-resolution application, Upscayl, to quickly enhance your images. By the end, you'll be able to give your photos a professional, crisp finish, making them poster-worthy in minutes.

Before you start

- **Install the free Upscayl application:** You'll need to download and install Upscayl on your computer. It's an open-source tool, which means it's developed by a community and freely available for use.
- **A computer with a dedicated graphics card (GPU) is recommended:** While Upscayl can run on most systems, a graphics card (like an NVIDIA, AMD, or Intel Arc GPU) will significantly speed up the processing of your images, making the enhancement much faster.
- **A photo to enhance:** Choose a digital image that's slightly blurry or low-resolution that you'd like to improve.
- **Time:** Allow around 5-7 minutes to follow the steps and enhance your first image. Upscayl is free to use with no hidden costs.

Step 1 — Launch Upscayl

First, open the Upscayl application on your computer. Look for its distinctive icon on your desktop, in your Start menu (Windows), Applications folder (macOS), or program launcher (Linux), and double-click it. When it launches, you'll typically see its main window, often featuring a clean interface with a central area to load images and various options organized on the side. This is your starting point for transforming your photos. **If it looks different:** If it's your very first time launching Upscayl, you might see a quick setup or welcome screen, possibly asking to confirm default file paths or to download initial AI models. Follow any clear on-screen prompts to complete this setup before proceeding. **You'll know it worked when:** The Upscayl application window fully appears on your screen, indicating it's ready for you to select an image and begin the enhancement process.

Step 2 — Load Your Image

Now it's time to bring in the photo you want to improve. In the Upscayl window, you'll usually find a prominent button labeled "Open Image" or a clear "Drag and Drop" area in the center of the screen. Click the "Open Image" button, and your computer's file browser window (like File

Explorer on Windows or Finder on macOS) will pop up. Navigate to where your blurry photo is stored on your computer, select it with a click, and then click "Open." Your chosen image will then appear in the main Upscayl window, usually displayed as a clear preview. **If it looks different:** If the image doesn't load immediately, an error appears, or the preview area remains empty, ensure your photo is in a common format like JPG, PNG, or WEBP. You might also try restarting the application and attempting to load the image again. **You'll know it worked when:** Your selected photo is clearly visible within the Upscayl application's main preview area, confirming it has been loaded successfully and is ready for enhancement. 💬 Example: Select a file named `old_family_photo.jpg` from your `Pictures` folder.

Step 3 — Choose an AI Model

With your photo loaded, you need to tell Upscayl which "recipe" or AI model to use for the enhancement. Think of an AI model as a specialized algorithm (a set of computational instructions) that the AI follows to intelligently reconstruct and add detail to your image. Each model is trained differently and can produce slightly varied results depending on the original image. On the left side of the Upscayl window, you'll see a "Models" section with a list of AI options like "Real-ESRGAN," "Remacri," or "Ultramix_Balanced." Click on a model to select it; for general photos, "Real-ESRGAN" or "Ultramix_Balanced" are excellent starting points known for good all-round performance. **If it looks different:** If you don't see many models listed, or options are greyed out, you might need to download additional models within Upscayl's settings. Look for a "Download Models" button or an option to manage models in the application's configuration or preferences menu. **You'll know it worked when:** The name of the model you clicked is clearly highlighted or appears as the active selection, indicating that Upscayl will use this specific AI algorithm for processing.

Step 4 — Select Upscale Factor

Next, you'll decide how much larger you want your image to become. The "upscale factor" determines how many times the image's width and height will be multiplied, significantly increasing its total number of pixels and its overall resolution. For example, a 2x factor makes the image twice as wide and twice as tall (quadrupling the total pixels), while a 4x factor makes it four times wider and four times taller (sixteen times the total pixels). Below the models section, you'll typically find options for "Upscale Factor," often presented as buttons like 2x, 3x, 4x, or sometimes a custom size. Select "4x" for a significant improvement in detail and resolution for most images, providing a good balance of quality and processing time. **If it looks different:** Some AI models are specifically designed for particular upscale factors, so if an option like "4x" is greyed out, try selecting "2x" or "3x" instead. Alternatively, you might need to switch to a different AI model (as described in Step 3) that supports a higher factor. **You'll know it**

worked when: The "4x" option (or your chosen upscale factor) is clearly selected or highlighted in the Upscayl interface, confirming your desired output size.

Step 5 — Start the Enhancement

Once your photo is loaded and you've selected your preferred AI model and upscale factor, it's time to let the AI do its work. Look for a prominent button, usually positioned at the bottom right of the Upscayl window, labeled "Upscayl" or "Enhance." Click this button to begin the super-resolution process. Depending on your computer's processing power (especially if you have a dedicated graphics card like an NVIDIA RTX GPU) and the original image size, this can take anywhere from a few seconds to a minute or two. You might see a progress bar, a spinning icon, or a message indicating that the AI is actively processing your image. **If it looks different:** If the "Upscayl" or "Enhance" button appears greyed out or inactive, double-check that you've successfully loaded an image and made definite selections for both the AI model and the upscale factor. All required options must be chosen to activate the enhancement button. **You'll know it worked when:** The progress indicator disappears, and Upscayl either displays the enhanced image in a side-by-side view or automatically prompts you to save it.

Step 6 — Review and Save Your Enhanced Image

After the AI has finished processing, you'll see your newly enhanced image. Take a moment to compare it with the original; Upscayl often provides a side-by-side comparison or an easy toggle to switch between the two. Notice the improved sharpness, clarity, and detail in areas that were once blurry or pixelated. If you're pleased with the result, Upscayl will usually prompt you with a "Save As" dialogue box, which is your computer's standard file explorer. Choose a location on your computer to save your crisp new poster-ready photo, and give it a descriptive name so you can easily find it later. **If it looks different:** If a "Save As" prompt doesn't automatically appear, look for a "Save" or "Export" button within the Upscayl application itself. Some versions might also automatically save the enhanced image to a default folder (like "Pictures" or "Downloads") if no prompt is explicitly given. **You'll know it worked when:** You have a new, sharper, and higher-resolution version of your photo saved as a separate file on your computer, ready for use. While Upscayl can significantly enhance your photos, it cannot perfectly invent details from an extremely blurry or tiny original image that were never present in the first place, but it can significantly improve clarity.

Common mistakes

- **Not downloading AI models:** Upscayl relies on various specialized AI models to perform its magic, and not all are included in the initial application installation.
    - **The fix:** If you see few or no options in the "Models" section (Step 3), go into Upscayl's settings or preferences and look for a "Download Models" or "Manage

Models" option to acquire more AI algorithms.

- **Choosing the wrong AI model:** Different models are optimized for different image types (e.g., general photos, illustrations, specific art styles). Using a model designed for, say, anime on a family portrait might give unnatural or distorted results.
    - **The fix:** Experiment! If an image looks strange after processing, try re-loading the original and selecting a different model. "Real-ESRGAN" or "Ultramix_Balanced" are generally good choices for everyday photos and realistic images.
- **Expecting miracles from extremely low-quality sources:** While AI super-resolution is powerful, it cannot invent information that simply isn't there in the original image. An image that is completely out of focus, heavily pixelated, or tiny from the outset will still have significant limitations even after enhancement.
    - **The fix:** Aim for input photos that are only moderately blurry or low-resolution. The better the starting material (even if just slightly), the more impressive and natural the AI's final result will be.

What this means for you

- **In everyday life:** You can rescue cherished old photos that are a bit blurry, prepare family pictures for larger prints or digital display, or make your social media images stand out with professional-level sharpness and detail.
- **For your business or work:** Enhance product photos for e-commerce listings, sharpen images for presentations or marketing materials, or improve scans of documents and graphics to make them clearer and more legible for your audience and colleagues.
- **If you're just getting started:** Pick one of your favorite slightly blurry photos – perhaps a holiday snap or an old family picture – and follow this guide. Start by selecting the "Real-ESRGAN" model and a 4x upscale factor. Seeing the transformation firsthand is the best way to understand the power of AI image enhancement.

Try it now

Grab a photo that needs a little boost, launch the free Upscayl application, and follow these steps to give it a fresh, crisp look in the next few minutes.

How to create and edit images in ChatGPT

6 steps · Read online

By the end of this guide, you'll be able to type a simple description and get a unique image from ChatGPT, then adjust it to suit your needs – all without needing any complicated design software. This is perfect for anyone who wants a quick visual for a blog post, a social media banner, or a personal creative project. Just remember, while ChatGPT can generate fantastic images, it can't perform complex photo manipulation like a dedicated graphic design program.

Before you start

- **An active OpenAI account:** You'll need to sign up at `openai.com` if you don't already have one.
- **Access to ChatGPT:** You can use the ChatGPT web interface (in your browser) or the official ChatGPT mobile or desktop app. Image generation is available to both free and paid users, though free users may have usage limits.
- **A basic internet connection:** ChatGPT runs online, so no extra software downloads are required, but a stable internet connection is essential.
- **Rough total time:** This guide should take about 5-10 minutes to follow.

Step 1 — Start a new conversation for image generation

Open ChatGPT and begin a fresh conversation. Image generation is built directly into ChatGPT's main interface, so you don't need to access any special mode or tool — simply type your image request in the chat.

What you see: The standard ChatGPT chat interface, with a text input box at the bottom of the screen. **Action:** Either start a new chat (click the "New chat" button) or use an existing conversation. There's no need to find a special image mode or custom GPT — just type your description directly. **What happens next:** ChatGPT will recognize when you want to generate an image based on your prompt and will process your request. **If it looks different:** If ChatGPT responds with text instead of an image, make sure your prompt clearly describes what you want to see (e.g., "Generate an image of..." or "Create a picture showing..."). The more specific you are, the better. 💬 **Example:** "You click 'New chat' and immediately type: 'Generate an image of a peaceful cottage garden.'" **You'll know it worked when:** You see a message like "Generating image..." or a progress indicator, followed by the image appearing in the chat.

Step 2 — Write your image description (prompt)

Now it's time to tell ChatGPT what picture you want. The instruction you give the AI is called a **prompt**, which is like giving a briefing to a human artist. Be specific about the main subject, desired style, colors, and any particular details.

What you see: A clear text input field. **Action:** Type a detailed description of the image you want ChatGPT to create. Mention key elements, the overall mood, color palette, and artistic style. You can start with phrases like "Generate an image of..." or "Create a picture showing..." **What happens next:** Once you press Enter or send your message, ChatGPT will process your request and begin generating the image. This usually takes a few seconds. **If it looks different:** If ChatGPT returns text instead of attempting to create an image, double-check that your prompt is clear and explicitly asks for an image. Try starting with "Generate an image of..." **Example:** `Generate an image of a peaceful, sun-drenched cottage garden in the American countryside, filled with native wildflowers, a small birdbath, and a maple tree in the background, painted in an impressionistic art style with soft, warm colours.` **You'll know it worked when:** You see a message like "Generating image..." or a progress bar indicating that the AI is working on your visual.

Step 3 — Review the generated image(s)

After a short wait, ChatGPT will present the image(s) it created based on your prompt. Take a moment to look at them closely.

What you see: One or more small images (thumbnails) displayed directly within your chat conversation. **Action:** Click on an image to view it in full size and examine the details. Think about what you like and what could be improved. **What happens next:** The image will expand, usually showing more detail. You might also see options to download or provide feedback. **If it looks different:** Sometimes ChatGPT might generate a single image instead of multiple options, or occasionally, if your prompt was unclear, it might offer text suggestions for refinement instead of images. If the image seems completely off, it might be worth going back to Step 2 and rewriting your prompt entirely. **Example:** "You notice the maple tree looks more like an oak. You also see a small download icon, usually a downward arrow or three dots, when you hover over the image." **You'll know it worked when:** Your chat window displays one or more visual results of your prompt, and you can click them to view them larger.

Step 4 — Refine with an edit request

It's rare to get a perfect image on the first try. Luckily, you can easily ask ChatGPT to make adjustments. You don't need to rewrite the whole prompt; just tell it what to change.

What you see: The previously generated image(s) still visible in the chat, and the active text input box below. **Action:** Type a new, concise prompt requesting specific changes to the

previous image. Start with a phrase like "Edit the image to..." or "Make the tree taller..." **What happens next:** ChatGPT understands you want to modify the existing image and will generate a new version incorporating your feedback, usually keeping the overall composition while applying your tweaks. **If it looks different:** If ChatGPT generates an entirely *new* image that doesn't seem related to the previous one, your refining prompt might have been too broad, or it might have lost the context. Try being more direct, e.g., "Adjust *the last image* by..." **Example:** `Edit the image to make the maple tree in the background taller and add a small, rusty tin shed near the cottage.` **You'll know it worked when:** A new image appears in the chat, showing the requested adjustments applied to the original concept.

Step 5 — Save your final version

Once you're happy with the results, it's time to save your unique creation to your device.

What you see: The final, refined image displayed in the chat. **Action:** Hover your mouse over the image (on desktop) or tap the image (on mobile) to reveal a download icon. This is often a downward-pointing arrow or a three-dot menu icon that opens more options, including 'Download'. Click or tap this icon. **What happens next:** Your browser or app will prompt you to save the image file to your computer or phone's downloads folder. You might be able to rename the file before saving. **If it looks different:** Some versions might require a right-click (desktop) or a long-press (mobile) on the image itself to bring up a context menu with a "Save image" or "Download" option. **Example:** "You click the downward-pointing arrow in the top-right corner of the image, and a file named 'image.png' appears in your browser's download bar." **You'll know it worked when:** The image file is successfully saved to your device, and you can open it with your usual photo viewer.

Step 6 — Understand usage guidelines

As you create images, it's important to be aware of OpenAI's usage policies. These rules help keep the platform safe and responsible.

What you see: No specific visual change, but your actions adhere to the platform's terms. **Action:** Be mindful of the content you ask ChatGPT to create or edit. Avoid prompts for anything that could be considered explicit, harmful, misleading, or illegal. **What happens next:** By following these guidelines, you ensure your images are generated successfully and your account remains in good standing. **If it looks different:** If your prompt violates the policy, ChatGPT will usually refuse to generate the image and might give you a message explaining why. Sometimes it might just generate a blank or highly abstract image. **Example:** "You avoid asking for images of specific copyrighted characters or content that promotes hate speech." **You'll know it worked when:** Your image is generated without any warning messages or content refusal from ChatGPT, allowing you to use it freely (within its licensing terms).

Common mistakes

- **Vague or short prompts:** Providing too little detail often leads to generic or unexpected results that don't match your vision.
    - **Fix:** Treat ChatGPT like an artist's brief. Include the subject, style, colors, mood, and any specific elements. For example, instead of "a cat," try "a fluffy ginger cat sleeping on a sunny windowsill, drawn in a cartoon style with vibrant colors."
- **Ignoring policy warnings:** Attempting to bypass content filters can lead to image generation being blocked or, in persistent cases, account warnings.
    - **Fix:** Always respect the content policies. If an image is flagged, try rephrasing your request to be more appropriate, or choose a different subject.
- **Not being specific enough about edits:** When refining an image, vague requests like "make it better" may not produce the desired result.
    - **Fix:** Be specific about what you want to change — for example, "make the background darker" or "add more flowers in the foreground."

What this means for you

- **In everyday life:** Easily create custom invitations for a family gathering, design a unique background for your phone, or illustrate a personal story without needing to be an artist.
- **For your business or work:** Quickly generate social media graphics, blog post headers, unique presentation slides, or marketing visuals, saving time and potentially design costs.
- **If you're just getting started:** The easiest first move is to open ChatGPT and type a one-sentence description of something you love – like "Generate an image of a cute koala hugging a eucalyptus leaf, in a colorful watercolour style."

Try it now

Spend the next two minutes opening ChatGPT, and typing a simple prompt such as **"Generate an image of a cup of coffee on a wooden table, sketch style"**. Download the result and see how fast a quality picture can appear.

How to Sort and Clean Up a Messy Phone Photo Library with Gemini

6 steps · Read online

> **In short:** Gemini can look through photos you upload, flag near-duplicates, group similar shots, and pull readable text out of old screenshots — all without touching your library. You stay in control: Gemini suggests, you decide what to delete.

By the end of this guide, you'll know how to use Gemini as a careful second pair of eyes for your photo mess — the kind where you have seven nearly-identical sunset shots and a folder full of screenshots you can never search. This is for anyone whose phone gallery has gotten out of hand and wants a low-risk way to start sorting through it.

Before you start

- A free Google account (any Gmail address works). The free Gemini tier is enough for small batches; a paid Google AI plan lets you upload more photos at once and keeps longer conversations in memory — check Google's official site for current plan options.
- The Gemini app on your phone, **or** a browser pointed at the Gemini website.
- About 20–30 minutes for your first batch.
- A folder of photos you actually want to clean up. Pick one small group to start with — don't try the whole library on day one.

Step 1 — Open Gemini and start a fresh chat

Tap the Gemini app icon, or visit the Gemini website in your browser. If you're not signed in, sign in with your Google account. Look for a **"New chat"** button — usually a plus sign or a small pencil icon at the top of the chat list — and tap it. A blank conversation opens. Starting fresh keeps the AI focused on the photos you're about to share, not on a leftover question from yesterday's weather check.

If you don't see "New chat," try the **three-line menu** (sometimes called a hamburger icon) in the top corner — the option to start over is usually inside it. On some devices the same button looks like a small **square with a pencil** inside instead.

💬 Example: tap **New chat**, then wait for the empty text box to appear.

You'll know it worked when you see a clean text box with no messages above it and a paperclip or upload icon somewhere nearby.

Step 2 — Gather a small batch of photos to upload

Don't try your whole library on day one. Pick **10 to 20 photos** that feel like they might overlap — the seven sunset shots from your last trip, the burst-mode photos from your kid's birthday, the screenshots you saved last month. In your phone's gallery, tap **Select** (or hold one photo until checkboxes appear), choose the photos, then tap **Share**, and pick **Gemini** from the list of apps. On a computer, drag the photos straight into the chat box.

If the share menu doesn't show Gemini, copy the photos to your clipboard and paste them into the chat. On older Android phones the upload icon looks like a **paperclip** 📎; on iPhones it's often a small **plus sign** + at the bottom of the text box.

💬 Example: open your gallery, select the 8 beach photos from last weekend, tap **Share** → **Gemini**.

You'll know it worked when the photos appear as small thumbnails in your message and Gemini replies with something like "Got it — I can see 8 photos."

Step 3 — Ask Gemini to flag duplicates and similar shots

Now comes the useful part. In the same chat, type a short, clear question about the photos you just shared. The clearer your instruction, the better the answer. You can ask Gemini to find photos that look almost the same, group them by scene, or just describe what makes each one different. It can read the image content directly — the colors, the framing, who's in the shot — and reason about which ones are worth keeping.

If Gemini's response is too generic, follow up with a more specific question. After a vague "they all look similar," try "Which one is sharpest?" or "Which one best shows both people smiling?" If the free tier says it's reached a limit, work in even smaller batches, or look into a paid plan for larger uploads.

💬 Example: "Here are 8 photos from the same sunset. Can you group any that look like near-duplicates, and tell me which one of each group is the sharpest or best composed?"

You'll know it worked when Gemini gives you a numbered list or short table, like "Group A: photos 1, 3, and 5 are nearly identical — keep #3."

Step 4 — Pull searchable text out of an old screenshot

Screenshots are useless until you can search them. Take a single screenshot — a recipe, an address, a confirmation number, a Wi-Fi password you saved months ago — and share it with Gemini the same way you did in Step 2. Then ask Gemini to read the text out of it and rewrite it in a clean, searchable form. This works for receipts, WhatsApp messages, store hours, anything that ended up trapped inside an image.

The text won't be added back to your photo library automatically — you'll copy it out yourself in Step 6. And Gemini can misread small or blurry text, so for anything important like medical info or legal documents, always double-check the original screenshot.

💬 Example: "Please transcribe the text in this screenshot. Format it as: store name, date, total amount, and any reference number."

You'll know it worked when Gemini replies with readable text you can copy and paste, and that text matches what you actually see in the screenshot.

Step 5 — Ask for a simple cleanup plan

Once Gemini has analyzed a batch, ask it to summarize what it found into a short, action-friendly list. You're looking for a "do this, then this, then this" answer — not a long essay. The point is to turn Gemini's observations into a checklist you can actually use in your phone's gallery. This is also where Gemini's "context window" (the amount of text and images it can keep track of in one conversation) matters: if your batch is too big, it may start forgetting earlier photos.

If the plan is too long, ask for the top three actions only. If it's missing something obvious, ask a follow-up like "Did you see any photos that look like receipts or notes I should transcribe?"

💬 Example: "Based on what you saw, give me a short checklist of which photos I can safely delete and which I should keep. Keep it under 5 bullet points."

You'll know it worked when you have a short, scannable list that feels doable in under five minutes.

Step 6 — Take the plan into your Photos app

Now switch to your phone's gallery or Google Photos. Open the same batch of photos. Use Gemini's checklist as your guide: delete the ones it flagged as duplicates, keep the sharpest of each group, and save the transcribed text somewhere searchable — a notes app, a contact entry, or a dedicated "Screenshots I've read" album. This is the step where the cleanup actually happens, and it's the step only you can do.

For an extra safety net, turn on your gallery's recently-deleted folder so any mistake can be undone within 30 days. On Google Photos it's usually at **Library** → **Trash**; on iPhone it's at **Albums** → **Recently Deleted**.

💬 Example: open the 8 beach photos, delete 5 based on Gemini's notes, star the best one, and paste the transcribed receipt text into your notes app.

You'll know it worked when your batch is smaller, the photos you kept are the ones you actually want, and the old screenshots have become searchable text.

Common mistakes

- **Uploading your whole library at once.** Gemini gets less accurate with very large batches, the conversation becomes hard to follow, and you may hit free-tier limits fast. Work in groups of 10–20.
- **Letting Gemini "decide" what to delete.** Gemini can suggest, but it can't actually delete anything from your phone — and it shouldn't. Always make the final call yourself, in your own gallery, with your finger on the button.
- **Trusting transcribed text blindly.** Especially for numbers, dates, and addresses — Gemini reads most of it right, but a single wrong digit on a confirmation code can cost you a refund. Always glance back at the original screenshot.

What this means for you

- **In everyday life:** That folder of 200 mystery screenshots from the last three years can finally become a few searchable notes — addresses you can find, receipts you can expense, articles you can actually read.
- **For your business or work:** If you snap whiteboards, business cards, or reference material on the go, this same flow turns them into text you can paste into documents, spreadsheets, or your CRM.
- **If you're just getting started:** Pick **one** folder — say, screenshots from last month — and run the full six steps on just that folder. Twenty minutes later, you'll have a feel for the process before touching anything bigger.

Try it now

Open Gemini right now, share it the last five photos you took, and ask: "Which of these is the best one, and why?" You'll get a tiny but real taste of how Gemini sees your library — and you'll have deleted nothing by accident.

Safely Rename a Batch of Photos with a GitHub Copilot Script

6 steps · Read online

> **In short:** With GitHub Copilot inside VS Code, you can describe a renaming task in plain English and get a working Python script back — no coding required. The trick is to test on copies first so your originals stay safe.

By the end of this short guide, you'll have asked Copilot to write a small Python script, tested it on copies of your photos, and renamed an entire folder without typing a line of code yourself. If you have a folder of vacation shots named `IMG_4829.jpg`, `IMG_4830.jpg`, and so on, this is for you.

Before you start

- **A computer** running Windows, macOS, or Linux.
- **VS Code installed** — the free code editor from Microsoft (download from `code.visualstudio.com`).
- **The GitHub Copilot extension** installed in VS Code, with a Copilot account signed in. Copilot offers a free tier and paid plans; check the official GitHub page for current availability and which features each tier includes.
- **Python installed** on your computer (download from `python.org`). Python is the language your script will be written in — think of it as plain English instructions for your computer.
- **A folder of photos** you'd like to rename, plus enough free disk space to make a copy. The whole task takes about **15 minutes**.

Step 1 — Make a working copy of your photos

Right-click your photo folder in your file manager, choose **Copy**, then **Paste** somewhere easy to find, like your Desktop. Rename the copy to something like `photos-test`. This is the folder you'll experiment on.

On macOS, you can also select the folder and press Option-drag to duplicate it. On Windows, hold Ctrl while dragging. You'll see a second folder appear next to your originals, with all the same files inside.

If it looks different: If your photos live in cloud storage like Google Drive or iCloud, the same Copy/Paste idea works — just make sure the copy finishes syncing before you start renaming anything. Some file managers hide the "Paste" option inside an Edit menu; look there if right-click doesn't show it.

You'll know it worked when you see two folders: your untouched originals and a clearly named test copy.

Step 2 — Open the test folder in VS Code

Open VS Code, then go to **File → Open Folder…** and pick your new `photos-test` folder. VS Code's left sidebar (called the **Explorer**) now lists every photo in the folder.

A workspace in VS Code is simply the folder you opened — think of it as putting that folder "on your desk" so the editor can see everything inside it. Nothing is moved or copied; VS Code is just looking.

If it looks different: On the VS Code welcome screen, there's a big **Open Folder** button you can click instead. If the Explorer sidebar is hidden, use **View → Explorer** to bring it back.

You'll know it worked when the names of your test photos appear in the left sidebar.

Step 3 — Open the Copilot chat panel

In VS Code, open the **Copilot Chat** panel. The keyboard shortcut is **Ctrl+Alt+I** on Windows and Linux, or **Cmd+Shift+I** on macOS. A chat window appears — usually on the side of the window.

Copilot is an AI coding assistant: you describe what you want in plain language, and it suggests code. You don't have to understand the code right now — that's the whole point of having an AI helper.

If it looks different: If the shortcut doesn't open anything, look for a small chat-bubble icon in the VS Code sidebar, or use the menu **View → Copilot Chat**. Some older versions put Copilot under **File → Preferences** instead.

You'll know it worked when a chat box is visible with a blinking text cursor, ready for your message.

Step 4 — Describe your renaming task in plain English

Click into the chat box and type what you want. Be specific about the folder, the file types, and the new naming pattern. Copilot reads your message and proposes a Python script in the chat.

The more clearly you describe the rule, the better the script will be. If your rule is "rename to `vacation-001.jpg`, `vacation-002.jpg`, …", say that exactly. If you want the date taken included, ask for that too.

💬 Example: `Write a Python script that renames every .jpg and .png file in the current folder to "vacation-2024-001.jpg", "vacation-2024-002.jpg", and so on. Skip files that already match this pattern.`

If it looks different: Copilot might ask a follow-up question like "What folder should I use?" — just answer in the same chat. If it offers several code blocks, pick the one that looks closest to what you asked for and ignore the rest.

You'll know it worked when a code block appears in the chat with a Python script inside it.

Step 5 — Save the script and run it on the test folder

When Copilot shows a script, look for a small **"Insert into new file"** button above the code (or a similar option like **"Apply"** or **"Copy"**). Click it, and VS Code creates a new file. Save this file inside your test folder as `rename.py`.

Now open VS Code's built-in terminal from the top menu: **Terminal → New Terminal**. A text panel slides up from the bottom. Type `python rename.py` and press Enter.

The terminal is a text-based way to talk to your computer directly — instead of clicking icons, you type commands and read text replies.

If it looks different: If `python` isn't recognized, try `python3 rename.py` instead — that's more common on macOS and Linux. On Windows, you may need to install Python from `python.org` and tick **"Add Python to PATH"** during setup for the command to work.

You'll know it worked when the terminal shows your photos being renamed one by one, and the file list in the Explorer updates.

Step 6 — Verify, then run the same script on your real folder

Open the test folder in your regular file viewer (Finder, Explorer, or Files) and check the new names. Open a few photos to make sure none are broken — the picture files themselves shouldn't change, only their names.

If the names look right, you have two safe options: copy your `rename.py` script into your real photo folder and run it there, or ask Copilot to adjust the script so it points to a different folder. Either way, your originals stay untouched.

If it looks different: If a photo won't open, the file may have been partially affected — restore it from your test copy or your originals. If the names look wrong, go back to Step 4 and ask Copilot to fix the rule, then re-test on the test folder before touching real files.

You'll know it worked when your real photo folder shows the clean new names, and every photo still opens normally.

Common mistakes

- **Skipping the test folder.** It's tempting to point the script straight at your real photos, but one wrong rule can scramble hundreds of filenames at once. Always run the script on copies first, then on the real folder only after you're sure the names are right.
- **Running the script twice without changing the names.** The second run will keep building on your filename pattern (`vacation-2024-001.jpg` could become `vacation-2024-001-001.jpg`). Either move the renamed files out, or include the "skip if already named" rule in your prompt.
- **Trusting Copilot's first answer without reading.** Copilot is helpful but not perfect. If the script does something you didn't ask for — like deleting files or changing the file extension — rewrite your prompt to be more specific and try again on the test folder.

What this means for you

- **In everyday life:** A folder of 200 vacation photos can go from chaos to `paris-day1-001.jpg`, `paris-day1-002.jpg`, and so on in under a minute. No more guessing which `IMG_4829.jpg` was the Eiffel Tower shot.
- **For your work or business:** Product photos, invoice scans, and event pictures often arrive with cryptic camera names. A two-minute script saves the headache of clicking through 80 files one by one, and the same approach scales to thousands of files.
- **If you're just getting started:** Don't worry about learning Python itself — you're learning how to ask an AI for a script. The skill that matters here is describing your rule clearly, not writing the code by hand.

Try it now

Open your test folder in VS Code, open the Copilot chat, and paste this exact message.

💬 `Write a Python script that lists every .jpg file in the current folder and renames them to "sample-001.jpg", "sample-002.jpg", in the order they appear. Print the new names before doing it.`

Run it on a tiny folder of 3 to 5 test files. Seeing those photos rename themselves is the moment this stops feeling like magic and starts feeling like a tool you control.

How to Turn a Long Podcast Transcript Into Shareable Notes With Claude

6 steps · Read online

> **In short:** You can paste a podcast transcript into Claude and ask for a 5-bullet summary, the best quotes, and an action list — usually in under a minute. You'll need the transcript as text (Claude can't listen to audio), and very long ones may need a paid plan.

A podcast you actually enjoyed drops, and now you want to share the best parts without making your friend sit through 90 minutes. Or you covered it for work and need notes by tomorrow. Here's how to use Claude to turn that wall of text into something a busy person will actually read.

Before you start

- A Claude account (the free tier works for shorter transcripts; very long ones may need a paid plan)
- Access to **claude.ai** on a web browser or the Claude mobile app
- A transcript of the podcast as plain text — either from the show's website, a transcription service, or your own export
- About 10 minutes total

Step 1 — Get the transcript as plain text

Before Claude can help, you need the transcript as plain text — words you can copy and paste. Many podcasts publish transcripts on their website or in the show notes. If yours doesn't, run the audio through a transcription service (there are free and paid options) and export the result. A simple .txt or .docx file works best; if you only have a .vtt file with timestamps, that's fine too — Claude can skip past the timecodes. If you only have the audio file and no transcript yet, pause here and get one first; Claude reads text, not sound. 💬 Example: a clean opening line — "Host: Today we're talking about how small teams ship faster than big ones." You'll know it worked when you have the spoken words sitting in a document or on your clipboard, ready to paste in the next step.

Step 2 — Open a fresh chat in Claude

Go to **claude.ai** on your computer, or open the Claude app on your phone. Look for a button that starts a brand-new conversation — it usually looks like a pencil icon, a plus sign, or carries a "New chat" label, often near the top-left of the screen. Click or tap it once, and the page clears to an empty conversation with a text box at the bottom. A fresh chat matters because mixing your

podcast transcript with old questions can muddle the summary later. If you don't see a clear "new chat" option, try logging out and back in, or closing the browser tab and reopening it. Example: no typing yet — you're just clearing the deck. You'll know it worked when the chat shows an empty message thread and a blinking cursor in the input box.

Step 3 — Paste the transcript and tell Claude what it is

Click into the empty text box, paste the full transcript with your keyboard (Ctrl+V on Windows, Cmd+V on Mac), and then add one short line underneath that explains what you're sharing. That line matters — without it, Claude doesn't know it's looking at a podcast and may try to edit the text as if it were your draft. If your transcript is very long and the paste doesn't fit, look for a paperclip or attachment icon near the text box and upload the file instead; files count toward the same content Claude can read. Example: after pasting, add this on a new line — "The text above is a podcast transcript. Please read it carefully before I ask my next questions." You'll know it worked when Claude replies with something like "Got it — I've read the transcript. What would you like to know?" or simply waits for your next instruction.

Step 4 — Ask for a 5-bullet summary

Stay in the same chat and type a follow-up request for a summary. Ask for a specific number of bullets and a specific audience — vague prompts get vague summaries. Claude usually answers in under a minute for a one-hour transcript, though very long ones may be too much for Claude to read all at once on the free tier. If the answer stops mid-sentence, just reply "keep going" and Claude will resume from where it left off. Example: "Now please summarize this podcast in exactly 5 bullets for a busy teammate who hasn't listened. Lead with the single most useful idea." You'll know it worked when you see five short, self-contained bullets that someone could read in 60 seconds and actually understand the episode.

Step 5 — Pull the best quotes and an action list

Ask for two more things in the same chat: the most quotable lines from the transcript, and a short list of action items — concrete things the listener might actually do after hearing the episode. Ask for them as separate, clearly labeled sections so you can copy each one cleanly. If a quote feels bland, ask for a more surprising or specific one. If the action list sounds generic, ask Claude to tie each item back to a moment in the transcript. Example: "Now give me two more sections: (1) the 3 most quotable lines, with who said them; (2) a 4-item action list for someone who wants to apply this tomorrow. Label each section clearly." You'll know it worked when you have three distinct blocks — summary, quotes, action list — that you can copy out separately.

Step 6 — Save and share the notes

Copy the three blocks out of the chat and paste them into your notes app, an email, or a Slack message. One honest note: every quote Claude pulls is something the speaker actually said in the transcript you gave it — but if you paraphrase a quote before sending, double-check it against the source, because Claude can occasionally mis-attribute a line. If you're on a paid plan, look for a "Share" option that gives your recipient a link to the whole conversation instead. Either approach works — pick the one your teammate will actually open. 💬 Example: "Summary: [paste 5 bullets]. Best quotes: [paste 3 lines]. Action list: [paste 4 items]. Source: [episode name and link]." You'll know it worked when you hit send and your recipient replies with something other than "what is this?"

Common mistakes

- **Pasting without telling Claude what the text is.** If you just drop in a transcript and ask "summarize this," Claude may treat it as your own writing and try to edit it. Always add one short line identifying it as a podcast transcript first.
- **Asking for too much in one prompt.** "Summarize this, pull quotes, give me five tweet ideas, and write a LinkedIn post" usually produces shallow output for everything. One clear request per turn gets much better results.
- **Trusting the quotes without checking.** Claude is reading your transcript, so the quotes are real — but if you rewrite them in your own words before sending, you can accidentally change the meaning. For important quotes, copy them verbatim and verify.

What this means for you

- **In everyday life:** instead of telling a friend "you have to listen to this episode," you can send the three best quotes and one sentence on why it matters — and they will actually read it.
- **For your business or work:** a 90-minute interview becomes a 60-second read for your team in roughly 10 minutes of work, transcribing included.
- **If you're just getting started:** this is one of the easiest first projects with Claude. The transcript gives the assistant a clear job to do, and you'll see useful output within minutes.

Try it now

Find the transcript of one podcast episode you've already listened to (anywhere from 20 minutes to an hour long), paste it into a fresh Claude chat, and ask for the 5-bullet summary using the prompt from Step 4. That single action is the whole first run — and it usually takes less time than making a cup of coffee.

How to Summarize a Long YouTube Video in 60 Seconds with Any AI

6 steps · Read online

> **In short:** Paste a YouTube link into any AI assistant and ask for a summary — in about a minute you get the main points, key quotes, and a one-sentence takeaway without watching. The catch: not every assistant reads YouTube links directly, so we cover that fallback too.

Found a one-hour interview you don't have time to watch? A 40-minute lecture you need for work tomorrow? You don't have to skip it. Any modern AI assistant can read the page a YouTube link points to and pull out the important parts. Here's the whole workflow, start to finish.

Before you start

- **A YouTube link** — the URL of the video you want summarized. It works on desktop, mobile, or in the YouTube app.
- **An AI assistant** you already have access to — ChatGPT, Gemini, Claude, Copilot, Le Chat, or any other mainstream option. Free or paid plans both work, though some assistants may need a paid plan to read external links reliably.
- **Roughly 2–3 minutes** for the whole thing, including reading the summary.

Step 1 — Copy the YouTube video link

Open the YouTube video you want summarized and look at the address bar at the top of your browser (or the link shown under the video title in the YouTube app). A YouTube link starts with **youtube.com** or **youtu.be** and ends with a jumble of letters and numbers — that's the video's unique ID. Tap and hold (on phone) or click and drag (on computer) to highlight the whole address, then copy it. On the YouTube mobile app, tap **Share** under the video and choose **Copy link** instead — it's faster. If you're on a TV or game console and don't see a copy option, open the video in a phone or computer browser to grab the link from there. 💬 Example: Your copied link will look something like youtube.com/watch?v=AbCdEfGh123 — the letters and numbers at the end are the video's ID. You'll know it worked when the address is highlighted or your phone briefly says "Copied."

Step 2 — Open your AI assistant

Open the website or app for the AI you normally use. Most readers already have at least one of these: ChatGPT, Gemini, Claude, Copilot, or Le Chat. Open a new chat — look for a button like **New chat**, **New conversation**, or a + or pencil icon in the corner of the screen. Starting fresh

matters because a long previous chat can confuse the assistant about what you want right now. You'll know it worked when you see an empty message box, usually at the bottom of the screen, with a prompt that says something like "Message..." or "Ask anything" — your exact wording may be different.

Step 3 — Paste the link and ask for a summary

Click in the empty message box and paste the YouTube link you copied. Then type your request. A short, clear instruction works best — tell the assistant what you want and how long the answer should be. Press **Enter** (or tap the send arrow). The assistant will fetch the video's public page, read its title and description, and (on most assistants) the transcript too. A short summary appears in under a minute. Example: "Here's a YouTube link: [paste your link]. Summarize this video in 5 bullet points and keep it under 150 words." You'll know it worked when the assistant replies with bullet points or a short paragraph that names the video's main topic and the people or ideas involved.

Step 4 — Pull out a few key quotes

The summary gives you the gist, but sometimes you want the actual words someone said — a striking line, a statistic, or a memorable phrase. Ask for those next. Quotes are useful when you want to share the video with a friend or quote it in your own work. Type a follow-up question that specifically asks for verbatim lines. The assistant will return them with quotation marks and usually say who said them. Example: "Now pull out 3 key quotes from that video, with the speaker's name for each one. Keep each quote to one sentence." You'll know it worked when the reply contains lines wrapped in quotation marks and the speaker is named beside each one.

Step 5 — Get a one-sentence takeaway

A bullet-point summary is great for studying, but a single sentence is easier to remember and share. Ask for the takeaway last. The assistant will compress everything it just read into one line you can use as a reference or send to a colleague. If you have a specific audience in mind — your team, your teacher, your family — say so, because the assistant will word the takeaway differently for each. Example: "In one sentence, what's the single most important takeaway from that video, written so I can send it to my manager?" You'll know it worked when the reply is a single complete sentence (not a paragraph) that captures the main idea.

Step 6 — Save or share your summary

Now that you have the points, quotes, and takeaway, decide where they should live. Most assistants have a **Copy** button at the top or bottom of each reply — tap it to grab the text and paste it into your notes app, email, or document. Some assistants also have a **Share** button that

creates a link you can send to someone. If you don't see either button, you can manually select the text and copy it the same way you copied the YouTube link in Step 1. You'll know it worked when the summary text is sitting in another app — your notes, an email draft, or a message — and you can read it again without opening the AI chat.

Common mistakes

- **Pasting the link without any instructions.** The assistant may reply with "What would you like me to do with this?" — always tell it you want a summary, in what length, and for whom.
- **Assuming every assistant reads YouTube links directly.** A few assistants only read the video's title and description, not its full transcript. If your summary feels thin or vague, ask the assistant to look up the video's transcript, or copy the transcript yourself from YouTube (look for an option like **Show transcript** under the video description) and paste that text instead of the link.
- **Trusting the summary on a topic you know nothing about.** Summaries are usually accurate, but AI can still mix up names, dates, or numbers. For important decisions — anything medical, legal, financial, or that affects other people — verify the quote or fact by checking the original video yourself.

What this means for you

- **In everyday life:** Next time someone sends you a 30-minute tutorial or a long podcast episode on YouTube, decide in 60 seconds whether it's actually worth your time. Skip the noise, keep the useful stuff.
- **For your business or work:** Before a meeting, paste the link to a recorded presentation or webinar and ask for the five main points plus the speaker's recommendations. You walk in prepared without sitting through the whole recording.
- **If you're just getting started:** Pick one YouTube video you've been putting off, follow the six steps above, and see how much you actually retain from the summary versus the full video. Most people are surprised how often the bullet points are enough.

Try it now

Open your AI assistant right now, paste the link to any YouTube video you've been meaning to watch, and type: "Summarize this video in 5 bullet points and pull out 2 short quotes." You'll have a usable summary in about a minute — that's the whole skill.

Create Royalty-Free AI Music Loops for Podcasts and Videos

5 steps · Read online

Imagine having custom background music for your podcast, YouTube video, or even a presentation, without needing a huge budget or any musical training. By the end of this guide, you'll know how to use Artificial Intelligence (AI) to generate unique, royalty-free music loops that perfectly match the mood of your content. This guide is for anyone who wants to add a professional touch to their projects without the hassle. It's important to know that while AI can create fantastic music, it cannot read your mind or fully understand complex musical theory, so clear instructions are key!

Before you start

- **A device with internet access:** You'll need a computer, laptop, or tablet with a stable internet connection. Most AI music tools, like the one we'll use, operate through a web browser.
- **A Suno.ai account:** This guide uses **Suno.ai** as a practical example. Visit the Suno.ai website and sign up using a Google, Discord, or email account. This process is typically quick.
- **About 10-15 minutes of time:** This is a rough estimate for following the steps and generating your first few music loops.
- **A creative idea:** Have a general sense of the mood, style, or genre of music you're looking for. Do you need something calm, energetic, mysterious, or upbeat?

Step 1 — Sign up or log in to Suno.ai

First, you need to access an **AI music generator**, which is a special online tool that uses **Artificial Intelligence (AI)** — computer programs that can learn and solve problems — to compose original music based on your descriptions. For this guide, we'll use **Suno.ai**, a popular and user-friendly option. Open your web browser (like Chrome, Firefox, or Safari) and navigate to the Suno.ai website (usually by typing "suno.ai" into your browser's address bar). You'll typically see prominent buttons to "Sign up" or "Log in." Choose the option that suits you, often by linking an existing Google or Discord account, or by creating a new account with your email.

If it looks different: If you're on a mobile device, the layout might be slightly condensed, but you'll still find clear sign-up or log-in options. If you're using a different AI music generator, the exact buttons might vary, but look for similar "Get Started" or "Join" prompts. **You'll know it worked when** you are successfully redirected to a main dashboard or a prominent creation screen, indicating you are logged in and ready to begin generating music.

Step 2 — Find the creation area and understand the prompt

Once you're signed in to Suno.ai, you'll typically land on a page that features a large text box, often labeled "Describe your music here" or "What do you want to hear?". This is your **prompt box** – a designated area where you'll type your instructions for the AI. A **prompt** is simply the text description you give to the AI, telling it exactly what kind of music you want it to create. Think of it like giving a detailed brief to a composer. Look around this area for other settings like a checkbox for "Instrumental" or a "Generate" button.

If it looks different: Some AI music generators might require you to click a "Create" or "Generate" button first to reveal the prompt box. Regardless of the tool, your goal is to locate the main text input field where you can describe your desired music, usually a long, empty box. **You'll know it worked when** you see a clear text input field on your screen, indicating that this is where you'll type your musical vision for the AI.

Step 3 — Craft your music idea (your 'prompt')

Now it's time to tell the AI what to compose! In the prompt box, type a clear and descriptive explanation of the music you're looking for. Consider these elements: the **genre** (like folk, electronic, rock), **mood** (happy, reflective, energetic), **tempo** (how fast or slow the music should be), and **instrumentation** (which instruments you want to hear, such as piano, guitar, or drums). The more specific you are, the better the AI can match your vision. Remember, we're aiming for a **music loop**, which is a short piece of music designed to play seamlessly over and over, so clarity is crucial.

💬 **Example:** `A calm, ambient electronic loop for a meditation podcast, featuring soft synth pads, a subtle bassline, and a gentle, pulsing rhythm. Focus on a peaceful and reflective mood.` **If it looks different:** Some AI music tools might offer specific dropdown menus for genre, mood, or instrumentation instead of requiring you to type them into the main prompt. If these dedicated options exist, it's generally best to use them as they can guide the AI more precisely. **You'll know it worked when** you have typed your detailed musical description into the prompt box, ready for the AI to process.

Step 4 — Generate your music and listen

After you've carefully written your prompt, look for the "Generate," "Create," or "Make Music" button, usually located below or next to the prompt box. Click this button to send your instructions to the AI. The tool will then process your request, using its understanding of musical patterns and your description to produce sound clips. This process typically takes a few seconds to a minute, depending on the tool and the complexity of your request. The tool may generate one or more variations of your music.

If it looks different: While the AI is generating, you might see a progress bar, a spinning icon, or a message indicating that the music is "processing." If you are using a free plan, you might only get one generation per prompt or have a daily limit on the total number of generations you can make. **You'll know it worked when** one or more generated music tracks appear on your screen, each typically accompanied by a play button. Press the play button on each track to listen and determine if it matches your original vision.

Step 5 — Refine and download your chosen loop

It's common not to get the perfect result on the very first try, so don't hesitate to experiment! If the music isn't quite right, go back to your prompt (from Step 3) and make small, precise adjustments. For instance, you could add details like "add a subtle orchestral swell" or change a mood word from "calm" to "serene." Once you find a music loop you love, look for a download button, often represented by an arrow pointing downwards, a cloud icon, or the word "Download." Click this to save your music file to your device. Most tools allow you to save the file as an **MP3** (a common, compressed audio format that balances quality and file size) or **WAV** (a higher-quality, uncompressed audio format ideal for professional use).

If it looks different: The download icon might be found within a three-dot menu (...) next to the track, or there might be a specific "Export" option. On a free tier, some tools may restrict downloads or offer lower-quality files. Always review the tool's terms regarding "royalty-free" usage to ensure you're allowed to use the music for your specific purpose, especially if it's for commercial content like monetised videos; some may require attribution (crediting the AI). **You'll know it worked when** the music file appears in your device's downloads folder or a designated location, ready for you to use in your podcast, video, or other creative projects.

Common mistakes

- **Being too vague in your prompt:** Just saying "happy music" doesn't give the AI enough detail to understand your specific needs.
    - **The fix:** Be much more specific. Add details about instruments ("upbeat acoustic guitar, light drums, simple piano melody"), tempo ("fast tempo"), and exact mood ("joyful and bright, like a morning sunrise").
- **Forgetting about "loop" functionality:** If the tool has a specific "loop" setting, or if you don't explicitly ask for it, you might get a song with a distinct ending that won't repeat seamlessly.
    - **The fix:** Always include "loop" or "seamless loop" in your prompt, e.g., "a seamless electronic loop." If the tool has a dedicated "loop" checkbox or setting, ensure it's ticked before generating.
- **Not checking usage rights:** Assuming all AI-generated music is free for any use can lead to unexpected copyright issues, especially if you plan to monetise your content.

    - **The fix:** Always confirm the AI music generator's terms and conditions, particularly for "royalty-free" usage and commercial licensing. "Royalty-free" usually means you don't owe ongoing payments, but some tools might still require you to credit the AI (attribution).

What this means for you

- **In everyday life:** You can easily create unique background music for your personal videos, social media stories, or even just for studying or relaxing. This allows you to personalize your digital content in a snap, making your home movies feel more professional or crafting the perfect atmosphere for your focus sessions.
- **For your business or work:** You can quickly produce bespoke, original music for marketing videos, online courses, corporate presentations, or business podcasts. This saves significant time and helps you avoid the complex licensing fees and legalities often associated with traditional stock music libraries.
- **If you're just getting started:** Begin with very simple prompts, like "upbeat acoustic guitar loop" or "calm piano instrumental." Listen to what the AI produces and then gradually add more detail or experiment with different genres.

Try it now

Open Suno.ai on your device, log in, and type a simple prompt like "energetic electronic loop for a workout video" into the prompt box. Click 'Generate' and listen to the results!

Remove Distractions from Photos with Google Photos Magic Eraser

5 steps · Read online

Have you ever taken a fantastic photo, only to have it ruined by a random person in the background, an overflowing bin, or an unsightly power cord? Google Photos' Magic Eraser feature uses a bit of clever Artificial Intelligence (AI) to help you instantly remove those distractions, leaving your pictures looking clean and professional. This guide will walk you through each simple step to make your photos shine.

Before you start

- You'll need the **Google Photos app** installed on your smartphone or tablet.
- You'll need a **Google Account** signed in to the app.
- Pick a photo from your library that has a clear object or person you'd like to remove.
- This feature is generally **free and available** on most modern devices via the Google Photos app, though very heavy or advanced use might have limitations on some free accounts.
- This whole process should take about **5 minutes**.

Step 1 — Open Google Photos and choose a photo

To begin, find and tap the **Google Photos app icon** on your device. It usually looks like a colorful pinwheel. Once the app is open, you'll see your photo library, often displayed as a grid of small picture previews, also known as thumbnails. Tap on the specific photo you want to edit.

What you see: Your device's home screen or app drawer, then your photo library in Google Photos, followed by your chosen photo appearing in full-screen view.

What happens next: The selected photo will open up, filling your screen, ready for you to view or edit.

If it looks different: If your photos aren't showing in a grid, look for a 'Library' or 'Photos' tab at the bottom of the screen. Tap it to see all your pictures.

You'll know it worked when your chosen photo fills the screen, showing all its details.

Step 2 — Find the Magic Eraser tool

With your photo open, you need to access the editing tools. Look for an **'Edit' button** at the bottom of the screen – it often looks like a pencil icon or a set of sliders. Tap this button to enter the editing mode. Once there, you'll see a row of different editing options like 'Crop', 'Adjust', or

'Filters'. Swipe through these options from right to left until you find and tap the **'Tools' section**. Inside 'Tools', you'll then see the 'Magic Eraser'.

What you see: Your photo with various editing icons or tabs appearing along the bottom. After tapping 'Edit', you'll see a new set of options, and then a menu specifically for 'Tools' will slide into view, revealing 'Magic Eraser'.

What happens next: The app transitions from viewing your photo to its comprehensive editing interface. Tapping 'Tools' brings up a selection of smart editing features, including the Magic Eraser.

If it looks different: The 'Edit' button might be a pencil or a 'Settings' icon. If you don't immediately see 'Tools' in the main editing bar, try looking under an 'Adjust' tab or swiping further along the bottom options.

You'll know it worked when you see 'Magic Eraser' clearly listed among the other editing tools, ready for you to select.

Step 3 — Let the AI suggest removals or draw your own

Now that you've selected 'Magic Eraser', the fun begins! The Artificial Intelligence (AI), which is like a clever computer program designed to think and learn, will immediately start analyzing your photo. It's looking for common distractions like people, power lines, or objects that stand out from the background. The AI will then present "suggestions" by highlighting these areas with a subtle, shimmering outline.

What you see: Your photo with the Magic Eraser interface active. You might notice flashing outlines around objects that the AI thinks you might want to remove. Below the image, you'll also see options to 'Erase' (using a brush) or 'Undo'.

What happens next: The AI has completed its scan. If it found anything, the suggested distractions will be highlighted on your screen. If no suggestions appear, that's perfectly fine – you can simply proceed to manually draw over any unwanted areas yourself.

If it looks different: If the highlights don't appear, or the AI doesn't suggest anything, don't worry. This just means you'll guide it yourself. The brush tool will still be active, ready for you to use.

You'll know it worked when the screen shows your photo with potential distractions outlined, or the brush tool is visible, indicating Magic Eraser is ready for input.

Step 4 — Erase unwanted elements

With Magic Eraser active, you have two ways to remove things. If the AI highlighted an object, simply **tap on the highlighted area**. The AI will then try to intelligently remove it and fill the gap with what it thinks the background should look like. For anything the AI didn't highlight, or for more precise control, use your finger to **draw a line or doodle directly over the unwanted object** on the screen. Try to cover the entire object you wish to remove.

What you see: As you tap or draw, the selected part of your photo will disappear almost instantly. The app will attempt to blend the surrounding pixels (the tiny dots that make up your image) to create a natural-looking background where the object once was.

What happens next: The unwanted element vanishes from your photo. You'll see the cleaned-up version right away, giving you a clearer, distraction-free image.

If it looks different: If the result isn't perfect, or if it leaves a blurry patch, don't panic. This tool works best on clear, distinct objects. Simply tap the 'Undo' arrow (often pointing left) at the bottom or top of the screen to reverse the last change and try again.

You'll know it worked when the unwanted object is gone, and the area where it used to be is filled in, making your photo look much tidier.

Step 5 — Refine and save your work

After erasing, take a moment to look closely at your photo. Sometimes, a first pass might leave a slight blur or an incomplete edge. If you need to make further adjustments, you can zoom in by pinching out on the screen with two fingers, and then use the brush tool again for finer detail. You can also use the 'Undo' arrow if you change your mind about a particular erasure. Once you're happy with your edits, look for a **'Done' button** or a tick icon, usually at the bottom-right, and tap it. Then, to save your masterpiece, tap **'Save copy'** (or sometimes just 'Save') at the bottom of the screen.

What you see: Your edited photo still on screen, possibly with 'Done' and 'Save copy' buttons. After tapping 'Save copy', you might see a brief 'Saving' message.

What happens next: Tapping 'Done' exits the Magic Eraser tool, taking you back to the main editing screen. Tapping 'Save copy' processes your changes and creates a brand-new version of your photo in your Google Photos library, leaving your original photo untouched.

If it looks different: The 'Done' button might be a simple checkmark. 'Save copy' might simply be labeled 'Save', or it could appear as a menu option after you tap 'Done'.

You'll know it worked when you see a confirmation message like 'Saved' or 'Copy saved', and you can then find the new, cleaned-up version of your photo alongside the original in your Google Photos library.

Common mistakes

- **Over-erasing**: Accidentally drawing over too much of your photo, or parts you actually wanted to keep, can make the result look unnatural.
    - **The fix:** Always use the 'Undo' arrow immediately if you make a mistake. It will revert your last action, allowing you to try again with more precision.
- **Unnatural background fill**: Sometimes, the AI can't perfectly guess what should be behind an object, leaving a blurry or repetitive pattern.
    - **The fix:** This often happens with complex backgrounds. Try erasing in smaller, more precise sections. If it still doesn't look right, remember that Magic Eraser works best on less complicated backgrounds.
- **Forgetting to save the copy**: After spending time perfecting your photo, it's easy to tap back out of the app without saving, losing all your hard work.
    - **The fix:** Always remember to tap 'Done' and then 'Save copy' (or 'Save') once you're satisfied. Google Photos will usually prompt you if you try to exit without saving, but it's good practice to make it a habit.

What this means for you

- **In everyday life:** Instantly clean up your holiday snaps by removing stray tourists, make family photos picture-perfect by getting rid of background clutter, or simply tidy up a shot of your pet for social media.
- **For your business or work:** Enhance product photos for online listings by removing dust or distracting elements, clear up real estate images by erasing power lines or bins, or prepare professional-looking profile pictures for your business's social media.
- **If you're just getting started:** Look for a photo with a single, clear, distinct object against a relatively uniform background – like a small sign on a wall, or a distant person against the sky. This will give you the best chance of a perfect 'erase' on your first try!

Try it now

Find a photo in your Google Photos library that has a simple, easy-to-remove object, like a small rubbish bin in the distance or a loose cable on the floor. Follow the steps above to magically erase it and see the difference!

How to use Kling AI video templates in Picsart Flow

5 steps · Read online

> **In short:** Picsart Flow includes Kling-powered video templates that turn a still photo, a character reference image, or a short script-plus-footage combo into a short AI-generated clip. You can try them without editing experience, though longer videos and higher resolution usually need a paid plan.

Want to make a short video but don't have filming skills, a camera, or editing software? Kling is a video-generation model that creates short clips from a still photo, a character image, or a written script combined with your existing footage. Picsart Flow hosts it as ready-made templates, so you mostly just upload, type a sentence, and wait. This guide walks through all three template styles one at a time. None of this requires a video background — only a clear idea of what kind of clip you'd like to end up with.

Before you start

- A free Picsart account (you can sign up with email, Apple, or Google — check the latest sign-up options on the Picsart login page).
- The Picsart app on your phone, or the Picsart web editor on a laptop. Picsart Flow tools may live in different spots in each, so this guide uses language like "look for" rather than exact menu names.
- A few source images saved to your device. For best results, use clear, well-lit photos where the subject is centered.
- Around 15–30 minutes for your first full run-through (most of that is waiting for the AI to render).
- A realistic expectation: AI video is impressive but imperfect. People may look slightly off, hands can blur, and motion can wobble — that's normal for current tools.

Step 1 — Open Picsart Flow and find the Kling templates

Tap the Picsart app to open it, then look along the bottom or side menu for a section called something like **Flow**, **AI Tools**, or **Video**. The exact label depends on your version and platform, so if you don't see "Flow", search inside the app for "Kling" or "AI video" — that usually leads you to the same area. Once you land on the Kling page, scroll until you see a row of template tiles. You'll typically find three styles: one that animates a single photo, one built around a character reference image, and one that combines your footage with a script.

Example: search terms to try inside the app — `Kling`, `AI video`, `photo to video`.

You'll know it worked when you see a screen with several preview tiles, each labeled with the input it needs (photo, image, or footage).

Step 2 — Use the photo-to-video template

Pick the **photo-to-video** template and tap it. The app will ask you to upload one image from your camera roll. Choose a clear photo where the main subject (a person, a pet, a product) is easy to see. After upload, a small text box appears asking what should happen in the video — that's called a **prompt** (the instruction you give the AI). Type one short sentence describing the motion you want, then tap the generate button. The app will show a progress indicator while Kling works. This step usually takes 30 seconds to a few minutes, depending on clip length and server load.

💬 Example: `The woman slowly turns her head to the right and smiles, with soft sunlight on her face.`

You'll know it worked when a short video preview appears in place of the still photo. If it looks wrong, tap regenerate or tweak the prompt — small wording changes make a big difference.

Step 3 — Try the character reference sheet template

Now pick the **character reference** template, sometimes labeled "character sheet" or "consistent character." This one is designed for keeping the same face or outfit across multiple clips — useful for short stories, social media series, or product mascots. The app will ask you to upload one or more images of the same character from different angles. Add a short prompt describing the scene, then generate. Kling tries to keep facial features and clothing consistent, which is harder than it sounds for AI, so your first result may need two or three tries.

💬 Example: `A young woman with short black hair, wearing a yellow raincoat, walks through a misty forest at dawn.`

You'll know it worked when the generated clip shows a person who actually resembles your reference images, in roughly the same clothes and setting. If the face drifts, regenerate or add more reference images — more input usually helps consistency.

Step 4 — Combine footage with a script

The third template, often called **video-to-video** or **script + footage**, lets you feed Kling a short video clip you already have plus a written description of what should change. Think of it as a soft redraw: the AI keeps the general shape and motion of your clip, but re-styles it. Upload a short clip (a few seconds is enough), then type a script describing the new mood, setting, or style. This is the template most people use for stylized social posts — turning a daytime clip into a moody

evening scene, for example. Generation takes longer here because the model has more to process.

💬 Example: `Change the setting to a neon-lit Tokyo street at night, keep the same walking motion.`

You'll know it worked when your uploaded clip now looks visually different (new lighting, color, or setting) while the original motion still feels familiar.

Step 5 — Preview, trim, and export

Once any template finishes, tap the result to open a preview screen. Most flows let you trim the start and end by dragging handles on a small timeline bar. If the clip is too long, shorten it; if the opening is awkward, trim the first second. When it looks right, look for an **Export**, **Save**, or **Download** button — usually at the top or bottom of the preview. On the free plan, you may get a watermark or a lower resolution; higher resolutions and watermark-free exports are typically behind a paid plan. The exact limits change often, so check Picsart's current pricing page if this matters for your project.

💬 Example: trim suggestion — `keep only seconds 2 through 6` (drag the handles until those seconds are highlighted).

You'll know it worked when a finished video file appears in your phone's photo library or downloads folder.

Common mistakes

- **Uploading a blurry or busy photo.** Kling struggles with tiny faces and cluttered backgrounds. Use the clearest, simplest image you have.
- **Writing prompts that are too long.** Two sentences is plenty. Long paragraphs confuse the model and often produce muddled motion.
- **Expecting perfect hands or text.** AI video still fumbles fingers, writing, and small details. If those matter, plan to crop them out of frame or pick a different shot.

What this means for you

- **In everyday life:** you can turn a favorite vacation photo into a short clip to share with family, without learning video software.
- **For your business or work:** small businesses can produce quick product clips, social teasers, or stylized B-roll (extra footage used to support a story) without hiring a video editor — useful for one-person shops and busy marketing teams.

- **If you're just getting started:** begin with the photo-to-video template on a single clear selfie or pet photo. It's the fastest way to see what Kling can do in under five minutes.

Try it now

Open Picsart, search for the photo-to-video template, and upload one clear photo of someone smiling. Type: `The person gently blinks and turns to look out a window, soft natural light.` Tap generate, wait a minute, and watch your first Kling clip come to life.

How to create images with AI: a beginner's first try

6 steps · Read online

Imagine typing a simple sentence and watching a brand-new, unique image appear on your screen in seconds. It's like having a digital artist ready to bring any idea to life, no drawing skills required. Whether you want to design a custom greeting card, create a visual for a work presentation, or just visualise a creative thought, **AI image generators** make it possible. By the end of this guide, you will know how to choose a free tool, write a description that the AI understands, and be aware of how you can use your new creations.

Before you start

- **A computer or smartphone** with an internet connection and a web browser.
- **A free account** with an AI image generation platform. Many trusted companies offer free ways to create images — for this guide, we suggest exploring **Microsoft Copilot** (at copilot.microsoft.com) or **Canva** (at canva.com), as both have options to create images without needing a paid subscription. You'll need to sign up for a free account if you don't already have one.
- **Roughly 10-15 minutes** of time to follow along with the steps and create your first images.

Step 1 — Choose and open a free tool

Your first action is to pick a free online tool and open it in your web browser. You don't need to spend money to try this out, as many trusted companies offer free **AI image generators** — these are smart computer programs that act like a digital artist, translating your written words into pictures. We'll use Microsoft Copilot as our primary example.

Open your web browser and go to `copilot.microsoft.com`. If you're not logged in, you'll see a screen asking you to sign in with a Microsoft account or create one. Once you're signed in, you'll see the Copilot interface. **If it looks different:** If you're using Canva, open `canva.com` and log in. Then look for AI-related features, often found in the design creation flow or through a search function. **You'll know it worked when** you see the main interface of your chosen tool, ready for you to input commands or search for tools.

Step 2 — Find the image generator

Now that you're in the tool, you need to tell it you want to create an image. In Microsoft Copilot, look for the image generation feature — this may be a dedicated button, a mode toggle, or you can simply describe what you want in the chat. Type a simple instruction into the input area. When you do this, you're giving a **prompt** to the AI – this is just a fancy word for the

instruction you write to tell the AI what you want to see. **If it looks different:** In Canva, explore the AI image creation options, which may be labeled differently or located in various places within the design interface. **You'll know it worked when** the tool responds by acknowledging your image request, or when a dedicated image generation interface appears with a specific text box for your image description.

💬 Example:

> "Create an image of"

Step 3 — Write your first detailed "prompt"

To get a great result, your **prompt** (the instruction) needs to be more than just one or two words. Think of it like giving a detailed recipe to a chef – the more specific you are, the better the dish will be! A good prompt usually has three main ingredients:

1. **The Subject:** Who or what is the main focus of the picture (e.g., "a fluffy ginger cat").
2. **The Details:** What they are doing, the colors, and the background (e.g., "wearing a blue knitted jumper, sitting on a wooden park bench, bright sunny day").
3. **The Style:** Do you want a realistic photo, a watercolour painting, a 3D cartoon, or something else (e.g., "soft watercolour painting style").

Now, type your complete, detailed description into the text box and send it. The AI will then start processing your request. **If it looks different:** If the AI seems confused, try rephrasing your request slightly, such as "Generate a picture that shows..." instead of "Create an image of...". **You'll know it worked when** you see a message like "Generating your image..." or a loading indicator, confirming the AI is working on your request.

💬 Example:

> "A fluffy ginger cat wearing a blue knitted jumper, sitting on a wooden park bench, bright sunny day, soft watercolour painting style."

Step 4 — Generate and review your images

After you send your detailed prompt, the AI will get to work. It usually creates a few different image options based on your description. This process can take anywhere from a few seconds to about half a minute, depending on the tool and how busy it is.

Once the AI has finished, you'll see a collection of images appear on your screen, usually displayed side-by-side or in a grid. Take a moment to look at them. Don't worry if they're not absolutely perfect yet – this is just the first draft! This step shows you the AI's interpretation of your words. **If it looks different:** Some tools might generate only one image at a time, or

show a progress bar as they work. If the images look completely wrong, double-check your prompt for any typos or unclear instructions. **You'll know it worked when** you see your newly generated images displayed clearly on your screen, all inspired by the description you provided.

Step 5 — Refine or download your favorite

You don't have to settle for the first set of images. Most AI tools allow you to refine your results, just like talking to a helpful assistant. You can ask for changes to the images you see or generate completely new ones with a tweaked prompt.

If you want to make changes, simply type a follow-up instruction. If you're happy with an image, look for a "Download" or "Save" button, which is often an arrow pointing downwards or a disk icon, usually found near the individual image or at the bottom of the collection. Clicking this will save the image to your device. **If it looks different:** The download button might be hidden behind a "..." (three dots) menu or a right-click option. If you're refining, some tools let you select a specific image to refine, while others apply changes to the whole set. **You'll know it worked when** you either see a new set of images based on your refinements, or you've successfully saved an image file to your computer or phone.

Example:

> "Change the jumper color to yellow, and add a small friendly magpie sitting on the bench next to the cat."

Step 6 — Review your image for watermarks and usage

Before you share your AI masterpiece, it's helpful to understand two quick concepts: **watermarks** and **usage rights**. Take a close look at the images you've generated.

A **watermark** is a faint, semi-transparent logo or symbol placed in the corner of an image to show which tool created it. Many free tools include these automatically to keep the images easily recognizable as AI-made. Next, think about **usage rights**, which are the rules for how you can use the images. As a general rule, images you make on free accounts are perfect for personal fun, sharing with friends, or school projects. However, if you plan to use them to make money (like on a commercial website or product packaging), you must check the tool's terms of service (their official list of rules) as some free tiers do not allow commercial use. **If it looks different:** Some tools might offer a paid plan to remove watermarks, or their terms of service might be difficult to find. Look for a "Terms" or "Legal" link, usually at the bottom of the website. **You'll know it worked when** you can identify if your image has a watermark, and you have a general understanding of whether your chosen free tool allows personal or commercial use of your creation.

Common mistakes

- **Being too vague:** Typing "a dog" gives the AI too much room to guess. You might end up with a cartoon pug when you wanted a realistic golden retriever.
    - **The fix:** Be specific! Add details about the breed, color, action, and setting (e.g., "A golden retriever puppy frolicking in a field of sunflowers, hyperrealistic photo style").
- **Expecting perfect text spelling:** AI image generators are brilliant with shapes and colors, but they often struggle to spell words inside images. If you ask for a sign that says "Welcome", the AI might print scrambled letters or strange symbols.
    - **The fix:** It's usually best to avoid asking the AI to generate text within your images. Create the image first, then add text later using a separate image editing tool.
- **Forgetting that AI does not know real people:** For safety and privacy reasons, free tools will not let you generate realistic images of real public figures, celebrities, or your friends.
    - **The fix:** Stick to fictional characters, general subjects, or broad descriptions (e.g., "a young woman reading a book" rather than "Sarah reading a book").

What this means for you

- **In everyday life:** You can instantly create custom birthday invitations, organize your home with beautifully styled labels, or make a unique phone wallpaper that no one else has – all with just a few words.
- **For your business or work:** You can quickly mockup visual ideas for presentations, create colorful backdrops for social media posts, or brainstorm logo concepts without spending hours searching through generic stock photo libraries.
- **If you're just getting started:** Remember that experimenting is half the fun. There are no wrong answers, and the more you practice describing your ideas, the better and more accurate your results will be.

Try it now

Let's get you started right now. Open Microsoft Copilot or Canva in a new tab, find the image generator function, and type this simple prompt:

*"A friendly koala wearing a green scarf, holding a warm mug of tea, soft digital illustration style"

How to Edit Your Phone Photos with Simple Text Instructions in Meta AI

6 steps · Read online

> **In short:** Send Meta AI a photo, type what you want changed ("remove the person on the left", "make the sky look like sunset"), and it returns a new version. Simple edits usually work on the first try; complex changes may need a couple of attempts.

Ever take a great photo and then notice a stranger walking through the background? Or snap a sunset shot where the sky came out dull and gray? You don't have to learn a photo editor to fix these — Meta AI can do small edits just by typing what you want. This guide walks through how, on your phone, in a few minutes.

Before you start

- A phone with **Meta AI** available — built into WhatsApp, Instagram, Facebook, or Messenger. If you don't see it in any of those, the same chat works at **meta.ai** in your phone's browser.
- A Meta account (the same one you use for Facebook or Instagram works).
- One photo saved on your phone that you'd like to edit — landscape, portrait, or a quick snapshot, any of them.
- A working internet connection.
- About **5 minutes** for your first edit, and a little patience if the first attempt isn't quite right.

> **One honest note:** Meta AI is great at small, focused edits — sky color, brightness, removing a clear object. It is not a replacement for a full photo editor like Photoshop. If your request is very complex, the result may surprise you, and that's normal.

Step 1 — Open a chat with Meta AI

Open WhatsApp, Instagram, Facebook, or Messenger, then find the contact or chat labeled **Meta AI**. It's usually marked with a little blue or purple ring around its profile picture. In WhatsApp, tap the green new-chat button and look for "Meta AI" near the top of your contacts. In Instagram, start a new direct message and search for "Meta AI". A one-on-one chat opens, with a text box at the bottom and a paperclip or + button on the left for attaching things. If you can't find Meta AI in any app, open your phone's browser and go to **meta.ai** — sign in with your Facebook or Instagram account if asked, and you'll land in the same chat. 💬 Example: Open WhatsApp → tap the new-chat icon → tap "Meta AI" in the contacts list. You'll know it worked when a chat thread opens with "Meta AI" at the top and a text box you can type into.

Step 2 — Send the photo you want to edit

Tap the paperclip, +, or camera icon next to the text box, then pick **Photo** or **Gallery** from the menu. Your phone's photo library opens. Select the photo you want — just one to start, so the AI has a clear focus. The photo appears as a small preview in the chat; tap **Send** (the arrow button) to deliver it. The image uploads in a few seconds, and Meta AI now "sees" it. If you're using **meta.ai** in the browser instead, look for an upload icon — usually an arrow pointing up or a small picture frame — on the left side of the message box. Tap it, pick the photo, and send the same way. 💬 Example: One landscape photo where the sky looks flat and gray — perfect for your first edit. You'll know it worked when your selected photo shows up in the chat as a sent message.

Step 3 — Type what you want changed

In the text box right below the photo, write a short, plain sentence describing the edit. Tap send. You'll see three small dots animating while Meta AI thinks — that means it's working. Within a few seconds, sometimes up to a minute, it replies with a new image, usually with a one-sentence acknowledgment like "Here's the edited version." The new photo lands in the chat as another message bubble. If the reply is **only text** with no image, your instruction was probably unclear — try again with the wording in the example below. 💬 Example: "Make the sky look like a bright sunset with orange and pink clouds" — or simpler, "Remove the person on the right" or "Brighten the photo so it doesn't look so dark." You'll know it worked when a new image appears in the chat that looks like your photo, but with the change you asked for.

Step 4 — Look closely at the edit

Tap the new image to open it full-screen. Pinch out with two fingers to zoom in on the area you asked to change. Now look at it honestly: did the edit land where you wanted? Sometimes Meta AI nails it. Other times the change is subtle, or it shifted something you didn't ask to change — a face looks slightly different, or a building in the background got smudged. That happens, and it doesn't mean you did anything wrong. If the image is small and won't zoom in the chat, take a screenshot of your phone (side button + volume up on most phones) — that captures the full-size version for closer inspection. 💬 Example: Ask yourself — is the person really gone, or is there a faint shadow where they stood? Is the sky the color you wanted, or does it look painted on? You'll know it worked when you can decide, honestly, whether this version is good enough to keep or whether you'd like to try again.

Step 5 — Ask for a redo if it's not quite right

If the edit missed the mark, reply in the same chat with a clearer or more focused instruction. Keep the conversation going — Meta AI keeps the photo and your request in context, so you can

iterate. Useful phrases: "Try again", "Make it stronger", "Keep everything else exactly the same and only change the sky", or pointing out what went wrong, like "The person is still there — try once more". If the AI keeps misunderstanding after several tries, start a brand new chat, send just the photo again, and write your best, clearest instruction first. Long threads with many edits sometimes confuse the model. 💬 Example: First try didn't remove the person cleanly? Send: "Try again — remove the person on the right side, but keep everything else exactly the same." You'll know it worked when a new version finally looks like what you wanted — or when you've accepted that this particular edit isn't possible and you'd rather try a different photo.

Step 6 — Save the photo you like

When you're happy with an edited version, save it to your phone. Tap and hold the image in the chat — a small menu appears with options like **Save**, **Save to Photos**, or **Share**. Tap "Save" and the photo lands in your camera roll or gallery, just like one you took yourself. In some apps you'll see a small downward-arrow icon on the image — tap that instead. In **meta.ai** in the browser, long-press the image (or right-click on a computer) and choose "Save image as". Only save once you're sure — once it's in your camera roll, you can't easily get back to the older versions still sitting in the chat. 💬 Example: Save your favorite version, then close the chat. Your original photo is untouched on your phone, and the edited version is now saved next to it. You'll know it worked when you open your phone's Photos or Gallery app and see the edited photo sitting there.

Common mistakes

- **Asking for too many changes at once.** "Remove the person, fix the sky, brighten everything, and add a reflection in the water" is too much for one message. The AI may handle one or two of those, badly, while ignoring the rest. **Fix:** ask for one change at a time, and send the photo back through Step 3 each time.
- **Writing vague instructions.** "Make it better", "fix this", or "improve it" leaves the AI guessing — usually wrong. **Fix:** describe exactly what to change in plain words. "Make the sky bluer" works; "nicer sky" doesn't.
- **Not zooming in to check the result.** On a phone screen, a small image can look fine until you look closely — then you notice a smudged face or a ghostly outline where someone was removed. **Fix:** always open the image full-screen and zoom in on the edited area before saving.
- **Saving the first version without comparing.** The first attempt is rarely the best one. **Fix:** wait until you've tried once or twice and you've got the version you actually want, then save that one.

What this means for you

- **In everyday life:** vacation photos where someone walked into your shot, birthday candles that came out dim, a beach picture where the sky was gray — all fixable in under a minute with one short instruction.
- **For your business or work:** quick product photos, social media posts, or marketplace listings where the background needs cleanup. Type what you want removed, send it through, save the result. It won't replace a designer, but it gets you 80% of the way in seconds.
- **If you're just getting started:** pick a photo that's been bothering you — maybe a great group shot with a stranger in the corner — and try removing them first. It's the easiest edit to see if Meta AI is working for you.

Try it now

Open Meta AI right now in WhatsApp, Instagram, or **meta.ai** in your browser. Send one photo you've taken recently — anything from this week. Then type one short instruction, such as "Make the colors look brighter" or "Remove the trash can in the corner". The whole thing takes under two minutes, and you'll see exactly how the tool responds to your wording — which makes the next edit easier too.

How to Generate Images with Google Gemini: A Beginner's Guide

6 steps · Read online

Imagine being able to paint a picture just by describing it out loud. Whether you need a quick illustration for a work presentation or a fun graphic for a party invitation, Google Gemini lets you turn your words into custom images in a matter of seconds. This guide will walk you through the simple steps to create your own unique visuals.

Before you start

- **A Google Account:** You will need to be signed into a standard Google account. If you use Gmail or YouTube, you already have one.
- **Access to Gemini:** Open your web browser and go to `gemini.google.com`, or open the Gemini app on your mobile device.
- **Device:** This guide works on any computer, tablet, or smartphone with internet access.
- **Cost:** Image generation in Gemini is currently available for free, though usage limits may apply or change over time with AI tools.
- **Time:** This guide will take approximately 5-7 minutes to follow.

Step 1 — Open Gemini and start a new chat

First, make sure you're ready to talk to Gemini. When you open `gemini.google.com` in your web browser, you'll see the main Gemini interface, which typically features a large text box at the bottom of the screen and a "New chat" button somewhere near the top. This is your blank canvas for a fresh conversation.

When you click or tap the "New chat" button, the text box at the bottom will clear, indicating Gemini is ready for your first instruction. This prepares the AI for a new request, ensuring it doesn't get confused with previous topics.

If you don't see a "New chat" button, look for a plus sign (+) icon or a menu option like "Start new conversation" in the top-left corner.

You'll know it worked when the text input area at the bottom of the screen is empty and ready for you to type.

Step 2 — Tell Gemini to generate an image

To create an image, you need to write a **prompt** (the written instruction you give to an AI, acting like a friendly request to a helpful assistant). Gemini needs to know right away that you

want a picture, not just text. Always start your prompt with a clear command for image creation. Type your command into the text box at the bottom of the screen.

After typing your command, press the "Enter" key on your keyboard or click the send icon (often a paper airplane or arrow button) next to the text box. Gemini will then process your request, and you'll typically see a "Generating..." message as it works.

If you accidentally just type a description without the command, Gemini might try to answer you with text. Just clarify in your next message, "Please make an image of that."

Example:

> "Create an image of..."

You'll know it worked when Gemini starts to show a message like "Generating image(s)..." or "I'm working on that for you."

Step 3 — Describe your subject with details

An AI **image generator** (a tool that turns written text into digital pictures) needs details to work its magic. Think of it like describing a scene to a friend who has their eyes closed. After your initial command from Step 2, continue to type in the main text box, adding specific information about what you want to see.

When you add more detail to your prompt and press Enter or the send button, Gemini will process this enriched request. It will use these clues to build a more accurate and vibrant picture, displaying the results directly in the chat window.

If the images don't quite match your vision, it often means your description was a little too vague. Try adding more adjectives or focusing on specific objects, colors, or actions.

Example:

> "Create an image of a fluffy golden retriever puppy wearing a bright red collar, sitting on a sun-drenched green lawn next to a yellow tennis ball."

You'll know it worked when Gemini presents you with usually two to four image options that directly reflect the details you provided in your prompt.

Step 4 — Choose your visual style

Gemini can create images in many different artistic styles. If you don't specify a style, it will usually default to a realistic photo. To change the mood or appearance, add style keywords to

your prompt in the same text box where you've been typing. This instructs Gemini on the artistic direction for the image.

When you include a style in your prompt and submit it, Gemini will generate new images that incorporate that specific artistic look. You'll see the images appear within the chat interface, often in a grid or gallery format. This allows you to easily compare different styles or refinements.

If the style doesn't look quite right, try using a slightly different keyword (e.g., "cartoon" instead of "3D cartoon") or being more specific (e.g., "impressionist oil painting").

Example:

> "Generate a watercolour painting of a quiet beach at sunset, with soft purple and orange colors in the sky."

You'll know it worked when the new images Gemini presents clearly show the artistic style you requested, such as brushstrokes for a "watercolour painting" or distinct lines for "line art."

Step 5 — Refine your generated images

Once Gemini generates your images, you might want to make some tweaks. You don't need to start from scratch. You can simply talk to Gemini like a person, typing a new instruction in the text box that refers to the images it just made. This tells Gemini to adjust the previous output rather than starting a completely new image set.

After you send your refinement request, Gemini will consider your feedback and generate updated versions of the images, incorporating the changes you asked for. These new images will appear in the chat, usually below your request.

If Gemini seems confused by your request, try being more specific about which image you want to change ("the last one," "the first image") or what exact element needs altering ("the background," "the animal").

Example:

> "Can you make the sky in that last picture a bit more blue, and add a single sailboat in the background?"

You'll know it worked when Gemini displays new images that show the specific adjustments you asked for, like a bluer sky or a newly added sailboat.

Step 6 — Save or share your image

Once you have an image you like, you'll want to save it or share it. Look for the image within the chat window. Usually, if you click or tap on an image, it will open in a larger view, where you'll find options to download or share.

When you select the download option (often an arrow pointing down, or a "Download" button), the image will be saved to your device's downloads folder or photo gallery. If you choose to share, a menu might appear allowing you to copy the image or share it directly to other applications.

If you don't see a clear download button, try right-clicking (on a computer) or long-pressing (on a mobile device) the image itself; this usually brings up a context menu with "Save image as..." or "Download image" options.

You'll know it worked when the image file appears in your device's downloads or photo library, ready for you to use or share.

Common mistakes

- **Being too brief:** Writing "a cat" will give you a random result. Instead, give Gemini specific details.
    - **The fix:** Writing "a sleepy tabby cat curled up on a knitted blue blanket in front of a warm fireplace" gives Gemini the clues it needs to make something beautiful and specific to your idea.
- **Asking for readable text:** AI image tools sometimes struggle to spell words correctly or make them look natural inside pictures.
    - **The fix:** It is usually best to generate the image without any text, and then add your words later using a simple design app or photo editor.
- **Using negative words:** AI models respond better to what *should* be in the picture, not what *shouldn't*.
    - **The fix:** Instead of saying "a kitchen with no clutter," try saying "a clean, tidy, minimalist kitchen with polished benchtops." Always frame your request in a positive, descriptive way.

What this means for you

- **In everyday life:** You can design personalized birthday cards, create unique social media posts, or even generate custom phone wallpapers to match your mood.
- **For your business or work:** You can quickly create eye-catching graphics for slide decks, blog posts, or social media updates without needing a massive budget for stock photos or a professional designer.

- **If you're just getting started:** Don't worry about making your prompts perfect. Just play around, try different styles, and see what the AI creates. The best way to learn is by experimenting!

Try it now

Let’s make your very first image. Copy and paste the prompt below into Google Gemini right now to see what happens:

> **"Create a simple, cheerful cartoon illustration of a happy koala holding a warm cup of tea."**

Create Your First AI Image with Midjourney: A Beginner Walkthrough

6 steps · Read online

> **In short:** Midjourney turns a written description into a picture in roughly a minute. You'll need a Midjourney subscription, about ten minutes, and a clear idea of what you want to see. The finished image is yours to download and use.

In the next ten minutes, you'll go from a blank screen to a finished picture saved on your device. No art background needed — if you can describe what you want, you can do this. We'll skip the jargon and focus on the buttons you'll actually press.

Before you start

- **A Midjourney account** — sign up on the official Midjourney website using your email, a Google account, or a Discord login. Generating images requires a paid subscription; current plans and pricing are listed on Midjourney's official pricing page, so check there before you commit.
- **A web browser or the Midjourney web app** — you can work on a computer, tablet, or phone. Discord also works as an alternative entry point if you already have an account there.
- **About 10–15 minutes** of your time, plus another minute or so while the AI draws.
- **One clear idea** of a picture you'd like to see. "A cat" works as a starter; "a calico cat asleep on a stack of books, soft morning light" gives a much better result.

Step 1 — Create your Midjourney account

Open your web browser and go to the official Midjourney homepage. Look for a sign-in or sign-up button — it's usually in the top-right corner, often labeled something like **"Sign In"** or **"Get Started."** Click it, then sign in with your email, a Google account, or a Discord login.

After signing in, the page will guide you through choosing a subscription plan. Pick whichever option feels right for a first try — you can change or cancel later. Once payment is confirmed, the page reloads and drops you into your personal Midjourney workspace.

If you'd rather start in Discord, the same tool runs inside a chat window there. Either path leads to the same image-quality results, so don't worry about choosing wrong.

💬 Example: when asked for your email, use the one you check most often — that's where receipts and account notices will go.

You'll know it worked when you see the Midjourney workspace, with a place to type your picture description and your account name visible somewhere on screen.

Step 2 — Find the create screen

Once you're in your workspace, look for the area where you can start a new image. On the Midjourney web app, this is usually a large text box near the middle of the screen, sometimes paired with a sidebar that lists past images or quick-start options. On mobile, the input area may fill the whole screen.

Click anywhere inside that text box. Your cursor should start blinking there. Don't type yet — just confirm you've found it. This box is where you'll describe your picture, so it's worth a moment to make sure you're in the right place.

If your version looks more like a Discord chat, type **/imagine** in any message field and Discord will offer to autocomplete it. Either way, you're heading to the same destination: a box waiting for your words.

If you don't see a text box at all, look for a button that says something like **"Create"** or **"Imagine,"** or shows a + symbol — clicking it usually opens the input area.

You'll know it worked when your cursor is blinking inside an empty text box and no error message is showing.

Step 3 — Write your first prompt

Now describe one picture in plain words. Imagine texting a friend who paints for you — tell them the subject, the mood, the setting, and the style. A **prompt** is just the name for this description; nothing fancy.

Be specific. *"A flower"* gives you a generic flower. *"A single red poppy in a glass jar, morning sunlight, photo-realistic"* gives you something that actually looks the way you pictured it. Short sentences are fine. So is a comma-separated list.

When your description feels right, look for a button nearby — usually a paper-plane icon, an arrow, or the word **"Generate"** — and click it once. Don't click twice; each click spends a credit and starts a fresh image, so patience saves you money.

💬 Example: `a cozy reading nook by a window, warm lamp light, watercolor illustration, soft colors`

You'll know it worked when the button greys out, a small progress message appears, and within about a minute the first signs of your picture start forming on screen.

Step 4 — Review your generated image

When generation finishes, you'll typically see your image appear on screen. Depending on your settings and the version being used, this might be a single image or a grid of variations. Either way, take a moment to look at what Midjourney created from your description.

If you're looking at a grid and none of the images feels right, look for a button like **"Re-roll"** or a circular ↻ arrow somewhere near the grid. That tells Midjourney to try again with the same prompt and produce new attempts.

You'll know it worked when you can clearly see your generated image(s) and at least one of them makes you think *"ooh, that's nice."*

Step 5 — Upscale your favorite

If you're working with a grid, near it you'll typically find a row of small buttons labeled something like **U1, U2, U3, U4** — one under each image. The **U** stands for **upscale**, which simply means turning the small preview into a full-size, sharp picture. There are usually also **V1–V4** buttons, which create **variations** on that specific image.

Click the **U** button under the picture you like most. Within a few seconds, a single full-resolution version of that image appears. This is the one you'll keep.

If the upscaled image has a small detail you'd like changed — say the cat's eyes are the wrong color — you can run a new variation on it, or simply upscale a different image from the grid instead.

You'll know it worked when a single, larger, sharper image appears alongside or in place of the small grid preview.

Step 6 — Download the image to your device

Click the upscaled image to open it full-size. Look for a download control — this might be a small downward arrow ⬇, a button that says something like **"Save"** or **"Download,"** or an option inside a right-click or menu. One tap or click should save a PNG or JPG file straight to your device.

On a phone, the image often opens in a new tab. From there, look in your browser's menu (often three dots or a share icon) for an option to save the picture to your Photos or Gallery app. On a computer, the file usually lands in your Downloads folder.

The file is now yours. You can post it, print it, set it as your wallpaper, or send it to a friend. Midjourney's standard paid plans deliver images without a watermark, so the picture is ready to

share as soon as it's saved.

You'll know it worked when the file appears in your **Downloads** folder (computer) or **Photos / Gallery** app (phone), and you can open it again from there.

Common mistakes

- **Writing a one-word prompt.** "Cat" gives you a random cat. Always include at least one detail about mood, style, or setting — Step 3 has an example you can copy and tweak.
- **Clicking Generate twice in a hurry.** Once is enough. Every extra click spends a credit and starts a fresh image. If the first result isn't ready yet, just wait — a progress indicator should be visible.
- **Asking the AI to draw readable text.** Midjourney struggles with words inside images — a sign that says "OPEN" often comes out misspelled or jumbled. Use a real text tool or a graphic editor for readable words, and let Midjourney handle the rest of the picture.

What this means for you

- **In everyday life:** your first image is great for a phone wallpaper, a birthday card, or a custom avatar — small personal touches that feel one-of-a-kind without paying a designer's fee.
- **For your business or work:** a single Midjourney picture can become a blog header, a slide background, or a quick mockup for a client conversation. It's not a replacement for a designer on real projects, but it's a fast way to test an idea visually.
- **If you're just getting started:** the easiest first move is to re-run the same prompt with one word changed each time — *"watercolor"* becomes *"oil painting"* becomes *"pencil sketch"* — and watch how the style shifts while the subject stays the same.

Try it now

Take one minute and write down a single sentence describing a picture you actually want to see. Then open Midjourney, paste it into the prompt box, and hit Generate. By the time you finish reading this paragraph, your first image will probably be ready to look at.

How to Create AI Logos with Recraft V4.1 Vector

6 steps · Read online

> **In short:** Recraft describes a two-stage AI logo workflow — generate lots of concepts on a raster model first to explore ideas, then switch to Recraft V4.1 Vector for the final mark. The vector output is an SVG file that stays crisp from a 16-pixel browser tab icon to a printed sign.

Designing a logo used to mean either learning vector design software yourself or paying a designer. Today, AI tools like Recraft can sketch dozens of ideas in seconds. The catch is the file format: a regular image — a "raster" picture made of tiny colored squares called pixels — can look fuzzy when you blow it up or shrink it down. A vector file, built from mathematical lines instead of pixels, stays sharp at any size.

This guide walks through Recraft's two-stage approach: explore concepts as raster images first, then generate the final logo in Recraft V4.1 Vector. The final file is an SVG (Scalable Vector Graphics) you can use on a tiny favicon, a website header, or a printed sign without any blurriness.

Before you start

- A free Recraft account (sign up at recraft.ai with an email address or a Google account — check the site for the current plan options, since free and paid tiers exist)
- A recent web browser (Chrome, Firefox, Safari, or Edge)
- A rough idea of your brand: name, vibe, and one or two colors you like
- About 20–30 minutes for your first attempt

Step 1 — Sign in and start a new project

Open recraft.ai in your browser and sign in with the account you just made. Look for a button that says something like "New design" or "Create" — on most creative tools it sits near the top-left of the canvas. After you click it, a blank workspace opens. You'll see a text box where you type what you want and some style controls to the side. If your screen looks crowded instead of clean, look for a "Blank canvas" option in a menu — it's the simplest starting point.

💬 Example: type a one-line description like "warm coffee shop, hand-lettered feel, soft cream and forest green" into the prompt box.

You'll know it worked when the canvas is empty and your cursor is blinking in the prompt field, ready for instructions.

Step 2 — Generate a batch of concepts on a raster model

In the style controls (often a sidebar or a row of icons above the prompt box), pick the option that produces raster-style images — sometimes labeled "Raster," "Photo," or shown with a small photo icon rather than a vector-shape icon. Recraft describes starting with raster because it's fast and easy to compare many ideas side by side. Then describe your logo in plain words and click the button that generates results — it usually looks like a small arrow, a sparkle ✨, or the word "Generate." Expect to wait about 10–20 seconds.

💬 Example prompt: "minimalist logo for a yoga studio, lotus flower, calm pastel palette, no text in the image."

You'll know it worked when several image options appear on the canvas within half a minute. If nothing appears, look for a small "Retry" or "Generate again" link — sometimes the first attempt stalls.

Step 3 — Pick a direction and copy what works

Scroll through the generated options and click the one that feels closest to what you want. Most AI image tools let you "favorite" or "select" an image — look for a heart icon, a star, or a "Use this image" button when you hover over it. You're not done yet: this is just your starting point. Read the prompt you typed in Step 2 and tweak the words that mattered most — shape, mood, colors — so your next round focuses on what you liked. This is the part of Recraft's workflow where you "lock in" the direction.

💬 Example: if you loved the lotus shape but the colors felt too cold, change "calm pastel palette" to "warm peach and sand palette" before you continue.

You'll know it worked when you can point to the one image you want to keep, and you have a tighter prompt that names its strongest traits.

Step 4 — Switch to Recraft V4.1 Vector for the final mark

Now go back to the style controls and switch from the raster setting to the vector setting. Recraft describes V4.1 Vector as the model that returns a clean SVG instead of a pixel image. SVG is just a file format where shapes are stored as mathematical instructions — like a recipe for drawing a circle — rather than a fixed grid of colored dots. Look for a label like "Vector," "V4.1 Vector," or an icon that looks like geometric shapes rather than a photograph. Make sure your refined prompt is still in the text box, then generate again.

💬 Example: paste your tightened prompt from Step 3 — for instance, "minimalist lotus logo, warm peach and sand palette, simple shapes only, no text."

You'll know it worked when your result looks like a clean illustration with crisp edges — you can zoom in really close without it turning blurry — and the file is described somewhere in the

interface as an SVG.

Step 5 — Export the SVG and verify it at small sizes

Find the export or download control — usually an icon that looks like a downward arrow ⬇, a share icon, or a button labeled "Export" or "Download." If asked to choose a format, pick SVG (not PNG or JPG — those are raster formats that go soft when scaled). Save the file somewhere easy to find, like your Desktop. Then comes the honest test: open the SVG on your phone or in any image viewer and shrink it to about 16 pixels tall — that's the size of a browser tab favicon. If your logo still reads clearly at that size, it works. If the details turn to mush, you may need to simplify the design.

💬 Example: open the file in a browser tab, then zoom out until it looks tiny. If the lotus shape is still recognizable, you're good.

You'll know it worked when the same SVG file looks sharp both as a 16-pixel favicon and at the size of a website header — no blurriness at either end.

Step 6 — Save variations and name your files well

Before you call it done, generate one or two more versions with small tweaks — different colors, a simpler version, or one with your business name typed in. Most tools let you download each result separately. Give every file a clear name that describes what's inside — your future self will thank you when you have six SVGs on your desktop and can't tell them apart. A simple naming pattern works: `brandname-logo-style-color.svg`, something like `lotus-logo-minimal-peach.svg`.

💬 Example: `lotusyoga-logo-minimal-peach.svg`, `lotusyoga-logo-simple-mono.svg`, `lotusyoga-logo-with-text.svg`.

You'll know it worked when you can open any file two weeks from now and instantly know which version it is, without having to open them all.

Common mistakes

- **Skipping the raster stage and going straight for vector.** The vector model is great for the final mark, but it's slower to compare ideas. Generate lots of cheap raster options first, then commit to vector once you know what you want — that two-step rhythm is the whole point of Recraft's described workflow.
- **Exporting as PNG or JPG instead of SVG.** Raster files look fine on screen but turn blurry the moment you scale them up or shrink them down. Always pick SVG for the logo you plan to use everywhere.

- **Adding too much detail.** Tiny text, thin lines, and busy patterns can vanish when the logo is shrunk to favicon size. If your design falls apart at 16 pixels, simplify it — a strong logo usually has one or two clear shapes, not ten.

What this means for you

- **In everyday life:** you can finally make a logo for your side project, bakery, garage band, or dog-walking business without paying a designer or learning Illustrator. Twenty minutes of experimenting is enough for a first version you can actually use.
- **For your business or work:** a single SVG file travels anywhere — your website header, email signature, business card, social media profile, even the tiny favicon in the browser tab — and looks identical at every size. That means one design, no redesigns.
- **If you're just getting started:** don't aim for perfect on your first try. Generate a batch, pick the one that makes you smile, then refine it. AI logo tools reward quick iteration, not overthinking.

Try it now

Open recraft.ai, type "friendly minimalist logo for a [your project name], [two colors you like]" into the prompt box, and generate a raster batch. Spend five minutes just looking at the results and noticing which shapes and colors pull you in — that's your style direction, and you'll use it as the seed for your vector version.

How to Create Custom Images with Copilot for Invites and Social Posts

6 steps · Read online

> **In short:** Copilot can turn a plain sentence into a custom image you can use on invitations, Instagram, or a personal vision board. Sign in with a Microsoft account, open the Designer tool, describe what you want in plain words, and download what comes out.

You don't need Photoshop, Canva, or any design background to get a one-of-a-kind picture for your next event, post, or vision board. Copilot's built-in image generator — Microsoft Designer — takes a sentence from you and turns it into a finished image in roughly 30 seconds. This walkthrough shows the whole flow, including what to write so the picture actually comes out the way you imagined.

Before you start

- A **free Microsoft account** (the same one you use for Outlook, OneDrive, or Xbox). If you don't have one, signing up takes about two minutes.
- A **browser** (Chrome, Edge, Safari, or Firefox) or the **Copilot mobile app**. Either works.
- Roughly **10–15 minutes** for your first picture — faster once you've done it a few times.
- The image generator is free to use, but the free plan has a **daily limit** on how many images you can create. If you want a lot more images or faster generation, look into the paid **Microsoft Copilot Pro** plan on Microsoft's official pricing page for the current rate.

Step 1 — Open Copilot and switch to image mode

Go to **copilot.microsoft.com** in your browser, or open the Copilot app on your phone. Sign in with your Microsoft account if you aren't already. Look for a chat box near the bottom of the screen that says something like **"Ask me anything..."** — and right next to it, find a small **image or picture icon** 🖼️, or a toggle labeled **"Designer"** or **"Create an image."** Tap or click that option. The chat box changes to invite an image description instead of a question, and a small note about Microsoft's content policy may pop up the first time. 💬 Example: don't type anything yet — just make sure you can see the image-creation prompt box. You'll know it worked when the chat box now asks you to **describe an image** instead of asking a question.

Step 2 — Describe the picture you actually want

Click inside the new prompt box and write a sentence or two describing your image. Be specific: include the **subject**, the **style** (watercolor, photo-realistic, cartoon, flat design), the **mood** (cheerful, calm, dreamy), and the **colors** if they matter. The more detail you give, the closer the

result matches what you imagined. If your screen looks different, look for a separate **"Designer" link** in the top menu, or a button labeled something like **"Image creator"** — Microsoft moves these around as they update the site. Example: "A watercolor birthday invitation background with soft pink peonies, gold confetti, and the words 'Emma turns 30' written in elegant script at the top." You'll know it worked when you hit Enter and Copilot starts **generating four different versions** of your image — this usually takes 15–30 seconds.

Step 3 — Review the four versions and pick your favorite

When the images finish loading, you'll see a grid of **four square images**, all based on your description. Don't rush — open each one in full size by clicking it, and check the small details that matter for your project: is the text readable, are the colors right, does the mood fit the occasion? If none of them are quite right, you don't need to start over — go back to the prompt box and describe what you'd change. Copilot can run another round based on your new description. Example: "Same image but make the peonies larger, and put the text 'Emma turns 30 — Saturday June 14' in the center." You'll know it worked when at least one of the new four images is **close enough to use**, or when you've nailed down exactly which change to ask for next.

Step 4 — Polish the prompt if the text in the image looks wrong

Here's an honest limitation: image generators — Copilot included — are **still hit-or-miss at writing words inside pictures**. Names, dates, and quotes can come out misspelled, scrambled, or warped. If your invite or vision board needs accurate text (like a name or a date), the safest workflow is to **leave the text out of the prompt entirely**, then add it yourself later in a free tool like Canva, PowerPoint, or even your phone's photo editor. This is one extra step, but it guarantees your text actually says what it should. Example: for the invite, ask for "A watercolor birthday invitation background with soft pink peonies and gold confetti, no text." Then you add "Emma turns 30 — Saturday June 14" yourself. You'll know it worked when you have a beautiful background image **without any text problems to fix**.

Step 5 — Download the image you want to use

Hover your cursor over the image you like best. A small toolbar appears underneath — look for a **download icon** (a downward arrow into a tray) or a button labeled **"Download"**. Click it once. Your browser saves the image to your normal Downloads folder, usually as a PNG or JPG file. On the phone app, the download button might say **"Save"** or look like a square with an arrow — if you can't find it, take a screenshot instead, which works on every phone. Example: click the **Download** button below your favorite peony image and wait for the file to appear in your Downloads folder. You'll know it worked when you can open the file on your computer or phone and see the finished picture, ready to drop into your design.

Step 6 — Drop the picture into your project

Open whichever tool you normally use — Canva, PowerPoint, Google Slides, Instagram, your phone's photo editor — and **import or insert** the image you just downloaded. Resize it to fit your canvas: square (1:1) for Instagram, portrait (9:16) for stories, landscape (16:9) for invitations printed on letter-size paper. If you want a different shape than the square Copilot gave you, you can either crop with the same tools or describe a different aspect ratio in Step 2 next time (some Copilot versions accept hints like "wide landscape format" or "tall poster format"). Example: in Canva, click **"Uploads"**, drag in your downloaded peony image, drop it onto a square Instagram template, and add the text "Emma turns 30 — details in bio" on top. You'll know it worked when your finished design **looks unique to you** and is ready to send, post, or print.

Common mistakes

- **Asking for text inside the image.** Image generators often misspell names and dates. Strip the text out of the prompt and add it yourself in Canva or your editor — it'll be sharper and faster.
- **Vague prompts.** "A nice background" gives you something generic. **"Watercolor peonies, pink and gold, birthday mood, square"** gives you something usable.
- **Stopping after the first round.** The first four images are a starting point, not the finish line. If they're close, ask for one specific change. The second round is usually the keeper.

What this means for you

- **In everyday life:** the next birthday invite, save-the-date, or holiday card can be made in the time it takes to brew coffee — no software to install, no design skills needed.
- **For your business or work:** social posts, internal event flyers, or vision-board slides for team workshops can be created on the fly. If you batch three or four prompts in one sitting, you have a week's worth of visuals before lunch.
- **If you're just getting started:** open Copilot today and try the birthday invite example from Step 2. Once you've made one, every future prompt is easier.

Try it now

Open **copilot.microsoft.com**, switch to image mode, and type: *"A dreamy watercolor vision board image with a calm beach at sunrise, a notebook, a cup of coffee, and soft pastel colors, no text."* Save the result you like best. That's your first custom picture — made in under five minutes.