

Everyday AI: Understand anything

12 step-by-step guides for real, everyday tasks

Understand anything

AI World HQ · <https://aiworldhq.com>

Free edition · CC BY-NC 4.0 · free to share with credit

Before you start — what this is, honestly

This collection is free, and every guide in it is also free on our website. This file just gathers them in one place, in reading order, so you can print it, keep it on your phone, or read it offline. Our website has no paywall either — every guide links back to its online page, which is always the newer version.

AI Every guide in this pack was written by AI and checked by automated fact and quality gates. **No human editor reviewed the text.** Every guide page on <https://aiworldhq.com> carries the same notice — we say it openly because you deserve to know who wrote what you are reading.

Apps change. Buttons move and menus get renamed. If a step does not match what you see, the guide's online version is the newer one — the link is under every title.

Free to share. This collection is licensed under Creative Commons Attribution-NonCommercial 4.0 (CC BY-NC 4.0): you may copy, print and share it — with your class, your library or your team — as long as you credit AI World HQ with a link to aiworldhq.com and do not sell it.

What's inside

1. Understand anything

1. How to Ask GitHub Copilot to Explain a Confusing Error Message
2. Compare two job offers side-by-side with Le Chat in plain English
3. Quickly Research Any Topic and Summarize Web Pages with Microsoft Copilot in Edge
4. Fact-Check Online Claims with Perplexity's Sources
5. How to Summarize a Long Article or PDF into Key Bullet Points
6. Translate a product description into Spanish and French with Cohere
7. How to compare two rental listings side by side with weighted pros and cons
8. How to use Gemini Deep Research to plan a big purchase decision
9. Catch Up Fast: Summarize Long Chats in WhatsApp and Messenger with Meta AI
10. Summarize a long PDF into 5 bullet points with Claude
11. Summarize Long Articles into Key Bullet Points with Grok on X
12. Translate and Explain Any Restaurant Menu with Qwen on Your Phone

Understand anything

How to Ask GitHub Copilot to Explain a Confusing Error Message

7 steps · [Read online](#)

In short: GitHub Copilot's chat can turn a confusing error message into plain English and a short list of things to check — you paste the error into the chat, then ask a simple question. It works inside editors like VS Code and on github.com; some chat features may require a paid plan, so check the official pricing page for the latest details.

You don't need to be a developer to get value from this. Plenty of people run into error messages while learning to code, tweaking a website, or running a script at work. GitHub Copilot is an AI coding assistant made by GitHub — think of it as a patient tutor sitting next to you, reading the error out loud in normal English. Here's the beginner-friendly way to use it.

Before you start

- A free GitHub account, signed in. Some chat features may require a paid plan — if a feature asks you to upgrade, that's why.
- The error message itself, copied from wherever you saw it (your terminal, your browser console, a log file, or your code editor).
- A code editor with the Copilot extension installed (VS Code is the most common), **or** the Copilot chat panel on github.com.
- Roughly 5 minutes.

Step 1 — Install the Copilot extension in your editor (if you haven't yet)

Open VS Code and look at the left sidebar — the narrow vertical strip with several icons. Click the square **Extensions** icon (it usually looks like four small squares, with one square separated from the others). A search panel opens on the left. Type **GitHub Copilot** into the search box at the top of that panel. The top result is published by GitHub itself — click **Install** on that row. After a few seconds, a new icon appears in your sidebar (often a chat-bubble shape). If you use a different editor like JetBrains or Visual Studio, the steps are similar — open your editor's Extensions or Plugins menu, search for "GitHub Copilot," and click Install. On github.com, you can skip this step entirely and go straight to Step 2. 💡 Example: (no typing needed in this step.) You'll know it worked when you see the Copilot icon in your sidebar, or when github.com shows a "Copilot Chat" option in its top menu.

Step 2 — Open the Copilot chat panel

Look for the Copilot icon you just installed — it's usually in the left or right sidebar, near the top. Click it once. A new side panel opens up, with a text box at the bottom and a label like "Ask Copilot" or "Send a message" above it. If you don't see the icon, try the keyboard shortcut: on Windows or Linux, press **Ctrl + Shift + I**; on macOS, press **Cmd + Shift + I**. The chat panel pops up the same way. Some setups hide the icon inside a **⋮ (three dots)** menu instead — open that menu and pick "Copilot Chat." 🗨 Example: (no typing needed in this step.) You'll know it worked when you see a blinking cursor inside a text box labeled something like "Ask Copilot" or "Send a message".

Step 3 — Paste the full error message

Click inside that text box and paste your error. To paste, use **Ctrl + V** on Windows/Linux or **Cmd + V** on macOS. The error text appears in the box, usually in red or in a monospaced font (where every letter takes the same width, like a typewriter). Try to copy the whole thing, not just the first line — most errors include a **stack trace** (a list of which files and lines were involved, like a breadcrumb trail showing where the problem happened) right under the main message, and that context is what helps Copilot pinpoint the problem. If the error is in your browser, open the **Console** tab first (usually under the browser's "More tools" or "Developer" menu), then copy from there. 🗨 Example: `TypeError: Cannot read properties of undefined (reading 'map') at App.render (App.jsx:14)` You'll know it worked when the error text sits in the chat box, exactly as it appeared on screen.

Step 4 — Ask for a plain-language explanation

Now type one short sentence *after* the error, asking Copilot to translate it. This part matters: the question you ask shapes the answer you get. Keep it simple and direct — no jargon required. Press **Enter** to send. Within a few seconds, Copilot replies with a friendly explanation of what the error means, in everyday words. If you don't see a response right away, look for a small **stop** or **regenerate** icon below your message and click it. If nothing comes back at all, your plan may not include chat — the reply box will usually show a small note asking you to upgrade. 🗨 Example: `Explain this error to me like I'm new to coding. What is it actually telling me to fix?` You'll know it worked when the reply describes the problem in normal English — something like "your code tried to use a list that doesn't exist yet" — instead of repeating the original error back at you.

Step 5 — Ask for a checklist of likely fixes

Don't stop at the explanation. Right below Copilot's reply, type a follow-up asking for a short list of the most likely causes and what to check first. Copilot usually responds with a numbered list — each item is a hypothesis to test against your own code. Treat the list as a **checklist**, not a verdict: not every suggestion will fit your situation, and that's normal. If the answer feels too

generic, add one line about what you're building (for example: *"I'm using React"* or *"I'm running Python on Windows"*) and ask again. 🗨 Example: `List the 3-5 most likely causes of this error. For each one, tell me what to look at in my code.` You'll know it worked when you get a numbered list with at least three concrete things to check, like "the variable might be empty" or "the function might be called before the data loads".

Step 6 — Test the suggestions against your own code

Open the file mentioned in the error (in the example, `App.jsx` on line 14). Read the line Copilot pointed to and see which suggestion fits. If none of them seem to match, paste a small chunk of your actual code into the chat and ask Copilot which suggestion is most likely. This is where the AI becomes useful — instead of guessing, you're now comparing each idea to the real code in front of you. One honest note: Copilot can be confidently wrong, especially for unusual setups — its suggestions are a starting point, not a final answer. 🗨 Example: `Here's the function that produces the error. Which of your suggestions fits this code? function render() { return items.map(...) }` You'll know it worked when at least one suggestion clicks — you think "oh, that's probably it" — and you have a clear next step to try.

Step 7 — Ask a follow-up when something still doesn't make sense

If a word in the explanation or the list is unfamiliar, just ask. Tell Copilot exactly which term confused you, and it'll rephrase that one part. This is the secret most beginners miss: the chat is a conversation, not a one-shot answer. Each follow-up makes the next answer tighter and more useful for your specific situation. If the answer goes off-track, click the small **refresh** or **try again** icon to get a new draft. 🗨 Example: `You mentioned "undefined". I don't fully get what that means in this context – can you explain just that word?` You'll know it worked when you can say the idea back in your own words, and you have at least one concrete thing to change in your code.

Common mistakes

- **Pasting only the first line of the error.** The lines below it (the stack trace) usually point to the file and line where things went wrong. Copy the whole block.
- **Forgetting to say what language or library you're using.** "Null" in JavaScript and "NULL" in a database mean similar but different things. Add one short line: *"I'm using React."*
- **Trusting the very first answer as the final answer.** Copilot can be confidently wrong, especially for unusual setups. If two or three ideas don't match your code, don't force the fourth — go back to the error message itself and ask again.

What this means for you

- **In everyday life:** when any tool gives you a wall of red text — a script, a build tool, a deployment — the same trick works. Paste, ask for a plain-language read, then ask for a checklist.
- **For your business or work:** faster debugging, fewer interruptions to teammates, and a growing habit of reading errors instead of fearing them. Over time you'll spot patterns and need the chat less.
- **If you're just getting started:** pick the next error you see today and use it as practice. One error a day, explained, adds up quickly.

Try it now

Open Copilot chat, paste the last error message you ran into — even if it's from yesterday — and send this single line:

💬 Example: Explain this error to me like I'm new to coding, then list the 3 most likely things to check.

You'll have a clear answer in under a minute.

Compare two job offers side-by-side with Le Chat in plain English

7 steps · [Read online](#)

By the end of this guide you'll be able to paste the key points of two job offers into Le Chat and receive a tidy, easy-to-understand table that highlights the differences—no legal jargon required. This is for anyone juggling two offers and wants confidence before signing.

Before you start

- A free (or paid) Mistral account so you can log into Le Chat.
- Access to the two job offers – either as PDF files you can copy from, or as plain-text emails.
- A computer or smartphone with an internet browser.
- About 10 minutes total – the guide walks you through every click and prompt.

Step 1 — Open Le Chat and start a fresh chat

Open your web browser and go to the Le Chat website (chat.mistral.ai). Sign in with your Mistral email and password. You'll see a clean white chat window with a text box at the bottom and a "New chat" button near the top-right corner. Click the "New chat" button; a fresh conversation pane slides in and the cursor blinks in the text box. If the button looks different on your screen (for example, it might be a plus-sign icon or simply the words "Start new chat"), look for the same meaning. 🟡 Example: *"Hello Le Chat, I need help comparing two job offers."* You'll know it worked when the chat area shows a greeting like "How can I help you today?" or "New conversation started."

Step 2 — Copy the key details from each offer into plain text

Open the first job offer in your email or PDF viewer. Select only the most important sections—typically base salary, bonuses, work hours, leave entitlements, notice period, and any special perks. Copy the selection (Ctrl+C on a computer, long-press → Copy on a phone). Repeat for the second offer, keeping the two chunks separate on your clipboard. Back in Le Chat you'll see the empty text box ready for your first prompt. If you cannot copy from a PDF, use a free PDF-to-text converter or type the details manually. 🟡 Example: *"Offer A: \$95,000 base salary; 15 days annual leave; 1 month notice; \$5,000 signing bonus. Offer B: \$92,000 base salary; 20 days annual leave; 2 weeks notice; 3% profit share."* You'll know it worked when the copied text appears in the box after you paste (Ctrl+V).

Step 3 — Ask Le Chat to build a side-by-side comparison table

With the two offer snippets pasted into the text box, type a clear instruction: *“Create a side-by-side comparison table in plain English for Offer A and Offer B. Include salary, annual leave days, notice period, bonuses, and any other notable perks. Avoid legal jargon.”* Press Enter or tap the send arrow. Le Chat will start typing a response that usually begins with a markdown-style table and then bullet points summarizing the contrast. If the assistant replies with a single paragraph instead of a table, simply add “Show it as a table” to the prompt. 🗨️ Example: *“Compare these two offers side-by-side and explain the pros and cons of each in simple language.”* You’ll know it worked when you see a two-column table with headings like “Offer A” and “Offer B” and simple sentences underneath each row.

Step 4 — Review the table and ask for missing rows

Look at the table Le Chat has generated. Each row should correspond to a contract element (salary, leave, notice, bonuses, etc.) and each column should contain the figure or phrase from the respective offer. If a row is missing (for example, “profit share”), type a follow-up such as *“Add a row for profit share details for both offers.”* The assistant will update the table accordingly. On some devices the table may appear as plain text rather than a rendered grid; if you see vertical bars (“|”) and dashes (“—”) instead of a visual grid, that’s still the correct markdown format. 🗨️ Example: *“Can you explain in everyday terms what the different notice periods mean for my life?”* You’ll know it worked when the assistant adds the requested row or provides a short, jargon-free explanation beneath the table.

Step 5 — Copy the final comparison into your notes

When the table looks complete, click the three-dot menu (or right-click) on the response and choose “Copy message” or “Download as text” depending on your browser. Paste the copied result into a notes app, email to yourself, or print it out. If the copy option isn’t visible, you can manually select the whole response and press Ctrl+C (or long-press → Copy on a phone). 🗨️ Example: *No extra prompt needed – just use the copy command.* You’ll know it worked when the entire table appears in your notes or email, ready for you to review side-by-side with a coffee.

Step 6 — Double-check the numbers against the original offers

Open the original offer documents again and compare the figures in your notes with the source. If anything looks different, paste the discrepancy into Le Chat and ask for clarification: *“The original Offer A PDF shows \$95,000 base salary, but your table shows \$92,000. Which is correct?”* Le Chat will usually point to the exact line in the text you pasted earlier. If the mismatch remains, paste the relevant paragraph again and ask for a corrected table. You’ll know it worked when the numbers in your notes match the original documents.

Step 7 — Decide with confidence

With the plain-English table in hand and the numbers verified, you now have a clear snapshot of the two offers. Take five minutes to weigh salary, leave, notice, bonuses, and any other perks against your personal priorities. You'll know you're ready to decide when you can confidently say which offer better supports your lifestyle and goals.

Common mistakes

- **Copying the entire PDF instead of the key points** – the assistant may get overwhelmed and miss important details. Fix: only copy salary, bonuses, leave, notice, and perks.
- **Using legal-sounding wording in the prompt** – Le Chat might mirror that jargon. Fix: ask explicitly for “plain English” or “simple language” in your first instruction.
- **Forgetting to ask for a table** – you may receive a paragraph of text. Fix: add “show it as a table” or “format as a two-column chart” in a follow-up prompt.
- **Not verifying the numbers against the originals** – a small typo can change the outcome. Fix: always compare the final table with the source documents before deciding.

What this means for you

- **In everyday life:** You can instantly see which offer gives you more vacation days or a higher take-home pay, without needing a lawyer or calculator.
- **For your business or work:** HR professionals can use the same approach to benchmark internal promotions against external market offers in minutes.
- **If you're just getting started:** The easiest first move is to copy the salary and leave details from one offer and ask Le Chat for a quick plain-English summary.

Try it now

Open Le Chat, paste the salary and leave figures from one of your current job offers, and type: *“Summarize these numbers in simple language.”* In under two minutes you’ll see how Le Chat turns raw numbers into an easy-to-understand sentence.

Quickly Research Any Topic and Summarize Web Pages with Microsoft Copilot in Edge

7 steps · [Read online](#)

Imagine reading a long article or research page and getting a clear, concise summary in seconds—without leaving your browser. With **Microsoft Copilot in Edge**, you can do exactly that. Whether you're studying, working, or just curious, this tool helps you understand web pages quickly and answer your questions on the spot. And it's free.

Before you start

- A **computer, tablet, or phone** with internet access.
- The **Microsoft Edge browser** installed (free from microsoft.com/edge).
- A **Microsoft account** (free to create at account.microsoft.com).
- About **5 minutes** of your time.

 **Tip:** You don't need to sign in to use basic features, but signing in lets you save your chats and access them across devices.

Step 1 — Open Copilot in Edge

Open the **Microsoft Edge** browser on your device. Look at the top-right corner of the screen. You'll see a small **blue Copilot icon** (it looks like a sparkle or a chat bubble with a sparkle inside).

Tap or click the **Copilot icon**. A sidebar will slide open from the right side of your screen. This is where Copilot lives—it's like a helpful assistant inside your browser.

You'll know it worked when ... the Copilot sidebar appears with a message like “Ask me anything” or “What can I help you with?”

 **Example:** Just click the sparkle icon in the top-right corner of Edge.

Step 2 — Choose “Summarize” from the menu

Inside the Copilot sidebar, look for a dropdown menu at the top. It might say “Chat” or “Compose” by default. Click the dropdown and select “**Summarize**”.

Now, Copilot is ready to summarize the web page you're viewing. It will read the content of the page and create a short, clear summary for you.

You'll know it worked when ... the sidebar shows a message like "Summarizing this page..." and then displays a bullet-point summary below.

💡 Example: In the Copilot sidebar, click the dropdown and choose "Summarize".

⚠️ If you don't see "Summarize" in the dropdown, make sure you're on a web page with text (not a video or image-only page). Also, look for a button like "📄 Summarize" or "📄 Summarize page".

Step 3 — Read the AI summary

After a few seconds, Copilot will show you a **short summary** of the web page. It might include:

- The main points or arguments
- Key facts or data
- A conclusion or takeaway

You can scroll through the summary to get the gist of the page without reading every word.

You'll know it worked when ... you see a clear, bulleted list or short paragraphs that capture the most important information from the page.

💡 Example: You'll see something like:

- "This page explains how to compost at home."
 - "It lists 5 easy steps and common mistakes to avoid."
 - "It also includes a list of local composting services."
-

Step 4 — Ask follow-up questions

After the summary appears, you can ask Copilot to **clarify or expand** on any point. Just type your question in the chat box at the bottom of the sidebar.

For example:

- "Can you explain point 3 in more detail?"
- "What are the benefits of composting?"
- "Give me a shorter version."

Copilot will respond based on the content of the page. This is like having a conversation with the article itself.

You'll know it worked when ... Copilot answers your question using information from the page, and you see a new message appear in the chat.

💡 Example: Type: "What are the top 3 benefits of composting mentioned here?"

⚠️ If Copilot gives an answer that doesn't match the page, try rephrasing your question or check if the page content changed. Copilot only knows what's on the current page.

Step 5 — Summarize a new page without leaving your tab

To summarize a **different web page**, just navigate to it in your browser. Copilot will automatically detect the new page and ask if you want to summarize it.

You can say "Yes" or click the "Summarize" button that appears in the sidebar. Copilot will generate a new summary for the new page.

You'll know it worked when ... the sidebar updates with a new summary of the page you're now viewing.

💡 Example: Open a new article or research page, and Copilot will prompt: "Want me to summarize this page?"

Step 6 — Save or share the summary

If you want to save the summary for later, you can:

- **Copy the text** and paste it into a document or note app.
- **Take a screenshot** of the summary.
- **Ask Copilot to save it** by typing: "Save this summary to my notes."

You'll know it worked when ... you have the summary text in your notes, or you see a confirmation like "Summary saved."

💡 Example: Type: "Save this summary to my notes" to store it in your Copilot chat history.

⚠️ On the free plan, your chat history is saved for a limited time. If you want long-term access, sign in with your Microsoft account.

Step 7 — Use Copilot on mobile (optional)

If you're using a phone or tablet, open the **Edge app**. Tap the **Copilot icon** (sparkle) in the bottom toolbar. Then tap the "**Summarize**" button at the top of the Copilot panel.

Copilot will summarize the page you're viewing. You can scroll, ask questions, or copy the summary—just like on desktop.

You'll know it worked when ... the summary appears in the Copilot panel on your mobile device.

💡 Example: In the Edge mobile app, tap the Copilot icon, then choose “Summarize”.

⚠️ On mobile, the interface may look slightly different. Look for a button like “ Summarize” or “Summarize page”.

Common mistakes

- **Mistake:** Trying to summarize a video or image-only page. **Fix:** Copilot only works on pages with text. Make sure the page has readable content.
 - **Mistake:** Expecting Copilot to understand pages you haven't opened. **Fix:** Copilot only sees the page you're currently viewing. Open the page first, then ask for a summary.
 - **Mistake:** Not signing in and losing your chat history. **Fix:** Sign in with your Microsoft account to save your summaries and access them later across devices.
-

What this means for you

- **In everyday life:** Quickly understand long articles, research papers, or product pages without reading every word. Great for students, parents, or anyone short on time.
 - **For your business or work:** Save hours reviewing reports, competitor websites, or industry updates. Ask Copilot to pull out key data or trends from any page.
 - **If you're just getting started:** Start by summarizing one article you're already reading. Try asking a follow-up question to see how it works.
-

Try it now

Open **Microsoft Edge**, go to any article or research page, and click the **Copilot icon** in the top-right corner. Choose “Summarize” from the dropdown, and see your first AI summary appear in seconds.

Fact-Check Online Claims with Perplexity's Sources

6 steps · [Read online](#)

In our busy online world, information flies around faster than ever. Sometimes, a post on social media might sound convincing but isn't quite right. This guide will show you how to use Perplexity, an AI assistant, to quickly check the facts and understand where information truly comes from. By the end, you'll be able to confidently verify claims, helping you decide what's reliable to share. While Perplexity is excellent at finding and compiling information, remember it doesn't perform human critical analysis or judgment; it's always up to you to evaluate the original sources.

Before you start

- A device with internet access (computer, tablet, or smartphone).
- Access to Perplexity, either through its website (perplexity.ai) or its mobile app. A free account is sufficient for this guide.
- Roughly 5-7 minutes of your time.

Step 1 — Open Perplexity

First, let's get Perplexity ready. Open your web browser and type `perplexity.ai` into the address bar at the top of the screen, then press your keyboard's Enter key. If you're on a mobile device, find and launch the Perplexity app. Once loaded, you'll typically see a clean, simple screen featuring a prominent search bar right in the middle, often displaying placeholder text like "Ask anything..." or "Search the web". This is your main input area, waiting for your questions.

If it looks different: The specific layout might vary slightly depending on whether you're using the website or the app, or due to recent updates. However, you should always find a clear, large text box or search bar positioned centrally, as it's the most obvious place to type your query.

You'll know it worked when you see the Perplexity interface with a clear text input area, ready for you to type your question.

Step 2 — Ask about your claim

Now, let's ask Perplexity to investigate a social media claim. Carefully type the specific statement or question you want to check directly into that main search bar. Try to phrase your query clearly and directly, just as you would ask a person for information. For example, if you saw a post about a new animal, you might ask "Is the 'blue-spotted tree frog' a real animal?". Once you've finished typing your question, press the Enter key on your keyboard or tap the "Ask"

button (which might appear as an arrow or a magnifying glass icon) usually located next to your input.

If it looks different: If there isn't a clearly visible "Ask" button, simply pressing Enter after typing your question usually works on most devices. Make sure your question is specific enough so Perplexity understands exactly what you're trying to verify.

💡 Example: Is it true that taking a cold shower every morning can cure the common cold?

You'll know it worked when Perplexity starts to process your request, often showing a loading animation, a spinning icon, or a "Generating answer..." message as it looks for information.

Step 3 — Review the summary and its citations

After a brief moment, Perplexity, acting as an **AI assistant** (a computer program designed to help with tasks and provide information), will present you with a **summary answer**. This is the AI's best attempt to directly address your question, drawing information from various online sources. Read through this summary carefully. You'll typically find a concise answer, along with key facts or explanations. Notice that within the summary, there will be small **numbered references**, like [1], [2], or [3], next to specific pieces of information. These numbers are crucial because they point directly to the original sources Perplexity used.

If it looks different: The summary might be formatted as a series of bullet points, a few longer paragraphs, or a mix of both. Regardless of the exact appearance, the core information and the numbered references should be clearly visible, usually appearing immediately below your initial question.

You'll know it worked when you see a generated answer to your query, complete with small, numbered citations embedded within the text.

Step 4 — Examine the source list

This is where Perplexity truly shines for transparency and fact-checking. Scroll down slightly below the main summary answer. You'll find a dedicated section, typically labeled "Sources" or "References." Here, you'll see a list of the actual websites or articles that Perplexity consulted to formulate its answer, corresponding to the numbers you saw in the summary. For each source, there will usually be a title, a brief description of the content, and the website's address. This **transparency** (showing you exactly where the information came from) is vital because it lets you see *where* the information originated, not just *what* the AI says.

If it looks different: The sources might be displayed as small, clickable cards, or as a simple list. On a mobile device, you might need to tap an "Expand Sources" or a similar button to reveal the full list. Just look for a distinct section that clearly lists external website links.

You'll know it worked when you see a list of external website links, each labeled with a number that corresponds to the citations in the AI's summary.

Step 5 — Dive into an original source

To truly verify a claim and understand its context, it's essential to go directly to the original source. Click on one of the numbered source links from the "Sources" section that seems most relevant to the claim you're checking or appears to be from a **reputable** (trustworthy and respected) organization. This action will open that original webpage in a new browser tab or sometimes within Perplexity itself (depending on your device settings). Now, you can read the information in its original context. Check if the information presented by Perplexity accurately reflects what the source says, look for publication dates, the author, and whether the website itself appears trustworthy (e.g., is it a known news organization, a research institution, or a personal blog?). Comparing the AI's summary with the actual source helps you confirm its accuracy.

If it looks different: If clicking a link doesn't open a new tab, check your browser's pop-up blocker settings, as they can sometimes interfere. You might also have the option to right-click (on a computer) or long-press (on a phone) the link to copy its address, then paste it directly into a new browser tab yourself.

You'll know it worked when you are actively viewing the actual website or article that Perplexity used, allowing you to read the original information.

Step 6 — Refine your query or ask follow-ups

Sometimes, your first question might not give you the complete picture, or the sources provided might be unclear. If you need more information or a different angle, you can ask a follow-up question. Below Perplexity's initial answer, you'll usually see an "Ask a follow-up" prompt or a new blank input box. Type your refined question there, focusing on clarifying details or asking for more specific types of sources. For example, if the initial sources were mostly opinion pieces, you might ask "Are there any scientific studies on this topic?". This iterative process helps you narrow down your search and get more precise, verifiable information.

If it looks different: Instead of an explicit "Ask a follow-up" button, you might simply see another blank input box directly below the answer, allowing you to continue the conversation. If you want to start a brand-new search entirely, look for a "New Thread" or "New Search" button, usually found near the top of the screen or within a side menu.

 Example: Can you find government health organisation statements about cold showers and colds?

You'll know it worked when Perplexity processes your follow-up query, generating a new answer and set of sources based on your refinement.

Common mistakes

- **Not checking the sources thoroughly:** It's tempting to just read Perplexity's summary and assume it's entirely correct. The AI compiles information, but it can sometimes misinterpret or give too much weight to less reliable sources. Always click through to the original sources to verify the context and content yourself.
- **Being too vague with your question:** Asking broad questions like "Tell me about this news" might not yield precise fact-checks. Instead, provide the specific claim or statement you want to investigate. Being specific helps Perplexity pull more relevant information.
- **Ignoring source reliability:** While Perplexity provides sources, it doesn't always rate their trustworthiness. A quick check of a source's website (e.g., is it a well-known news organization, a research institution, or just a personal blog?) and its publication date is crucial for deciding if the information is truly reliable.

What this means for you

- **In everyday life:** You can quickly become a more informed social media user. Before sharing a sensational post or falling for misinformation, you can take a minute to verify it, protecting yourself and your friends from spreading inaccurate content.
- **For your business or work:** If you work in communications, marketing, or any field where accurate information is key, Perplexity can be a fast way to cross-check industry news, claims from competitors, or statistics before you use them in a report or presentation.
- **If you're just getting started:** The easiest first move is to pick one social media post you're curious about today and run its exact claim through Perplexity. Focus on just reading the summary and then clicking through one source link to see the original content.

Try it now

Find a recent claim or statistic on your social media feed that caught your eye, perhaps something surprising or highly shared. Type that exact claim into Perplexity's search bar and hit Enter. Once you have the summary, scroll down and click on the first source link provided to see where the information originally came from.

How to Summarize a Long Article or PDF into Key Bullet Points

6 steps · [Read online](#)

In short: Paste the text (or upload the PDF) into any AI assistant, then use a short "summarize into bullet points" prompt. You'll get the main argument and the key takeaways in under five minutes — works in ChatGPT, Gemini, Claude, Copilot, Le Chat or DeepSeek.

You've just opened a 20-page report, a long news feature, or a 40-page PDF someone sent you. You don't have time to read all of it, but you do need to know what it actually says. This guide walks you through a simple workflow that turns a long document into a tight set of bullet points you can scan in a minute, using a prompt you can reuse on any AI assistant.

Before you start

- **An AI assistant** you can type into — ChatGPT, Gemini, Claude, Copilot, Le Chat or DeepSeek all work. Pick the one you already have an account with.
- **A free plan is fine** for short articles. Uploading large PDFs (say, over 50 pages) may need a paid plan or a paid tier with bigger file limits — check the assistant's current limits if your file is very long.
- **The article or PDF you want summarized.** Have it open in another tab, or saved on your device.
- **Roughly 5 minutes** from opening the assistant to having your bullet-point summary ready.

One honest note before you start: an AI summary is a starting point, not a replacement for reading. If the document is something important (a contract, a medical paper, a policy that affects you), treat the summary as a map of the territory, not the territory itself. Use it to decide what to read carefully.

Step 1 — Open your AI assistant and find the chat box

Open the website or app of your chosen AI assistant in a browser, or launch the mobile app. Look for the chat box — the long text field at the bottom of the screen, usually with a placeholder line like *"Message [assistant name]..."* or *"Ask anything..."*. Tap or click inside it so your cursor is blinking in the field. If you're on a phone and don't see a text field, look for a round + or **new chat** button on the home screen, then start a fresh conversation. A new, blank chat is ideal — older messages in the thread can confuse the summary, and the assistant only has a limited "memory" of what you paste in (this is called the **context window** — think of it as the assistant's short-term working memory, the amount of text it can actively read in one go).  Example: type nothing yet — just confirm the chat is open and ready. You'll know it worked

when you see an empty text field with a send button (usually a paper-plane icon or an up-arrow) next to it.

Step 2 — Get the text you want summarized ready

Open the article or PDF in another window. You have two options, and which one is easier depends on the assistant and the file type. **Option A (easiest for articles):** select the text in the article, copy it (Ctrl+C on Windows, Cmd+C on Mac, or long-press and "Copy" on a phone), then switch back to the chat. **Option B (easiest for PDFs):** look for a **paperclip icon** , an **attachment icon**, or a + button near the chat box — many assistants let you upload a PDF directly. If you don't see a paperclip, look for a menu labeled "Add files", "Upload" or similar. If neither option works, fall back to opening the PDF in a reader, selecting the text, copying it, and pasting it into the chat as plain text. 💡 Example: highlight the first few paragraphs of a long article, copy them, and have your clipboard ready. You'll know it worked when you either have text highlighted and copied, or a file ready to upload in the chat.

Step 3 — Paste the text or upload the PDF

Click into the chat box and paste the text (Ctrl+V / Cmd+V, or long-press and "Paste" on a phone). The text will appear inside the chat field as a block of words. If you're uploading a PDF, tap the paperclip or + button, choose your file from the list, and wait for the upload progress bar to finish — a small file thumbnail or filename usually appears in the chat field once it's uploaded. If the assistant says the file is "too large", try copying just the first half of the text and pasting it as plain text instead. On the free plan, file size limits are common, and they change often. 💡 Example: paste a 3,000-word article into the chat, or upload a 12-page PDF called "*Q3-strategy-review.pdf*". You'll know it worked when you can see the pasted text block (or the uploaded file name) sitting in the chat field, ready to send.

Step 4 — Add a "summarize into bullet points" prompt

This is the most important step. The quality of your summary depends almost entirely on the instruction you give the assistant. A vague request like "summarize this" usually gives you a waffly paragraph. A specific request gives you clean, scannable bullets. Type the following prompt directly under the pasted text or uploaded file, then send the message (press the paper-plane or up-arrow button). The assistant will start "reading" the content, and within a few seconds you'll see bullet points appear. 💡 Example: *"Please summarize this into bullet points. I want: 1) the main argument or thesis in one sentence, 2) the 5–7 key takeaways as short bullets, 3) any important numbers, dates or names mentioned, 4) one sentence on what the author is asking the reader to do or believe. Keep it under 250 words."* You'll know it worked when the assistant replies with a structured list of bullets (not a long paragraph), and you can identify the main point within the first three bullets.

Step 5 — Check the summary and ask for a tweak

Read the bullet points. Ask yourself: do I now know what this document is actually about? If the summary is too long, too vague, or missing the main point, send a short follow-up. The assistant remembers the original document in the same chat, so you can refine without re-pasting anything. If the summary missed something, name what you wanted. If it's too long, ask for it shorter. If you want a different angle — say, the parts that affect your job — ask for that view.

💡 Example: *"That's good. Now give me just the 3 bullets most relevant to a small business owner."* or *"Cut this down to 5 bullets maximum."* You'll know it worked when the second reply gives you a summary that you can read in under a minute and confidently explain to someone else.

Step 6 — Save or share the result

Once you're happy with the summary, copy it out of the chat (select the assistant's reply, copy, and paste into a note, email or document). If your assistant has a **share** or **copy** button near the reply, use that — it's usually faster. Saving the bullets somewhere means you don't have to come back to the chat later, and you can paste them into a meeting agenda, a study note, or a message to a colleague. Over time, building a small folder of "bullet summaries" is a quiet superpower for keeping up with reading you don't have time to do properly. 💡 Example: paste the bullets into a note called *"Reading — June 2026"*, with a one-line note above them saying where they came from. You'll know it worked when the summary is saved in a place you can find again, and you can skim it in under a minute the next time you need it.

Common mistakes

- **Asking a vague question.** "Summarize this" usually gets you a wobbly paragraph. Always specify the format: bullet points, a number of bullets, and what you want covered (main argument, key takeaways, numbers). The more specific the instruction, the better the output.
- **Pasting a 50-page PDF as plain text.** If the document is very long and you paste it all in, the assistant may quietly drop the middle of the text because of the context window limit. For long files, use the PDF upload option if your assistant offers one, or paste just the introduction and the conclusion first to get a quick orientation.
- **Trusting the summary without a sanity check.** AI assistants sometimes invent a number, a name, or a quote that isn't in the original. If a bullet contains a specific fact that matters to you, scroll back to the original document and confirm it's there. Treat surprising claims as a prompt to verify, not as confirmed truth.

What this means for you

- **In everyday life:** Got a long newsletter, a council report, or a terms-of-service update you don't want to read in full? Run it through this workflow in five minutes and you'll know whether it's worth your time.
- **For your business or work:** Use it on meeting prep notes, industry reports, competitor blog posts, or long client briefs. A bullet summary you can re-read in 60 seconds is far more useful than a document you read once and never open again.
- **If you're just getting started:** Pick one document you've been meaning to read — an article someone sent you, a report you bookmarked — and run it through Steps 1–6 today. The first time feels slow; the second time is genuinely five minutes.

Try it now

Open the assistant you use most, find a long article you've been putting off, and run it through this exact prompt: *"Please summarize this into bullet points. I want: 1) the main argument in one sentence, 2) the 5–7 key takeaways as short bullets, 3) any important numbers or dates, 4) one sentence on what the author wants the reader to do or believe. Keep it under 250 words."* Five minutes, one document, one useful habit started.

Translate a product description into Spanish and French with Cohere

6 steps · [Read online](#)

In short: Use Cohere's free Playground to turn one English product description into Spanish and French in a single prompt. You'll need a Cohere account (free to start), and the result is a strong first draft — you should still read it through before publishing.

By the end of this guide, you'll have two ready-to-use product descriptions — one in Spanish, one in French — sitting in your store or in a backup doc. This walkthrough is for small business owners and marketers who want a fast, good-enough first draft without hiring a translator for every product.

Before you start

- **A free Cohere account.** Sign-up takes about 2 minutes at the official Cohere site.
- **Your English product description.** Copy it from your store or product sheet and have it ready to paste.
- **Roughly 10 minutes.** Most of that is reading and double-checking the AI's answer.
- **A free plan is enough** for this task. Cohere may limit how many messages you can send per minute, but a single product description easily fits inside the free tier.

A quick note on what Cohere **cannot** do for this task: it won't catch every specialist word (a fabric name, a regulatory phrase, a brand-specific term) and it isn't a substitute for a sworn legal translation. Treat its output as a polished first draft, not a final one.

Step 1 — Create your free Cohere account

Open your browser and go to Cohere's homepage. Click the button in the top-right corner — it's usually labeled "Sign up", "Get started", or "Try for free". You'll land on a sign-up page. You can register with your email and a password, or click the Google or GitHub button to skip the form. After signing up, Cohere may send a confirmation email; open it and click the verification link inside before you continue.

💡 Example: use your real work email so future receipts, API keys, and account-recovery messages all reach you.

If the button wording looks different ("Get API key" or "Start building"), click that instead — it leads to the same dashboard. Cohere updates its labels often.

You'll know it worked when you see a screen that greets you by name, or a dashboard page that lists "Playground" as one of the options.

Step 2 — Open the Playground

Once you're signed in, look at the left-hand navigation menu and click the item labeled "Playground" (it may also be called "Try it now" or sit under a "Tools" header). A new screen opens with two big text boxes side by side: a prompt box on the left where you type your request, and a response box on the right where the AI's answer will appear. At the top of the panel there's a dropdown menu for choosing which AI model to use — leave it on the default for now.

💡 Example: Cohere's "Command" family of models handles everyday translation tasks well; no need to change it for a product description.

If you don't see a left-hand menu, look in the top-left corner for a "hamburger" button — three horizontal lines — and click it. On a smaller screen or a narrow window, the menu hides behind that icon.

You'll know it worked when you see two empty boxes labeled something like "Prompt" (left) and "Response" or "Output" (right), with a "Generate" button near the bottom of the prompt box.

Step 3 — Paste your English copy and write a clear instruction

Click inside the prompt box on the left. Press Ctrl+V (Windows) or Cmd+V (Mac) to paste your English product description directly into the box — don't retype it, that introduces typos. Now add a short instruction above or below the pasted text, telling the AI exactly what you want: which languages, what tone, and how to label the answer.

💡 Example prompt: "Translate the following product description into natural, shopper-friendly Spanish and natural, shopper-friendly French. Label each version clearly. Match the warm, persuasive tone of the original. Here is the description: [your English copy here]"

If your description has unusual characters (curly quotes, em dashes), don't worry — they paste fine. If the box only shows part of your text, that's normal; the AI can still read the whole input.

You'll know it worked when the pasted English text shows up inside the prompt box, with the same line breaks as in your original.

Step 4 — Generate both translations

Click the button at the bottom of the prompt box — look for "Generate", or a play-arrow icon (▶). The right-hand box will start filling with the AI's answer. Most product descriptions finish in under 30 seconds; longer ones may take up to a minute. When it's done, you'll see two clearly separated blocks: one Spanish version, one French version, each labeled.

💡 Example of what a good output looks like: "Spanish: ... [paragraph] ... French: ... [paragraph] ..."

If nothing appears after a minute, click the button once more — the first click sometimes doesn't register. If you only get one language, the prompt may have been cut off; see Step 5.

You'll know it worked when you can read two separate paragraphs, one Spanish and one French, each with its own label.

Step 5 — Ask for a redo if anything feels off

If a translation sounds too stiff, too long, or just doesn't read like something your customers would say, you don't have to start over. In the same playground panel, click at the end of the previous answer (or in a new prompt box that appears below it) and type what you'd like changed. The model treats the new message as a follow-up, so it knows what you're reacting to.

💡 Example follow-up: "Make the Spanish version about 20% shorter and a bit warmer — like a friend recommending the product. Keep the French version unchanged."

If the playground doesn't keep a chat history, paste the previous answer back into the prompt box and ask for the rewrite from scratch — slower, but it still works.

You'll know it worked when the new reply shows the changed version, and any language you asked to keep untouched is still there in the right language.

Step 6 — Copy each version into your store

Highlight the Spanish text in the response box, copy it (Ctrl+C or Cmd+C), and paste it into your store's Spanish-language product page or into a translations spreadsheet. Then go back and do the same with the French text. Save the originals somewhere you'll find again — a Google Doc, a Notion page, or a spreadsheet with one row per product and one column per language.

💡 Example spreadsheet columns: English | Español | Français — one row per product, so any future edit stays in sync across all three languages.

If your store only has a single language field, save the translations in a notes app or a Google Doc titled with the product name, so you can find them later.

You'll know it worked when both Spanish and French versions sit in your store (or in your backup doc), readable, with no garbled characters, no English words left over, and a clear label on each.

Common mistakes

- **Pasting the description alone.** If you only paste the description with no instruction, you may get stiff, word-for-word translations. Always tell the AI the tone (friendly, professional, playful) and who's reading (shoppers, parents, professionals).
- **Trusting the output blindly.** Cohere is strong at everyday language but can mistranslate specialist words (a specific fabric, a regulatory term, a brand name). Read both versions out loud — anything you wouldn't say yourself is worth fixing.
- **Forgetting to save your English source.** If you close the playground and come back later, the original prompt may be gone. Keep your best English wording in a doc alongside the translations, so future edits have a single source of truth.
- **Mixing up the labels.** The AI sometimes merges Spanish and French into one block, or swaps the order. Always scan for the "Spanish:" and "French:" headers before you copy — that's how you know which version is which.

What this means for you

- **In everyday life:** If you ever need to email a French-speaking friend or read a Spanish recipe at a friend's house, the same Playground prompt works — paste the text, ask for a translation, done.
- **For your business or work:** A team of one can keep a small product catalog in three languages without paying a translator per product. You stay in control of the wording, and you ship faster.
- **If you're just getting started:** Pick your best-selling product today. Translate just that one. See how the result reads, then decide whether to roll the workflow out across the rest of your catalog.

Try it now

Open Cohere's Playground, paste this exact prompt, and replace the bracketed text with your own description: "Translate the following product description into natural Spanish and natural French. Label each version clearly. Here is the description: [paste your English product copy here]." Hit Generate, read the result out loud, and save the two versions somewhere you'll find them again.

How to compare two rental listings side by side with weighted pros and cons

6 steps · [Read online](#)

In short: Open any AI assistant, paste both listings, list your priorities with weights, and ask for a scored comparison table. The assistant won't visit the apartment for you, but it will turn a wall of listing text into a clear trade-off table in minutes.

When you're choosing between two rentals, the listings blur together. This five-minute workflow works in any AI assistant — ChatGPT, Gemini, Claude, Copilot, Le Chat, DeepSeek, or whichever you already have — and scores both places against the priorities you set yourself.

Before you start

- An account with any AI assistant you already use (free or paid — both work for this).
- The full text of both rental listings, copied from the listing page or app. Screenshots alone won't work — only pasted text.
- A short list of what actually matters to you (commute, noise, utilities included, lease length, pets, parking — three to six items is enough).
- About 15 minutes.

Step 1 — Open a new chat and paste both listings

Start a fresh conversation — don't drop this into an old chat, or the assistant will weigh your rentals against whatever you were talking about before. In the message box at the bottom of the screen, paste the first listing, leave a blank line, then paste the second one. End the message with a one-line note telling the assistant what it's looking at. 💡 Example: "Listing A — 2-bed in Eastside, \$X/month, 12-month lease. — Listing B — 1-bed + den in Northgate, \$Y/month, month-to-month. Compare these two rentals." If the listing is mostly images or a PDF, you'll need to type out the key facts by hand — assistants can't reliably read screenshots unless your plan includes file upload or vision, and even then the result varies. You'll know it worked when you can scroll up and see both listings sitting in the chat, and the assistant acknowledges them in its first reply.

Step 2 — Tell the assistant your priorities and how much each matters

Send a second message that lists the things you care about and gives each one a weight — how much it should count in the final score. Weights are how you stop the assistant from treating "has a dishwasher" as equally important as "20-minute commute." The format that works best is a numbered list with a 1–5 or 1–10 weight next to each priority. 💡 Example: "Here are my

priorities, weighted 1–5: 1) Commute under 30 min by transit (5). 2) Quiet street, no main road (4). 3) Utilities included in rent (3). 4) Allows my cat (5). 5) In-unit laundry (2). Score both listings against these." If your assistant has a "memory" or "custom instructions" setting (different apps call it different things), you can save the weights there once and reuse them for future comparisons. You'll know it worked when the assistant restates your priorities back to you without being asked — that means it actually loaded them into the conversation.

Step 3 — Ask for a clean side-by-side table first, before any scores

Now ask for a comparison table with one row per feature and one column per listing. Don't ask for scores yet — you want a neutral snapshot of what each listing actually says, so you can spot places where the assistant fills in gaps with guesses. The clearer this table is, the more honest the scoring will be in the next step. 💡 Example: "Make a table with these rows: rent, lease length, square feet, utilities included, parking, laundry, pet policy, commute time to downtown, nearest noise source. One column for Listing A, one for Listing B, one for 'verified or missing.' If a fact isn't in the listing, write 'not stated.'" If the table renders as plain text instead of a real grid (some mobile apps do this), ask the assistant to "rebuild it with pipes and dashes" and it will produce a text-only version that still reads clearly. You'll know it worked when every row is filled with text from the listing — or says "not stated" — instead of the assistant inventing something to fill the gap.

Step 4 — Add weighted scores and a recommendation

Send a third message asking the assistant to score each listing on your weighted priorities and add a one-line recommendation at the bottom. This is the step that turns the table into a decision. The math is simple: multiply each priority's weight by the listing's score on that priority, then add them up. The listing with the higher total wins — *on your priorities*, not on someone else's. 💡 Example: "Now score both listings 1–10 on each priority from Step 2, multiply by the weight, and show the totals. Add a two-sentence recommendation naming which one wins on my priorities and why." If the assistant hedges with "it really depends on you," that's a sign it didn't read your weights carefully — nudge it with "use the weights literally, don't substitute your own." You'll know it worked when you see a final number for each listing and the recommendation names a clear winner with a stated reason.

Step 5 — Push back on vague answers and list what to verify

Before you act on the recommendation, ask the assistant to flag anything it had to guess at, plus every "not stated" cell from the table. This step is the safety net — assistants can sound confident about facts that aren't actually in the listing, and a lease you signed based on a wrong assumption is an expensive lesson. The honest answer often involves several things you need to

ask the landlord directly. 🗨️ Example: "List every fact in the table you marked 'not stated' and anything you had to infer. For each one, write a one-line question I can email the landlord." If the assistant invents a detail (like "utilities are included" when the listing didn't say so), that's a hallucination — when an AI confidently states something that isn't actually there. Call it out with: "You made that up; the listing didn't say it. Replace with 'unverified, ask landlord.'" You'll know it worked when you have a short, concrete list of questions to send before you sign anything.

Step 6 — Save the prompt structure as a reusable template

Copy the three prompts you sent (Steps 1, 2, and 4) into a notes app and label it "Rental comparison template." Next time you're weighing two places, open a new chat, paste the template, swap in the new listings and weights, and you'll get the same scored comparison in minutes — no setup, no re-thinking the structure. 🗨️ Example: "PROMPT 1: [paste listings + intro line]. PROMPT 2: [paste weighted priorities]. PROMPT 3: [paste score-and-recommend request]." If your assistant supports a way to save reusable prompts (different apps call this Projects, Gems, custom instructions, or similar — check the menus if you're not sure), you can store the template there so it loads automatically next time. You'll know it worked when the template is sitting in your notes ready to go the next time a listing site or a landlord's email hands you two options.

Common mistakes

- **Letting the assistant pick your priorities.** It will default to generic things like "location" and "price." Insist on *your* weights, even if they feel quirky — a quiet street might matter more to you than square footage, and that's the whole point of doing this yourself.
- **Pasting only the headline summary.** The first paragraph of a listing is marketing. The fine print about utilities, parking, and pet fees is usually further down — paste the whole thing, or you'll be comparing ad copy instead of facts.
- **Trusting the score without checking the table.** A 9/10 on "utilities included" means nothing if the listing never actually said that. Always read the table before you read the score, and send the "not stated" list to the landlord before you commit.

What this means for you

- **In everyday life:** When you're weighing two rentals (or two cars, or two phone plans), this same three-prompt pattern turns a messy comparison into a clear trade-off table in minutes — without you having to argue with yourself about which factor matters more.
- **For your work or business:** The same structure works for choosing between vendors, job offers, software plans, or any two options where you want *your* priorities to drive the

score — not the assistant's defaults.

- **If you want to go further:** Try extending the prompt to three listings instead of two, or add a "dealbreaker" rule that automatically disqualifies any option missing a non-negotiable (like "no cat-friendly place counts, period").

Try it now

Pick one priority that matters to you (commute, pets, quiet, anything), find two short rental listings — even made-up ones work for practice — and run Steps 1 to 4 in a fresh chat. You'll see the decision become obvious in under five minutes.

How to use Gemini Deep Research to plan a big purchase decision

6 steps · [Read online](#)

In short: Gemini Deep Research browses dozens of websites for you and returns a structured report with source links, so comparing big purchases like cars, appliances, or laptops stops being a full weekend of manual Googling.

Buying a new laptop, a refrigerator, or a used car usually means opening forty browser tabs and reading half of them. Gemini's **Deep Research** feature does that scrolling for you: it reads many sites, summarizes what it found, and hands you a tidy report with links back to where each claim came from. This guide walks you through using it, step by step, so the report actually answers *your* question.

Before you start

- **A Google account** signed in on the Gemini website or the Gemini app.
- **Access to Deep Research.** It may be limited to paid Google AI plans on some accounts, so check the plan page for the latest availability before relying on it.
- **A rough idea of what you're shopping for** (a category, a budget range, and 2–3 must-haves work best — more on this in Step 2).
- **About 15–25 minutes.** Deep Research usually takes a few minutes to run, and you'll want time to read the report it returns.
- An honest expectation: the report gives you a strong starting point, but it does **not** buy the thing for you, and prices or availability can be outdated by the time you read them.

Step 1 — Open Gemini and pick the Deep Research mode

Go to the Gemini website or open the Gemini app, then look in the chat box for a model or tool picker — usually a small dropdown above the text field. From that picker, choose the option labeled "**Deep Research.**" On some accounts the option appears as a separate **Tools** button or a dedicated tab at the top of the chat. If you don't see it, it might not be enabled on your plan yet, or it could live behind a different name like "**Research**" or "**Pro.**"

 Example: when the picker appears, click the dropdown that currently shows "Fast" or "Pro" and switch it to **Deep Research**.

You'll know it worked when the text field gains a small brief outline or "Plan" panel, and the placeholder text inside the box mentions research or sources.

Step 2 — Describe your purchase in concrete detail

This is the most important step: the more specific your prompt, the more useful the report. Type a short paragraph that names the **item, your budget range, your must-haves, your deal-breakers, and how you'll use it**. Think of it like a one-page brief you'd give a friend who's shopping on your behalf — vague briefs produce vague reports. Avoid jargon the AI won't understand; instead, describe the thing in plain words.

💡 Example: "I need a quiet, energy-efficient dishwasher for a small kitchen, budget roughly \$700–\$1,000 USD, must fit a 24-inch opening and run under 45 dB. I cook most nights for two, sometimes four. Compare the top 5 models available in my country, with pros, cons, and where each one tends to be cheapest."

You'll know it worked when you see a clear "Research plan" or outline appear below your message — Gemini usually shows what it intends to search for before it starts.

Step 3 — Approve the plan and let it run

Below the research plan, look for a button like "**Start Research**" or "**Begin**." Click it, then step away. Deep Research browses many websites in the background and writes as it goes, so the response can take **several minutes** (sometimes five, sometimes fifteen) — far longer than a normal chat reply. A progress bar or running indicator usually appears. **Do not close the tab** while it's running, or you'll lose the work.

💡 Example: hit the "Start Research" button, then make a coffee; the first draft of the report appears inline as it finishes each section.

You'll know it worked when a long, structured report lands in the chat, with section headings, bullet points, and small numbered or linked citations you can click.

Step 4 — Read the report and click at least three sources

Don't trust any AI research report blindly. Scroll to the end of the report first and look for a **Sources** section, or for inline footnote-style numbers next to specific claims. Click at least three of those links and read the original page for a minute each, especially anything related to price, warranty, or safety. The sources tell you whether the report is repeating marketing copy or actually weighing real reviews.

💡 Example: open the report, jump to the "Sources" list at the bottom, and click the first three links — if any of them are dead, irrelevant, or from a brand's own press release, treat the matching claim with caution.

You'll know it worked when you can point to a specific source link and say, out loud, "this is where that fact came from."

Step 5 — Ask two or three follow-up questions to narrow the field

Once the first report lands, treat it like a starting point, not the final answer. Ask Gemini to compare your top two candidates head-to-head, to look for known reliability complaints, or to translate the report into a simple "if you care most about X, choose Y; if you care about Z, choose W" decision guide. One question per message works best. If an answer feels too generic, push back: ask for specific model numbers, specific dates, or specific failure modes.

💡 Example: "Now compare models A and B from your report on noise, energy use, and reliability over five years — which one wins for a household of two who cooks daily?"

You'll know it worked when Gemini's reply is sharper and more specific than the original report, and ideally points back to the same source list.

Step 6 — Turn the report into a personal decision checklist

Before you close the chat, ask Gemini to summarize the entire research session into a short **decision checklist** you can actually use: 3–5 finalist products, the deciding criteria for each, and the questions you still need to answer before you spend money. Copy that checklist into a note on your phone or print it. The point of the report is to leave your head clear, not crammed with new jargon.

💡 Example: "Summarize this whole conversation into a checklist: top 3 dishwashers I should see in a store, what to test in person, and what to double-check on the manufacturer's site before I buy."

You'll know it worked when you have a one-page summary in your hand that you could read in two minutes in a showroom.

Common mistakes

- **Mistake:** Treating the report as a buy recommendation. **Fix:** It's a starting brief — the final choice is still yours, and any purchase should be confirmed against the official product page and the live price.
- **Mistake:** Skipping the sources. **Fix:** Always click at least three source links. If you can't find them, or they look unreliable, ask Gemini to redo the research with more trusted sites.
- **Mistake:** Closing the tab while research is running. **Fix:** Deep Research takes several minutes; leave the tab open until the full report appears, or you'll lose the work.
- **Mistake:** Asking a vague prompt like "best laptop under \$1,000." **Fix:** Add your must-haves (size, weight, battery life, ports) and your main use (work, gaming, school) — the more detail, the better the report.

What this means for you

- **In everyday life:** A big purchase — a car, a fridge, a laptop — stops being a weekend of forty open tabs. You get a sourced brief instead, then spend your time on the decisions that actually need a human.
- **For your business or work:** The same workflow works for buying equipment, software subscriptions, or vendor comparisons at work: ask Deep Research to compare three vendors against your checklist, then read the sources yourself before signing anything.
- **If you're just getting started:** Pick one purchase you've been putting off, write a two-sentence brief in Step 2, and run the report end to end. The first run is the slowest; the second one feels easy.

Try it now

Open Gemini, switch to **Deep Research**, and paste this single sentence as your first prompt:

"I'm buying a [item] with a budget of [low]–[high], and my must-haves are [X, Y, Z].

Compare the top 5 options and end with a short decision checklist." Fill in the brackets, hit Start Research, and read the sources before you trust any of it.

Catch Up Fast: Summarize Long Chats in WhatsApp and Messenger with Meta AI

6 steps · [Read online](#)

By the end of this guide you'll know how to ask Meta AI (a built-in helper inside WhatsApp and Messenger) to read through a long chat and give you a short, useful summary in plain language. This is for anyone who has ever scrolled back through hundreds of messages thinking "*what did I miss?*" — whether it's a family group chat or a busy work thread.

Before you start

- A phone with **WhatsApp** or **Messenger** installed, updated to the latest version (Meta AI is rolled out in stages, so update to be sure).
- A personal Meta account already logged in on that app.
- Meta AI available in your country or region — if you cannot see it after updating, it may not have reached your area yet.
- Around **5 to 10 minutes** of your time, plus the chat you want summarized open in another window or conversation.

One honest note up front: Meta AI **cannot read old messages you have not shared with it**. You will need to either paste the messages into a chat with Meta AI or @-mention it in a recent thread. It is free to use, but anything you send is processed by Meta — so avoid pasting very sensitive personal details.

Step 1 — Open WhatsApp or Messenger and look for Meta AI

Open the app you usually use (WhatsApp or Messenger) and look at the top of your chat list or your contacts. You should see Meta AI listed as an option, sitting alongside your usual contacts. Tap it once. A new chat window opens, just like starting a conversation with a friend. You'll know it worked when you see a message from Meta AI indicating it's ready to help. **If it looks different:** some users see a different icon or a "New chat" option with an AI symbol, or a search bar at the top — any of these can lead to Meta AI. If you see nothing related to AI, your app may need an update, or Meta AI may not yet be available in your region.

Step 2 — Open the chat you want summarized, side by side

Leave the Meta AI chat open, then open the long chat you want to catch up on. On most phones you can switch between apps with the square "recent apps" button, or use your phone's split-screen feature if available. You will see the full message history of your long chat on one screen and the Meta AI chat on the other. This makes copying easy. **If it looks different:** on tablets or

split-screen devices you can arrange both windows side by side so both sit next to each other at the same time.

Step 3 — Copy the messages you want summarized

In the long chat, press and hold the first message you want Meta AI to read, then choose **Copy**. For a range of messages, keep tapping each one until you have selected the section you missed (often the last 20 to 50 messages is plenty). Then tap **Copy** again. The text now sits on your phone's clipboard, ready to paste — like copying a paragraph from a website. **If it looks different:** some Android phones show a small menu with "Copy", "Forward" and "Share"; on iPhone you'll see "Copy" appear above the selected message. If selecting one-by-one is too slow, try **Forward** instead, then paste into the Meta AI chat in one go.

💡 Example: *"I copied messages from a 40-person family group covering about a day's worth of chat — roughly 30 messages with photos, links and a few voice notes."*

You'll know it worked when your phone briefly shows a small "Copied" label or a tick next to the selected messages.

Step 4 — Paste the messages into Meta AI and ask for a summary

Switch back to the Meta AI chat. Tap inside the message box, hold for a moment, and choose **Paste**. The copied messages appear as one long block. On a new line underneath, type a clear, kind request that tells Meta AI exactly what you want — think of this as giving instructions to a helpful assistant, not just dumping text on it. Then tap the send arrow. After a few seconds, Meta AI replies with a summary. **If it looks different:** if the paste does not appear, try tapping once inside the text box first, then holding your finger down again. If Meta AI replies in a language you did not expect, add "Please reply in English" at the end of your message.

💡 Example: *"Here are the messages from my family group from yesterday. Please summarize them in 5 short bullet points, focus on any decisions or plans people made, and ignore the photos. Reply in plain English."*

You'll know it worked when Meta AI replies with a summary rather than echoing the messages back at you.

Step 5 — Refine the summary if it is too long or off-target

Read Meta AI's reply. If the summary is too vague, too long, or misses what you actually care about, reply with a short follow-up message asking for a change. The assistant remembers the conversation, so you can be casual about it. Each new request usually takes a few seconds. **If it**

looks different: if Meta AI seems stuck or repeats the same answer, try starting a brand-new chat and pasting the messages again with a clearer instruction — sometimes a fresh start produces better results.

💡 Example: *"Great, but please make it 3 bullet points instead of 5, and add a one-line 'action items' section at the end showing what I might need to reply to."*

You'll know it worked when the new summary matches what you actually wanted to know — short, focused and useful.

Step 6 — Save or share the summary

Once you are happy with the summary, you can save it for later or send it to someone else. Tap and hold the assistant's message, then choose **Copy** to put it on your clipboard, or look for the forward option to send it to another chat or contact. Many phones also let you tap **Share** and save it straight to Notes, email or a document. **If it looks different:** if you don't see forward options, look for a menu icon or arrow next to the message. From there you can usually forward, share or save the message.

💡 Example: *"I forward the summary to a private chat with myself called 'Catch-up notes' so I can search it later without scrolling through the original group."*

You'll know it worked when the summary now lives somewhere easy to find — your notes app, a saved message, or a chat you control.

Common mistakes

- **Pasting months of messages at once.** Meta AI has a limit to how much text it can read in one go, and very long pastes often get cut off or produce a generic answer. **Fix:** paste the last 20 to 50 messages only, or split a longer chat into two requests.
- **Sharing sensitive personal information.** Anything you paste into Meta AI is processed by Meta's servers, so names, addresses, phone numbers or financial details should be removed first. **Fix:** change real names to "Alex", "Sam" and "the host" before pasting, especially for group chats with people you do not know well.
- **Expecting Meta AI to know your whole chat history without pasting.** Unlike asking a friend, the assistant only sees what you have shared with it in that conversation. **Fix:** always paste the messages you want summarized, or @-mention Meta AI inside a recent thread so it can see the surrounding context.

What this means for you

- **In everyday life:** catch up on noisy family or friend group chats before replying, instead of scrolling through 200 messages on the train.
- **For your business or work:** turn long WhatsApp or Messenger threads with clients, teammates or event groups into a short briefing you can skim in 30 seconds and forward to a colleague.
- **If you're just getting started:** pick one group chat you fall behind on, copy the last 20 messages, paste them into Meta AI and ask for 3 bullets — that single step is the whole skill in miniature.

Try it now

Open your busiest group chat right now, copy the last 20 messages, paste them into a Meta AI chat, and ask: *"Summarize these in 3 bullet points and list anything I need to reply to."* It takes under two minutes and you'll have the skill for life.

Summarize a long PDF into 5 bullet points with Claude

6 steps · [Read online](#)

Ever faced a huge PDF document and wished you could get the gist of it without reading every single page? This guide will show you exactly how to use Claude, an AI assistant, to summarize any PDF file into just five easy-to-read bullet points. By the end, you'll be able to quickly understand important documents for work, study, or personal life.

Before you start

- You'll need a free [Claude account](#). If you don't have one, it only takes a few minutes to sign up with an email address or Google account.
- Access to a computer or mobile device with internet.
- A PDF file you want to summarize.
- This guide assumes you are using the free version of Claude, which allows ample usage for this task.
- **Rough total time:** 6 minutes.

Step 1 — Get Ready for Claude

First, you need to open Claude in your web browser. Claude is an **AI assistant** – think of it like a very smart digital helper you can chat with, designed to understand your requests and provide helpful responses.

Open your web browser (like Chrome, Firefox, or Safari) and go to **claude.ai**. You'll see either a login screen or a chat interface if you're already logged in. If prompted, either sign in with your email and password or use a linked account like Google.

Once you're successfully logged in, you'll see Claude's main chat interface, which usually has a clear area to type messages and a history of past conversations on the left. You'll know it worked when you see the main Claude chat screen, ready for a new conversation.

Step 2 — Start a Fresh Conversation

To get the best results, it's always a good idea to start a new, clean conversation for each task. This ensures Claude focuses solely on your current request without being influenced by previous discussions.

Look for a button that says **"Start a new chat"** or **"New Chat"**. This is usually found in the top-left corner of the screen, often next to a Claude logo. Click on this button.

After clicking, the main chat area will clear, and you'll see a fresh, empty conversation window. A text box at the bottom will say something like "Message Claude..." or "Type your message here...". If it looks different, look for an icon like a plus symbol (+) or a speech bubble to begin a new discussion.

You'll know it worked when the chat panel is empty, ready for your first message.

Step 3 — Attach Your PDF Document

Now it's time to give Claude the document you want summarized. A **PDF** (Portable Document Format) is a common file type for documents, reports, and books, designed to look the same on any device.

In the message box at the bottom of the Claude chat window, look for an **attachment icon**. This usually looks like a **paperclip** or a **plus symbol (+)**. Click on this icon. A file explorer window will pop up from your computer. Navigate to where your PDF document is saved, select the PDF file, and then click "**Open**" or "**Choose**".

You'll see the name of your PDF file appear above the message box, sometimes with a small preview or an "X" icon next to it, indicating it's ready to be sent. If you can't find the paperclip, try dragging and dropping your PDF file directly into the message box. You'll know it worked when your PDF's filename appears within the Claude chat interface, ready for your instruction.

Step 4 — Tell Claude What You Need

This is where you give Claude your **prompt** – which is simply the instruction you want it to follow. The clearer your prompt, the better Claude's response will be. You'll ask it to summarize the document into bullet points.

In the message box, where it says "Message Claude...", type your request. Make sure to specify the format (bullet points) and the number of points you want.

 **Example:** Summarize this document into 5 key bullet points. Please use simple, plain English.

Once you've typed your prompt, double-check it to make sure it's clear and asks for exactly what you want. You'll know it worked when your text appears in the message box, alongside the attached PDF file, ready to be sent.

Step 5 — Send Your Request to Claude

With your PDF attached and your prompt typed, the next step is to send it to Claude so it can start working its magic.

Look for the **send button**, which is typically on the right side of the message box. It often looks like a **paper airplane** icon or a **right-pointing arrow**. Click on this button to send both your PDF and your instruction to Claude.

Claude will then begin processing your request. You'll usually see a message like "Claude is thinking..." or "Generating response..." while it works. This might take a few moments, especially for longer documents. If the send button seems unresponsive, ensure you have both the PDF attached and text in the message box. You'll know it worked when Claude starts typing out a response in the chat window, directly below your sent message and the attached PDF.

Step 6 — Review and Refine Claude's Summary

Claude will now present you with the summary based on your prompt. It's important to review this response to ensure it captures the essence of your document accurately.

Read through the 5 bullet points Claude has provided. Check if they cover the main ideas and if the language is indeed simple and plain as you requested. While Claude is very capable, it can sometimes **hallucinate** (confidently make up information that isn't in the document) or miss a subtle nuance, especially with very complex or highly technical texts.

If the summary isn't quite right, or if you'd like more detail on a specific point, you can ask Claude to refine it. For example, you could type: `Can you expand on the third bullet point?` or `Please rephrase the first point to be even simpler.` You'll know it worked when Claude presents a clear list of 5 bullet points in response, ready for your review or further instructions.

Common mistakes

- **Forgetting to attach the PDF or sending a blank message:** Claude can only summarize what you provide. Always make sure the document name appears above the message box before you hit send. If you sent a blank message, simply attach the PDF and send your prompt in a new message.
- **Using a vague prompt:** If you just say "Summarize this," Claude might give you a very general overview or a different format than you wanted. Be specific about the number of points and the desired language, like "Summarize into 5 bullet points, in simple English."
- **Expecting perfect accuracy without review:** While Claude is powerful, it's an AI and not infallible. It might occasionally misinterpret information or omit crucial details. Always read the summary carefully and compare it to your understanding of the document, especially for critical information.

What this means for you

- **In everyday life:** Instead of spending hours reading lengthy terms and conditions, user manuals, or community notices, you can quickly get the essential information in a few bullet points, saving you time and effort.
- **For your business or work:** Instantly summarize research papers, market reports, client proposals, or meeting minutes. This helps you and your team quickly grasp key takeaways, make faster decisions, and stay informed without getting bogged down in dense text.
- **If you're just getting started:** The easiest first move is to try it with a PDF you've already read. This way, you can easily compare Claude's summary to your own understanding and see how well it works.

Try it now

Find a PDF document on your computer – perhaps an old instruction manual, a news article you saved, or a school assignment. Head over to **claude.ai**, start a new chat, attach your document, and use the example prompt to get your first 5-point summary! You'll be amazed at how quickly you can distil information.

Summarize Long Articles into Key Bullet Points with Grok on X

6 steps · [Read online](#)

In today's fast-paced world, it's easy to feel overwhelmed by the sheer volume of information available online. This guide will show you how to use Grok, an AI assistant integrated into the X platform, to quickly condense long news articles or reports into concise bullet points. By the end, you'll be able to get the essential information from any article in moments, letting you stay informed without needing to read every word. While Grok can summarize the content of an article, it cannot verify the accuracy of the original article's claims or perform actions like sending emails for you.

Before you start

- An active X Premium+ subscription, as Grok is only available to subscribers of X's highest tier.
- The X app on your mobile device or access to the X website on a computer.
- Roughly 5-10 minutes of your time.

Step 1 — Access Grok on X

First, open the X app on your phone or visit the X website on your computer. Once the X feed has loaded, look for a special icon that represents Grok. On mobile, this is usually a diamond-shaped icon found in the bottom navigation bar. On a computer, you might find a "Grok" option listed in the left-hand sidebar menu, often near your profile or notifications. Tap or click on this icon to enter the Grok interface. If it looks different: The X interface can change, so if you don't see a diamond icon or a clear "Grok" label, look for an icon that stands out, or check the 'More' menu (often three lines or dots) for an entry labeled "Grok". You'll know it worked when you see Grok's chat interface, which usually includes a welcome message or a list of previous conversations.

Step 2 — Copy the Article Link

Now, you need to find the online article you want to summarize. Open your web browser (like Chrome, Safari, or Edge) and navigate to the news article, blog post, or report you're interested in. Once you're on the page, look for the web address (URL) at the top of the browser window. Tap or click on this address to highlight it, then use your device's copy function (e.g., right-click and 'Copy' on a computer, or long-press and 'Copy' on a phone) to save the link to your clipboard. If it looks different: Some apps have a 'Share' button (often an arrow pointing upwards or a three-dot menu) that includes a 'Copy Link' option. If you can't easily highlight the

URL, look for that share option. You'll know it worked when the link is saved and ready to be pasted elsewhere.

Step 3 — Start a New Chat with Grok

With your article link copied, return to the Grok interface within the X app or website. To begin a fresh summarisation task, you'll need to start a new conversation. Look for a button that typically indicates starting something new, like a '+' symbol, a pencil icon, or a "New Chat" button, usually located at the top-right or bottom-right corner of the Grok chat screen. Tapping this will clear any previous messages and give you a clean slate. If it looks different: If you don't see an obvious 'new chat' icon, you can often just type your request directly into the existing message box. Grok is usually smart enough to treat a new request as a new topic. You'll know it worked when the chat area is empty and a message input box is clearly visible at the bottom of the screen, ready for you to type.

Step 4 — Paste Your Link and Request a Summary

Now it's time to tell Grok what you want. Click or tap inside the message input box at the bottom of the Grok chat screen. Paste the article link you copied in Step 2 into the box. After the link, type a clear instruction asking Grok to summarize the content. Be specific about the format you prefer, like asking for bullet points, to help Grok deliver exactly what you need. If it looks different: If pasting doesn't work, try typing the full URL manually if it's short, or check your device's clipboard settings to ensure the link was copied correctly. Grok might indicate if it's having trouble accessing the link. 🗨️ Example: Please summarise this news article into 5 clear bullet points, focusing on the main findings: [paste your article link here] You'll know it worked when Grok displays a message like "Generating response..." or "Thinking...", followed shortly by the summary appearing in the chat.

Step 5 — Review Grok's Summary

Once Grok has finished processing your request, its summary will appear in the chat window, presented in the bullet-point format you asked for. Read through these points carefully. This is your chance to quickly understand the core message, key arguments, and important facts from the article without having to read every sentence of the original text. Take a moment to ensure the summary captures what you expected and is easy to understand. If it looks different: Sometimes, Grok might ask a clarifying question if the article is too long, behind a paywall (which Grok cannot bypass), or contains ambiguous content. If this happens, try rephrasing your request or finding a more accessible article. You'll know it worked when you have a concise list of key information that helps you quickly grasp the article's main content.

Step 6 — Refine the Summary (Optional)

Grok's initial summary might be great, but you can always ask for more details or a different perspective if needed. If a particular bullet point sparked your interest, or if you feel something important was missed, simply type a follow-up question into the message box. You can ask Grok to expand on a specific detail, explain a complex term, or even rephrase the summary for a different audience (e.g., "Summarize this for a high school student"). This allows you to dig deeper without returning to the original lengthy article. If it looks different: If Grok says it can't find the information, it might be that the detail wasn't explicitly in the article, or it's too specific for a general summary. Try rephrasing your follow-up in simpler terms. 🗨️ Example: Can you elaborate on the economic implications mentioned in bullet point 2? You'll know it worked when Grok provides a more detailed or adjusted response, showing it's continuing to process the article's information based on your new prompt.

Common mistakes

- **No Grok access:** Grok is a premium feature, so if you can't find it or get an error, ensure your X subscription is the highest tier (X Premium+). The fix is to upgrade your X subscription.
- **Pasting a broken or inaccessible link:** If Grok says it can't access the article, check that the link is correct and that the article isn't behind a paywall or login that Grok can't bypass. The fix is to find a publicly accessible link or try a different article.
- **Vague summary request:** If Grok gives a less-than-useful summary, your request might have been too general. The fix is to be more specific in your prompt, like "Summarize into 5 bullet points" or "Focus on the environmental impact."

What this means for you

- **In everyday life:** Quickly catch up on news headlines or understand long reports from community groups without spending hours reading. It's perfect for when you're short on time but want to stay informed.
- **For your business or work:** Efficiently digest industry news, competitor analysis, or lengthy internal reports, allowing you to extract critical insights faster and make more informed decisions.
- **If you're just getting started:** Practice by picking a news article you'd normally skim, then use Grok to get a bullet-point summary. Compare it to what you would have highlighted yourself.

Try it now

Open the X app, navigate to Grok, and paste a link to any recent news article. Ask Grok to summarize it for you in three bullet points. See how quickly you can get the main gist!

Translate and Explain Any Restaurant Menu with Qwen on Your Phone

6 steps · [Read online](#)

In short: Snap a photo of a foreign-language menu and Qwen will translate each dish into English, list the likely ingredients, and flag common allergens. The free version handles photos and text, but it cannot guarantee allergen safety — always confirm with staff if you have a serious allergy.

Here's how to turn a confusing wall of unfamiliar words into a clear, ordered list of dishes you can actually understand and order from, even when you don't speak the language.

Before you start

- A phone with a working camera (iPhone or Android).
- The free Qwen app installed, or a browser open to Qwen's website if you'd rather not install anything. No paid plan is required for basic menu translation.
- A clear, well-lit photo of the menu, or the ability to type or paste the menu text. Total time: about 5 minutes for a single menu.

Step 1 — Open Qwen on your phone

Tap the Qwen app icon on your home screen. If you haven't installed it yet, open your phone's app store, search for "**Qwen**", and download the official app from the developer listed as Alibaba. When the app opens, you'll see a chat screen with a text box near the bottom and a short welcome message above it. If you're using a browser instead, the same chat layout appears once the page loads. If you don't see a chat box, look for a button labeled "**New chat**" or a + icon in the corner of the screen. You'll know it worked when the chat box is visible, empty, and ready for you to type or attach something.

Step 2 — Start a fresh chat

Tap "**New chat**" (or the + / pencil icon) so you begin with a clean conversation. A blank screen appears with just the input field at the bottom. This matters because earlier messages in a long chat can confuse the assistant when you switch topics. If you already have an old chat open, look for a small + or "**New**" label near the chat history list, sometimes hidden inside a side menu (a ≡ or ≡ icon). A fresh chat also means the response won't be colored by whatever you asked Qwen yesterday. You'll know it worked when the screen shows no previous messages and only an empty input field at the bottom.

Step 3 — Add the menu (photo or text)

Tap the **paperclip / attachment / + icon** next to the chat box and choose "**Camera**" to snap a fresh photo of the menu, or "**Photos**" to pick one you already took. Hold the phone steady, frame the dish list, and capture one section at a time if the menu is long — sharp photos translate far better than blurry ones. If you have the menu as text (maybe you copied it from a website), tap the chat box, **long-press**, and choose **Paste**. The image preview or pasted text appears inside the message bubble before you send. If you don't see a paperclip, look for an icon that looks like a **plus sign** or a **picture frame**. 🗨️ Example prompt to use after pasting a text menu: "Here's a menu from a small restaurant in Lyon, France." You'll know it worked when the photo or text shows up inside your message bubble, ready to send.

Step 4 — Ask for translation with ingredients and allergen notes

Tap the text field and type your request, then send it with the arrow or paper-plane button. In a few seconds, Qwen replies with a translated, dish-by-dish breakdown. Be specific in your ask: a one-line "translate this" usually gives less useful answers than a request that names what you want to know. If the response is too brief, tap "**Regenerate**" or send a follow-up asking for more detail. 🗨️ Example: "Translate this menu into English. For each dish, give a one-sentence description, list the main ingredients, and flag any common allergens like dairy, gluten, eggs, nuts, shellfish, soy, or pork." You'll know it worked when Qwen returns a clear list where each dish has a name, a description, ingredients, and an allergen note.

Step 5 — Ask follow-up questions about specific dishes

Read through the translation and pick anything that sounds unusual or risky for you. Tap the chat box, type a question about a single dish, and send it. Qwen keeps the menu in memory, so you don't need to re-upload or re-explain. This is the moment to ask about taste, preparation style, or whether something is vegetarian, spicy, or contains a hidden ingredient. If the answer is fuzzy, ask the same question again with sharper wording — for example, instead of "is this safe?", try "does this dish contain any peanut products?" 🗨️ Example: "What does 'confit de canard' usually taste like, and is it ever made with something other than duck?" You'll know it worked when Qwen gives a confident, specific answer that helps you decide whether to order it.

Step 6 — Save or share the translation

When the translation is complete and you have what you need, keep a copy for the table. On most phones you can long-press a Qwen reply and choose **Copy**, or take a regular screenshot by pressing the side buttons together. You can also tap "**Share**" (often a small arrow or three dots near the reply) and send the text to a travel companion, your own notes app, or a messaging group. If you'd like the menu translated into another language too (for example, to show a server in their own tongue), ask Qwen — it can switch languages without re-uploading the

photo. You'll know it worked when you can scroll through the translation offline and show a single dish to your server without any guesswork.

Common mistakes

- **Forgetting to mention allergens.** If you only ask for a translation, Qwen won't flag allergens by default. Always include the line "and flag any common allergens" in your prompt, and if you have a serious allergy, double-check with restaurant staff — Qwen's guesses are helpful, not medical advice.
- **Sending a blurry photo.** Small, dim menu photos produce incomplete or wrong translations. Tap to focus before you shoot, hold steady, and capture one section at a time instead of the whole spread.
- **Asking for "the best" dish.** Qwen doesn't know your taste, your budget, or what your travel companions feel like eating. Ask for short descriptions instead, then decide yourself.

What this means for you

- **In everyday life:** the same trick works for grocery labels, medicine packaging instructions, and street signs when you travel. Once you've done a menu, the workflow is identical.
- **For your business or work:** if you host international guests or clients, you can use the same steps to translate a sample menu, brochure, or schedule from a foreign partner — and ask Qwen to flag anything culturally sensitive, like dishes that include alcohol or non-halal meat.
- **If you're just getting started:** don't worry about the prompt being perfect. Even a rough "translate this menu and tell me what's in each dish" gets you surprisingly far. You can always add detail in a second message.

Try it now

Open Qwen, snap a photo of the first menu you see today — even an English one — and ask it to list the ingredients of one dish you've never tried. You'll get a feel for the workflow in under two minutes, and you'll be ready for the real thing on your next trip.