Everyday AI: Understand anything

20 step-by-step guides for real, everyday tasks

Understand anything

Updated: November 8, 2026

AI World HQ · https://aiworldhq.com

Free edition · CC BY-NC 4.0 · free to share with credit

Before you start — what this is, honestly

This is an updated edition. It was rebuilt from our current guides on the date shown on the cover: more guides per topic, retired guides removed, and every step taken from the latest version of each guide.

This collection is free, and every guide in it is also free on our website. This file just gathers them in one place, in reading order, so you can print it, keep it on your phone, or read it offline. Our website has no paywall either — every guide links back to its online page, which is always the newer version.

AI **Every guide in this pack was written by AI** and checked by automated fact and quality gates. **Our editor spot-checks published guides every day, but not every text is read by a person.** Every guide page on https://aiworldhq.com carries the same notice — we say it openly because you deserve to know who wrote what you are reading.

Apps change. Buttons move and menus get renamed. If a step does not match what you see, the guide's online version is the newer one — the link is under every title.

Free to share. This collection is licensed under Creative Commons Attribution-NonCommercial 4.0 (CC BY-NC 4.0): you may copy, print and share it — with your class, your library or your team — as long as you credit AI World HQ with a link to aiworldhq.com and do not sell it.

What's inside

1. **Understand anything**
    1. Fact-Check Online Claims with Perplexity's Sources
    2. How to Ask Gemini to Explain Your Bank Statement in Plain English
    3. How to Fact-Check a Suspicious Text or Link Before Forwarding It
    4. How to compare two rental listings side by side with weighted pros and cons
    5. Catch Up Fast: Summarize Long Chats in WhatsApp and Messenger with Meta AI
    6. Summarize Long Articles into Key Bullet Points with Grok on X
    7. Summarize a long chat thread into 3 bullet points with Copilot
    8. Ask Copilot to compare two product specs and list the differences
    9. Summarize Long Research PDFs with Cohere’s Clever AI
    10. Use Claude to Pull Out the Key Points of a Long Research Paper
    11. Use Qwen to Translate a Short Message for a Friend
    12. How to use Gemini Deep Research to plan a big purchase decision
    13. How to Ask ChatGPT to Explain Any AI News Headline in Plain English
    14. Research and Draft Short Social Media Captions with Perplexity
    15. Compare Products with Perplexity Before You Buy
    16. How to Compare Phone Plans and Pick the Right One Without the Jargon
    17. How to ask Alexa+ to translate a short phrase on your Echo
    18. How to Compare Two Job Offers Side by Side Using AI Safely
    19. Use Grok to Explain a Confusing News Article in Plain English
    20. Explain a tricky school topic with Grok in five minutes

Understand anything

Fact-Check Online Claims with Perplexity's Sources

6 steps · Read online

In our busy online world, information flies around faster than ever. Sometimes, a post on social media might sound convincing but isn't quite right. This guide will show you how to use Perplexity, an AI assistant, to quickly check the facts and understand where information truly comes from. By the end, you'll be able to confidently verify claims, helping you decide what's reliable to share. While Perplexity is excellent at finding and compiling information, remember it doesn't perform human critical analysis or judgment; it's always up to you to evaluate the original sources.

Before you start

- A device with internet access (computer, tablet, or smartphone).
- Access to Perplexity, either through its website (perplexity.ai) or its mobile app. A free account is sufficient for this guide.
- Roughly 5-7 minutes of your time.

Step 1 — Open Perplexity

First, let's get Perplexity ready. Open your web browser and type `perplexity.ai` into the address bar at the top of the screen, then press your keyboard's Enter key. If you're on a mobile device, find and launch the Perplexity app. Once loaded, you'll typically see a clean, simple screen featuring a prominent search bar right in the middle, often displaying placeholder text like "Ask anything..." or "Search the web". This is your main input area, waiting for your questions.

If it looks different: The specific layout might vary slightly depending on whether you're using the website or the app, or due to recent updates. However, you should always find a clear, large text box or search bar positioned centrally, as it's the most obvious place to type your query.

You'll know it worked when you see the Perplexity interface with a clear text input area, ready for you to type your question.

Step 2 — Ask about your claim

Now, let's ask Perplexity to investigate a social media claim. Carefully type the specific statement or question you want to check directly into that main search bar. Try to phrase your query clearly and directly, just as you would ask a person for information. For example, if you saw a post about a new animal, you might ask "Is the 'blue-spotted tree frog' a real animal?". Once you've finished typing your question, press the Enter key on your keyboard or tap the "Ask"

button (which might appear as an arrow or a magnifying glass icon) usually located next to your input.

If it looks different: If there isn't a clearly visible "Ask" button, simply pressing Enter after typing your question usually works on most devices. Make sure your question is specific enough so Perplexity understands exactly what you're trying to verify.

Example: `Is it true that taking a cold shower every morning can cure the common cold?`

You'll know it worked when Perplexity starts to process your request, often showing a loading animation, a spinning icon, or a "Generating answer..." message as it looks for information.

Step 3 — Review the summary and its citations

After a brief moment, Perplexity, acting as an **AI assistant** (a computer program designed to help with tasks and provide information), will present you with a **summary answer**. This is the AI's best attempt to directly address your question, drawing information from various online sources. Read through this summary carefully. You'll typically find a concise answer, along with key facts or explanations. Notice that within the summary, there will be small **numbered references**, like [1], [2], or [3], next to specific pieces of information. These numbers are crucial because they point directly to the original sources Perplexity used.

If it looks different: The summary might be formatted as a series of bullet points, a few longer paragraphs, or a mix of both. Regardless of the exact appearance, the core information and the numbered references should be clearly visible, usually appearing immediately below your initial question.

You'll know it worked when you see a generated answer to your query, complete with small, numbered citations embedded within the text.

Step 4 — Examine the source list

This is where Perplexity truly shines for transparency and fact-checking. Scroll down slightly below the main summary answer. You'll find a dedicated section, typically labeled "Sources" or "References." Here, you'll see a list of the actual websites or articles that Perplexity consulted to formulate its answer, corresponding to the numbers you saw in the summary. For each source, there will usually be a title, a brief description of the content, and the website's address. This **transparency** (showing you exactly where the information came from) is vital because it lets you see *where* the information originated, not just *what* the AI says.

If it looks different: The sources might be displayed as small, clickable cards, or as a simple list. On a mobile device, you might need to tap an "Expand Sources" or a similar button to reveal the full list. Just look for a distinct section that clearly lists external website links.

You'll know it worked when you see a list of external website links, each labeled with a number that corresponds to the citations in the AI's summary.

Step 5 — Dive into an original source

To truly verify a claim and understand its context, it's essential to go directly to the original source. Click on one of the numbered source links from the "Sources" section that seems most relevant to the claim you're checking or appears to be from a **reputable** (trustworthy and respected) organization. This action will open that original webpage in a new browser tab or sometimes within Perplexity itself (depending on your device settings). Now, you can read the information in its original context. Check if the information presented by Perplexity accurately reflects what the source says, look for publication dates, the author, and whether the website itself appears trustworthy (e.g., is it a known news organization, a research institution, or a personal blog?). Comparing the AI's summary with the actual source helps you confirm its accuracy.

If it looks different: If clicking a link doesn't open a new tab, check your browser's pop-up blocker settings, as they can sometimes interfere. You might also have the option to right-click (on a computer) or long-press (on a phone) the link to copy its address, then paste it directly into a new browser tab yourself.

You'll know it worked when you are actively viewing the actual website or article that Perplexity used, allowing you to read the original information.

Step 6 — Refine your query or ask follow-ups

Sometimes, your first question might not give you the complete picture, or the sources provided might be unclear. If you need more information or a different angle, you can ask a follow-up question. Below Perplexity's initial answer, you'll usually see an "Ask a follow-up" prompt or a new blank input box. Type your refined question there, focusing on clarifying details or asking for more specific types of sources. For example, if the initial sources were mostly opinion pieces, you might ask "Are there any scientific studies on this topic?". This iterative process helps you narrow down your search and get more precise, verifiable information.

If it looks different: Instead of an explicit "Ask a follow-up" button, you might simply see another blank input box directly below the answer, allowing you to continue the conversation. If you want to start a brand-new search entirely, look for a "New Thread" or "New Search" button, usually found near the top of the screen or within a side menu.

💬 Example: `Can you find government health organisation statements about cold showers and colds?`

You'll know it worked when Perplexity processes your follow-up query, generating a new answer and set of sources based on your refinement.

Common mistakes

- **Not checking the sources thoroughly**: It's tempting to just read Perplexity's summary and assume it's entirely correct. The AI compiles information, but it can sometimes misinterpret or give too much weight to less reliable sources. Always click through to the original sources to verify the context and content yourself.
- **Being too vague with your question**: Asking broad questions like "Tell me about this news" might not yield precise fact-checks. Instead, provide the specific claim or statement you want to investigate. Being specific helps Perplexity pull more relevant information.
- **Ignoring source reliability**: While Perplexity provides sources, it doesn't always rate their trustworthiness. A quick check of a source's website (e.g., is it a well-known news organization, a research institution, or just a personal blog?) and its publication date is crucial for deciding if the information is truly reliable.

What this means for you

- **In everyday life:** You can quickly become a more informed social media user. Before sharing a sensational post or falling for misinformation, you can take a minute to verify it, protecting yourself and your friends from spreading inaccurate content.
- **For your business or work:** If you work in communications, marketing, or any field where accurate information is key, Perplexity can be a fast way to cross-check industry news, claims from competitors, or statistics before you use them in a report or presentation.
- **If you're just getting started:** The easiest first move is to pick one social media post you're curious about today and run its exact claim through Perplexity. Focus on just reading the summary and then clicking through one source link to see the original content.

Try it now

Find a recent claim or statistic on your social media feed that caught your eye, perhaps something surprising or highly shared. Type that exact claim into Perplexity's search bar and hit Enter. Once you have the summary, scroll down and click on the first source link provided to see where the information originally came from.

How to Ask Gemini to Explain Your Bank Statement in Plain English

6 steps · Read online

> **In short:** You can upload a bank statement into Google's Gemini and ask plain-English questions to surface real spending, recurring charges and hidden fees. Gemini can't read your bank account — think of it as a translator for the numbers on your statement.

A monthly bank statement is full of small numbers, merchant codes and fee names that rarely mean anything at first glance. With Google's Gemini, you can hand it that document and ask, in normal sentences, "what am I actually paying for?" — and get a readable answer back. This guide walks you through it, one step at a time.

Before you start

- A free Google account (the same one you use for Gmail or YouTube).
- Access to Gemini in a browser (gemini.google.com) or the Gemini mobile app — both work for this. The free version is enough for reading a single statement.
- A copy of your monthly bank statement, saved as a PDF or a screenshot — most banks let you download it from their website or app.
- About 10 minutes the first time; once you've done it once, the same questions work on next month's statement.

Step 1 — Open Gemini and start a fresh chat

Go to **gemini.google.com** in your browser, or open the **Gemini app** on your phone. If you are not signed in, the page will ask you to sign in with your Google account — the same email and password you use for Gmail. Once you are in, look for a button or link along the left side or top of the screen that says something like **"New chat"** or **"Start a new chat"** and click or tap it. A blank conversation opens. If you don't see that exact label, look for a **pencil icon**, a **"+" sign**, or a text box that says **"Ask Gemini"** — any of those will start a fresh conversation.
Example: open the Gemini app, tap "New chat," and you should see an empty text box with the prompt "Ask Gemini" at the bottom. You'll know it worked when you see an empty text field with no previous messages above it.

Step 2 — Upload your bank statement

Look for a small **paperclip icon**, a **"+" button**, or an **"Upload"** option near the text box — every version of Gemini has a way to attach a file. Click it, then choose the **PDF** of your statement from your computer or phone. If your statement is a stack of screenshots instead, you

can either upload them one at a time or, on the desktop site, drag and drop them all into the chat window together. The file name appears in your message once it uploads. If you can't find an upload option at all, your version may only accept **pasted text** — open your PDF, select all the text (Ctrl+A or Cmd+A), copy it (Ctrl+C or Cmd+C), and paste it (Ctrl+V or Cmd+V) directly into the chat box instead. Example: "I'm pasting my bank statement below — it's about 3 pages of transactions." You'll know it worked when a small thumbnail or file name shows up next to your message, and Gemini replies with something like "Got it, I'm reading the statement now."

Step 3 — Ask for a plain-language summary

Now type a first question in the same chat. Don't ask for anything fancy yet — just ask Gemini to read the statement and tell you, in normal words, what happened in the month. Think of it as asking a friend who is good with numbers to glance at the page and explain it. Be specific about the period, and tell Gemini how you'd like the answer — short bullets, full sentences, whatever you prefer. If Gemini answers with generic budgeting advice instead of reading your actual statement, it probably didn't see the file — re-upload it or paste the text directly into the chat. Example: "Please read the statement I just uploaded and give me a plain-English summary of what happened in November. List the top 3 things I spent money on and any fees I was charged. Keep it under 200 words." You'll know it worked when Gemini replies with a readable summary that mentions real names and amounts from your statement — not a generic answer about budgeting.

Step 4 — Ask about recurring charges and small fees

Now go one level deeper. Small monthly fees — a $4.99 streaming service, a $12 "account maintenance" charge, a $2 ATM surcharge — are easy to miss when you scroll quickly. Ask Gemini to look specifically for anything that appears more than once, or anything that looks like a fee rather than a purchase. This is the step where most people discover charges they've been paying for months without noticing. If Gemini says it can't see any recurring charges, the file may have been only partially read — ask it to look at every page, or paste the full transaction list into the chat. Example: "Now look only at the charges that happened more than once this month, and anything that looks like a fee or service charge. For each one, tell me the name, the amount, and how many times it appeared." You'll know it worked when you get a short list of repeated merchants and fee names you can scan in 30 seconds — and probably one or two that surprise you.

Step 5 — Ask about your real spending patterns

The next question is about categories. Gemini can group your transactions by type — food, transport, subscriptions, cash withdrawals — and show you where most of the money went. This

is helpful even if your statement already has a "category" column, because Gemini's version often groups things more sensibly for everyday reading, and you can ask follow-up questions about any category that looks too big or too small. If Gemini lumps almost everything into "other," your statement may have lots of unfamiliar merchant names — ask it to list the biggest "other" items so you can see what got hidden. Example: "Group all my spending from this statement into simple categories like food, transport, subscriptions, bills and other. Show me the total for each category and the percentage of my total spending. Then tell me which category grew the most compared to a normal month." You'll know it worked when you see a small table or list with named categories, dollar totals and clear percentages — something you could read once and remember.

Step 6 — Save your insights or turn them into questions for your bank

The last step is making sure you don't lose what you just learned. If something in the summary or the fee list looked wrong, copy Gemini's exact wording of that charge — the merchant name, the date, the amount — and use it next time you contact your bank. You can also ask Gemini to rewrite the whole summary into an email or a checklist you can save. To come back to this conversation later, look in the left sidebar for your recent chats — your version of Gemini should keep a running list you can reopen. If you can't find it, bookmark the page in your browser before you close the tab. Example: "Turn the summary from Step 3 and the fee list from Step 4 into a short email I can send to my bank asking about the $12 monthly fee. Keep the tone polite and include the date and amount for each charge." You'll know it worked when you have either a saved conversation, a copyable email draft, or a clear list of questions to bring to your bank's customer service.

Common mistakes

- **Uploading a statement from a bank that masks real names.** Some banks replace merchant names with codes (like "POS DEBIT 8842"). Gemini can't decode those on its own — you'll get a vague answer. If this happens, ask your bank for an itemized version.
- **Asking one giant question instead of several small ones.** A request like "tell me everything about my statement" usually produces a fuzzy answer. Three or four short, focused questions (like in Steps 3, 4 and 5) get sharper results.
- **Forgetting to remove sensitive info if you share the chat.** A bank statement has your account number, sometimes your address and always your full name. If you ever screenshot a Gemini chat to share it, blur those parts first.
- **Treating Gemini's answer as a final verdict.** It can misread a number, miss a charge or guess wrong about a fee. Always check the original statement before calling your bank — Gemini helps you ask better questions, it doesn't replace them.

What this means for you

- **In everyday life:** the next time a charge looks unfamiliar on your statement, you can paste it into Gemini and ask, in one sentence, "what is this $9.50 charge from XYZ?" before you spend 20 minutes on hold with customer service.
- **For your business or work:** if you run a small business and review statements monthly, this same set of questions can become a 5-minute weekly check instead of a 90-minute pile of receipts.
- **If you're just getting started:** try it on a statement you already understand well — your last month, perhaps. That way you can quickly see whether Gemini's summary is accurate, before you trust it on a more confusing one.

Try it now

Open Gemini, upload this month's bank statement, and paste this single prompt: "Read this statement and list every charge I don't recognize, with the date and amount. Keep it to 10 items or fewer." Ninety seconds, one question — and you'll likely learn at least one thing you didn't know you were paying for.

How to Fact-Check a Suspicious Text or Link Before Forwarding It

6 steps · Read online

> **In short:** Paste a suspicious message into any AI assistant to list its warning signs, then verify the claim yourself through a trusted channel. The AI cannot confirm that a link is safe or that a message is true, so treat its answer as a second pair of eyes.

A message that says "forward this to everyone now!" can make anyone hit send. This guide shows you how to use an AI assistant as a careful first reader: you'll build a reusable checklist, get a clear list of red flags, and finish with a plan for what to check before you share. It's written for anyone who already uses a chat assistant and wants a more reliable routine.

Before you start

- An account on any AI assistant you already use. You know how to start a chat and send a message, so we skip that part.
- The suspicious text or link, copied into a notes app on your phone or computer.
- A free plan is usually enough for this. Very long pastes can be cut short on some plans, and limits can change, so check the reply is complete.
- About 10 minutes.

Step 1 — Save a reusable checklist

Open your notes app and create a note called something like "Scam check template." Paste in the instructions below, so you never have to rewrite them. Most people find this helpful because the same checklist gets used every time, which keeps your review consistent.

Once saved, the template is simply copied into a new chat each time. You don't need any special settings or paid features for this; it's just a saved piece of text you control.

If your assistant has a feature for saved instructions or custom profiles, you can paste the template there instead. Look for a settings area with words like "instructions" or "custom." If you can't find one, the notes app works just as well.

💬 Example: "You are helping me check a message before I forward it. Look for these warning signs: urgency or pressure, requests for money, gift cards or passwords, a sender I can't verify, links that don't match the company they claim to be, and wording that feels off. Then list what you cannot tell from the text alone."

You'll know it worked when your note holds the full template and you can copy it in one tap.

Step 2 — Start a fresh chat with the message

Open your AI assistant and start a new chat, so earlier conversations don't mix into this check. Copy the template from your note and paste it into the message box. Then add the suspicious message below it. Most assistants accept both in one message, and you should see your text appear in the chat box before you send.

Tell the assistant who sent it, if you know. That detail matters. If the message says it came from a friend, a bank, or a delivery company, mention that, but don't assume it's true.

If the box won't take a long paste, try pasting in two parts, with the template first and the message second in a new message. If the reply only covers part of the text, ask it to continue.

Example: "Here is the message I received from [who sent it]: [paste the message here]"

You'll know it worked when the assistant replies with a structured list or a set of observations about the message, rather than just repeating your text back to you.

Step 3 — Ask for the evidence behind each flag

A red flag is only useful if you know why it's there. Read the assistant's list, then send a follow-up asking it to show its reasons. This builds on the first answer, which is a simple but powerful way to get a more useful response.

Look for each flag to point to exact words from the message. If the assistant claims something is suspicious but can't point to a specific sentence, be skeptical of that flag and ask again.

If the follow-up answer is vague, rephrase your request with more specifics, such as "only list flags you can tie to a quoted sentence." Keep the chat open so the assistant remembers the message you pasted.

Example: "For each warning sign you found, quote the exact words from the message and explain in one sentence why it could be a concern."

You'll know it worked when every flag comes with a quoted phrase and a short reason, so you can check the reasoning yourself in the original text.

Step 4 — Inspect a link without tapping it

If the message contains a web address, don't tap it yet. Copy the address as text, or long-press it to copy, and paste it into the chat. The AI can break the address into parts so you can read it more clearly.

Here's what to look for. In an address like `https://usps-tracking.example.com/pay-now`, the name that matters most is the part right before the main ending, which here is `example.com`. The words in front of it can be anything, so a famous company's name at the start proves very little. The AI can help you spot that pattern, but it cannot visit the site to tell you whether it's live or dangerous.

Be cautious of shortened links and links that ask you to log in or pay right away. Those are reasons to slow down, not to panic.

💬 Example: "Break down this web address into its parts and tell me which part is the actual site name: [paste the address here]"

You'll know it worked when the assistant explains which part of the address belongs to the site and points out any mismatch with the company it claims to be.

Step 5 — Confirm through a channel you trust

This is the most important step, and no AI can do it for you. Ask the assistant which claims in the message you would need to confirm directly. Then check each one through a source you already trust: type the company's address yourself, use the official app, or call the phone number printed on a bill or a card you own. Don't use a number or link that came inside the message.

If the message claims a friend is in trouble or needs money, call that person on a number you already have saved. Scammers often rely on urgency to stop people from checking.

If you can't reach anyone to confirm, the safe choice is to hold off on sharing it. Waiting costs almost nothing.

💬 Example: "Based on this message, what are the three most important claims I should confirm directly with the company before I share it?"

You'll know it worked when you have a short list of specific things to verify and a plan for how you'll check each one outside the message.

Step 6 — Draft a calm heads-up for family

Once you know what you're going to do, ask the assistant to help you write a short message to the people you'd otherwise have forwarded to. A calm, specific note is more useful than an all-caps warning, and it helps your family learn the checking habit too.

Keep the draft short and kind. Mention what you found and what you're doing about it, without blaming the person who sent it to you. Read the draft before sending, because you know your family better than the assistant does.

💬 Example: "Write a short, calm message I can send to [family member name] saying I'm not sure this text is real, that we should check it with [trusted contact] first, and that they shouldn't share it yet."

You'll know it worked when you have a short, friendly draft that you can edit and send in your own voice.

Common mistakes

- **Trusting the AI as the final judge.** The assistant reads the text you give it and can be wrong. The fix is to always confirm the key claims through a source you trust, as in Step 5.
- **Clicking the link to "see what it says."** Even a quick tap can load a page you didn't mean to visit. The fix is to paste the address into the chat as text, as in Step 4, and check it through the official site you type yourself.
- **Forwarding first and checking later.** Once a message is sent, you can't pull it back. The fix is to make Steps 2 to 5 a habit: pause, paste, check, then decide.

What this means for you

- **In everyday life:** When a worrying text arrives about a package, a bill, or a relative in trouble, paste it into your checklist and follow the steps before anyone in your family acts on it.
- **For your business or work:** Keep the template in a shared note so coworkers can use the same checklist before forwarding a suspicious invoice or request to the team.
- **If you want to go further:** Try building a short library of templates for different situations, such as payment requests, delivery notices, or account alerts. Keep each one short, and test it on a message you already know is legitimate so you learn what a normal message looks like.

Try it now

Open your notes app and save the template from Step 1. Then take one message you received recently, paste it into a new chat with the template, and read the first list of flags. You'll know this worked when you have a saved template and a short, quoted list of observations for a real message.

How to compare two rental listings side by side with weighted pros and cons

6 steps · Read online

> **In short:** Open any AI assistant, paste both listings, list your priorities with weights, and ask for a scored comparison table. The assistant won't visit the apartment for you, but it will turn a wall of listing text into a clear trade-off table in minutes.

When you're choosing between two rentals, the listings blur together. This five-minute workflow works in any AI assistant — ChatGPT, Gemini, Claude, Copilot, Le Chat, DeepSeek, or whichever you already have — and scores both places against the priorities you set yourself.

Before you start

- An account with any AI assistant you already use (free or paid — both work for this).
- The full text of both rental listings, copied from the listing page or app. Screenshots alone won't work — only pasted text.
- A short list of what actually matters to you (commute, noise, utilities included, lease length, pets, parking — three to six items is enough).
- About 15 minutes.

Step 1 — Open a new chat and paste both listings

Start a fresh conversation — don't drop this into an old chat, or the assistant will weigh your rentals against whatever you were talking about before. In the message box at the bottom of the screen, paste the first listing, leave a blank line, then paste the second one. End the message with a one-line note telling the assistant what it's looking at. Example: "Listing A — 2-bed in Eastside, $X/month, 12-month lease. — — Listing B — 1-bed + den in Northgate, $Y/month, month-to-month. Compare these two rentals." If the listing is mostly images or a PDF, you'll need to type out the key facts by hand — assistants can't reliably read screenshots unless your plan includes file upload or vision, and even then the result varies. You'll know it worked when you can scroll up and see both listings sitting in the chat, and the assistant acknowledges them in its first reply.

Step 2 — Tell the assistant your priorities and how much each matters

Send a second message that lists the things you care about and gives each one a weight — how much it should count in the final score. Weights are how you stop the assistant from treating "has a dishwasher" as equally important as "20-minute commute." The format that works best is a numbered list with a 1–5 or 1–10 weight next to each priority. Example: "Here are my

priorities, weighted 1–5: 1) Commute under 30 min by transit (5). 2) Quiet street, no main road (4). 3) Utilities included in rent (3). 4) Allows my cat (5). 5) In-unit laundry (2). Score both listings against these." If your assistant has a "memory" or "custom instructions" setting (different apps call it different things), you can save the weights there once and reuse them for future comparisons. You'll know it worked when the assistant restates your priorities back to you without being asked — that means it actually loaded them into the conversation.

Step 3 — Ask for a clean side-by-side table first, before any scores

Now ask for a comparison table with one row per feature and one column per listing. Don't ask for scores yet — you want a neutral snapshot of what each listing actually says, so you can spot places where the assistant fills in gaps with guesses. The clearer this table is, the more honest the scoring will be in the next step. Example: "Make a table with these rows: rent, lease length, square feet, utilities included, parking, laundry, pet policy, commute time to downtown, nearest noise source. One column for Listing A, one for Listing B, one for 'verified or missing.' If a fact isn't in the listing, write 'not stated.'" If the table renders as plain text instead of a real grid (some mobile apps do this), ask the assistant to "rebuild it with pipes and dashes" and it will produce a text-only version that still reads clearly. You'll know it worked when every row is filled with text from the listing — or says "not stated" — instead of the assistant inventing something to fill the gap.

Step 4 — Add weighted scores and a recommendation

Send a third message asking the assistant to score each listing on your weighted priorities and add a one-line recommendation at the bottom. This is the step that turns the table into a decision. The math is simple: multiply each priority's weight by the listing's score on that priority, then add them up. The listing with the higher total wins — *on your priorities*, not on someone else's. Example: "Now score both listings 1–10 on each priority from Step 2, multiply by the weight, and show the totals. Add a two-sentence recommendation naming which one wins on my priorities and why." If the assistant hedges with "it really depends on you," that's a sign it didn't read your weights carefully — nudge it with "use the weights literally, don't substitute your own." You'll know it worked when you see a final number for each listing and the recommendation names a clear winner with a stated reason.

Step 5 — Push back on vague answers and list what to verify

Before you act on the recommendation, ask the assistant to flag anything it had to guess at, plus every "not stated" cell from the table. This step is the safety net — assistants can sound confident about facts that aren't actually in the listing, and a lease you signed based on a wrong assumption is an expensive lesson. The honest answer often involves several things you need to

ask the landlord directly. 💬 Example: "List every fact in the table you marked 'not stated' and anything you had to infer. For each one, write a one-line question I can email the landlord." If the assistant invents a detail (like "utilities are included" when the listing didn't say so), that's a hallucination — when an AI confidently states something that isn't actually there. Call it out with: "You made that up; the listing didn't say it. Replace with 'unverified, ask landlord.'" You'll know it worked when you have a short, concrete list of questions to send before you sign anything.

Step 6 — Save the prompt structure as a reusable template

Copy the three prompts you sent (Steps 1, 2, and 4) into a notes app and label it "Rental comparison template." Next time you're weighing two places, open a new chat, paste the template, swap in the new listings and weights, and you'll get the same scored comparison in minutes — no setup, no re-thinking the structure. 💬 Example: "PROMPT 1: [paste listings + intro line]. PROMPT 2: [paste weighted priorities]. PROMPT 3: [paste score-and-recommend request]." If your assistant supports a way to save reusable prompts (different apps call this Projects, Gems, custom instructions, or similar — check the menus if you're not sure), you can store the template there so it loads automatically next time. You'll know it worked when the template is sitting in your notes ready to go the next time a listing site or a landlord's email hands you two options.

Common mistakes

- **Letting the assistant pick your priorities.** It will default to generic things like "location" and "price." Insist on *your* weights, even if they feel quirky — a quiet street might matter more to you than square footage, and that's the whole point of doing this yourself.
- **Pasting only the headline summary.** The first paragraph of a listing is marketing. The fine print about utilities, parking, and pet fees is usually further down — paste the whole thing, or you'll be comparing ad copy instead of facts.
- **Trusting the score without checking the table.** A 9/10 on "utilities included" means nothing if the listing never actually said that. Always read the table before you read the score, and send the "not stated" list to the landlord before you commit.

What this means for you

- **In everyday life:** When you're weighing two rentals (or two cars, or two phone plans), this same three-prompt pattern turns a messy comparison into a clear trade-off table in minutes — without you having to argue with yourself about which factor matters more.
- **For your work or business:** The same structure works for choosing between vendors, job offers, software plans, or any two options where you want *your* priorities to drive the

score — not the assistant's defaults.

- **If you want to go further:** Try extending the prompt to three listings instead of two, or add a "dealbreaker" rule that automatically disqualifies any option missing a non-negotiable (like "no cat-friendly place counts, period").

Try it now

Pick one priority that matters to you (commute, pets, quiet, anything), find two short rental listings — even made-up ones work for practice — and run Steps 1 to 4 in a fresh chat. You'll see the decision become obvious in under five minutes.

Catch Up Fast: Summarize Long Chats in WhatsApp and Messenger with Meta AI

6 steps · Read online

By the end of this guide you'll know how to ask Meta AI (a built-in helper inside WhatsApp and Messenger) to read through a long chat and give you a short, useful summary in plain language. This is for anyone who has ever scrolled back through hundreds of messages thinking *"what did I miss?"* — whether it's a family group chat or a busy work thread.

Before you start

- A phone with **WhatsApp** or **Messenger** installed, updated to the latest version (Meta AI is rolled out in stages, so update to be sure).
- A personal Meta account already logged in on that app.
- Meta AI available in your country or region — if you cannot see it after updating, it may not have reached your area yet.
- Around **5 to 10 minutes** of your time, plus the chat you want summarized open in another window or conversation.

One honest note up front: Meta AI **cannot read old messages you have not shared with it**. You will need to either paste the messages into a chat with Meta AI or @-mention it in a recent thread. It is free to use, but anything you send is processed by Meta — so avoid pasting very sensitive personal details.

Step 1 — Open WhatsApp or Messenger and look for Meta AI

Open the app you usually use (WhatsApp or Messenger) and look at the top of your chat list or your contacts. You should see Meta AI listed as an option, sitting alongside your usual contacts. Tap it once. A new chat window opens, just like starting a conversation with a friend. You'll know it worked when you see a message from Meta AI indicating it's ready to help. **If it looks different:** some users see a different icon or a "New chat" option with an AI symbol, or a search bar at the top — any of these can lead to Meta AI. If you see nothing related to AI, your app may need an update, or Meta AI may not yet be available in your region.

Step 2 — Open the chat you want summarized, side by side

Leave the Meta AI chat open, then open the long chat you want to catch up on. On most phones you can switch between apps with the square "recent apps" button, or use your phone's split-screen feature if available. You will see the full message history of your long chat on one screen and the Meta AI chat on the other. This makes copying easy. **If it looks different:** on tablets or

split-screen devices you can arrange both windows side by side so both sit next to each other at the same time.

Step 3 — Copy the messages you want summarized

In the long chat, press and hold the first message you want Meta AI to read, then choose **Copy**. For a range of messages, keep tapping each one until you have selected the section you missed (often the last 20 to 50 messages is plenty). Then tap **Copy** again. The text now sits on your phone's clipboard, ready to paste — like copying a paragraph from a website. **If it looks different:** some Android phones show a small menu with "Copy", "Forward" and "Share"; on iPhone you'll see "Copy" appear above the selected message. If selecting one-by-one is too slow, try **Forward** instead, then paste into the Meta AI chat in one go.

Example: *"I copied messages from a 40-person family group covering about a day's worth of chat — roughly 30 messages with photos, links and a few voice notes."*

You'll know it worked when your phone briefly shows a small "Copied" label or a tick next to the selected messages.

Step 4 — Paste the messages into Meta AI and ask for a summary

Switch back to the Meta AI chat. Tap inside the message box, hold for a moment, and choose **Paste**. The copied messages appear as one long block. On a new line underneath, type a clear, kind request that tells Meta AI exactly what you want — think of this as giving instructions to a helpful assistant, not just dumping text on it. Then tap the send arrow. After a few seconds, Meta AI replies with a summary. **If it looks different:** if the paste does not appear, try tapping once inside the text box first, then holding your finger down again. If Meta AI replies in a language you did not expect, add "Please reply in English" at the end of your message.

Example: *"Here are the messages from my family group from yesterday. Please summarize them in 5 short bullet points, focus on any decisions or plans people made, and ignore the photos. Reply in plain English."*

You'll know it worked when Meta AI replies with a summary rather than echoing the messages back at you.

Step 5 — Refine the summary if it is too long or off-target

Read Meta AI's reply. If the summary is too vague, too long, or misses what you actually care about, reply with a short follow-up message asking for a change. The assistant remembers the conversation, so you can be casual about it. Each new request usually takes a few seconds. **If it**

looks different: if Meta AI seems stuck or repeats the same answer, try starting a brand-new chat and pasting the messages again with a clearer instruction — sometimes a fresh start produces better results.

💬 Example: *"Great, but please make it 3 bullet points instead of 5, and add a one-line 'action items' section at the end showing what I might need to reply to."*

You'll know it worked when the new summary matches what you actually wanted to know — short, focused and useful.

Step 6 — Save or share the summary

Once you are happy with the summary, you can save it for later or send it to someone else. Tap and hold the assistant's message, then choose **Copy** to put it on your clipboard, or look for the forward option to send it to another chat or contact. Many phones also let you tap **Share** and save it straight to Notes, email or a document. **If it looks different:** if you don't see forward options, look for a menu icon or arrow next to the message. From there you can usually forward, share or save the message.

💬 Example: *"I forward the summary to a private chat with myself called 'Catch-up notes' so I can search it later without scrolling through the original group."*

You'll know it worked when the summary now lives somewhere easy to find — your notes app, a saved message, or a chat you control.

Common mistakes

- **Pasting months of messages at once.** Meta AI has a limit to how much text it can read in one go, and very long pastes often get cut off or produce a generic answer. **Fix:** paste the last 20 to 50 messages only, or split a longer chat into two requests.
- **Sharing sensitive personal information.** Anything you paste into Meta AI is processed by Meta's servers, so names, addresses, phone numbers or financial details should be removed first. **Fix:** change real names to "Alex", "Sam" and "the host" before pasting, especially for group chats with people you do not know well.
- **Expecting Meta AI to know your whole chat history without pasting.** Unlike asking a friend, the assistant only sees what you have shared with it in that conversation. **Fix:** always paste the messages you want summarized, or @-mention Meta AI inside a recent thread so it can see the surrounding context.

What this means for you

- **In everyday life:** catch up on noisy family or friend group chats before replying, instead of scrolling through 200 messages on the train.
- **For your business or work:** turn long WhatsApp or Messenger threads with clients, teammates or event groups into a short briefing you can skim in 30 seconds and forward to a colleague.
- **If you're just getting started:** pick one group chat you fall behind on, copy the last 20 messages, paste them into Meta AI and ask for 3 bullets — that single step is the whole skill in miniature.

Try it now

Open your busiest group chat right now, copy the last 20 messages, paste them into a Meta AI chat, and ask: *"Summarize these in 3 bullet points and list anything I need to reply to."* It takes under two minutes and you'll have the skill for life.

Summarize Long Articles into Key Bullet Points with Grok on X

6 steps · Read online

In today's fast-paced world, it's easy to feel overwhelmed by the sheer volume of information available online. This guide will show you how to use Grok, an AI assistant integrated into the X platform, to quickly condense long news articles or reports into concise bullet points. By the end, you'll be able to get the essential information from any article in moments, letting you stay informed without needing to read every word. While Grok can summarize the content of an article, it cannot verify the accuracy of the original article's claims or perform actions like sending emails for you.

Before you start

- An active X Premium+ subscription, as Grok is only available to subscribers of X's highest tier.
- The X app on your mobile device or access to the X website on a computer.
- Roughly 5-10 minutes of your time.

Step 1 — Access Grok on X

First, open the X app on your phone or visit the X website on your computer. Once the X feed has loaded, look for a special icon that represents Grok. On mobile, this is usually a diamond-shaped icon found in the bottom navigation bar. On a computer, you might find a "Grok" option listed in the left-hand sidebar menu, often near your profile or notifications. Tap or click on this icon to enter the Grok interface. If it looks different: The X interface can change, so if you don't see a diamond icon or a clear "Grok" label, look for an icon that stands out, or check the 'More' menu (often three lines or dots) for an entry labeled "Grok". You'll know it worked when you see Grok's chat interface, which usually includes a welcome message or a list of previous conversations.

Step 2 — Copy the Article Link

Now, you need to find the online article you want to summarize. Open your web browser (like Chrome, Safari, or Edge) and navigate to the news article, blog post, or report you're interested in. Once you're on the page, look for the web address (URL) at the top of the browser window. Tap or click on this address to highlight it, then use your device's copy function (e.g., right-click and 'Copy' on a computer, or long-press and 'Copy' on a phone) to save the link to your clipboard. If it looks different: Some apps have a 'Share' button (often an arrow pointing upwards or a three-dot menu) that includes a 'Copy Link' option. If you can't easily highlight the

URL, look for that share option. You'll know it worked when the link is saved and ready to be pasted elsewhere.

Step 3 — Start a New Chat with Grok

With your article link copied, return to the Grok interface within the X app or website. To begin a fresh summarisation task, you'll need to start a new conversation. Look for a button that typically indicates starting something new, like a '+' symbol, a pencil icon, or a "New Chat" button, usually located at the top-right or bottom-right corner of the Grok chat screen. Tapping this will clear any previous messages and give you a clean slate. If it looks different: If you don't see an obvious 'new chat' icon, you can often just type your request directly into the existing message box. Grok is usually smart enough to treat a new request as a new topic. You'll know it worked when the chat area is empty and a message input box is clearly visible at the bottom of the screen, ready for you to type.

Step 4 — Paste Your Link and Request a Summary

Now it's time to tell Grok what you want. Click or tap inside the message input box at the bottom of the Grok chat screen. Paste the article link you copied in Step 2 into the box. After the link, type a clear instruction asking Grok to summarize the content. Be specific about the format you prefer, like asking for bullet points, to help Grok deliver exactly what you need. If it looks different: If pasting doesn't work, try typing the full URL manually if it's short, or check your device's clipboard settings to ensure the link was copied correctly. Grok might indicate if it's having trouble accessing the link. 💬 Example: `Please summarise this news article into 5 clear bullet points, focusing on the main findings: [paste your article link here]` You'll know it worked when Grok displays a message like "Generating response..." or "Thinking...", followed shortly by the summary appearing in the chat.

Step 5 — Review Grok's Summary

Once Grok has finished processing your request, its summary will appear in the chat window, presented in the bullet-point format you asked for. Read through these points carefully. This is your chance to quickly understand the core message, key arguments, and important facts from the article without having to read every sentence of the original text. Take a moment to ensure the summary captures what you expected and is easy to understand. If it looks different: Sometimes, Grok might ask a clarifying question if the article is too long, behind a paywall (which Grok cannot bypass), or contains ambiguous content. If this happens, try rephrasing your request or finding a more accessible article. You'll know it worked when you have a concise list of key information that helps you quickly grasp the article's main content.

Step 6 — Refine the Summary (Optional)

Grok's initial summary might be great, but you can always ask for more details or a different perspective if needed. If a particular bullet point sparked your interest, or if you feel something important was missed, simply type a follow-up question into the message box. You can ask Grok to expand on a specific detail, explain a complex term, or even rephrase the summary for a different audience (e.g., "Summarize this for a high school student"). This allows you to dig deeper without returning to the original lengthy article. If it looks different: If Grok says it can't find the information, it might be that the detail wasn't explicitly in the article, or it's too specific for a general summary. Try rephrasing your follow-up in simpler terms. 💬 Example: `Can you elaborate on the economic implications mentioned in bullet point 2?` You'll know it worked when Grok provides a more detailed or adjusted response, showing it's continuing to process the article's information based on your new prompt.

Common mistakes

- **No Grok access:** Grok is a premium feature, so if you can't find it or get an error, ensure your X subscription is the highest tier (X Premium+). The fix is to upgrade your X subscription.
- **Pasting a broken or inaccessible link:** If Grok says it can't access the article, check that the link is correct and that the article isn't behind a paywall or login that Grok can't bypass. The fix is to find a publicly accessible link or try a different article.
- **Vague summary request:** If Grok gives a less-than-useful summary, your request might have been too general. The fix is to be more specific in your prompt, like "Summarize into 5 bullet points" or "Focus on the environmental impact."

What this means for you

- **In everyday life:** Quickly catch up on news headlines or understand long reports from community groups without spending hours reading. It's perfect for when you're short on time but want to stay informed.
- **For your business or work:** Efficiently digest industry news, competitor analysis, or lengthy internal reports, allowing you to extract critical insights faster and make more informed decisions.
- **If you're just getting started:** Practice by picking a news article you'd normally skim, then use Grok to get a bullet-point summary. Compare it to what you would have highlighted yourself.

Try it now

Open the X app, navigate to Grok, and paste a link to any recent news article. Ask Grok to summarize it for you in three bullet points. See how quickly you can get the main gist!

Summarize a long chat thread into 3 bullet points with Copilot

6 steps · Read online

Ever feel lost trying to find the key takeaways from a really long email chain or group chat? This guide is for anyone who wants to quickly cut through the noise and get straight to the important decisions and actions. By the end, you'll be able to use Microsoft Copilot to summarize any long text into three clear bullet points, making it easy to catch up or share concise updates.

Before you start

- A free Microsoft account (you can sign up for one easily if you don't have it).
- A web browser like Microsoft Edge, Google Chrome, or Mozilla Firefox.
- A long piece of text you want to summarize, such as an email conversation or a chat history.
- Roughly 5-7 minutes to follow the steps.

Step 1 — Open Copilot in your browser

Your first move is to open Copilot in your preferred web browser. Simply type `copilot.microsoft.com` into your browser's address bar at the top of the window, then press `Enter` on your keyboard. The Copilot homepage should load, presenting you with a clear chat input box usually located at the bottom of the screen, waiting for your commands. If the page looks a little different, don't worry – websites update often. Just look for a prominent chat area or a button that says "Start chat" or "Ask me anything." You'll know it worked when you see a clear input box at the bottom of the screen, probably saying "Ask me anything" or similar.

Step 2 — Sign in to your Microsoft account

To get the most out of Copilot, it's best to sign in. Look for a "Sign in" button, typically positioned in the top-right corner of the Copilot page, and click it. A new screen will appear, prompting you to enter your Microsoft account details, like your email address or phone number, and then your password. If you're already signed in, you'll see your profile picture or initials in that corner and can simply skip this step. If it offers choices like "work or school account," usually you'll select "personal account" for this task. You'll know it worked when your profile picture or initials appear in the top-right corner, indicating you are successfully logged in.

Step 3 — Craft your summary request

Now it's time to tell Copilot exactly what you want it to do. In the chat input box at the bottom of the screen, type out your specific instruction. Be clear that you want a summary, how many bullet points you need, and what you want it to focus on – in this case, decisions and action

items. This "instruction" is often called a **prompt** (which is just the message you send to an AI). Your typed message will appear in the input box, ready for you to add the content you want summarized. 💬 Example: `Summarize the following chat thread into exactly three short bullet points, focusing on key decisions and action items:` You'll know it worked when your complete instruction is clearly visible in the chat box, ready for the content you're about to paste.

Step 4 — Paste your long text

Next, you need to provide Copilot with the text to summarize. Go to the original source of your long chat thread or email chain (like your email client or messaging app), select all the text you want, and copy it. Then, click directly into the Copilot chat input box, making sure your cursor is positioned right after the instruction you typed in Step 3. Now, paste the copied text. The entire block of text will appear in the input box after your instruction. Copilot has a "context window" (the amount of text it can "remember" and process in one conversation), but for most chat threads, it will fit easily. You'll know it worked when all your copied text is visible in the input box, ready to be sent.

Step 5 — Send your request

With your instruction and the long text in place, it's time to send it to Copilot. Press the `Enter` key on your keyboard, or click the "Send" button that looks like a paper plane icon, usually located next to the input box. Once sent, your entire message will move from the input box into the main chat history. Copilot will then show a "thinking" indicator, meaning it's processing your request. After a short moment, it will generate its response, providing your summary. If the "Send" icon looks a little different, just know it's always near the input box. You'll know it worked when Copilot starts typing its response, and you see the summary appear as three bullet points in the chat window.

Step 6 — Review and refine your summary

Read through the three bullet points Copilot has generated. Does it capture the essence of the conversation? Does it highlight the key decisions and actions as you intended? If it's not quite perfect, you can ask Copilot to refine it. Simply type a follow-up request in the input box, explaining what changes you'd like. Copilot will then generate a revised summary based on your new feedback. This helps you get precisely what you need from the conversation. 💬 Example: `Please make the bullet points even shorter, just focusing on who needs to do what.` You'll know it worked when you have a concise, useful three-point summary of your original long text that meets your needs.

Common mistakes

- **Pasting text without a clear instruction:** If you just paste a long thread into Copilot without telling it what to do, it might simply repeat parts of the text or give a generic response. Always start with a specific instruction, like "Summarize this into three bullet points."
- **Assuming the AI knows context:** Copilot only "knows" what you provide in the current conversation. If your chat thread refers to external documents or past discussions, Copilot won't have that information unless you paste it in. Make sure all necessary information is included for the best summary.
- **Too much text for one go:** While Copilot has a large capacity, extremely long documents (tens of thousands of words) might exceed its "context window" (the amount of text it can process at once). If Copilot seems to miss the end of your content, try breaking the original text into two or three smaller chunks and summarizing them separately.

What this means for you

- **In everyday life:** Imagine catching up on a flurry of family messages about a holiday plan or quickly understanding the core points of a long community group discussion. Copilot helps you get the gist without having to read every single message, saving you precious time.
- **For your business or work:** Instantly extract the crucial decisions, assigned tasks, and next steps from lengthy meeting transcripts, project update emails, or internal team chats. This ensures you're always informed and can provide concise updates to colleagues without sifting through pages of text.
- **If you're just getting started:** A great way to practice is to take a long news article or a detailed product review you're interested in, copy the text, and ask Copilot to summarize it into three bullet points. This helps you see how it condenses information.

Try it now

Find a recent long email or a lengthy article online that you've been meaning to read. Copy the main body of text. Then, open Copilot, paste your copied text, and simply add the instruction: "Summarize this into three clear bullet points." See how quickly you can get the core information!

Ask Copilot to compare two product specs and list the differences

6 steps · Read online

You've got two product pages open and need to decide which one fits your needs. Instead of reading both from top to bottom, give Copilot the key specs and let it list the differences for you. In minutes, you'll have a clear side-by-side summary—no PhD in spreadsheets required.

Before you start

- A **GitHub account** (free is fine)
- A **computer or phone** with a web browser
- Two product pages open in separate tabs (for example, two software pricing pages or two laptop spec sheets)
- Rough time needed: **5 minutes**

Step 1 — Open Copilot in your browser

Go to github.com/github-copilot and sign in with your GitHub account. Alternatively, you can find Copilot through the GitHub homepage by clicking the Copilot option in the navigation menu. You'll see a chat box ready for your question.

If you already use Copilot in another way (like in your IDE or on Windows), you can also access the web-based chat from the main GitHub interface.

Example: Type or paste: "Hello"

You'll know it worked when Copilot replies with something like "Hi there! How can I help you today?"

Step 2 — Share the first product's specs

Since Copilot can't automatically read external web pages, you'll need to copy the key specs from the product page yourself. Go to your first product page, select and copy the main specifications (like processor, RAM, storage, screen size, weight, and price), then paste them into the chat box.

Example: Paste: "Here are the specs for the first laptop: Processor: Intel Core i7-13700H, RAM: 16GB LPDDR5, Storage: 1TB SSD, Screen: 14" 2.8K OLED 120Hz, Weight: 1.4kg, Price: $2,199"

You'll know it worked when Copilot acknowledges the specs and confirms it's ready to compare.

Step 3 — Ask Copilot to format the specs

Now, tell Copilot how you want the specs organized. Type: "Format these specs as a clean bullet-point list."

Copilot will return a neat list of the specs you shared, making them easy to compare with the second product.

Example: Type: "Format these specs as a clean bullet-point list."

You'll know it worked when you see a clear list like: • Processor: Intel Core i7-13700H • RAM: 16GB LPDDR5 • Storage: 1TB SSD • Screen: 14" 2.8K OLED 120Hz • Weight: 1.4kg • Price: $2,199

Step 4 — Repeat for the second product

Go back to your browser and copy the key specs from the second product page. Return to Copilot and paste those specs into the chat.

Example: Paste: "Now for the second laptop: Processor: AMD Ryzen 9 7940HS, RAM: 32GB LPDDR5, Storage: 512GB SSD, Screen: 16" 2.5K IPS 165Hz, Weight: 1.8kg, Price: $1,799"

You'll know it worked when Copilot acknowledges the second set of specs.

Step 5 — Ask Copilot to compare the two lists

Now ask Copilot to compare the two sets of specs. Type: "Compare these two products and list the differences in a clear table or bullet list."

Copilot will take the two lists and highlight what's different—like RAM upgrades, storage options, screen refresh rates, or price changes.

Example: Type: "Compare these two products and list the differences in a clear table or bullet list."

You'll know it worked when you see a concise comparison like: • RAM: 16GB vs 32GB (second product has more) • Storage: 1TB vs 512GB (first product has more) • Price: $2,199 vs $1,799 (second product is cheaper) • Screen: 14" vs 16" (second product has a larger screen)

Step 6 — Ask for a recommendation (optional)

If you're still unsure, ask Copilot to give a simple recommendation based on your needs. Type: "Which of these two products is better for video editing on a budget?"

Copilot will weigh the specs against your stated use case and suggest the better option.

Example: Type: "Which of these two products is better for video editing on a budget?"

You'll know it worked when Copilot gives a concise recommendation with reasons.

Common mistakes

- **Pasting the wrong specs**: Double-check what you're copying from the product page. If you include the wrong specs, just paste the correct information and ask Copilot to start over.
- **Not including enough detail**: If you only paste partial specs, Copilot can't give you a full comparison. Include the key specs: processor, RAM, storage, screen, weight, and price.
- **Assuming Copilot remembers**: Copilot doesn't keep memory between chats. If you open a new chat, you'll need to paste both sets of specs again. Bookmark your chat or keep it open to avoid redoing steps.

What this means for you

- **In everyday life**: Next time you're choosing between two gadgets, appliances, or software tools, use Copilot to do the heavy lifting. Copy the specs, paste them in, and get a clear comparison in seconds—no more squinting at tiny font on a phone screen.
- **For your business or work**: When your team debates which tool to adopt, run this comparison in a shared chat. Everyone sees the same facts, and you avoid long email threads full of conflicting opinions.
- **If you're just getting started**: Start with two simple product pages you already have open. Try the steps above, and you'll see how easy it is to turn messy specs into a clean comparison.

Try it now

Open two product pages in your browser right now. Copy the first set of specs, go to github.com/github-copilot, paste them in, and ask Copilot to format them. Do the same for the second page, then ask it to compare them. You'll have your answer in under five minutes—no spreadsheets, no PhD required.

Summarize Long Research PDFs with Cohere's Clever AI

5 steps · Read online

Imagine turning a dense, fifty-page research paper or a lengthy report into a crisp, one-page summary in under ten seconds. Whether you're a student tackling academic journals or a professional keeping up with industry reports, this guide will show you how to use Cohere to extract key insights without the headache, using a simple online tool.

Before you start

- **A Cohere account:** You'll need a free developer account to access the tools. Sign up on Cohere's website.
- **A research PDF:** Have a PDF document ready on your computer that contains selectable text (not scanned images).
- **A web browser and computer:** You'll use the Cohere website on a desktop or laptop, as the interface might look different on a mobile device.
- **Time needed:** This guide should take about 10-15 minutes to follow.

Step 1 — Log in and find your API Key

Before we can ask the AI to read our document, we need a digital passkey to access its brain. This passkey is called an **API Key (Application Programming Interface Key)**, which is a unique password that securely connects your computer or web browser directly to the AI engine.

Once you log into your new Cohere dashboard, you'll usually see a menu on the left side of the screen. Look for a tab or section labeled "API Keys" or "Developers." Click on this to reveal your keys. Your active "Trial" API key will be displayed there; make sure to copy it. Keep this key private, just like you would with a bank password, to protect your account. You'll know it worked when you see a long string of letters and numbers (your unique key) next to "Trial API Key" or similar wording on the screen.

💬 Example: Your key will look like a long string of random letters and numbers, such as `a1B2c3D4e5...`

Step 2 — Open the Cohere Playground

Writing code can be intimidating if you're just starting out, but Cohere provides a **Playground** (an online sandbox or testing space) where you can experiment with AI settings using simple buttons instead of writing code.

From your Cohere dashboard, look for a prominent "Playground" button, which is often found at the top or in the sidebar. Clicking this will take you to a page where you can choose different

AI "templates." Select the "Summarize" option from the list of templates. After selecting, a new interface will load, typically showing a large text box on the left for your input and various setting sliders on the right. If you don't immediately see "Summarize," look for a section like "Templates" or "Use Cases." You'll know it worked when you see a main text area on the left and settings options on the right, with "Summarize" active or prominently displayed.

Step 3 — Extract and paste your PDF text

Currently, most AI tools work best when you feed them raw text directly. We need to copy the text out of our PDF and give it to the tool. Remember, the Cohere Summarize tool processes text you provide; it can't directly *open* your PDF file itself.

Open your research PDF on your computer using any PDF reader. Carefully highlight the main sections of the paper you want summarized – the introduction, key findings, and conclusion are usually the most important for a comprehensive overview – and copy them (using Ctrl+C on Windows or Command+C on Mac). Switch back to your Cohere Playground window. You'll see a large input box on the left, usually labeled "Input" or "Text to Summarize." Paste the text you copied from your PDF directly into this box. Don't worry if the formatting looks a bit messy after pasting; the AI is smart enough to read past odd spacing and line breaks. You'll know it worked when your text appears in the input box, ready for the AI to process.

💬 Example: Paste the text you copied from your PDF into the large input box on the left side of the Playground.

Step 4 — Customize your summary settings

On the right side of the screen, you'll see a panel of options. These sliders and dropdowns let you tell the AI exactly how you want your information presented, giving you control over the summary's output.

You'll typically find settings for "Length," "Format," and "Extractiveness."

1. **Length:** Choose between "Short," "Medium," or "Long." For a quick overview of a dense paper, "Short" is usually best.
2. **Format:** Select "Bullets" (perfect for quick scanning of key points) or "Paragraph" (better for a narrative, flowing summary).
3. **Extractiveness:** This slider controls how much the summary sticks to the exact wording of your original text versus rephrasing it in simpler, more general terms. Set it to "Low" if you want a plain-English translation that simplifies academic jargon, or "High" if you need the exact academic terms and phrases preserved.

If these options aren't immediately visible, look for a "Settings" or "Options" panel on the right side of the screen that you might need to expand. You'll know it worked when your chosen settings (e.g., "Short," "Bullets," "Low") are displayed next to their respective labels.

Step 5 — Generate your summary

Now it's time to let the AI do the heavy lifting and transform your long text into a concise summary.

Once you've adjusted your settings, scroll to the bottom of the right-hand panel or the Playground interface. Look for a button labeled "Generate" or "Run." Click this button to begin the AI summarisation process. Within a few seconds, a polished, easy-to-read summary will appear in a separate output window or section on the right side of the screen. If the "Generate" button is greyed out, double-check that you have pasted some text into the input area. You'll know it worked when a new section on the right-hand side of the screen populates with your AI-generated summary, which you can then copy and save for your study notes or work files.

💬 Example output:

> **Key Findings:**
>
> - The study tested 500 participants over twelve months.
> - Using AI tools reduced administrative tasks by 40%.
> - Participants reported higher satisfaction when using clear, visual interfaces.

Common mistakes

- **Pasting scanned images instead of text:** If your PDF is a scanned photocopy of an old book, you won't be able to highlight and copy the text.
    - **The fix:** Ensure you use "selectable" PDFs, where you can actually highlight individual words. If you only have a scanned image, you'll need to run the file through an Optical Character Recognition (OCR) tool first to convert the images into real text before copying.
- **Hitting the text limit with very long documents:** While Cohere can handle substantial text, free accounts often have limits on the total number of words (or "tokens") you can process in one go. If you paste an entire book, it might fail or only summarize a portion.
    - **The fix:** For extremely long documents, try breaking them into key sections (e.g., summarize chapter by chapter) or focus on only the most critical parts like the introduction, methods, results, and conclusion.
- **Leaving your API key visible:** Your API key is like a personal password; sharing it publicly could allow others to use your Cohere account, potentially incurring costs or

misusing the service.

    - **The fix:** Never share screenshots that show your API key, and never paste it into public forums or messages. If you accidentally share it, simply click "Regenerate Key" in your Cohere dashboard to disable the old one and get a new, secure key.

What this means for you

- **In everyday life:** You can use this method to quickly understand long product manuals, complicated rental agreements, or lengthy council planning reports before making big decisions, saving you hours of reading.
- **For your business or work:** You can easily digest competitor research, market studies, or complex industry regulations without spending hours reading every page, helping you stay informed and make quicker decisions.
- **If you're just getting started:** Try starting with a short news article (about 500-1000 words). Copying and pasting a single page of text into the Playground is the gentlest way to see how the AI structures its thoughts before tackling bigger documents.

Try it now

Open a browser tab, sign up for a free Cohere account, and copy just three paragraphs from any online news article. Paste it into the Cohere Summarize Playground, hit **Generate**, and watch how quickly it extracts the main point for you!

Use Claude to Pull Out the Key Points of a Long Research Paper

5 steps · Read online

> **In short:** Upload a research paper PDF to Claude and ask for a plain-English summary in five short sections. You'll need a Claude account and the PDF file; the whole thing takes about ten minutes.

Research papers are dense by design — they're meant to be read carefully and argued with, not skimmed on a Tuesday afternoon. Claude, the chat assistant from Anthropic, is unusually good at reading long documents and pulling out the bones of what's actually being said. This guide walks you through turning any academic PDF into a clean, structured summary you can actually use, whether you're a student trying to understand a paper for class, a journalist scanning sources, or a professional keeping up with research in your field.

Before you start

- A free Claude account at **claude.ai** (a paid plan unlocks more file uploads and longer documents — check Anthropic's site for current plan details).
- The research paper saved as a PDF on your computer or phone. Most journal articles have a free author copy on the author's university page or on repositories like PubMed Central.
- About 10 minutes.

Step 1 — Open Claude and start a fresh chat

Open **claude.ai** in your browser, or launch the Claude mobile app, and sign in. If you don't have an account yet, the sign-up button is right at the top of the page — you'll need an email address and a password. Once you're in, look for a button labeled **"New chat"** (or a **+ icon**) and click it. A blank conversation screen opens with an empty message box at the bottom. 💬 Example: N/A — no typing in this step. You'll know it worked when you see a clean, empty chat with no previous messages and a text box ready at the bottom of the screen.

Step 2 — Attach the PDF

Click the **paperclip icon** 📎 next to the message box. A file picker opens; browse to your PDF, select it, and confirm. On mobile, the same paperclip opens your phone's file or photos app — pick **"Files"** or **"Documents,"** not photos. If you can't spot a paperclip, look for a **+ plus icon** or a button that says **"Add file"** or **"Upload"** — Claude offers the same feature under slightly different labels depending on the device and platform. After you select the file, its filename appears as a small chip just above the text box. 💬 Example: N/A — no typing in this step. You'll

know it worked when the PDF's filename is visible as a label or chip above the message box, before you've typed anything.

Step 3 — Ask for a structured five-section summary

Click into the message box and type a clear, specific prompt asking for the summary in five labeled parts. Then send it — press Enter on a computer, or tap the upward arrow on mobile. You'll see a small "thinking" indicator for a few seconds while Claude reads the PDF, and then its reply appears. The reply should be longer than a chat message — usually several short paragraphs with headings or numbered points. 💬 Example: "Please summarize this paper in plain English using five short sections: (1) the main question the research tries to answer, (2) the method they used, including who or what they studied, (3) the key findings, (4) what the authors conclude, and (5) the limits of the study the authors mention. Keep each section to 2–4 sentences." You'll know it worked when Claude's reply breaks into labeled sections that match what you asked for — usually with bold headings, numbered items, or clear paragraph breaks.

Step 4 — Ask follow-up questions on the parts that confuse you

Stay in the same chat and type a question about something in the summary that didn't quite land — a method, a term, a number, anything. Press Enter. Claude replies with a focused answer because it still has the full PDF in memory and the previous summary right in front of it. This is where Claude really earns its keep: you don't have to re-upload or re-explain. If the first answer is still fuzzy, ask a sharper follow-up, like "Can you give me a concrete example?" or "What does that mean in everyday language?" 💬 Example: "In section 2 you mentioned 'a randomized controlled trial with 240 participants' — can you explain in plain words what that actually means they did, step by step?" You'll know it worked when Claude's reply addresses that one specific question and references details from the paper rather than answering in general terms.

Step 5 — Copy or export the summary for later

Hover over Claude's reply and look for a small **copy icon** (usually two overlapping squares) in the corner of the message — click it and the entire reply lands on your clipboard, ready to paste into a note, an email, or a Google Doc. If you don't see a copy icon, click and drag to highlight the text manually, then press Ctrl+C (or Cmd+C on a Mac). For the full conversation history, look for a **download or share** option in the chat's menu — Claude can export the whole thread as a text or PDF file. One honest note: Claude can summarize well, but it can miss nuances, misread a graph caption, or smooth over a contradiction the authors are actually arguing about, so use the summary as a map, not the territory. 💬 Example: N/A — no typing in this step. You'll know it worked when you can switch to another app and paste the summary cleanly, or when a downloaded file appears in your downloads folder.

Common mistakes

- **Asking for "a summary" without specifying the format.** Claude will give you something, but it will often be one wall of paragraph text. Always tell it the structure you want — sections, length, audience — because it's dramatically better at filling a template than guessing what you meant.
- **Forgetting to check that the PDF actually uploaded.** If you typed your prompt and sent it without seeing a file chip above the box, Claude is answering from its general knowledge, not from the paper. Scroll up, look for the filename, and re-attach if it's missing.
- **Trusting the summary as a substitute for the paper.** For anything important, spot-check the original section Claude points you to. The summary is a map; the paper is the territory, and the map can drop landmarks.

What this means for you

- **In everyday life:** When a news article cites a study and you want to know what it actually said — without paying $40 for the journal — drop the free PDF into Claude and get a five-section summary in two minutes, enough to decide whether it's worth your evening.
- **For your business or work:** If you review research, write literature reviews, or build products grounded in evidence, this turns an afternoon of skimming into a quick first pass. You still need to read the key papers yourself, but Claude handles the long tail of "is this worth my time?" decisions.
- **If you're just getting started:** Open Claude right now and pick the shortest paper you can find — a two-page report or even a one-paragraph abstract works for practice. Ask for the same five-section summary and see how Claude structures the answer.

Try it now

Find any open-access research paper (try Google Scholar with the PDF filter, or browse PubMed Central), upload it to Claude, and paste the example prompt from Step 3. You'll have a structured summary in under two minutes — and you'll know whether the paper deserves a closer read.

Use Qwen to Translate a Short Message for a Friend

5 steps · Read online

> **In short:** You can paste any short message into Qwen (Alibaba's AI assistant), name the language and tone you want, and get a ready-to-send reply in seconds. It's free to start, no setup beyond a quick sign-up, and works in your browser.

Got a message in another language — a sweet note from a relative overseas, a WhatsApp from a friend abroad, or a school update you can't quite read? Qwen can translate it for you, and even rewrite it in the tone you want: friendly, formal, or casual. Here's how to do it in five minutes, even if you've never used an AI tool before.

Before you start

- A free Qwen account (sign up with email, Google, or Apple — takes about a minute).
- A phone, tablet, or computer with a browser. Qwen works on all three.
- The message you want translated — copy it from your messaging app, email, or screenshot.
- About five minutes total. No payment needed for the basics, though some heavy users may eventually want a paid plan for faster replies and higher usage limits. Check the official Qwen site for current plan details.

Step 1 — Open Qwen and sign in

Open your browser and go to the Qwen website. On the homepage, look for a button labeled something like **Sign in** or **Log in**, usually at the top-right corner. Tap it, then choose how you want to sign in — email, Google account, or Apple ID all work. If you've never used Qwen before, you'll be asked to create a free account; the form just asks for an email and a password, then sends a quick verification code to your inbox.

If your screen looks different, don't worry. Some users land on a regional version of the site; the sign-in button may say **Get started** or **Try free** instead. Look for any button near the top of the page that invites you in. 💬 Example: sign in with your usual Gmail address. You'll know it worked when you see a chat window with a text box at the bottom and a short greeting from the assistant.

Step 2 — Paste the message you want translated

Click inside the text box at the bottom of the screen — it often shows placeholder text like **Message Qwen** or **Ask anything**, though your version may phrase it slightly differently. Tap or click inside the box, then paste your message. You can paste with the keyboard shortcut **Ctrl+V**

on a computer or **long-press → Paste** on a phone. Don't worry if the original message is messy or has emojis; Qwen reads through all of it.

If you only have a photo of the message (a screenshot, say), look for a small **paperclip icon** or **image icon** next to the text box. Tap it, choose the photo from your library, and Qwen will read the words in the image. Example: paste this short Spanish line — *"Hola María, ¿puedes venir a cenar el sábado? Te echo de menos."* You'll know it worked when your pasted text appears inside the chat box, ready for the next step.

Step 3 — Tell Qwen the language and tone you want

Now write a short instruction below your pasted message — in the same text box, just press **Enter** to add a new line, or send the message first and then ask for the translation. The instruction is called a **prompt** (think of it like a request slip you hand to a helpful assistant). Keep it simple: name the target language and the feeling you want.

If the message comes from a relative and you want it to sound warm, ask for a **friendly** tone. If it's for a teacher or business contact, ask for **formal**. If it's for a close friend, **casual** is fine. You can also add things like "keep it short" or "include the original text too" if that helps. Example: *"Please translate this to English in a friendly tone, as if I'm writing back to my aunt."* You'll know it worked when Qwen replies with a clear translation and a short note like "Here you go!" or "Sure, here's the translation:"

Step 4 — Check the translation and ask for tweaks if needed

Read Qwen's reply out loud in your head. Does it sound natural? If something feels off — maybe a phrase is too stiff, or the assistant kept a word that should sound softer — just ask for a change in the same chat. There's no need to start over. Type one more short sentence explaining what to adjust, and Qwen will rewrite the line for you.

If your screen looks different, the reply may appear as a longer message with the translation first and a short explanation second — that's normal. Some Qwen versions also show a small **copy icon** (two overlapping squares) in the corner of the reply; clicking it copies the translation so you can paste it straight into your messaging app. Example: *"Lovely, but can you make the second sentence a little warmer? She always signs off with love."* You'll know it worked when Qwen gives a fresh version that sounds closer to what you want, or clearly explains why it kept a certain phrase.

Step 5 — Send the translated message

Once the reply sounds right, copy it. On a computer, highlight the text and press **Ctrl+C** (or **Cmd+C** on a Mac). On a phone, long-press the translation and choose **Copy** from the menu

that pops up. Then switch to your messaging app, tap into the text field, and paste with **Ctrl+V** or long-press → **Paste**. Hit send, and you're done.

One honest note: Qwen can translate the words beautifully, but it can't read tone from a voice message or know inside jokes between you and your friend. If a line still feels awkward, trust your gut — rephrase it yourself, or add a smiley to soften it. 💬 Example: the final English reply — *"Hi María, can you come over for dinner on Saturday? I miss you."* You'll know it worked when your friend or family member replies in their language, smiling.

Common mistakes

- **Forgetting to specify the tone.** A literal translation can sound cold or robotic. Always add a word like *friendly*, *formal*, or *casual* — it changes the result more than you'd think.
- **Asking for a whole conversation at once.** Qwen handles short, single messages best. If you have a long thread, break it into one message at a time so the translation stays accurate.
- **Trusting names, numbers, or dates without checking.** Qwen can occasionally guess wrong on a name or misread a phone number from a photo. Always glance over those before hitting send.

What this means for you

- **In everyday life:** You can reply to a relative's message in their own language even if you only speak a few words. It's a small way to stay close across distance.
- **For your business or work:** If you chat with international clients or colleagues, Qwen can quickly turn a short business note into formal English, Spanish, French, or dozens of other languages — handy for a one-line email when you don't have time to draft from scratch.
- **If you're just getting started:** The easiest first move is exactly what you just did — pick one message you've been meaning to reply to, paste it in, and try a translation. Once you see how natural it feels, you'll find more uses on your own.

Try it now

Open Qwen, paste the last short message you received that you couldn't fully read, and ask: *"Translate this to English in a friendly tone."* You'll have your first translation in under a minute — and probably your second idea for how to use it before you've finished reading this.

How to use Gemini Deep Research to plan a big purchase decision

6 steps · Read online

> **In short:** Gemini Deep Research browses dozens of websites for you and returns a structured report with source links, so comparing big purchases like cars, appliances, or laptops stops being a full weekend of manual Googling.

Buying a new laptop, a refrigerator, or a used car usually means opening forty browser tabs and reading half of them. Gemini's **Deep Research** feature does that scrolling for you: it reads many sites, summarizes what it found, and hands you a tidy report with links back to where each claim came from. This guide walks you through using it, step by step, so the report actually answers *your* question.

Before you start

- **A Google account** signed in on the Gemini website or the Gemini app.
- **Access to Deep Research.** It may be limited to paid Google AI plans on some accounts, so check the plan page for the latest availability before relying on it.
- **A rough idea of what you're shopping for** (a category, a budget range, and 2–3 must-haves work best — more on this in Step 2).
- **About 15–25 minutes.** Deep Research usually takes a few minutes to run, and you'll want time to read the report it returns.
- An honest expectation: the report gives you a strong starting point, but it does **not** buy the thing for you, and prices or availability can be outdated by the time you read them.

Step 1 — Open Gemini and pick the Deep Research mode

Go to the Gemini website or open the Gemini app, then look in the chat box for a model or tool picker — usually a small dropdown above the text field. From that picker, choose the option labeled **"Deep Research."** On some accounts the option appears as a separate **Tools** button or a dedicated tab at the top of the chat. If you don't see it, it might not be enabled on your plan yet, or it could live behind a different name like **"Research"** or **"Pro."**

💬 Example: when the picker appears, click the dropdown that currently shows "Fast" or "Pro" and switch it to **Deep Research**.

You'll know it worked when the text field gains a small brief outline or "Plan" panel, and the placeholder text inside the box mentions research or sources.

Step 2 — Describe your purchase in concrete detail

This is the most important step: the more specific your prompt, the more useful the report. Type a short paragraph that names the **item, your budget range, your must-haves, your deal-breakers, and how you'll use it.** Think of it like a one-page brief you'd give a friend who's shopping on your behalf — vague briefs produce vague reports. Avoid jargon the AI won't understand; instead, describe the thing in plain words.

💬 Example: "I need a quiet, energy-efficient dishwasher for a small kitchen, budget roughly $700–$1,000 USD, must fit a 24-inch opening and run under 45 dB. I cook most nights for two, sometimes four. Compare the top 5 models available in my country, with pros, cons, and where each one tends to be cheapest."

You'll know it worked when you see a clear "Research plan" or outline appear below your message — Gemini usually shows what it intends to search for before it starts.

Step 3 — Approve the plan and let it run

Below the research plan, look for a button like **"Start Research"** or **"Begin."** Click it, then step away. Deep Research browses many websites in the background and writes as it goes, so the response can take **several minutes** (sometimes five, sometimes fifteen) — far longer than a normal chat reply. A progress bar or running indicator usually appears. **Do not close the tab** while it's running, or you'll lose the work.

💬 Example: hit the "Start Research" button, then make a coffee; the first draft of the report appears inline as it finishes each section.

You'll know it worked when a long, structured report lands in the chat, with section headings, bullet points, and small numbered or linked citations you can click.

Step 4 — Read the report and click at least three sources

Don't trust any AI research report blindly. Scroll to the end of the report first and look for a **Sources** section, or for inline footnote-style numbers next to specific claims. Click at least three of those links and read the original page for a minute each, especially anything related to price, warranty, or safety. The sources tell you whether the report is repeating marketing copy or actually weighing real reviews.

💬 Example: open the report, jump to the "Sources" list at the bottom, and click the first three links — if any of them are dead, irrelevant, or from a brand's own press release, treat the matching claim with caution.

You'll know it worked when you can point to a specific source link and say, out loud, "this is where that fact came from."

Step 5 — Ask two or three follow-up questions to narrow the field

Once the first report lands, treat it like a starting point, not the final answer. Ask Gemini to compare your top two candidates head-to-head, to look for known reliability complaints, or to translate the report into a simple "if you care most about X, choose Y; if you care about Z, choose W" decision guide. One question per message works best. If an answer feels too generic, push back: ask for specific model numbers, specific dates, or specific failure modes.

💬 Example: "Now compare models A and B from your report on noise, energy use, and reliability over five years — which one wins for a household of two who cooks daily?"

You'll know it worked when Gemini's reply is sharper and more specific than the original report, and ideally points back to the same source list.

Step 6 — Turn the report into a personal decision checklist

Before you close the chat, ask Gemini to summarize the entire research session into a short **decision checklist** you can actually use: 3–5 finalist products, the deciding criteria for each, and the questions you still need to answer before you spend money. Copy that checklist into a note on your phone or print it. The point of the report is to leave your head clear, not crammed with new jargon.

💬 Example: "Summarize this whole conversation into a checklist: top 3 dishwashers I should see in a store, what to test in person, and what to double-check on the manufacturer's site before I buy."

You'll know it worked when you have a one-page summary in your hand that you could read in two minutes in a showroom.

Common mistakes

- **Mistake:** Treating the report as a buy recommendation. **Fix:** It's a starting brief — the final choice is still yours, and any purchase should be confirmed against the official product page and the live price.
- **Mistake:** Skipping the sources. **Fix:** Always click at least three source links. If you can't find them, or they look unreliable, ask Gemini to redo the research with more trusted sites.
- **Mistake:** Closing the tab while research is running. **Fix:** Deep Research takes several minutes; leave the tab open until the full report appears, or you'll lose the work.
- **Mistake:** Asking a vague prompt like "best laptop under $1,000." **Fix:** Add your must-haves (size, weight, battery life, ports) and your main use (work, gaming, school) — the more detail, the better the report.

What this means for you

- **In everyday life:** A big purchase — a car, a fridge, a laptop — stops being a weekend of forty open tabs. You get a sourced brief instead, then spend your time on the decisions that actually need a human.
- **For your business or work:** The same workflow works for buying equipment, software subscriptions, or vendor comparisons at work: ask Deep Research to compare three vendors against your checklist, then read the sources yourself before signing anything.
- **If you're just getting started:** Pick one purchase you've been putting off, write a two-sentence brief in Step 2, and run the report end to end. The first run is the slowest; the second one feels easy.

Try it now

Open Gemini, switch to **Deep Research**, and paste this single sentence as your first prompt: *"I'm buying a [item] with a budget of [low]–[high], and my must-haves are [X, Y, Z]. Compare the top 5 options and end with a short decision checklist."* Fill in the brackets, hit Start Research, and read the sources before you trust any of it.

How to Ask ChatGPT to Explain Any AI News Headline in Plain English

6 steps · Read online

> **In short:** Paste any confusing AI headline into ChatGPT with a short instruction, and it usually returns a plain-English summary plus why the news matters to you. The trick is telling it exactly what you want — a short summary, no jargon, and a "why should I care" sentence.

Every week, AI companies release a flood of announcements packed with words like "multimodal," "fine-tuned," and "context window." Most headlines read like they were written for engineers, not for the rest of us. The good news: ChatGPT is unusually good at translating that jargon back into normal language — if you ask it the right way. Here's the simple routine.

Before you start

- A free or paid ChatGPT account (the free tier is enough for this task).
- Access to ChatGPT in a browser or in the app on your phone.
- One AI news headline that confused you — copied and ready to paste.
- About 5 minutes for the whole routine.

Step 1 — Pick one headline that's giving you trouble

Scroll your news feed, email, or a trusted news site until you find a single AI headline that made you stop and think, "wait, what does that actually mean?" Copy that exact headline to your clipboard, exactly as it was written. You are about to hand it to ChatGPT as raw material — the AI works best when it sees the original wording, not your summary of it. If nothing on your feed fits, any announcement from the past week works fine; you don't need breaking news. 💬 Example headline to try: *"Anthropic's Claude gets a 1M-token context window with extended thinking mode"* You'll know it worked when one specific headline sits in your clipboard, ready to paste.

Step 2 — Open ChatGPT and start a fresh chat

Open ChatGPT in your browser or in the app. Look for a button or option labeled something like **"New chat"** — it's usually near the top of the sidebar on desktop, or behind a small **compose** or + icon on mobile. Tap or click it. A blank chat window appears. Starting fresh matters here: an old conversation can confuse the AI, because it might try to answer in the style of whatever you were discussing before. A clean chat is a clean slate. 💬 Example: just open the app — no typing needed yet. You'll know it worked when you see an empty text box with a blinking cursor and no previous messages above it.

Step 3 — Paste the headline and ask for a plain-English summary

Click into the empty text box and paste your headline. Then, on a new line, type a short instruction telling ChatGPT exactly what you want. The simplest prompt that almost always works is: *"Explain this headline in plain English, like I'm not a tech expert. Keep it under 100 words."* Press **Enter** (or tap the send arrow). ChatGPT will start typing a reply within a few seconds. The reason this prompt works is that you told it three things at once: who the audience is (a non-expert), how long the answer should be (under 100 words), and what kind of answer you want (a plain summary, not a definition list). 💬 Example full prompt to paste:

> Anthropic's Claude gets a 1M-token context window with extended thinking mode
> Explain this headline in plain English, like I'm not a tech expert. Keep it under 100 words.

You'll know it worked when the reply uses everyday words — verbs like "remembers," "thinks longer," "reads at once" — instead of repeating the original jargon.

Step 4 — Add a second prompt: "Why does this matter to me?"

Most AI news only becomes useful when you know whether it actually changes your life. In the same chat (don't start a new one — ChatGPT remembers the conversation), click into the text box again and ask a follow-up. A reliable follow-up is: *"Now explain why this might matter to a regular person who isn't a developer or AI researcher. Give me 2-3 concrete examples."* Press **Enter**. ChatGPT will usually answer in a few short paragraphs with specific scenarios — like whether it changes how you write emails, search for information, or use your phone. If the answer still feels vague, ask one more time: *"Give me a real-world example I might notice this week."* 💬 Example follow-up to paste:

> Now explain why this might matter to a regular person who isn't a developer or AI researcher. Give me 2-3 concrete examples.

You'll know it worked when you can finish the sentence "This matters to me because…" without thinking.

Step 5 — Ask about any word you still don't understand

You will almost certainly hit one word in the reply that still sounds like gibberish — something like "token," "fine-tuning," or "multimodal." Don't skip past it. In the same chat, type: *"In one sentence, what does [word] mean? Use an analogy a 10-year-old would get."* ChatGPT will swap the technical term for a familiar picture, like comparing a token to a puzzle piece or a context window to a notebook page. This is the moment where a confusing headline finally

becomes something you actually understand — and you can repeat this step as many times as you need. 💬 Example follow-up to paste:

> In one sentence, what does "context window" mean? Use an analogy a 10-year-old would get.

You'll know it worked when you can explain the term to a friend in one breath, without looking anything up.

Step 6 — Save the chat so you can find it again

Once the explanation makes sense, keep it. In the ChatGPT sidebar, look for your chat's title — ChatGPT usually names it automatically after the first message. If the auto-title is something like *"New chat"*, click the title and rename it to something you'll recognize later, like *"Context window — what it actually means."* Now your explanation lives in your sidebar like a small personal glossary, ready the next time the same word pops up in a headline. Over time, this becomes your private AI dictionary, built one confusing headline at a time. 💬 Example new title: *"Claude context window — plain English"* You'll know it worked when you can scroll your sidebar a week from now, click the renamed chat, and instantly re-read the explanation.

Common mistakes

- **Pasting just a phrase instead of the full headline.** A snippet like "new reasoning model" gives ChatGPT nothing to explain. Always paste the whole headline as it was published.
- **Forgetting to specify the audience.** If you don't say "plain English" or "I'm not a tech expert," ChatGPT may answer for an engineer, which can be more confusing than the headline itself.
- **Starting a new chat for every follow-up question.** Each new chat makes ChatGPT forget what you already discussed. Keep all your questions in the same conversation so it can build on the explanation.
- **Trusting the explanation as the final word.** ChatGPT can misread a headline or fill in gaps that weren't there. If the answer sounds important, check the original source before acting on it — see Step 4 for the "real-world example" trick.
- **Skipping over a word you don't understand.** One unexplained term can quietly undo the whole summary. Use Step 5 until nothing in the reply feels like a foreign language.

What this means for you

- **In everyday life:** the next time a friend sends you an AI headline and asks "is this a big deal?", you'll have a 30-second way to find out instead of doom-scrolling for an hour.

- **For your business or work:** if you manage a team, marketing channel, or client communication, you can use this same routine to translate any technical announcement into a short brief your colleagues will actually read — and you can save the explanations as a shared reference.
- **If you're just getting started:** pick one headline from this week, paste it into ChatGPT, and ask for a plain-English summary. That's the whole skill — the rest of the routine grows from there.

Try it now

Open ChatGPT right now, paste one AI headline you saved recently, and send this exact prompt: *"Explain this headline in plain English, like I'm not a tech expert. Keep it under 100 words."* Read the reply once. If one word still feels unclear, ask ChatGPT to define it using a 10-year-old analogy. That's the routine — done in under two minutes.

Research and Draft Short Social Media Captions with Perplexity

6 steps · Read online

By the end of this guide, you can turn a fuzzy idea into a ready-to-post social media caption backed by real sources — whether you manage a personal account or post on behalf of a small business. This is for beginners who want to save time without sounding robotic.

Before you start

- **A free Perplexity account** — sign up at the Perplexity homepage with an email or social login. A free plan is enough for this guide; Perplexity Pro (paid) gives more daily searches and access to advanced AI models.
- **A device with a browser** — Perplexity works on desktop, mobile web, and the iOS/Android apps. Screens look similar across all three.
- **A topic in mind** — even something rough, like "sustainable coffee" or "weekend hiking near Nashville".
- **Rough total time: 15–20 minutes.**

> **One honest note:** Perplexity can research facts, summarize articles, and draft captions, but it can't post to social media for you — you'll still copy the caption into Instagram, LinkedIn, X, or your scheduler.

Step 1 — Open Perplexity and start a new thread

Open Perplexity in your browser or app, then sign in if prompted. Look for a **"New Thread"** button or a clean search box at the top of the screen — it usually has placeholder text like *"Ask anything..."* or *"What do you want to know?"*. Click or tap inside the box. The cursor blinks, ready for typing.

If your screen looks different, look for a **pencil icon**, a **"+" button**, or a **"Home"** tab that returns you to a blank search box. Each Perplexity conversation is called a "thread", so starting fresh keeps your research tidy.

You'll know it worked when the page reloads with an empty box and no previous answers above it.

Step 2 — Ask Perplexity to research trending angles for your topic

Type a clear, specific question into the search box. Beating your hands at the keyboard here won't help — give Perplexity a **topic**, an **audience**, and a **goal** in plain English. After you press Enter (or tap the arrow icon), Perplexity starts writing a short answer with **numbered**

footnotes — those tiny superscript numbers are links to the sources it used. Below the answer, you'll often see a list of **related questions** you can click for a deeper dig.

If the answer feels too generic, try the **"Focus"** dropdown (look for buttons labeled **Web**, **Social**, **Academic**, or **All**) — switching to **Social** tells Perplexity to lean on posts and trending discussions rather than formal articles.

Example: *"I'm writing a 3-sentence Instagram caption for a small American coffee roastery. What are 5 trending topics in sustainable coffee right now in the US, with sources from 2025 or 2026?"*

You'll know it worked when Perplexity returns an answer with at least 3 numbered footnote links you can tap to read the underlying articles.

Step 3 — Pick one angle and ask for supporting facts

Scroll through the answer and choose **one** angle that excites you — don't try to cover all five. Click the related question that matches your pick, or type a new follow-up question into the same box. A new answer appears below the first, still with its own numbered sources. The conversation history stays visible, so you can scroll up and compare.

If a footnote link looks interesting, click it — Perplexity opens the source in a new tab. Don't worry about reading every word; skim for a quote, a stat, or a phrase you might want to mention.

Example: *"Tell me more about the rise of compostable coffee pods in the US. Give me 3 specific statistics or quotes from American sources I could mention in a post."*

You'll know it worked when you have one sentence or stat from a source you'd be happy to credit in your caption.

Step 4 — Switch into "drafting mode" with a clear instruction

Now you've done the research — time to write. Stay in the same Perplexity thread and type a new prompt. This time, you're telling Perplexity to **write**, not research. Be specific about the **platform**, **length**, **tone**, and what facts to include. Perplexity drafts a caption right inside the chat.

If the tone feels off, try adding words like "warm", "playful", "matter-of-fact", or "professional" to your instruction. Perplexity works best when you tell it exactly how the caption should sound.

Example: *"Write a 3-sentence Instagram caption for our compostable coffee pods. Tone: warm and friendly, not salesy. Mention that 1 in 3 people now look for compostable packaging,*

according to a 2025 survey. End with a question to invite comments."

You'll know it worked when Perplexity produces a short caption (3 sentences or so) that includes your chosen fact and ends with a question or call-to-action.

Step 5 — Refine the draft in the same conversation

Read the caption out loud — does it sound like you? If a sentence feels clunky, type a quick refinement request right below the draft. Perplexity updates the answer in place, and the old version scrolls up so you can compare. This back-and-forth is normal and usually takes 2–3 quick rounds.

If Perplexity suggests adding emojis and you'd rather skip them, just say so. The same goes for hashtags — you can ask it to suggest 3–5 relevant ones or to leave them out entirely.

💬 Example: *"Make sentence 2 shorter and a bit punchier. Keep the question at the end. Suggest 5 relevant hashtags."*

You'll know it worked when the caption sounds natural when you read it aloud, includes your key fact, and fits the platform's vibe.

Step 6 — Save your sources and copy the caption

Before you leave Perplexity, capture the **footnote links** from Steps 2 and 3 — you may want to credit a source in your post, or come back to read more later. Look for the small numbered superscripts next to sentences; click them to copy the URL from your browser's address bar.

Then copy the caption itself. Look for a small **copy icon** (two overlapping squares) near the draft, or simply highlight the text with your mouse or finger and copy it. Paste it straight into Instagram, LinkedIn, X, or your scheduling tool.

If Perplexity's Pro plan would help you here, note that the **free plan has a daily cap on Pro searches**, but regular searches are unlimited for most users — you can complete this whole guide on the free tier.

You'll know it worked when you have the final caption ready to paste, plus 2–3 source links saved in case you want to reference them later.

Common mistakes

- **Mistake:** Asking one giant question like *"Write me a caption about coffee."* **Fix:** Always include the **topic**, the **audience**, the **platform**, and the **length** — specificity saves you a round-trip.

- **Mistake:** Trusting Perplexity's footnotes without checking one or two. **Fix:** Click at least one source to confirm the stat or quote actually says what Perplexity implied. AI can misread a source.
- **Mistake:** Forgetting to add a human touch. **Fix:** After copying the caption, change one small thing — a word, an emoji, your own opener — so it sounds like *you*, not a template.

What this means for you

- **In everyday life:** Posting about a hobby, a trip, or a family milestone becomes a 10-minute task instead of a 30-minute staring-at-the-screen session.
- **For your business or work:** You can keep a steady posting rhythm without paying an agency, and every caption can be backed by a real source — handy if you ever get asked "where did you get that number?"
- **If you're just getting started:** Treat the first caption as practice, not a masterpiece. The second one will already feel easier.

Try it now

Open Perplexity, type this exact prompt, and watch what comes back:

> *"I'm writing a 2-sentence LinkedIn post for a small accounting firm in Portland. Find 1 recent statistic about AI tools being used by American small businesses in 2025 or 2026, with a source. Then draft the post in a confident but friendly tone, ending with a question."*

You should have a researched, source-backed draft within 2 minutes.

Compare Products with Perplexity Before You Buy

6 steps · Read online

Want to make sure you're getting the best value or the right features before spending your hard-earned cash? This guide will show you how to use Perplexity, an AI-powered search engine, to quickly compare products, understand their differences, and even check user opinions, all in one place. By the end, you'll be able to compare anything from new phones to kitchen appliances, helping you shop smarter and more confidently.

> **In short:** Perplexity lets you compare products by asking direct questions, then shows you AI-generated summaries with source links — useful for research, though the free version has usage limits.

Before you start

- An internet-connected device (computer, tablet, or phone).
- Access to the Perplexity website (perplexity.ai) or the Perplexity app.
- Perplexity offers both free and paid tiers — the free tier has daily search limits, while paid plans (Pro, Enterprise) offer higher usage. This guide works with either, but you may hit limits on the free version.
- The whole process should take about 5-7 minutes.

Step 1 — Open Perplexity

Start by opening Perplexity in your web browser or launching the app on your device. When you first arrive, you'll see a clean interface designed for asking questions. The main feature is usually a large text box in the center, often labeled "Ask anything..." or "Search here." This is your gateway to getting information. If you're on a mobile device, the layout might be slightly more compact, but the central search bar will still be prominent. You'll know it worked when you see the main Perplexity screen with a clear input box, ready for your query.

Step 2 — Enter your comparison query

Now it's time to tell Perplexity what you want to compare. Click or tap inside the main search box and type a clear question asking for a comparison. Think about what specific products you're interested in and what aspects you want to know about. After typing, you might see a "Search" or "Ask" button appear, or you can simply press the 'Enter' key on your keyboard or the 'Go' button on your mobile device's keyboard. If it looks different, some versions might have a microphone icon if you prefer to speak your query. You'll know it worked when the AI starts processing your request, often showing a spinning icon or a "Generating answer" message.

Example: `Compare the Google Pixel 8 vs Samsung Galaxy S24 on camera quality and battery life.`

Step 3 — Review the AI's summary

Once Perplexity has finished processing your request, it will present a summary of the comparison. This usually appears as a concise, easy-to-read answer, often using bullet points, short paragraphs, or even a simple table to highlight key differences and similarities between the products. Below the summary, you'll typically see numbered source links – these are the websites and articles Perplexity used to gather its information. If it looks different, the exact layout can vary slightly, but always look for a clear answer section at the top of the page. You'll know it worked when you see a well-organized summary directly comparing the products you asked about.

Step 4 — Explore the sources

To ensure the information is reliable and to dig deeper into specific details, it's a good idea to check Perplexity's sources. Each piece of information in the summary often has a small, numbered link next to it. Click on these numbers, or find the "Sources" section usually located below the main answer. What happens next is that clicking a source link will open the original webpage in a new tab in your browser. If it looks different, sometimes the sources are just listed with their titles, and you click the title rather than a number. You'll know it worked when a new browser tab opens to an external website, letting you see where the AI got its information from, like a review site or a manufacturer's product page.

Step 5 — Ask follow-up questions

Perplexity isn't a one-shot tool; you can continue the conversation to get more detailed insights. Below the initial answer, you may find a section with suggested "Related Questions" or a new input box where you can type your own follow-up. This allows you to ask for clarification, or to explore other aspects of the products you're comparing. For example, if the initial summary didn't cover specific features you care about, now's your chance to ask. If it looks different, you might need to scroll down to find a new prompt area or simply type directly into the main chat box again. You'll know it worked when Perplexity provides a fresh answer, refining its previous information based on your new question. 💬 Example: `Which one is generally considered better for low-light photography according to user reviews?`

Step 6 — Refine your query for specific details

Sometimes, you need to go beyond a general comparison and focus on very specific features or scenarios. You can start a completely new query (by clicking "New Thread" or clearing the

search bar if available) or continue in your current conversation. For this step, try asking Perplexity a highly specific question about a particular product feature or how it performs in certain conditions. This is where Perplexity shines for deep dives, pulling targeted information from various sources. If it looks different, some versions might automatically clear the search bar for a new query, while others require you to click a "Start a new query" button. You'll know it worked when Perplexity gives you a highly focused answer, directly addressing your specific detail. 💬 Example: `What are the pros and cons of an OLED TV screen for watching sports, compared to a QLED screen?`

Common mistakes

- **Being too vague with your query:** If you just type "compare phones," Perplexity won't know what you're looking for.
    - **The fix:** Be specific! Name the products and the aspects you care about (e.g., "Compare iPhone 15 vs Samsung Galaxy S24 for camera and battery").
- **Not checking the sources:** Perplexity is excellent, but like any AI, it draws from various sources, and some might be outdated or biased.
    - **The fix:** Always glance at the numbered sources. If you see a review site you trust, or the manufacturer's official page, that's a good sign. If something seems off, check another source.
- **Expecting live pricing or stock information:** Perplexity excels at product specifications, features, and reviews, but it's not a live shopping tool.
    - **The fix:** Use Perplexity for product research, then check specific retailer websites for the latest prices, deals, and stock availability.
- **Hitting free tier limits:** If you're on the free version, you may reach daily search limits quickly.
    - **The fix:** If you need unlimited comparisons, consider upgrading to a paid plan. Otherwise, space out your queries or focus on the most important comparisons.

What this means for you

- **In everyday life:** Imagine you're buying a new fridge or a pair of running shoes. Instead of trawling through countless review sites, you can ask Perplexity to summarize the key differences, pros, and cons of your top choices. This saves you time and helps you feel more confident about your purchase.
- **For your business or work:** If you're deciding between two different project management software tools, or researching the best laptop for your team, Perplexity can provide a quick breakdown of features, pricing tiers (if publicly available), and user feedback, helping you make an informed decision for your budget and needs.

- **If you're just getting started:** Try comparing two items you know well, like different brands of coffee machines or popular car models, focusing on just one or two aspects you care about most.

Try it now

Open Perplexity and ask it to compare two household items you might consider buying, like "Compare features of the Dyson V15 Detect vs Samsung Bespoke Jet for pet hair removal." See how quickly you get a concise answer!

How to Compare Phone Plans and Pick the Right One Without the Jargon

6 steps · Read online

> **In short:** Pick a plan by looking at your real monthly data use, the all-in cost (not just the advertised price), and the coverage in your area — most people overpay because they pick by brand instead of fit.

Picking a phone plan can feel like decoding a foreign language. Carrier websites throw around words like "unlimited," "deprioritization," "throttling," and "MVNO" without explaining any of them. This guide walks you through a simple way to compare plans on your own terms: what you actually use, what you actually pay, and what you actually need.

Before you start

- Your last 2-3 phone bills (paper, email, or in your carrier's app)
- About 30-45 minutes of quiet time
- A notebook, spreadsheet, or a notes app
- A friend or family member on a different plan (optional, for coverage feedback)
- Roughly 7 days to think it over — no need to decide in one sitting

Step 1 — Look at your real monthly usage

Open your last 2-3 phone bills — paper, email, or the carrier's app. Find the "data used" line. Data means the internet your phone uses over the cell network, not your home Wi-Fi. If your bill just says "unlimited," check your phone's settings: most phones show usage under **Settings > Network & internet > Data usage**, and you'll see a number like "8.2 GB used this month." If you can't find it, menus change between phone brands — look for "Cellular" or "Mobile data" instead. No number at all? A rough rule: light users (email, maps, some video) use 5-10 GB; heavy users (lots of streaming, video calls) use 20+ GB. 💬 Example: "I used 12 GB last month — mostly maps, Spotify, and one Zoom call a day." You'll know it worked when you have one number written down that you trust.

Step 2 — Decide which features actually matter to you

Make a short list of must-haves and nice-to-haves. Must-haves are things you'd cancel the plan over if missing. Common features to consider: **hotspot** (using your phone as a Wi-Fi source for a laptop — some plans limit this to 5-10 GB even on "unlimited" tiers), international calling or roaming, talk and text (most plans include unlimited, but double-check), family or multi-line discounts, and 5G access. Don't trust the word "unlimited" without reading the fine print — many "unlimited" plans slow your speeds after 20-50 GB, a practice called **throttling**, or push

you behind other users in crowded areas, called **deprioritization**. 💬 Example list: "Must-haves: at least 15 GB, hotspot, 5G. Nice-to-haves: international roaming." You'll know it worked when you have a clear top-3 list of features you actually care about.

Step 3 — Learn the four main plan types

There are roughly four kinds of phone plans, and knowing which one you're looking at saves real money. **Postpaid** = you pay at the end of the month, usually through a big carrier, often with a phone installment plan attached. **Prepaid** = you pay upfront for 30 days, no contract. **MVNO (Mobile Virtual Network Operator)** = smaller brands (think Mint, Consumer Cellular, Cricket) that rent coverage from the big carriers (Verizon, AT&T, T-Mobile) and resell it cheaper. **Family plan** = multiple lines on one account, usually with a per-line discount. The catch: MVNOs are usually much cheaper but may run slightly slower in busy areas because the host carrier prioritizes its own customers first. 💬 Example: "I'll compare a Verizon postpaid plan, an AT&T-based MVNO, and a T-Mobile prepaid plan." You'll know it worked when you can tell which type a plan is just by reading its website.

Step 4 — Calculate the real total monthly cost

The advertised price is almost never what you actually pay. Add up: the plan's base price, taxes (often 10-20% extra, varies by region), any **device installment** (paying off a phone in monthly chunks — say $25 for a phone on a 24-month plan), and any "regulatory" or "administrative" fees. Then check for auto-pay or paperless billing discounts — many carriers knock off a small amount if you let them charge your card automatically. Also watch for **introductory pricing** that jumps up after 12 or 24 months. The real question isn't "what does it cost now" but "what will it cost in year two?" A $30 plan that becomes $70 after a year is the more expensive plan. 💬 Example: "$40 plan + $6 taxes + $25 phone = $71/month total. After promo ends: $95/month." You'll know it worked when you can see — in one number — what this plan will cost you a year from now.

Step 5 — Check coverage where you actually live and travel

A cheap plan is worthless if it doesn't work in your apartment, your office, or your parents' house. Coverage maps on carrier websites are optimistic — they're the best case. For a real check, ask friends and neighbors on different carriers how their service is at your home address. If you commute, ask about the route. If you travel, look for the carrier's international options or roaming add-ons. Some MVNOs have slightly weaker indoor signal in rural areas, even though they use the same towers as the host carrier. If you live in a city, almost any plan will work fine. If you live in a small town or rural area, this is the most important step in the whole guide. 💬 Example: "I asked my neighbor: 'How's your T-Mobile signal in our building?' He said: 'Two

bars, calls drop.'" You'll know it worked when you can name two people whose coverage matches your daily life.

Step 6 — Compare 2-3 plans side by side

Pick your top three plans from the previous steps. Make a simple table — paper, notes app, or spreadsheet — with these columns: plan name, monthly cost in year two (including everything), data allowance, your must-have features (yes/no), and a "gut feeling" line. The gut-feeling column matters: a plan that ticks every box but has awful customer service is a plan you'll hate in six months. Read recent reviews, but focus on patterns, not single angry comments. Once your table is filled in, the right plan usually jumps out: it hits your must-haves, stays affordable in year two, and has good coverage where you live. If two plans tie, pick the one with the friendlier cancellation policy — flexibility is worth real money. 💬 Example row: "Plan A: $55/mo, 25 GB, hotspot yes, support fair. Plan B: $45/mo, 15 GB, no hotspot, support good. Pick A — hotspot is a must-have." You'll know it worked when you can point to one plan and say, "That's the one — and here's why."

Common mistakes

- **Picking by brand loyalty instead of fit.** The cheapest "unlimited" plan is rarely the cheapest plan for you specifically.
- **Forgetting the year-two price.** Promotional rates that double after 12-24 months are the most common cause of bill shock.
- **Trusting "unlimited" without reading the fine print.** Most "unlimited" plans slow down your speeds or deprioritize your traffic after a set amount of data.

What this means for you

- **In everyday life:** a $20-40 monthly difference over two years is $500-1,000 — real money for a vacation, an emergency, or a small home repair.
- **For your work or business:** if your job depends on reliable calls or hotspot for a laptop, the cheapest plan is rarely the right one — downtime costs more than the savings.
- **If you're just getting started:** pick one plan from a major carrier or well-known MVNO with no contract, then reassess in 3 months once you know your real usage. No decision has to be permanent.

Try it now

Pull up your last phone bill right now. Find the data usage number. Write it down somewhere you can find again. That's the single most important number in choosing a plan — and you just got it in under a minute.

How to ask Alexa+ to translate a short phrase on your Echo

5 steps · Read online

> **In short:** Your Echo can translate short spoken phrases into many languages, hands-free, as long as you say the wake word first. Some languages and accents work better than others, and Alexa+ cannot reliably translate long conversations or read written text from a photo.

If you've ever frozen at a bakery counter abroad or waved at a relative who doesn't share your language, the Echo on your kitchen counter can quietly fix that moment. You don't need an app, a phone or a dictionary — just your voice, the wake word, and a short, clear sentence. This guide walks you through a translation workflow that goes beyond the basic one-off request: setting the target language once, asking for the slow spelling when a word matters, and recovering when Alexa+ gets it wrong.

Before you start

- An Echo device (any generation that supports Alexa+ — the newer the device, the smoother the experience, but older Echos still work for most translations).
- An Amazon account already linked to your Echo, and a working internet connection (translations run in the cloud, not on the device).
- The wake word turned on — by default it's "Alexa," but if you've changed it, use whatever you chose.
- A rough idea of the target language and how its name sounds in English (for example, "Spanish," "Japanese," "Brazilian Portuguese"). Alexa+ understands English language names; it does not usually need the exact locale code.
- About 5 minutes for the first run, then it gets faster the more you practice.

Step 1 — Say the wake word and a simple "translate" request

Stand within a few feet of the Echo so the microphone can pick you up cleanly — the seven-mic array on most Echos hears across a room, but a quiet kitchen beats a loud living room. Say the wake word, then a short sentence that names the source and target language. The light ring on top turns blue, then a spinning cyan or white ring while Alexa+ thinks, then it speaks the translation out loud.

Example: "Alexa, how do I say 'Where is the bathroom?' in French?"

If your Echo doesn't light up at all, the wake word may be off or the mic button on top is pressed in (red ring). Look for a small microphone or mic-off icon on top and tap it once to unmute. If you renamed the wake word to "Echo" or "Amazon," swap that in for "Alexa" in every example

below. You'll know it worked when the ring spins, you hear a short chime, and Alexa+ speaks the translated phrase clearly.

Step 2 — Set a target language so you don't repeat yourself

If you know you'll be asking for several phrases in the same language, you can pin it for the conversation. Say a short setup line first, then each phrase after that only needs the wake word. The light ring pulses briefly to acknowledge the mode, and every following translation goes into that language until you change it or stop talking for a minute or two.

💬 Example: "Alexa, speak Spanish for the next five minutes." Then: "Alexa, how do I say 'the check, please'?"

If "speak X for the next Y minutes" doesn't work on your device, try the simpler form: "Alexa, translate everything I say into Italian." It does the same thing in plain words. Some older Echos may not support timed modes — if nothing changes, fall back to including the target language in every request, like Step 1. You'll know it worked when your second request comes back in the same language without you naming it again.

Step 3 — Ask for the slow version when pronunciation matters

A spoken translation is only useful if a real person can understand you. When a phrase contains a tricky word — a city name, a food, a medication — follow up and ask Alexa+ to say it more slowly, or to spell it out letter by letter. The Echo will repeat the phrase at a slower pace, or say each letter with pauses, which you can copy down or repeat after.

💬 Example: "Alexa, can you say that more slowly?" — or — "Alexa, spell that in French."

If the slow version is still too fast, ask "even slower" or "one more time, slowly" — most Echoes will keep slowing down. If the spelling uses letters your language doesn't share (like the Spanish ñ or French accents), ask "how is that spelled with accents?" and Alexa+ usually spells it out phonetically. You'll know it worked when the second version is noticeably slower than the first, or when you hear each letter named individually.

Step 4 Translate what someone else just said to you

This is the move most people miss: you can have Alexa+ listen to a foreign-language phrase and bring it back in English (or another language you know). Hold the conversation close to the Echo and ask for a translation after the speaker finishes. The light ring stays blue while it's "listening" past the wake word, so you don't need to say the wake word again until Alexa+ replies.

💬 Example: "Alexa, translate that into English." — after someone has spoken in another language near the device.

If Alexa+ says it didn't catch anything, the room was probably too noisy, or the speaker paused too long and the Echo thought its turn was over. Try again, and this time say "Alexa, listen to this and translate it" before the other person speaks, so the listening window stays open. You'll know it worked when Alexa+ repeats the foreign phrase back to confirm what it heard, then gives you the English version.

Step 5 — Save the translation in your Alexa app for later

Sometimes the phrase matters — a hotel address, an allergy, a relative's name. Open the Alexa app on your phone, tap the **Activity** tab (look for a clock or list icon along the bottom), find the translation in the recent list, and tap the three dots or a share icon to send it to yourself by message, email, or notes. This way you don't have to remember it from memory or write it down on a napkin.

💬 Example: after a translation you want to keep, open the app and search the activity feed for the first word of the phrase.

If the Activity tab is empty or you can't find the entry, the translation may not have been logged — not every Echo logs every interaction. As a backup, just ask: "Alexa, what did you just say?" and it will repeat the translation, which you can then type into your phone. You'll know it worked when the phrase appears in the app's activity list with a timestamp, or when you can copy it from a message you sent yourself.

Common mistakes

- **Speaking too fast or mumbling.** The Echo's translation is only as good as what it heard. Slow down, pause between the wake word and the phrase, and face the device. The fix: treat it like talking to someone across a quiet table, not someone across a noisy street.
- **Asking for a language Alexa+ doesn't support well.** Most major languages — Spanish, French, German, Italian, Portuguese, Japanese, Mandarin, Hindi, Arabic, plus regional variants like Brazilian Portuguese and Mexican Spanish — translate cleanly. Smaller languages or dialects may come back patchy or romanized (written in Latin letters instead of the original script). The fix: check Amazon's current list of supported languages before you rely on it for a trip.
- **Trying to translate a full conversation.** Alexa+ handles short phrases; if you speak for 30 seconds straight, it often misses the second half. The fix: break the conversation into one- or two-sentence chunks, and let Alexa+ respond between them.

What this means for you

- **In everyday life:** the next time a neighbor, a delivery driver or a family member speaks a different language at the door, you can get a polite phrase in seconds without pulling out your phone or asking them to write it down.
- **For your travel or work:** before a trip, sit with the Echo for two minutes and learn ten core phrases — greetings, please and thank you, yes and no, "how much," "where is," and "I don't understand." The repetition is a built-in flashcard drill.
- **If you want to go further:** connect the Alexa app to a note-taking tool like a shared phone note or a simple spreadsheet, and build a personal phrasebook from your activity feed over a week — the phrases you'll actually use, in the order you'll need them.

Try it now

Pick the one phrase you'd most like to have in another language this week — "how much is this?", "thank you very much," or your name — and say it to your Echo right now. Ask for the slow version, then save it in the Alexa app. You'll have your first saved, rehearsed translation in under two minutes.

How to Compare Two Job Offers Side by Side Using AI Safely

5 steps · Read online

> **In short:** You can use any AI assistant — ChatGPT, Gemini, Claude, Copilot, Le Chat or DeepSeek — to build a side-by-side comparison of two job offers, as long as you scrub personal info first. The AI builds the table; you keep control of your privacy.

Two job offers can look almost identical on paper until you line them up next to each other. This guide shows you how to ask any AI assistant to build a comparison table from the offer details you already have, and how to keep sensitive information out of the chat. It is for anyone weighing two offers and feeling a little lost in the numbers.

Before you start

- A free account with at least one AI assistant (ChatGPT, Gemini, Claude, Copilot, Le Chat, DeepSeek — any one works).
- Both job offers in some written form: letters, emails, or your own notes from a phone call.
- About 15 minutes of quiet time, plus a blank note on your phone or computer where you will rewrite the offer facts by hand.

Step 1 — Rewrite the offer facts by hand, in your own notes

Open a fresh blank note on your phone or computer — the Notes app, Google Docs, Notion, or even paper — and copy the facts of each offer into it by hand. Type out only what would go on a comparison table: job title, base pay, bonus structure, location, working hours, vacation days, health benefits, start date. Leave out your name, home address, employee ID, references, signatures, and any greeting lines.

What you see: a plain page filling with short, neutral lines. What happens next: you now have an anonymized (stripped of any detail that could identify you) summary of each offer — clean enough to share, but stripped of identifying details.

If it looks different: any plain text editor works, even your phone's notes app. The point is a fresh page, not your offer email.

💬 Example: Offer A — "Title: Senior Marketer; Base: $X; Bonus: up to 10%; Location: 3 days office, 2 remote; Vacation: 18 days; Health: full coverage; Start: 2 weeks." Use the same shape for Offer B.

You'll know it worked when a stranger could read your note and not learn your name, your address, or anything beyond the jobs themselves.

Step 2 — Read your note once and remove anything sensitive

Before opening any AI assistant, read your note with fresh eyes and ask: is there anything here that could identify me, or that I would not want sitting on a third-party server (a company computer that stores your chat so the AI can process it)? Remove anything that answers yes.

What you see: a short checklist of risky items. What happens next: you strip your note down to the safest possible version before a single character is sent to the AI.

If it looks different: even assistants with strong privacy settings benefit from this habit. Building the muscle now means you will not accidentally share something sensitive the next time you paste text into a chat.

💬 Example: never paste your full name plus employer combo, home address, phone number, government ID numbers (Social Security, passport), bank details, names of current coworkers or your manager, references' contact info, NDA-protected clauses from your current employer, or your past salary history.

You'll know it worked when the only information left in your note is the kind you would feel comfortable writing on a public whiteboard.

Step 3 — Open the AI assistant and paste a clear prompt

Open the assistant of your choice in a fresh chat — never inside an existing conversation, because older messages could leak in. Paste your cleaned-up note, then ask the AI to build a side-by-side comparison table.

What you see: a chat box, usually empty or with a short greeting. What happens next: the assistant returns a formatted table with one column per offer and one row per category — pay, benefits, time off, and so on.

If it looks different: every assistant has its own layout. Look for a text box labeled something like "Message", "Ask anything" or "Send a message" — yours may use slightly different words, but the idea is the same: a blank area where you type or paste.

💬 Example: "Here are the facts of two job offers I'm comparing. Please build a clear side-by-side table with one column per offer and rows for: base salary, bonus, vacation, health benefits, location, working hours, and start date. At the bottom, add a short Pros and Cons summary for each." Then paste your note.

You'll know it worked when a structured table appears on screen with both offers in their own columns and the categories you asked for in the rows.

Step 4 — Add the things the AI cannot measure for you

The AI's table is great for numbers, but it cannot feel your gut, meet your future manager, or drive the commute at 5 p.m. Open your note again and add three rows the AI cannot fill: how the role feels, how the team treated you, and how daily life would actually work.

What you see: your note grows by a few lines. What happens next: you now have the AI's clean data side-by-side with your lived sense of each offer — the part no chatbot can replace.

If it looks different: if you only have written offers and never visited either workplace, write what you do know. Reading company reviews, looking at the office on a map, or checking the commute at peak hour all count as "your sense" of the offer.

Example: add three rows — "Team vibe — Offer A felt warm in the interview; Offer B felt formal"; "Commute — Offer A 25 min by train; Offer B 55 min by car at peak"; "Growth — Offer A clearer path to senior; Offer B broader role but less clear ladder".

You'll know it worked when your final comparison has both the hard numbers from the AI and the human notes only you can write.

Step 5 — Decide your next move and keep your privacy

Read the full table and ask which offer you keep coming back to. Then close the chat: most assistants let you delete a conversation from the side panel.

What you see: a delete or "remove" option in the chat menu. What happens next: the conversation is gone from your account; your offer details live only in your notes app, where you control them.

If it looks different: some assistants keep a short-term copy after deletion, and paid plans may save chat history by default. Look under Privacy or Data controls for a toggle like "Save conversations" or "Improve the model" — turning it off stops future chats being used to train the AI.

Example: in your assistant's settings, search for a toggle named something like "Save conversations" or "Improve the model for everyone" — switch it off if you want your offer details to stay private.

You'll know it worked when you have a clear favorite between the two offers, your note is the only place the details live, and the chat is either deleted or set to stay private.

Common mistakes

- **Pasting the whole email without cleaning it first.** Greetings, signatures, and "cc" lists can include names, phone numbers, or links back to your identity. Strip the note down before the AI ever sees it — see Step 2.
- **Trusting the AI to make the decision for you.** A table is a great mirror, but it cannot choose for you. If the AI's "Pros and Cons" feels off, ignore it and re-read your human notes from Step 4.
- **Forgetting the things no spreadsheet can measure.** Vacation days and salary are easy to compare; team vibe, manager style, commute, and growth path are not. If those rows are missing from your final table, the comparison is incomplete.

What this means for you

- **In everyday life:** the same habit — stripping personal info before pasting it into any chatbot — protects you whether you are comparing jobs, asking about medical symptoms, or sharing a tricky email you want help rewriting.
- **For your work or business:** if you ever help someone else compare offers (a partner, a friend, a mentee), walk them through the scrub-the-note-first step. It is the single most useful privacy habit to teach.
- **If you are just getting started with AI:** pick one task you already do by hand — comparing offers, drafting an email, summarizing a long article — and try it once with a free assistant. The skill is repetition, not genius.

Try it now

Open your notes app, find one offer letter or email, and rewrite its facts into a fresh note with names and address removed. That single act is the whole habit — and it takes under two minutes.

Use Grok to Explain a Confusing News Article in Plain English

5 steps · Read online

> **In short:** Paste a confusing news article into Grok with a short, specific request, and it returns a plain-English summary of who is involved, what happened, and why it matters. Paywalled articles usually have to be copied and pasted manually, since Grok often can't read them from a URL alone.

We've all clicked a headline that sounded important, then bounced off after two paragraphs of jargon. This guide walks you through turning one of those articles into something you can actually understand — using Grok, the AI assistant from xAI. By the end, you'll have a five-minute routine you can repeat whenever a story makes your eyes glaze over.

Before you start

- **An X account** (the same one you'd use to post on X / Twitter). It's free to create one if you don't have it.
- **A web browser** or the **X app** on your phone — Grok lives inside both.
- **One news article** you want explained. Have it open in another tab or app before you begin.
- **About 5 minutes.** That's the whole routine.
- A free X account works for basic chats. Heavier or longer use may need a paid plan; check Grok's official page for the latest details.

Step 1 — Open Grok and start a fresh chat

Open a browser and go to **grok.com**, or open the X app and find the **Grok** option. If you're not signed in, Grok will ask for your X account. Once inside, look for a text box near the bottom of the screen, usually labeled something like "Ask Grok" or "Message Grok." Tap inside it to start a new conversation. If older chats are showing, look for a **pencil icon** or a **+ button** to begin fresh; this keeps your new article separate from earlier context. If you're on the X mobile app, the Grok icon usually sits in the middle of the bottom navigation bar — tap it and the chat window slides up. Example: type a quick "hi" first, just to confirm everything is working, before pasting anything. You'll know it worked when you see an empty chat box with a blinking cursor, ready for your question.

Step 2 — Copy the article text

On the page with the article you want explained, select the text. On a computer, click at the start and drag your mouse to the end; on a phone, long-press the first word and drag the highlight bar to the last word. Then choose **"Copy"** — usually from a right-click menu on desktop, or a small

pop-up menu on mobile. The words now live on your device's clipboard, ready to paste anywhere. If copying doesn't work on some sites, try your browser's **Reader Mode** or print the page to a PDF, then copy from there. If the article is paywalled, you can usually only copy the visible preview — but that's still better than pasting only the link, because Grok often can't read paywalled pages from a URL alone. Example: copy the headline, byline, and the first three to five paragraphs — usually enough for a clear summary. You'll know it worked when you paste the text into any other app (like Notes) and the article words appear there.

Step 3 — Paste the article and ask for a plain-English summary

Click into Grok's text box and paste your copied text — **Ctrl+V** or **Cmd+V** on a computer, **long-press then "Paste"** on a phone. Right below the pasted text, add a clear instruction asking for a plain-English explanation. Hit **Enter** or tap the **send arrow** (usually a small upward icon next to the box). Grok starts typing a reply within a few seconds — you'll see words appear one by one or in chunks, finishing with a complete summary. If the reply is too short or misses the point, send a follow-up message in the same chat with a more specific prompt. The interface varies slightly between web and mobile, but the basic paste-and-send flow is the same everywhere. Example: paste the article, then on a new line type: *"Explain this like I'm smart but not an expert. Who is involved, what actually happened, and why does it matter?"* You'll know it worked when Grok replies with a short, organized summary written in everyday words.

Step 4 — Ask follow-up questions about the parts you still don't understand

Read Grok's reply. If anything is still fuzzy — a person you don't recognize, a phrase you don't know, or a number that doesn't quite make sense — just type another question in the same chat. Grok remembers the conversation, so you don't need to paste the article again. Watch for the new reply appearing just below your question, usually within seconds. If Grok doesn't have enough information to answer, it will usually say so honestly or ask you to clarify. If the chat window looks frozen, scroll down — sometimes the reply sits below where you're looking. This is where the conversation really starts to feel useful. Example: *"Who is [Name] you mentioned? Why are they important in this story?"* You'll know it worked when Grok answers your specific question in plain English, building on the earlier summary.

Step 5 — Ask for the "why does this matter" angle

This is the step most people skip — and it's the most useful. Ask Grok to connect the dots to your everyday life: why should a regular person care, what might happen next, or what to watch for. Type a short request and send it. Grok will reply with a brief take. Treat it as a starting point, not the final word — news is complex, and Grok's reading is just one view among many. If you want a different angle, simply ask — for example, how a small business owner might read the same

story, or how it could affect renters in a city. The free plan covers basic chats; very long sessions may eventually hit a limit, and exact limits change, so check Grok's official page for current details. 💬 Example: *"In two or three sentences, why should an ordinary person care about this story, and what's the most likely next development?"* You'll know it worked when you can explain the article to a friend in your own words — that's the real test.

Common mistakes

- **Pasting only a paywalled URL.** Grok can sometimes read public web pages from a link, but paywalled articles often block that — only a short preview loads, or nothing does. The fix: copy the visible text and paste it directly into the chat, even if it is only the first few paragraphs.
- **Asking "explain this" with no context.** A bare "what does this mean?" usually gets a vague reply. The fix: tell Grok who you are and what you want — for example, *"I'm a regular reader, not an expert — please give me the short version and the context."*
- **Trusting the summary without checking key facts.** Grok can sometimes mix up details (this is called a **hallucination** — when an AI confidently states something that isn't actually true). The fix: for anything that affects a real decision, double-check the names, numbers, and dates against the original article or another reliable source.

What this means for you

- **In everyday life:** the next time a confusing headline crosses your feed — about a new law, a court ruling, a market shift — you have a five-minute routine to actually understand it instead of scrolling past.
- **For your work:** when a colleague shares a dense industry article or a long report, you can pull out the key points quickly and decide whether the full read is worth your time.
- **If you're just getting started:** try this once with one short article before you need it. Build the habit on something low-stakes, so the routine feels natural the next time something genuinely confusing shows up.

Try it now

Pick one headline from today's news that made you feel a little lost. Copy the first three to five paragraphs, paste them into Grok, and type: *"Explain this like I'm smart but not an expert. Who is involved, what actually happened, and why should I care?"* Two minutes, start to finish.

Explain a tricky school topic with Grok in five minutes

5 steps · Read online

> **In short:** Open Grok, name the topic and your child's age, and ask it to retell the subject as a short story. You'll have a simple, kid-friendly explanation in about a minute, ready to read together at the kitchen table.

Homework time can go from calm to chaos in about ten minutes flat — usually the moment a tricky word appears in the textbook and nobody in the room can explain it. Grok is an AI assistant made by xAI, and one of the things it does well is take a confusing subject and rewrite it as a story a kid can actually follow. You don't need to be a teacher. You just need to know the topic, the age, and a good opening message. One honest note up front: Grok can rewrite the topic as a story, but it can't replace reading the textbook or doing the actual homework — your child still needs to know what the teacher is asking for.

Before you start

- A Grok account — you can sign up at the official Grok site or inside the X app. The basic chat is usually available on the free plan; some newer or more powerful models may need a paid tier, so check the current plan list if you want the strongest answer.
- A phone, tablet or computer with Grok open in a browser or app.
- The exact homework topic from your child's book, worksheet or class note — even a rough idea works.
- About **5–7 minutes** total.

Step 1 — Open Grok and start a new chat

Open the Grok app on your phone, or visit grok.com in a browser on your computer. Look for the **chat box near the bottom of the screen** — usually a rounded text field labeled something like "Ask anything" or "Message Grok." If you've used Grok before, your old chats appear in a sidebar on the web or behind a menu icon on mobile. Tap the **"New chat"** button or the **pencil / "+" icon** to start fresh, so the assistant doesn't get confused by an earlier conversation. 💬 Example: tap the "+" or pencil icon at the top of the chat list to open a blank chat. If your screen looks different, the new-chat option is often hidden behind a **☰ menu** at the top-left, or inside an account/profile icon at the top-right — open it and look for a "New chat" entry. You'll know it worked when you see a blank chat screen with a single text field and no old messages above it.

Step 2 — Name the topic, the age and the goal

Type one clear sentence into the chat box that includes three things: the subject (for example, "the water cycle" or "the French Revolution"), your child's age or grade, and what you want —

usually "as a short story" or "in simple words." Then press the **paper-plane icon ➤** or hit **Return / Enter** to send. Grok starts typing a reply within a few seconds; you'll see words appear one by one, the way a person would text. 💬 Example: *"My son is 10. Please explain the water cycle as a short story he can read on his own — about 150 words, no big words."* If your screen shows a **microphone icon** instead of a text box, tap the small **keyboard icon** to switch to typing first. Some mobile builds put the new-chat option behind the **☰ menu** — open it and choose "New chat" if you skipped Step 1. You'll know it worked when Grok starts writing a response — usually a title line, then a paragraph or two of story text.

Step 3 — Ask for one small example or scene

Stories land better when there's a concrete moment the child can picture in their head. After the first reply, send a short follow-up asking Grok to add **one tiny example, analogy or scene** — something a kid can see (a leaf, a king, a battery, a fish). This makes the idea stick. Type it the same way: into the text box, then tap the send arrow. Grok rewrites or extends the explanation in plain English. 💬 Example: *"Can you add one short scene with a leaf falling off a tree, so the 'evaporation' part makes sense?"* If the reply feels too long or too grown-up, just say so — *"A bit shorter, please, and use simpler words."* The assistant adjusts in the next message. If the text field won't send, check that you didn't accidentally tap a microphone, and that the message bubble actually appears above the chat box. You'll know it worked when the new reply contains a concrete image your child can repeat back to you — "a leaf falls into a puddle and floats up as steam."

Step 4 — Read it together and ask one question

Sit with your child and read the story out loud — slowly is fine. Pause at the end and ask **one** question to check the idea landed. You can type it into Grok first to see if the assistant agrees with the answer, or ask your child directly and then check together. This is where homework stops being a fight: your child is retelling, not being quizzed. 💬 Example: ask your child *"So what does the sun do to the water in the puddle?"* — then type *"Did I explain the water cycle correctly to a 10-year-old?"* into Grok for a sanity check. If Grok's answer comes back with new vocabulary your child doesn't know, ask it to **swap the hard word for a friendlier one** — *"Replace the word 'precipitation' with something a 10-year-old would actually say."* If your child loses interest halfway, ask Grok to **shrink the story** to about 80 words and try again. You'll know it worked when your child can summarize the topic in their own words without staring at the textbook.

Step 5 — Save the story for next time

If the topic is going to come up again — most school subjects do — keep the chat. Look for a **"Share" or "Copy" icon** near Grok's reply, often a small arrow pointing up or two overlapping

squares. Tap it to copy the text into a notes app, an email or a message to yourself. That way, tomorrow's homework — or next week's review — starts with a story your child already half-remembers. 💬 Example: tap the share icon, then paste the story into Apple Notes, Google Keep or a simple text message to yourself. If you don't see a share button, **long-press the reply text on mobile** (or **right-click on a computer**) and choose **"Copy"** — the text goes to your clipboard, ready to paste anywhere. Inside the X app, the chat lives within X itself, so saving usually means copying the story out before you close the tab. You'll know it worked when the story is safely in your notes app, ready to open at the kitchen table tomorrow.

Common mistakes

- **Asking for "an explanation" without naming the age.** Grok will write for an adult, and the child will tune out. Always include the age in the very first message.
- **Treating the first reply as the final one.** A great explanation usually takes two or three short follow-ups. If it feels off, ask Grok to try again in simpler words — don't give up after one round.
- **Quizzing instead of retelling.** Asking "did you understand?" puts a child on the spot. Asking "tell it back to me in your own words" is much kinder and shows whether the story actually landed.

What this means for you

- **In everyday life:** the next time a homework page makes everyone in the room sigh, you can turn it into a five-minute story instead of a thirty-minute argument.
- **For your business or work:** the same trick — *"explain X like I'm 10"* — is one of the fastest ways to break down a dense report, a contract clause or a long email you don't fully understand.
- **If you're just getting started:** start with one subject your child already half-likes — a favorite animal, a sport, a video game world — and ask Grok to explain it in story form. It's the easiest first win.

Try it now

Pick one topic from tonight's homework, open Grok, and paste this: *"My child is 10. Please explain [TOPIC] as a 150-word story with no big words. End with one tiny example."* Read the result together and see how it lands.