

Everyday AI: Staying safe

9 step-by-step guides for real, everyday tasks

Staying safe

AI World HQ · <https://aiworldhq.com>

Free edition · CC BY-NC 4.0 · free to share with credit

Before you start — what this is, honestly

This collection is free, and every guide in it is also free on our website. This file just gathers them in one place, in reading order, so you can print it, keep it on your phone, or read it offline. Our website has no paywall either — every guide links back to its online page, which is always the newer version.

AI Every guide in this pack was written by AI and checked by automated fact and quality gates. **No human editor reviewed the text.** Every guide page on <https://aiworldhq.com> carries the same notice — we say it openly because you deserve to know who wrote what you are reading.

Apps change. Buttons move and menus get renamed. If a step does not match what you see, the guide's online version is the newer one — the link is under every title.

Free to share. This collection is licensed under Creative Commons Attribution-NonCommercial 4.0 (CC BY-NC 4.0): you may copy, print and share it — with your class, your library or your team — as long as you credit AI World HQ with a link to aiworldhq.com and do not sell it.

What's inside

1. Staying safe

1. How to check if a suspicious message is a scam by pasting it into an AI
2. How to Spot an AI Scam Call, Fake Email, or Deepfake Voice Today
3. Turn Off AI Chat History and Data Saving in ChatGPT, Claude, and Gemini
4. What You Should Never Share With an AI Assistant: A Privacy Guide
5. How to use AI safely with your personal information
6. How to Spot a Phishing Email or Scam Text Before You Click
7. How to Spot AI Scams, Deepfakes, and Fake Content
8. How to Spot AI-Generated Scam Messages Before You Click or Reply
9. How to Spot a Deepfake Video or Voice Clone Before You Share It

Staying safe

How to check if a suspicious message is a scam by pasting it into an AI

6 steps · [Read online](#)

In short: Paste a sketchy message into any AI assistant — ChatGPT, Gemini, Claude, Copilot, or Le Chat — and ask it to flag the red flags. You'll get a second opinion in under a minute, but never paste real passwords, and always verify with the company first.

Scam messages have gotten good. They use real company logos, your real name, sometimes even your real address. Before you click anything, copy the message into an AI assistant and ask for a second pair of eyes. This guide walks you through that in six steps — even if you've never used an AI assistant before. One honest warning up front: the AI is a useful first check, not a verdict, and it cannot actually click the link or call the company for you. You still have to do that part.

Before you start

- An AI assistant you can sign into — ChatGPT, Gemini, Claude, Microsoft Copilot, and Le Chat all work, and their free plans are fine for this.
- The suspicious message itself, ready to copy (an email, a text, a chat screenshot, or a DM).
- About 5 minutes in total.
- Important: never paste real passwords, full card numbers, or government ID numbers, even if the message contains them — Step 2 shows you how to handle that.

Step 1 — Open a fresh chat in any AI assistant

Open the website or app of any AI assistant — ChatGPT, Gemini, Claude, Microsoft Copilot, Le Chat, or whichever you already use. Start a brand-new conversation rather than dropping the message into an older thread. On a phone, the app icon usually looks like a small chat bubble or a starry logo. On a computer, just go to the assistant's homepage and click into the empty chat box. A clean thread means the AI won't get distracted by older context and will focus entirely on your message. If you don't already have an account, the free sign-up usually takes under a minute and only needs an email address. 💡 Example: open ChatGPT, Gemini, Claude, Microsoft Copilot, or Le Chat — whichever one you already have an account on. You'll know it worked when you see an empty chat box with a placeholder like "Message..." or "Ask anything" and a send button (often an arrow or paper-plane icon).

Step 2 — Redact personal details before you paste

Before you copy the message, take 30 seconds to hide any sensitive bits it contains. Cover real passwords, full credit card numbers, your Social Security or national ID number, your home address, and any one-time login codes. A quick search-and-replace in any notes app works fine. Think of it like covering your hand on a debit-card PIN pad at the store — the AI doesn't need those details to spot a scam, and fewer real numbers in the chat mean less risk if your account is ever hacked. If the message has no personal details at all, you can skip ahead to Step 3. 🗨️ Example: turn "my SSN is 123-45-6789, please call 555-0123" into "my SSN is [REDACTED], please call [REDACTED]". You'll know it worked when the message still reads naturally but no real numbers, codes, or addresses remain.

Step 3 — Paste the message and ask the AI to flag red flags

Tap or click inside the empty chat box and paste the redacted message. Then add a short prompt asking the assistant to look at it as if it were a security expert. Press the send button (usually an arrow ↗ or paper-plane icon, but apps vary). After a few seconds the AI will start typing a reply — usually a short list of warning signs plus a one-line verdict like "this looks like a phishing attempt" (a fake message designed to steal your login or money) or "this appears legitimate, but verify the sender". If nothing appears after a few seconds, look for a small "stop" or square icon — the response sometimes needs a moment to begin. Apps change often, so if your screen looks a little different, just look for any button that submits your message. 🗨️ Example: "Here's a message I just received. Can you act as a security expert and tell me whether it looks like a scam? List the specific red flags you see, and tell me whether I should click any links in it. Message: [paste here]" You'll know it worked when you see a written response that names specific warning signs, such as urgency, suspicious links, an odd sender, or requests for personal info.

Step 4 — Read the AI's answer with a critical eye

Treat the AI's reply like a friend's opinion, not a doctor's diagnosis. Read each red flag it names and ask yourself: does this actually match the message I received? Many scams lean on urgency ("act in the next 24 hours"), impersonate a brand you trust, include a link that almost looks right, ask you to click or log in, or demand payment in gift cards or cryptocurrency. If the AI flags something that isn't actually in your message, ignore that point — assistants sometimes guess wrong, and you are still the expert on your own inbox. Conversely, if something feels off to you but the AI didn't mention it, trust your gut and move to the next step. Be especially wary if the AI says "this looks fine" but the message asks you to click a link, send money, or share a code — that's the moment to get a second human opinion. 🗨️ Example: if the AI says "the sender domain doesn't match the company", check that yourself — does the "from" address really end in the company's real website, or in something like "support-service.example.com"? You'll know it

worked when you can repeat back, in your own words, two or three specific reasons the AI gave for its verdict.

Step 5 — Verify any links or sender names yourself

Don't click anything yet. If the message contains a link, hover over it on a computer or long-press it on a phone — you'll see the real destination. Real companies use their own short domain (the brand's actual homepage), while scam links often hide behind words like "secure-login" or "verify-account" on a totally different address. The simple rule: if the link doesn't end on the brand's real site, it's almost certainly fake. If the message claims to be from a person you know, contact that person through a channel you already trust — a phone call or a separate chat — and never reply to the suspicious message itself. If it claims to be from a company like your bank, look up the official number on their real website (type the address yourself, don't use any number inside the message) and ask them directly. 💡 Example: a "USPS" tracking link points to `usps-tracking.example.com/pay-now` — that is NOT `usps.com`, so it's a fake. You'll know it worked when you have either confirmed the message is genuine through a trusted channel, or you've concluded it's a scam.

Step 6 — Decide what to do, and don't act in a rush

Now choose one of three actions. If it's a scam: delete it, block the sender, and — if it pretended to be a company — report it to that company through their real website. If it's genuine but unexpected: still don't click links inside it; go to the company's real site yourself and log in there. If you're not sure: leave it alone for an hour, then re-read it cold; urgency is the scammer's favorite tool, and a slow second look usually reveals more. Tell a friend or family member if you're worried — a second human pair of eyes still catches things AI misses, especially with voice or video scams. The whole point of using an AI here is to slow down, not to rush. 💡 Example: text a family member — "Hey, I just got a weird message about my Amazon account, does this sound legit to you? [paste]" You'll know it worked when you take one concrete action in the next few minutes — delete, report, or verify through a trusted channel — instead of leaving the message sitting unread in your inbox.

Common mistakes

- **Pasting real passwords or full card numbers.** The AI doesn't need them, and storing them in a chat creates risk if your account is ever hacked. Fix: redact first (Step 2), or describe the message in your own words instead of pasting it.
- **Trusting the AI's verdict without checking.** AI assistants sometimes miss obvious scams or flag clean messages by mistake. Fix: always cross-check the specific red flags it names against the actual message (Step 4) and verify with the real company (Step 5).

- **Clicking the suspicious link "just to see what happens".** That's exactly what the scammer wants. Fix: hover or long-press to read the real URL, and never log in through a link inside an unexpected message — go to the company's website yourself.

What this means for you

- **In everyday life:** the next time a "your package couldn't be delivered" text or a "your account will be locked" email lands, you can get a clear second opinion in under a minute instead of guessing whether to tap it.
- **For your business or work:** if you handle invoices or vendor payments, this same 5-minute check on unusual payment-change requests can stop a business email compromise (a scam where a criminal impersonates a vendor or boss to redirect a payment) — one of the most expensive scams going.
- **If you're just getting started:** pick the AI assistant you already use, bookmark this guide, and try it the next time something feels even slightly off. Practice on a junk email first if you want to build confidence.

Try it now

Find one message in your inbox you've been unsure about, redact any personal details, paste it into your AI assistant, and ask it to flag the red flags. Five minutes, one fresh chat, one safer decision.

How to Spot an AI Scam Call, Fake Email, or Deepfake Voice Today

6 steps · [Read online](#)

In short: AI scams now sound like your family, write like a real person, and spoof your bank's phone number. A five-second pause plus one private family word catches almost all of them — here's how to set both up today.

By the end of this guide you'll know how to recognize the three most common AI scams circulating right now, what one-second check stops each one, and what to do the moment you realize something is wrong. This is written for anyone who uses a phone and email — no tech background needed, and nothing here requires buying or installing anything.

Before you start

- A phone (mobile or landline) and an email account you check regularly.
- About 10 minutes total, mostly to set up one protection you can use for years.
- One or two trusted people — a family member or close friend — for the family safe-word step.

Step 1 — Listen for the giveaway in a "family emergency" voice call

When someone calls claiming to be a relative in trouble, the very first move is to **hang up** and dial that person's number yourself — the one already saved in your phone. What you'll usually see on screen is a familiar name on the caller ID, sometimes with a stranger's voice asking "guess who this is?" to make you fill in the blank yourself. After you hang up and redial, a real relative will answer or call you back; a scam line often disconnects, plays a generic voicemail, or doesn't ring at all. If your caller ID shows nothing or a random number, the same rule applies — never continue the call. Voice cloning (an AI that copies someone's voice from a few seconds of audio) is good enough now that even adult children of the "caller" can't tell the difference in the first few seconds. 🗨️ Example: "Grandma, it's me, I was in an accident, please send money" → hang up → dial your grandchild's saved number. You'll know it worked when the real person answers your second call and confirms they never phoned you.

Step 2 — Catch a fake customer-service agent

Any time someone calls YOU about a problem with your account — bank, phone, delivery, internet — hang up and call the company back using the number printed on your last statement or shown on the company's official site. What you'll usually see in the original call is an "agent" who already knows your name, the last four digits of your card, or recent transactions; the voice

sounds professional and there's often office noise in the background to feel legit. A real company will let you hang up and call back; a scammer will push you to "stay on the line" or refuse to give you a number you can verify. If the caller insists they can't take a callback because the case is "time-sensitive," that's actually the biggest giveaway. 🧠 Example: "Hi, this is Maria from your bank's fraud department" → hang up → look up the bank's number on your statement → dial it yourself. You'll know it worked when the real company's recorded greeting answers and the human you reach has no record of the previous call.

Step 3 — Spot the AI-written phishing email

Before clicking any link in a message that asks you to act, **hover over the sender's address and any link** to see where they actually point. What you'll see is an email that looks like it came from a delivery company, a streaming service, or your bank — the wording is fluent, friendly, and free of the typos that old scams used to have; there's often a small urgency like "act within 24 hours" to nudge you past thinking. Hovering reveals a different domain — something like `usps-tracking.example.com/pay-now` instead of a real postal address, or a long random-looking string of letters. On a phone you can long-press a link to preview it; if you can't tell where it goes, don't tap. AI-written phishing is convincing because the grammar is now perfect — the giveaway is no longer the spelling, it's the destination. 🧠 Example: "Your package is held at customs, pay a small fee to release it" → hover the link → see it points to a domain you've never heard of → delete. You'll know it worked when you can name the actual destination of the link before deciding whether to click.

Step 4 — Use the one-second check that works for all three

When any message or call feels urgent, scary, or too lucky to be true, pause for five seconds and ask yourself one question: **"Would I be in trouble if I did nothing for the next ten minutes?"** What you'll notice in those five seconds is your own hands wanting to act fast — and the caller or email pushing you to "do it now." Most scams collapse under a ten-minute pause: the "agent" hangs up and moves on, the email can sit unread overnight, and your "grandchild" is still safe wherever they actually are. Some scams add a fake countdown timer to fight this exact habit; remember that a real emergency from a real company will still be a real emergency ten minutes from now, and a real loved one will understand a short wait. 🧠 Example: "Pay this fee now or your account closes in five minutes" → pause → ask the ten-minute question → almost always the answer is no. You'll know it worked when the urgency melts away and you make your next move from calm instead of panic.

Step 5 — Set up a family safe word today

Pick one unusual word with one or two people you trust, and agree that any "emergency" call must include this word before any money moves. A short group chat, a text, or even an in-

person conversation is enough — no app or password manager needed. What this gives you is a private test that even a perfect voice clone can't pass: the scammer simply doesn't know your word, and your real relative does. If your family is scattered, agree on a private question with a private answer instead — "what did we name the cat at grandma's house?" works better than a single word when kids are involved. Write the word down once, share it with the people who need it, and don't post it anywhere online. 🗨️ Example: Text your sister: "If anyone calls saying they're me and needs money, ask them for our safe word. The word is platypus." You'll know it worked when at least two people in your circle can recite the word without looking it up.

Step 6 — Know what to do in the moment you realize it's a scam

Hang up, close the email, or stop replying — and then **report it**, even if you didn't lose a cent. Your heart will still be beating fast, the "agent" will be gone, and the email will still be sitting in your inbox; that's normal. Reporting helps block the next person, and in the US you can forward phishing emails to reportphishing@apwg.org and file a report at the FTC's reportfraud.ftc.gov, or with the FBI's Internet Crime Complaint Center at ic3.gov. If you DID send money or share a one-time code, call your bank's fraud line immediately — the number on the back of your card — and ask for a recall; speed matters more than embarrassment here. 🗨️ Example: Realized mid-call that the "bank agent" was fake → hung up → called the number on my debit card → told them what happened → account flagged for monitoring. You'll know it worked when the call is over, you've told one official channel, and your accounts and passwords are unchanged.

Common mistakes

- **Trusting caller ID.** Scammers can spoof any name or number. The display showing "Daughter" or "Your Bank" proves nothing — hang up and dial back on a number you saved yourself.
- **Acting on the "ten-minute countdown."** That timer is part of the scam. Real companies and real family emergencies don't disappear while you take a breath; pausing is always safe.
- **Sharing a one-time code "to verify yourself."** No legitimate company will ever ask for a code they just sent you. The instant someone asks for that string of numbers, the conversation is a scam — full stop.
- **Skipping the family safe word because "it'll never happen to us."** The setup takes one minute today and saves potentially life-changing money later. Treat it like a smoke detector: boring to install, priceless when it goes off.

What this means for you

- **In everyday life:** the next time your phone rings with a panicked voice, you'll know to hang up first and call back on a number you trust — instead of sending money in the first thirty seconds.
- **For your business or work:** any team that handles invoices, supplier emails, or wire transfers should adopt a private code word the same way families do; deepfake voice scams have already hit finance and HR departments, not just grandmothers.
- **If you're just getting started:** do just one thing today — text one person in your family the safe word from Step 5. That single message is more protection than any antivirus app on your phone.

Try it now

Right now, open your messaging app and text one trusted person: *"If I ever call you in a panic and ask for money or codes, ask me for our safe word. The word is [pick something unusual, like 'platypus' or 'silver-spoon']."* That one text takes under two minutes and is the single highest-value defense in this whole guide.

Turn Off AI Chat History and Data Saving in ChatGPT, Claude, and Gemini

6 steps · [Read online](#)

In short: You can stop ChatGPT, Claude, and Gemini from using your conversations to train future AI models by flipping a privacy toggle in each service's settings — but the exact menu names change often, and turning off training does not erase the conversations those companies have already saved.

By default, ChatGPT, Claude, and Gemini save your chats and may use them to train future versions of themselves. This guide shows you how to switch that off in each service — on web and mobile — and explains what each toggle actually changes, and what it quietly does not.

Before you start

- A free or paid **personal** account on each AI service you want to adjust. Work or school accounts are usually managed by your employer and may not let you change these settings.
- A web browser is the easiest place to do this. Phone apps work too, but menus shift around more often.
- Set aside about **15 minutes total** — roughly five minutes per service.
- Make sure you are signed in to each service before you begin, so the settings actually save.

Step 1 — Open ChatGPT's Data Controls

Go to **chatgpt.com** in a browser and sign in if you are not already. Click your **profile picture or initials** in the top-right corner — a small menu appears. Look for an option like **Settings** (sometimes shown as a gear icon ) and click it. In the Settings panel that opens on the left side of the page, look for a section called **Data Controls** and click that. If your screen looks different, try a section called **Privacy** instead — older versions used that name.

If you are on a phone, tap the **two lines or three dots** in the top corner of the ChatGPT app, then tap your name at the top of the menu that slides up. The same Settings → Data Controls path should appear.

You'll know it worked when you see a list of toggles — look for ones mentioning "Improve", "training", or "history".

Step 2 — Turn off ChatGPT's training toggle

On the Data Controls page, find the toggle labeled something like **"Improve model for everyone"** or **"Chat history & training"** — the exact label changes often, so read each

toggle's description. Click the toggle so it slides to the **off** position (it usually turns gray or loses its color). You may see a small popup asking you to confirm — click **Confirm** or **Save**. The page may refresh, or you may need to close and reopen it for the change to show.

If you do not see a training toggle at all, your account may be on a plan or in a region where the option is hidden. In that case, leave the page and check back in a few weeks — settings are added over time.

You'll know it worked when the toggle stays gray (off) after a page refresh and no longer shows an "on" indicator. Even after turning this off, OpenAI may still keep your most recent chats for up to 30 days for safety monitoring — so avoid pasting anything truly sensitive in the meantime.

Step 3 — Open Claude's privacy settings

Go to **claude.ai** and sign in. Click your **profile picture or initials** — usually the bottom-left on web or the top-right on mobile. From the menu that pops up, choose **Settings**. Inside Settings, look for a tab or section called **Privacy**, **Account**, or **Profile** — the exact name varies. Within that section, hunt for a toggle that mentions "**help improve Claude**" or "**use my chats for training**".

On the Claude mobile app, the same path usually works: tap your **profile icon**, then **Settings**, then look for the privacy section. If your app opens straight into a chat, look for a **gear icon**  in the top corner instead.

You'll know it worked when you see a privacy section with at least one on/off toggle related to training or improving the assistant.

Step 4 — Turn off Claude's training toggle

Click the toggle next to "**Help improve Claude**" (or whichever similar phrase appears on your screen) so it slides to the **off** position. Some users see a confirmation popup — accept it. The change usually saves instantly without a separate Save button, but if you see one, click it.

A few honest notes here. On some free Claude accounts, the option to opt out of training may be limited or unavailable — Anthropic has updated this policy more than once. If you do not see the toggle at all, check whether you are signed in with a **work or school email**; those accounts are managed by an admin and cannot be changed here. If you pay for Claude, you may see the toggle in a slightly different place — the same Settings area, but possibly under a **Privacy** sub-tab.

You'll know it worked when the toggle is gray and stays gray after you refresh the page. If the setting does not seem to stick, sign out, sign back in, and check again — sometimes it takes a moment to sync.

Step 5 — Open Gemini's activity controls

Gemini handles this a little differently because it lives inside your Google account. Go to **gemini.google.com** or open the Gemini app. Click your **profile picture** in the top corner. Look for **Settings** or **Activity** — on the web, you may see a direct link labeled "**Gemini Apps Activity**" that takes you straight into your Google account controls. Click it.

On mobile, tap your **profile picture** in the top corner of the Gemini app, then look for **Gemini Apps Activity** in the menu that appears. If it is not there, your app version may have moved it — try searching your phone's settings app for "**Gemini**" instead, or open **myactivity.google.com** in a browser.

You'll know it worked when a Google page opens showing a toggle labeled "**Gemini Apps Activity**" and a short description of what it does.

Step 6 — Turn off Gemini's Apps Activity

On the Gemini Apps Activity page, click the toggle at the top so it switches to **off**. A small confirmation panel may pop up explaining what will change — read it, then confirm. You may also see a separate option to **delete past activity**; if you want a clean slate, click that too and follow the prompts.

This single Google toggle covers the Gemini web app, the Gemini mobile app, and any other place Gemini saves your chats to your Google account. Disabling it does not delete your past conversations automatically — you need to do that separately if you want them gone. And remember: turning off Gemini Apps Activity does not affect your Google Search history or YouTube history, which are controlled by separate toggles.

You'll know it worked when the toggle is off and a line of text confirms activity is paused. You can verify by starting a new chat, then returning to this page — the activity should not appear there.

Common mistakes

- **Looking for a "delete my data" button.** None of these toggles erase what the company already has on file. If you want old chats removed, each service has a separate **Delete** or **Clear conversations** option — and even then, companies often keep copies for legal or safety reasons for a while.
- **Assuming one toggle fixes everything.** In Gemini especially, there are **multiple** Google activity controls. Turning off Gemini Apps Activity does not affect Search history, YouTube history, or other Google products. Check each one separately if you care about all of them.

- **Trusting the toggle forever.** Companies update their settings and policies regularly. The opt-out you flipped today may quietly disappear or change meaning in six months — set a calendar reminder every few months to come back and check.

What this means for you

- **In everyday life:** Turning these off is a good habit if you ever paste anything personal into an AI chat — names, addresses, health questions, work emails. It does not make you invisible, but it stops your words from being used to teach the next model.
- **For your business or work:** If you handle customer data, contracts, or anything confidential, do not rely on these toggles alone. Many companies explicitly say opt-outs do not cover legal requests or internal abuse detection. Use a separate AI account with synthetic data, or stick to enterprise plans that promise stronger privacy.
- **If you're just getting started:** Pick the service you use the most, open its settings, and turn off the training toggle right now. That single move is the most useful thing you can do today — you can come back for the other two later.

Try it now

Open the AI assistant you used most recently, click your profile picture, find the privacy or data settings, and turn off the training toggle. It takes about a minute, and you'll see exactly what your settings look like today.

What You Should Never Share With an AI Assistant: A Privacy Guide

6 steps · [Read online](#)

By the end of this guide you'll know the three buckets of personal data that belong strictly offline, how to spot any of those details in your chat history, and what to do if you've already let an AI see them. It's written for anyone who chats with tools like ChatGPT, Gemini, Claude or Copilot.

Before you start

- A device with internet access (phone, tablet or computer).
- An account on at least one AI assistant (e.g., ChatGPT, Gemini, Claude – the free tier is enough).
- About 5 minutes of uninterrupted time.

Step 1 — Spot the three privacy buckets

Open the AI assistant you use most and start a fresh conversation. On screen you'll see a blank text box waiting for your prompt, and a "Send" button on the right. When you type, the assistant will start generating a reply beneath your prompt. If the layout looks different, look for a field that says "Ask anything" or a similar text-input area. 💡 Example: "What types of personal information should I never share with you?" You'll know it worked when the assistant lists three groups such as (1) identifiers (name, address, ID numbers), (2) financial details (bank accounts, credit-card numbers), and (3) health or biometric data.

Step 2 — Scan your recent chats for those buckets

Go to the chat-history screen – usually accessed by tapping a clock-icon or a menu labeled "History". You'll see a list of past conversations, each summarized by its first line. Selecting a conversation opens the full exchange, where you can scroll through every message you sent. If you don't see a clock-icon, look for a three-line menu and choose "Recent chats" instead. 💡 Example: type "Show me my last five messages" (many assistants support a simple request like this). You'll know it worked when the assistant displays the recent prompts you entered, giving you a chance to review them.

Step 3 — Delete or redact any message that contains private info

In the opened conversation locate the message that holds an identifier, financial detail or health data. Most assistants let you tap and hold the message, revealing options such as "Delete" or "Remove". Choose "Delete" and confirm when a pop-up asks "Are you sure?". If the option isn't

obvious, look for a trash-can icon or a “Clear” button at the bottom of the chat. 🗨️ Example: after selecting the message, type “Delete this message”. You’ll know it worked when the message disappears from the chat view and the history no longer shows that content.

Step 4 — Turn off chat-history saving (if your plan allows)

Open the settings panel – usually reached by tapping your profile picture or a gear-shaped icon in the top-right corner. The settings screen will list toggles like “Save chat history” or “Conversation storage”. Switch the toggle off; a brief note may appear saying the change applies to future chats only. On a free plan some assistants only let you “clear history” rather than disable it permanently, so look for a button that says “Clear all chats”. 🗨️ Example: “Turn off chat history”. You’ll know it worked when the toggle changes color or moves to the “off” position, and the app confirms “Your chats will not be saved”.

Step 5 — Verify that nothing sensitive is being stored

Start a new conversation and type a harmless query such as “What’s the capital of the US?”. The assistant will reply below the prompt. Then open the history again and check that only this new, non-personal question appears. If older, private messages still show up, repeat Steps 2–4 for any that remain. 🗨️ Example: “Show my recent chats”. You’ll know it worked when the list only contains the recent, innocuous question you just asked.

Step 6 — Request a full data-deletion if you think something slipped through

If you suspect you shared something the assistant cannot delete yourself (for example a piece of text buried deep in a long conversation), look for a “Help” or “Support” link in the settings menu. Most providers have a contact form or email address labeled “Data-privacy request”. Send a short message stating “Please delete all personal data associated with my account [your email]”. On some platforms you may need to verify your identity first. 🗨️ Example: “Dear Support, please remove any personal identifiers from my account.” You’ll know it worked when you receive a confirmation email or a system message saying “Your request has been received”.

Common mistakes

- **Sharing full addresses or ID numbers:** If you copy-paste your home address, the assistant now knows a location you might not want stored. Fix it by deleting the message and turning off history as shown in Steps 3–4.
- **Posting bank or credit-card numbers:** Even a partial number can be used for fraud. Remove the message immediately and contact your bank to monitor for any unusual activity.

- **Assuming the assistant deletes data for you:** Deleting a message only removes it from your view; the provider may still keep a copy on its servers. Always use the built-in “Clear all chats” option or submit a formal deletion request (Step 6).

What this means for you

- **In everyday life:** You can safely ask an AI to brainstorm a grocery list without worrying that your phone number or home address is being stored forever.
- **For your business or work:** Confidential client numbers and contract figures stay out of the AI’s training data, reducing the risk of accidental leaks.
- **If you’re just getting started:** The easiest first move is to turn off chat-history saving – it’s a single toggle that protects all future conversations.

Try it now

Open the AI assistant you use most, find the settings gear, and toggle **off** the “Save chat history” option. In under two minutes you’ll have taken a solid step toward keeping your private information out of AI-generated archives.

How to use AI safely with your personal information

4 steps · [Read online](#)

Want to use smart AI tools without accidentally sharing your private life with the world? This guide will show you exactly how to protect your personal details while still getting the most out of AI. You'll finish knowing how to thoughtfully interact with AI, manage privacy settings, and clean up sensitive information.

Before you start

- **An active account** on any popular AI tool like ChatGPT, Claude, or Gemini. A free plan is usually sufficient for this guide.
- **A computer or smartphone** with internet access. The steps apply whether you're using a web browser or a mobile app.
- **Rough total time:** About 10-15 minutes to read through and apply the steps to your chosen AI tool.

Step 1 — Understand the "postcard rule"

Before typing anything, pause and consider your message. The core idea for using AI safely is simple: treat everything you type into a chat window like writing on a postcard. If you wouldn't want a stranger reading it at the post office, do not send it to an AI. Most AI systems use your conversations for **data training** – this is where the AI company uses past chats to teach their computer models how to write and think better. Think of it like letting a student read your old letters so they can learn to write their own, but on a massive, anonymised scale.

You'll know it worked when you instinctively rephrase potentially sensitive information before hitting send. This habit helps the AI tool to assist you effectively without needing your private data. While AI can draft an email, it cannot magically remove personal information you've already typed.

💡 Example: Instead of typing: *"My name is Jane Smith, I live at 24 Smith Street and need help budgeting my \$1,500 weekly salary. My email is jane.smith@example.com."* Rephrase it to: *"Help me budget a weekly salary of \$1,500 for a typical household in a major city."*

Step 2 — Organize your settings and opt out of training

Almost all major AI platforms allow you to turn off data training. When you switch this off, the company still processes your requests to give you an answer, but they promise not to use your chat history to teach future versions of their system. This setting usually applies to future chats, not ones you've already had.

To do this, open your AI tool. Look for a gear icon, your profile picture, or a menu with three horizontal lines, usually located in a corner of the screen (often bottom-left or top-right). Tap or click on this to open a menu. From the menu, find and select a button labeled "Settings", "Privacy", or "Data Controls". Within this section, look for a toggle switch or checkbox that controls how your data is used for "model improvement" or "training". Toggle this option *off*. What happens next is a confirmation that your setting has been saved, or the toggle will simply show it's now off.

If it looks different: App layouts vary between AI tools and change over time. If you don't see exact labels, explore sections related to 'Privacy & Safety', 'Data Management', or 'Chat History'. The key is to find an option that controls how your data is used for "model improvement" or "training". You might need to scroll down or look through a few sub-menus.

💡 Example: In ChatGPT, open Settings (tap your profile picture or the gear icon), then look for "Privacy" or "Data Controls" and toggle off any option related to using your conversations for model improvement.

You'll know it worked when the toggle switch clearly shows it is in the "off" position (it might turn gray, or have a circle on the left side).

Step 3 — Sanitise your text before copying and pasting

"Sanitising" means cleaning your text to remove any specific details like names, phone numbers, email addresses, or financial figures *before* you feed them to the AI. This is your strongest defense against data leaks and ensures the AI doesn't learn sensitive facts about you. This action happens in your mind or a separate document, before you ever type into the AI chat.

When you review your text and replace real details with generic placeholders, the AI will still understand the context perfectly and can help you with your task. What happens is your input becomes general, focusing on the problem rather than the specifics of your personal life.

If it looks different: This step is about your pre-AI preparation, so what you see will be your own document, notes app, or even just your thoughts. No AI screen will look different here, as you're doing the sanitisation yourself.

💡 Example: If you want the AI to polish a professional email, change the real names and figures to generic labels like this: Original: *"Hi Sarah, can you please transfer the \$450 refund for the project at 12 High St to my Commonwealth account by Friday?"* Sanitised version: *"Hi [Recipient Name], can you please transfer the [Refund Amount] refund for the project at [Project Address] to my [Bank Name] account by [Date]?"*

You'll know it worked when you read through your prepared text and can't find any specific detail that would directly identify you, another person, or a specific location or account number.

Step 4 — Use temporary chats for sensitive tasks

If you need to discuss a particularly sensitive topic — like brainstorming a tricky work situation, or sorting out a personal issue — look for a temporary chat option. Some tools call this "Temporary Chat", "Incognito Mode", or "Ephemeral Chat". It works like write-on, wipe-off ink: once you close the window, the conversation is deleted from your history and generally won't be used for model training, even if your main training setting is on.

To activate this, open your AI tool and look near the top of your chat window or near the "New Chat" button for an option to start a temporary or incognito chat. Tap or click this before you begin your conversation. What happens is a new chat session will open, often with a visual indicator (like a different background color or an "Incognito" label) to show it's temporary. This temporary chat is usually designed to be forgotten once you close it.

If it looks different: Not all AI tools offer a specific "temporary chat" feature, especially on free plans. If you can't find it, don't worry – simply ensure your main data training setting (from Step 2) is switched off, and always sanitise your information (Step 3).

💡 Example: In Claude, look for a "Temporary Chat" button next to "New Chat" in the sidebar. Click it to start a new, unsaved conversation.

You'll know it worked when you see a clear visual cue that your chat is temporary (e.g., a message like "This chat will not be saved," or a different interface color).

Common mistakes

- **Uploading complete PDF documents without checking:** Many people upload entire utility bills, tax documents, or medical records for the AI to summarize. These files often contain private tax file numbers, account passwords, or personal health information.
 - **The fix:** Always open and review documents yourself *before* uploading. If there's sensitive info, either redact it (black it out or delete it) or copy-paste only the specific, non-sensitive parts you need the AI to process.
- **Treating the AI like a diary:** Do not use the chat box to vent about specific people by name, or share highly emotional, identifiable personal secrets. Even if you've opted out of training, the raw content might still be viewable by human reviewers at the AI company for safety checks.
 - **The fix:** Use generic terms like "a friend," "a colleague," or "a family member" instead of names. Focus on the situation or problem rather than the specific individuals involved. If it's truly private, talk to a trusted person instead.

- **Assuming "deleted" means instant removal:** When you delete a chat from your history sidebar, the AI company may still keep a copy on their secure servers for a short period (days or weeks) to check for harmful content, or for legal compliance. It's not always an instant, permanent deletion across all systems.
 - **The fix:** The best way to protect sensitive data is to *never put it into the AI in the first place*. Rely on the postcard rule (Step 1) and sanitising your text (Step 3) to prevent the data from ever being recorded.

What this means for you

- **In everyday life:** You can safely write personal emails, plan holidays, or draft letters to your local council without worrying about your home address leaking, provided you've followed these steps.
- **For your business or work:** You can use AI to summarize meeting notes, brainstorm marketing plans, or draft reports, provided you swap out client names, proprietary business secrets, or identifiable employee information.
- **If you're just getting started:** The easiest first move is to build a "pause habit"—take three seconds to look over your prompt before hitting the send button. Ask yourself: "Would I be okay with this being on a postcard?"

Try it now

Open your favorite AI tool right now. Navigate to the **Settings** menu, find the **Data Controls** or **Privacy** tab, and toggle off the option for **Model Training** or **Improve the model for everyone**. This simple action takes less than a minute and instantly secures your future chats from being used for broad model training.

How to Spot a Phishing Email or Scam Text Before You Click

6 steps · [Read online](#)

In short: Copy any suspicious email or scam text into an AI assistant with a short "fraud detector" prompt, and it lists the red flags — fake urgency, off-brand sender addresses, sketchy links — so you can decide whether to delete, block, or report the message before you click anything.

That pang of "wait, something feels off" you get when a strange email or text arrives — the AI can confirm it (or calm you down) in about 30 seconds. This guide shows you exactly what to type, and how to read the answer so you don't miss the warning signs. It works with ChatGPT, Gemini, Claude, Copilot, Le Chat, DeepSeek, or any other assistant you already use. One honest note up front: an AI can spot a lot, but it can't actually click the link for you or verify whether a sender is real — the final "yes" or "no" is still your call.

Before you start

- An account on any AI assistant you already use (free versions work fine — ChatGPT, Gemini, Claude, Microsoft Copilot, Le Chat, or DeepSeek all do this).
- The suspicious message itself, opened on your screen but with no links tapped or replies sent.
- Roughly 2 minutes.

Step 1 — Read the message calmly before doing anything

Open the suspicious email or text, and do exactly one thing: read it top to bottom with your finger off the mouse. Glance for three things — a sender or number you don't recognize, a tone that pushes you to act NOW, and any link you weren't expecting. You'll see the message exactly as you normally would; nothing changes yet. 🗨️ Example: if it's "Your package is delayed, confirm here: bit.ly/x9zq", just read the words. Don't tap the link. If it looks different: in some messaging apps long-pressing a link shows a small URL preview — that's actually useful, but still don't open it. Reading is the goal here. You'll know it worked when you can describe the message in one short sentence without having clicked anything.

Step 2 — Open any AI assistant you already use

On your phone or laptop, open the AI assistant app or website where you already have an account. No special "fraud tool" is needed — every general-purpose assistant can do this job. You'll see an empty chat box, usually labeled "Message", "Ask anything", or similar. If it looks different: every assistant has its own empty state, but the chat box is almost always in the middle

of the screen on desktop or at the bottom on mobile. If you're stuck, start a "New chat". You'll know it worked when you see a blinking cursor and a send button, ready for you to type.

Step 3 — Paste in a "fraud detector" prompt first

Before pasting the suspicious message, give the AI a one-line role instruction. Type it into the chat box and hit send. The AI will respond with a short confirmation — usually something like "Sure, paste it in" — that primes it to focus on fraud patterns instead of generic advice. This single line is what turns a normal chat into a fraud check. 🗨️ Example: "Act as a fraud detector. Read the message I'm pasting next and list every red flag you can see, then tell me clearly whether it's safe to reply, click any link, or open any attachment. If you're not sure, say so." If it looks different: some assistants reply with a longer setup paragraph — that's fine; what matters is that it now knows to focus on scams. You'll know it worked when the AI answers something like "Got it — paste the message" or asks a clarifying question.

Step 4 — Paste the full suspicious message, then send

Highlight the entire message — sender name, subject line, body, and the link text — but leave out your own name, address, or any account numbers if they appear, since the AI doesn't need them to spot trouble. Paste it into the chat box and hit send. The AI now has the raw text and starts analyzing it. If it looks different: if the message is unusually long (a fake invoice, say), you can paste a trimmed version — but keep the opening sentence and any links exactly as written, because that's where the red flags usually live. 🗨️ Example of what to paste: "From: Amazon Support <> — Dear customer, your account has been suspended due to unusual activity. Click here to verify: <http://amazon-secure.example.com/login> — please respond within 24 hours." You'll know it worked when the AI starts replying with a bulleted list of specific things it spotted — usually 4 to 8 items.

Step 5 — Read the red flags and trust the obvious ones

The AI's reply is usually a checklist: things like "the sender domain doesn't match Amazon," "the URL uses a zero instead of an 'o'," "the message creates urgency with a 24-hour deadline," or "the greeting is generic (no real name)." Read each one and ask yourself, "Does this match what I'm seeing on my screen?" You'll start to see a pattern you can recognize next time without any help. 🗨️ Example of an AI answer: "Red flags: sender domain is amazon-support.example.com, not amazon.com; the link is 'http' (typo) and uses a zero in 'amazon'; urgency tactic ('within 24 hours'); generic greeting. Verdict: do NOT click, do NOT reply. Report and delete." If it looks different: the AI might use softer words ("likely a phishing attempt", "this looks suspicious"). That's the same meaning. If it ever says "I'm not certain", take that seriously and lean toward treating the message as a scam. You'll know it worked when you can name at least three specific reasons the message is suspicious and explain them out loud.

Step 6 — Decide, then report and delete

Once the AI gives its verdict, decide based on the strength of the red flags. If even one major red flag is present — a sketchy sender domain, pressure to act NOW, or a request for passwords or payment — don't click and don't reply. Use the email or messaging app's built-in "Report phishing" or "Report junk" option if you can find it, then delete the message. The dangerous message is gone from your screen and flagged to the platform. 🟡 Example of a good follow-up question to the AI: "Can you explain in one sentence what an 'urgency tactic' means, so I know what to look for next time?" If it looks different: Gmail, Outlook, Apple Mail, and most phone carriers all have a Report button, usually under a "..." three-dot menu or near "Mark as" / "Move to". If you can't find the report option, deleting is good enough — what matters is that you don't reply. You'll know it worked when the message is no longer in your inbox, you've flagged it to your email or carrier, and you haven't tapped the link.

Common mistakes

- Replying to the scammer just to ask "is this real?" — replying confirms your address is live, and more scam messages follow. Delete and report instead.
- Trusting the AI without reading the explanation — AI can miss things. If the message "feels" off but the AI says it's fine, trust your gut and don't click. The AI is a second opinion, not a final one.
- Pasting personal details along with the message — your name, address, and account numbers aren't needed to spot red flags; trim them out before pasting, because anything you send into a chat is processed by that service.

What this means for you

- **In everyday life:** A 30-second check before clicking unknown links saves you from stolen passwords, drained bank accounts, and long phone calls with your bank. It becomes a habit that protects you on autopilot.
- **For your work:** Make this your default first step before forwarding suspicious work emails to IT. You'll often catch obvious scams that IT would have flagged anyway — which saves everyone time and earns you a reputation as the careful one on the team.
- **If you're just getting started:** Try it on the very next weird message you receive. You don't need to be perfect — you only need to slow down for 30 seconds before you click.

Try it now

The next time an email or text feels off, copy it into your AI assistant with the prompt from Step 3, paste the full message from Step 4, and read the red flags from Step 5 before you tap anything. The whole check takes about half a minute.

How to Spot AI Scams, Deepfakes, and Fake Content

5 steps · [Read online](#)

In short: AI-generated content is getting more convincing, but you can still protect yourself by verifying through independent channels and staying calm when messages create urgency.

By the end of this guide, you will know how to identify AI-generated images, synthetic voices, and deceptive messages before they trick you. This practical guide is designed to help anyone navigate the internet safely, without needing a degree in computer science. While these tips are powerful, remember that AI technology is always evolving. No single method is 100% foolproof, so staying updated and critically evaluating content is your best defense. This guide teaches you to be observant, but it cannot make decisions for you about the trustworthiness of a source.

Before you start

- **A device with internet access:** This could be your mobile phone, tablet, or computer.
- **About 10-15 minutes of your time:** Enough to read through the guide and try a few observations.
- **An observant and calm mindset:** AI changes quickly, so ongoing vigilance and a cool head are your best tools against deception.
- **No special apps or accounts needed:** We'll focus entirely on using your natural observation skills and everyday verification methods.

Step 1 — Spot visual "glitches" in photos

AI image generators are computer programs that create pictures based on text descriptions, much like a highly skilled digital painter who has studied millions of real photos. While they're incredibly talented, older versions of these programs often struggle with fine details, and examining photos closely can sometimes reveal inconsistencies.

When you see a photo that feels "off," your first action should be to zoom in and closely examine the specific details. You'll usually see the image take up more of your screen as you zoom. Once magnified, check areas like hands, feet, accessories, clothing, and backgrounds for subtle errors. If your device doesn't have a simple pinch-to-zoom feature, look for a magnifying glass icon or a dedicated zoom option in your photo viewer. You'll know it worked when you can clearly see distorted features, strange shadows, or illogical elements in the image.

💡 **Example:** You see a photo of a well-known person in an unusual situation. Zoom in on their hands and face. If you notice anything unusual — extra fingers, distorted features, or inconsistent lighting — investigate further before sharing or believing it.

Note: Modern AI generators have improved dramatically, so these artifacts are becoming less common. When in doubt, verify the image through a reverse image search or other means.

Step 2 — Identify "deepfake" video clues

A **deepfake** is a video where a person's face or body has been digitally altered or even completely swapped using smart computer systems. Think of it like a sophisticated digital mask that moves in real time, trying to perfectly mimic a real person. While detection is getting harder as the technology improves, there are still some things to watch for.

Start by playing the suspicious video, then pause it frequently to get a good look at specific frames. As you play the video, you'll see the frames move, and then when you pause, a still image will appear, allowing you to scrutinise the facial areas without the distraction of constant motion. If your video player doesn't have an obvious pause button, look for a large play/pause icon, often in the center or bottom of the screen. You'll know it worked when you spot any inconsistencies in how the face blends with the hair, unusual lighting, or areas that look slightly blurred or distorted.

💡 Example: You receive an urgent video message from your company's director asking for an immediate fund transfer. Pause the video and look closely at their jawline, neck, and hairline. If you notice anything unusual — a fuzzy outline, unusual shadows, or lip movements that don't quite match the audio — treat it as suspicious and verify through another channel before acting.

Note: High-quality deepfakes are increasingly difficult to spot with the naked eye. If a video requests money or sensitive information, always verify through a separate, trusted method.

Step 3 — Listen for voice cloning issues

Voice cloning is technology that takes a short recording of a real person's voice and then uses a computer to make them say entirely new things. Imagine it like a high-tech text-to-speech tool that can perfectly imitate your friend, child, or colleague. Modern voice cloning can be very convincing, so the best protection is not relying on voice alone.

When you receive a suspicious call or audio message, listen intently to the speaker's voice, paying close attention to their tone and speech patterns. You'll hear the words being spoken and can analyze how natural or unnatural they sound. If you are on a call, try to get to a quiet place or use headphones if available to reduce background noise. If anything feels slightly off — an unusual tone, unexpected phrasing, or a request that feels out of character — take that as a signal to verify through other means.

💡 Example: You receive an urgent phone call from a friend asking for financial help. The voice sounds exactly like them, but the request is unusual or the tone feels strange in context. This is a

red flag. Hang up immediately and call them back on their usual, trusted number to verify.

Step 4 — Verify suspicious messages using a second channel

Deceptive messages often leverage AI to write highly convincing, personalized emails or text messages, making them a potent form of automated scamming. The golden rule of online safety is simple: never use the contact details provided *within* a suspicious message to verify its authenticity, as these will likely lead you directly to the scammer.

Your action is to deliberately step away from the original message and initiate contact through a different, independently verified method. You'll see the original message on your screen (email, SMS, social media), which might contain urgent requests or "updated" details. By using an independent contact, you avoid giving the scammer direct access. If the request was via email, for instance, don't reply to that email. Instead, find a trusted phone number from a previous bill or the company's official website. You'll know it worked when the person on the other end confirms (or denies) the request, using a contact method you initiated and trust, separate from the potentially fake one.

💡 Example: A supplier sends an email with "updated bank details" for a payment, urging you to use the new account. Do not reply to that email or call the phone number listed in its signature. Instead, find the phone number from your original paper contract or the supplier's official website and call them directly to confirm the change.

Step 5 — Look for generic or overly perfect text

AI-generated text can often be identified not by obvious errors, but by its *perfection* – or lack of human nuance. These **AI-written texts** might be overly formal, incredibly verbose, repetitive in their phrasing, or simply too generic to sound like genuine human communication. They lack the quirks, emotional inflections, and subtle imperfections that make human writing unique.

Your action is to read the message or article carefully, paying close attention to its style, flow, and word choice. As you read, you'll be evaluating the tone and identifying any unusual patterns. This applies whether you're looking at an email, a social media post, an online article, or even a review. If the text reads like it was written by a robot trying to sound human, it's a good clue. You'll know it worked when the writing feels unnaturally smooth, free of human quirks or slang, or uses overly formal language for a casual situation.

💡 Example: You receive an email from a supposed friend, but it uses overly formal language like "Furthermore, I require immediate assistance" rather than their usual casual style, and every sentence is perfectly structured without any contractions or slang. This unnatural perfection suggests AI involvement.

Common mistakes

- **Trusting "official" branding:** Scammers can easily use AI to copy company logos, official-looking fonts, and professional language. The mistake is to assume legitimacy based on appearance. **The fix:** Always focus on the *action* they are asking you to take (like clicking a suspicious link or sending money), rather than how professional or authentic the message looks.
- **Panicking and acting too fast:** AI scams are designed to create urgency or fear, making you feel pressured to act without thinking. **The fix:** If a message demands immediate action, take a deep breath and pause. Time is always on your side when it comes to verifying facts and protecting yourself.
- **Missing the full context:** Sometimes a small snippet of an image, video, or audio can be shared out of its original context, making it seem suspicious when it's actually harmless. **The fix:** Always try to find the original source or the wider story if something looks questionable. Search online for the image or video to see where it first appeared and what the original context was.

What this means for you

- **In everyday life:** You can protect your family from distressing phone scams by realizing that a familiar voice on the phone might not always belong to the person you expect. Using these observation skills helps you verify crucial information before making hasty decisions, keeping your loved ones safe from emotional and financial exploitation.
- **For your business or work:** Safeguarding your company's assets and reputation becomes easier. You'll be better equipped to detect phishing emails, fraudulent invoices, or deepfake video calls that could lead to significant financial loss or data breaches. Protecting your team and clients from sophisticated AI-powered deception is vital for maintaining trust and security.
- **If you're just getting started:** The most important takeaway is to cultivate a healthy dose of skepticism about content you encounter online, especially if it elicits a strong emotional response or demands urgent action. Start by simply pausing and asking yourself, "Does this feel real?" before you click, reply, or share.

Try it now

Next time you're scrolling through social media or checking your emails, pause when you see an image, video, or message that looks a little too good to be true, or surprisingly out of character. Apply one of the quick checks from this guide – perhaps zoom in on a photo, or read a message aloud to detect robotic tones.

Conclusion

Staying informed and vigilant is your strongest defense against the evolving landscape of AI-driven scams and misinformation. By learning to recognize potential red flags in AI-generated content, you empower yourself and those around you to navigate the digital world with greater confidence and security. Always remember to question, verify, and trust your instincts when something feels off.

How to Spot AI-Generated Scam Messages Before You Click or Reply

5 steps · [Read online](#)

In short: Modern scams often use AI to write surprisingly convincing messages, but you can paste the message into any AI assistant (ChatGPT, Gemini, Claude, Copilot, Le Chat, DeepSeek — whichever you already use) and ask it to flag the warning signs. This guide walks you through a simple five-step check before you click a link, reply, or send money.

Scammers used to write in clumsy English. Today, AI helps them write polished messages that look exactly like a real text from your bank, USPS, or a family member. The good news: the same AI assistants we use every day are also very good at spotting these tricks when you ask the right questions. In the next five minutes, you'll learn a repeatable check you can run on any suspicious message.

Before you start

- An account with at least one AI assistant you can chat with — ChatGPT, Gemini, Claude, Microsoft Copilot, Le Chat, DeepSeek, or similar. The free tier is enough for this.
- A phone or computer where you can copy the suspicious message and paste it into the chat.
- The original message itself (screenshot or text), plus the sender's number, email or username if you have it.
- Roughly **5 minutes** for your first check; faster once you've done it two or three times.

A quick honest note: an AI assistant can flag *likely* warning signs. It cannot guarantee a message is safe, and it cannot detect every new trick. When in doubt, treat the message as suspicious and contact the organization directly using a phone number from their official website.

Step 1 — Don't click anything yet; copy the message safely

The very first action is the most important: **do not tap any link, reply, or call any number inside the suspicious message.** Scam links can take you to convincing fake sites, and replying confirms your number is real. Instead, take a breath, then open the message and copy the full text. On most phones, press and hold the message bubble and look for a **Copy** option in the small menu that pops up. On a computer, highlight the text with your mouse and press **Ctrl+C** (or **Cmd+C** on a Mac). If the message includes an image (a screenshot of an invoice, for example), take a screenshot of that too — look for a button like **Share** or **Screenshot** on your device. 🗨️ Example: a message you might want to check — *"Hi Mom, I dropped my phone and*

I'm using a friend's number. Can you urgently transfer \$450 to this account for a new one? I'll pay you back tonight." You'll know it worked when the suspicious text is sitting in your phone or computer's clipboard, ready to paste, and you have not tapped any link inside it.

Step 2 — Open your AI assistant in a separate window

Open a fresh browser tab or your AI app — whichever you normally use — so you can see the suspicious message and the assistant side by side. The reason for a separate window: if the scam message came through email, replying inside your email app can be risky, and a fresh chat keeps things clean. Look for the chat box, which is usually a large text field at the bottom of the screen, sometimes with a placeholder like **"Message ChatGPT"** or **"Ask Gemini"**. If you don't see a chat box, look for a button like **New chat** or a + icon at the top-left. If your assistant is on a paid plan, this check works the same; free plans handle short messages easily. 💡 Example: open the assistant you already use, and start a brand new chat titled something like *"Check this message for scam signs"*. You'll know it worked when you can see an empty chat with a text field ready for typing, and the original message is still in your clipboard.

Step 3 — Paste the message and ask for a scam-spotting breakdown

Tap inside the chat field and paste the message with **Ctrl+V** or **Cmd+V** (or long-press and **Paste** on a phone). Then, on a new line under the pasted message, ask the assistant to look for scam warning signs. A good prompt is short and direct — you don't need any special wording. If the assistant offers extra features (like web search, voice mode, or "deep research"), you can leave them off for this task; the default chat is plenty. 💡 Example prompt to copy: *"I got this message and something feels off. Can you list any scam warning signs you notice? Please look at the tone, the request, any pressure to act fast, and anything that doesn't quite add up. Don't tell me what to do — I just want a list of red flags."* You'll know it worked when the assistant replies with a numbered or bulleted list pointing out specific things in the message (for example: *"1. Urgency — the word 'urgently' is a pressure tactic"*).

Step 4 — Ask the assistant to check the link and the sender

Now add the sender's details and any link, and ask the assistant to look at those too. The **From** field on an email, the phone number on a text, or the username on a DM can all reveal patterns — for example, a "PayPal" email sent from a Gmail address, or a "bank" link that doesn't start with the bank's real website. Paste any link in full, exactly as it appeared (don't tap it). Then ask the assistant three things: who the sender claims to be, whether the link looks like a real address for that organization, and whether the sender's address matches the company they say they represent. If a tool can browse the web, you can also ask it to summarize what the public link actually points to. 💡 Example: *"The sender claims to be from USPS. Their email is '[usps-](#)*

delivery-support@gmail.com' and the link is 'usps-tracking.example.com/pay-now'. Does that look like a real USPS address? Just point out anything suspicious — don't visit the link for me." You'll know it worked when the assistant flags the mismatch — a Gmail address for a big company, or a link that almost spells the brand but is slightly off.

Step 5 — Read the summary with your own judgement

The assistant's list is a **starting point**, not a verdict. Take a moment to read its answer and ask yourself: do these warning signs match my gut feeling about the message? If the assistant flagged several red flags (urgency, money request, weird sender, off-brand link), treat the message as a scam — don't reply, don't click, and don't send any money. If you want to verify, contact the real organization using a phone number from their **official website** (type the address into your browser yourself; don't use any number from the suspicious message). Delete the message afterwards. If the assistant found no red flags but the message still feels wrong, trust your feeling — scams evolve faster than any checklist, and no AI can promise a 100% clean bill of health. 🧠 Example: ask a follow-up like *"Based on the red flags you found, is this more likely a scam or a real message? Give me your best guess and a one-sentence reason."* You'll know it worked when you can clearly say to yourself either "I'm confident this is a scam, I'm deleting it" or "I'm confident enough to verify it through the official channel — and I still haven't clicked the link in the message."

Common mistakes

- **Tapping a link "just to look"**. A scam site can install harmful software the moment it opens, or quietly capture data. Always check links in the assistant's chat, never in your browser.
- **Asking the assistant to "visit the link and tell me what's there"**. Some assistants can browse, but pasting an unknown URL into a tool you don't fully understand adds risk. Ask for an analysis of the URL itself (the words and structure) instead.
- **Trusting the message because the writing sounds perfect**. Perfect grammar used to be a sign of a real company. With AI, polished writing is now a common sign of a scam. Polished ≠ safe.
- **Ignoring pressure words**. "Urgent", "today only", "act now", "or your account will be closed" are some of the strongest scam signals. If you see them, slow down — that is exactly what the scammer doesn't want you to do.
- **Replying to ask "is this real?"**. Replying tells scammers your number is active. If you genuinely want to verify, contact the organization through a channel you found yourself, not through the message.

What this means for you

- **In everyday life:** treat any unexpected message asking for money, passwords, codes, or quick action as suspicious until you've checked it with your AI assistant. This covers texts from "your bank", "USPS", "IRS", "CouriersPlease", and even messages that look like they're from a family member using a new number.
- **For your business or work:** scam messages often target work emails (fake invoices, fake login pages, fake boss messages). Run the same five-step check on any work message that asks for payment changes, gift card purchases, or password resets. When in doubt, talk to a colleague in person before acting.
- **If you're just getting started:** keep one AI assistant open in a tab during the day. The moment a message feels off, paste it in and run through the checklist. After two or three times, the whole check takes under a minute — and you'll start spotting red flags on your own before you even paste.

Try it now

Find one message in your phone or inbox that you've been unsure about (not from the last hour — something older is fine). Paste it into your usual AI assistant with the prompt from Step 3, and read the first three red flags it lists. That's your first five-minute check.

How to Spot a Deepfake Video or Voice Clone Before You Share It

5 steps · [Read online](#)

In short: Look at the eyes, mouth and face edges for small glitches; listen for flat, robotic speech; check who sent it and what it asks. If anything feels off, don't forward it — verify with the real person through a trusted channel first.

A short clip of a public figure saying something outrageous lands in your family group chat. A frantic voicemail from your "boss" asks for gift cards. None of these are obviously fake, and that's exactly the point. A **deepfake** is a video, image or audio clip made with AI to look or sound like a real person who never actually said or did the thing, and **voice cloning** (reproducing someone's voice convincingly from a short audio sample) is now easy and cheap. The good news: a few careful seconds of looking is usually enough to catch the most common giveaways. Even so, no single check is perfect, so your final safety net is always verifying with the real person before you share.

Before you start

- The clip itself (saved to your phone or laptop, or a link in a message)
- A web browser (for the reverse image check in Step 4)
- About 2–3 minutes per clip
- No account, no app, no paid plan needed

Step 1 — Watch the face for small glitches

Pause the video and study the face for about 10 seconds. Focus on the **eyes, the edges of the face, and the mouth**. Common giveaways: eyes that don't quite blink in time; glasses that seem to slide or warp when the head turns; teeth that blur into a single white block when the mouth moves; a face outline that looks soft or "painted" where it meets the hair and background; lighting on the face that doesn't match the rest of the scene. If you can't pause or zoom, slow playback to 0.5x and look again. On a phone, tap the screen to bring up the playback speed option — it's usually a small "1x" button near the timeline. You'll know it worked when you spot at least one thing that feels off — or, just as usefully, confirm the face looks natural.

Step 2 — Listen to the voice for flat delivery

Close your eyes for a moment and just listen. Real human speech has small imperfections: a quick breath, a half-stutter, a sentence that gets quieter when the speaker is unsure. AI-generated voices (TTS, short for text-to-speech — software that turns written words into spoken audio) tend to be too clean, too even in volume, and oddly flat in emotion. Listen for sentences

with no breath at all, words that all sit on the same pitch, or a tone that doesn't match the situation (someone supposedly panicking, but the voice sounds calm and steady). On a phone, plug in headphones so background hiss doesn't mask the small tells. You'll know it worked when you can tell whether the voice feels human or **synthesized** (made by a computer, not a real person).

Step 3 — Read the message around the clip

Before you forward anything, stop and read the context. Who sent it? Is it really them, or just a forwarded link from someone you trust? What does the clip ask you to do — donate money, click a link, send a code, stay quiet about something? Most scams lean on **urgency, secrecy or big emotional claims**: "send this in the next hour", "don't tell anyone", "this is huge." Those are the same red flags as any scam, and they're a strong reason to slow down. If you can't answer "who, what, and why" from the message alone, that's already a signal. You'll know it worked when you have a clear answer to those three questions — or a clear reason to be suspicious.

Step 4 — Run a quick source check

Take a screenshot of the person's face in the video (or use the platform's "save image" option). Open a reverse image search (commonly available on **Google Images**, **TinEye** or **Yandex Images** — no account needed) and upload the picture. Look for earlier appearances of the same face. If the face is supposed to be a public figure but the search shows nothing, or if older photos don't match the person in the clip, treat it as a likely fake. If the clip is supposedly a private person, send a quick message to that person on a channel you already had before the clip — not the one the clip came from — and ask if they really sent it. 💬 Example: "Hey, I got a voice note from someone claiming to be you asking about gift cards. Did you actually send it? Just checking." You'll know it worked when you can either confirm the source or decide it isn't worth the risk.

Step 5 — Decide what to do next

If any of the checks above raised doubt, the safest move is simple: **don't share it**. Reposting a fake — even with a "this might be fake" caveat — still spreads the original. Tell the person who sent it, politely, that you couldn't verify it, and ask them where they got it. If the clip claims to be a real person you know, give them a heads-up that someone may be using their face or voice. If the clip is a scam attempt, report it to the platform where you saw it, and if money or account access was involved, contact your bank or the company being impersonated directly. You'll know it worked when you stop and think before tapping "forward."

Common mistakes

- **Trusting the "share if you agree" framing.** If a clip is trying to get you to forward it before you've checked, that's the scam working. Slow down — the share button can wait 30 seconds.
- **Assuming a friend's name means a friend sent it.** Accounts get hacked and contacts get **spoofed** (a message made to look like it came from someone you know). Verify through a separate channel before trusting.
- **Relying on a single "AI detection" app.** No consumer app is reliable enough to be your only check, and some detection tools themselves are scams. Your eyes, ears and a quick source check beat any single app.

What this means for you

- **In everyday life:** Treat any unexpected video or voice message that asks for money, codes, or secrecy as a scam until proven otherwise. A two-minute check protects your bank account and your reputation.
- **For your work:** If a "colleague" sends an unusual request by voice or video, especially one involving wire transfers or credentials, confirm it through a known channel like a phone call or in-person chat before acting.
- **If you're just getting started:** Save this page and use it as a checklist the next time a clip makes you feel strongly — good or bad. Strong feelings are exactly what fake clips are designed to trigger.

Try it now

Pick one short video you saw recently that made you feel a strong reaction. Spend two minutes running Steps 1, 2 and 4 on it. Even if it turns out to be real, the practice makes the next check faster and easier.