Everyday AI: Work & email

20 step-by-step guides for real, everyday tasks

Work & email

Updated edition · October 19, 2026

AI World HQ · https://aiworldhq.com

Free edition · CC BY-NC 4.0 · free to share with credit

Before you start — what this is, honestly

This is an updated edition. It was rebuilt from our current guides on the date shown on the cover: more guides per topic, retired guides removed, and every step taken from the latest version of each guide.

This collection is free, and every guide in it is also free on our website. This file just gathers them in one place, in reading order, so you can print it, keep it on your phone, or read it offline. Our website has no paywall either — every guide links back to its online page, which is always the newer version.

AI **Every guide in this pack was written by AI** and checked by automated fact and quality gates. **Our editor spot-checks published guides every day, but not every text is read by a person.** Every guide page on https://aiworldhq.com carries the same notice — we say it openly because you deserve to know who wrote what you are reading.

Apps change. Buttons move and menus get renamed. If a step does not match what you see, the guide's online version is the newer one — the link is under every title.

Free to share. This collection is licensed under Creative Commons Attribution-NonCommercial 4.0 (CC BY-NC 4.0): you may copy, print and share it — with your class, your library or your team — as long as you credit AI World HQ with a link to aiworldhq.com and do not sell it.

What's inside

1. **Work & email**
    1. How to Get Quick Summaries from Long Documents with Claude
    2. Turn Messy Notes into a Polite Meeting Agenda in Five Minutes
    3. How to Clean Up a Flooded Email Inbox Using AI in One Afternoon
    4. Check your resume for ATS-friendly keywords using Hugging Face's Inference API
    5. Write a Polite Complaint Email to a Company with DeepSeek
    6. Drafting Professional Presentations from Scratch Using Microsoft Copilot in PowerPoint
    7. Digitize Handwritten Notes into Editable Text with AI
    8. How to turn a long PDF contract into a plain-English summary in 5 minutes
    9. Unlock Document Insights with Google NotebookLM
    10. Summarize Any Long Document or Article with AI
    11. How to Prepare for a Job Interview Using a Pasted Job Post
    12. Write a Resume From Scratch (Even With No Work Experience)
    13. How to Organize a Messy Email Inbox Into Simple Folders in One Sitting
    14. Practice job interview questions with Qwen before the real interview
    15. How to write a polite but firm complaint email that actually gets a reply
    16. Translate and Explain a Confusing Email or Letter with Perplexity
    17. Draft a Polite but Firm Refund Email with Le Chat
    18. How to Use Gemini Deep Research for Thorough Reports
    19. How to Draft a Polite Follow-up Email After a Job Interview
    20. How to practice for a job interview with an AI mock interviewer tonight

Work & email

How to Get Quick Summaries from Long Documents with Claude

7 steps · Read online

Ever feel swamped by lengthy reports, articles, or complex documents? This guide will show you how to use Claude, an **AI (Artificial Intelligence)**, which is like a very smart computer program designed to think and learn, to quickly get the main points from almost any text. By the end, you'll be able to upload a document and ask Claude to summarize it for you, giving you the essential information you need in a flash.

Before you start

- You'll need a free Claude account (or a paid plan if you prefer). This will give you access to Claude's website.
- Access to the Claude website through your computer's web browser.
- A digital document you want to summarize, such as a PDF or a plain text file. Note that Claude has limits on the file size and length it can process, but free plans are usually sufficient for most common documents.
- This whole process should take you about 5-7 minutes.

Step 1 — Access Claude's Website

Your first step is to open your web browser and navigate to the Claude website. Think of Claude as a powerful **Large Language Model (LLM)**, which is like the highly intelligent engine behind an AI, specifically designed to understand and generate human language. You'll usually see a clean, welcoming chat interface where you can begin interacting with the AI.

You'll know it worked when you see a main chat area, often with a text box at the bottom ready for you to type in.

Step 2 — Start a New Conversation

To ensure you're starting fresh, it's a good idea to begin a new conversation with Claude. Look for a "New Chat" button, which is commonly found on the left-hand side of the screen or near the top. Clicking this will give you a clean slate, meaning Claude won't carry over any previous topics or instructions from old chats, so your summary request will be clear.

If it looks different on your screen, don't worry! Apps change often. You might see a small plus (+) icon, or a menu option like "Start new conversation." Simply click whatever seems to open a fresh chat.

You'll know it worked when the main chat window is clear, with no previous messages, and the message box at the bottom is empty, waiting for your input.

Step 3 — Upload Your Document

Now it's time to get your document into Claude. Look for a paperclip icon, usually found near the message box where you type. This icon is universally used for attachments. Clicking it will open your computer's file explorer, allowing you to browse and select the PDF or text file you wish to summarize. Claude can handle various formats, but PDFs and plain text are generally the most reliable for summarisation. Remember, very large files might not upload on a free plan, as Claude, like all AIs, has a **context window** – this is like the AI's short-term memory, limiting how much text it can 'remember' at once.

If you don't see a paperclip, look for an "Attach file" text link or a similar icon that suggests adding a file.

You'll know it worked when the name of your uploaded document appears in the message box, usually next to your cursor, indicating it's ready to be sent.

Step 4 — Craft Your Summary Request

With your document attached, it's time to tell Claude what you want. This instruction is called a **prompt**, which is simply the command you give to the AI. In the message box, type out a clear request for a summary. Being specific helps Claude understand exactly what you're looking for, rather than just giving a generic overview. For instance, you could ask for specific details or a summary tailored to a particular audience.

If your message box is hard to find, it’s usually a long, horizontal bar at the very bottom of the chat window.

You'll know it worked when your complete summary request appears in the message box, ready to be sent to Claude. 💬 Example: "Please summarize this document for me, highlighting the main points and key takeaways relevant to small businesses."

Step 5 — Send Your Request

Once your document is attached and your summary request is typed, the next step is to send it to Claude. Look for a "Send" button or a paper plane icon, usually located at the right end of the message box. Clicking this will submit both your document and your prompt to the AI. Claude will then begin processing the information, reading through your document and generating a summary based on your instructions. This process usually happens quite quickly.

If the button looks different, it might be an arrow or just the word "Send". Just click the button that looks like it will submit your message.

You'll know it worked when your message disappears from the box, and Claude's response starts appearing in the main chat window, beginning with your requested summary.

Step 6 — Review and Refine the Summary

After Claude generates the summary, take a moment to read through it. AI responses are usually very helpful, but sometimes they might not hit every single point you need, or you might want more detail on a specific section. This is where you can interact further. If you need more information or a different perspective, just type another prompt in the message box, referring back to the document or the summary Claude just gave you. For example, you could ask Claude to elaborate on a particular section or rephrase the summary for a different audience.

If the answer isn't quite what you expected, don't worry, that's normal with AIs. Just try rephrasing your question or adding more specific details in your next prompt.

You'll know it worked when you've either received a summary that fully meets your needs or have asked follow-up questions that yield the specific additional information you were looking for. 💬 Example: "Can you list the top three challenges mentioned in the document and suggest potential solutions?"

Step 7 — Save or Copy Your Summary

Once you're happy with Claude's summary, you'll likely want to save or copy it for your own use. You can typically highlight the text of Claude's response using your mouse, then right-click and select "Copy." You can then paste this summary into a document, email, or wherever you need it. Some AI tools also offer specific "Copy" buttons next to their responses or even export options, though these vary.

If you don't see a direct "copy" button, highlighting the text and using your computer's standard copy (Ctrl+C on Windows, Command+C on Mac) and paste (Ctrl+V or Command+V) functions will always work.

You'll know it worked when you've successfully pasted the summary into another application or saved it to a file, making it ready for your tasks.

Common mistakes

- **Uploading unsupported files or files that are too large:** Claude works best with common document types like PDFs, Word documents, and text files. If you try to upload

an image or a video, or a document that's many hundreds of pages long, it might not work.
    - **The fix:** Convert your file to a supported format like PDF or plain text. If it's very long, you might need to split it into smaller sections or use a paid plan with a larger context window.
- **Vague summary requests:** Just saying "Summarize this" might give you a general overview, but it might not be what you actually need. Claude is good, but it can't read your mind!
    - **The fix:** Be specific! Ask for "three key takeaways," "a summary for someone who knows nothing about this topic," or "the pros and cons of this proposal." The more detail you give in your **prompt**, the better the result.
- **Not verifying critical information:** While Claude is powerful, it's an AI and can sometimes **hallucinate**, which is when an AI confidently presents false or inaccurate information as if it were true. It generates text based on patterns, not absolute truth.
    - **The fix:** Always cross-check any critical facts, figures, or sensitive information from the summary with the original document or another reliable source, especially if you plan to use the summary for important decisions. Claude can draft an excellent summary, but it can't verify the truthfulness of *every single detail* for you, especially in complex or niche topics.

What this means for you

- **In everyday life:** Imagine quickly understanding the long terms and conditions for a new app, getting the gist of a lengthy news article, or grasping the main points of a new recipe or instruction manual without reading every single word. Claude can save you precious time.
- **For your business or work:** This skill is invaluable for professionals. You can rapidly review lengthy reports, condense meeting minutes, extract key findings from research papers, or quickly understand client briefs, allowing you to focus on decision-making rather than sifting through pages of text.
- **If you're just getting started:** Pick a news article you're interested in online. Copy the main text, paste it into a new Claude conversation, and ask, "Summarize this for me in three sentences." It's a quick, easy way to get a feel for how it works!

Try it now

Find a short PDF document (perhaps a product manual or a short article) or copy-paste the text from a recent online news story into Claude, and then ask it to "Give me the top three main points from this text." See how quickly you can get to the core information!

Turn Messy Notes into a Polite Meeting Agenda in Five Minutes

7 steps · Read online

> **In short:** Paste your messy email thread or bullet list into any AI assistant with a clear instruction about timing, owners, and tone, then iterate. You'll get a polite, structured agenda in under five minutes — though you'll still need to send the invites yourself.

A 30-minute meeting with seven bullet points, three owners, and zero structure usually ends with 20 minutes wasted on "wait, who's covering that?" — and another 10 on a follow-up email nobody reads. This guide shows how to feed that mess into any AI assistant and walk out with a timed, owner-clear agenda you can actually send. It's written for people who already use AI for everyday tasks and want one more workflow that saves real time at work.

Before you start

- An account on any AI assistant you already use — ChatGPT, Claude, Gemini, Copilot, Le Chat, or DeepSeek all handle this pattern. No new signup needed.
- The messy source material within reach: an email thread to copy from, a notes app with bullet points, or a transcribed voice memo.
- About five minutes once you've done it once; a little longer the first time as you build your reusable template.

Step 1 — Gather the messy material in one place

Open your AI assistant in a browser or app and start a fresh chat. Paste your raw material — the email thread, the bullet list, the voice-note transcript — straight into the message box, and add two short sentences of context before you hit send: what the meeting is for, how long it is, and who's coming. You're not asking for the agenda yet; you're staging the inputs.

💬 Example: "Below is the email thread about our Q3 launch. Goal: 30-minute sync to unblock the design freeze. Attendees: Maria, John, Sarah."

You'll see your pasted text sitting in the input area, with the send button below it. After you send, the assistant confirms receipt and begins working. If your assistant has a "Projects" or "custom instructions" feature, you can save the context sentence there for next time — but a plain chat works fine without it.

You'll know it worked when the assistant's first reply acknowledges the inputs you provided and starts producing structured output rather than asking what you meant.

Step 2 — Ask for a first draft with explicit structure

Send a follow-up instruction in the same chat telling the assistant how to shape the draft: numbered topics, a duration in minutes next to each, a named owner per topic, a polite opener, and a small buffer between items. The clearer your structure request, the less you'll need to fix later.

Example: "Now produce the draft. Use a polite opener, list each topic with a duration in minutes and a single owner, leave a 2-minute buffer between topics, and end with a one-line confirmation request."

You'll see a numbered agenda appear, with minutes and names attached to each line. If your assistant formats the response as a table by default — some do — reply with "please rewrite as a numbered list with bold topic names" and the next version arrives in the shape you wanted.

You'll know it worked when your chaotic items are reordered into clear blocks with names and minutes, and you can read the whole meeting in one glance.

Step 3 — Calibrate the politeness

Reply in the same chat asking for tone adjustments. The first draft is almost never the right tone — it usually sounds either too stiff ("Please find below the agenda…") or too casual for a workplace you share with senior people. Tell the assistant which direction to move.

Example: "The opener feels stiff. Loosen it. Our team uses first names and avoids phrases like 'circling back'. Keep the same structure and timings."

You'll see a revised version where the wording changes but the structure stays. If the assistant doesn't pick up your tone signal, paste one sentence of how you'd actually open the meeting and ask it to match that register — that's the fastest shortcut to a natural voice.

You'll know it worked when the opener reads like something you'd say out loud without wincing.

Step 4 — Lock down the owners and timing

Send another short instruction asking the assistant to verify two things: each topic has exactly one named owner, and the minutes add up to the meeting length. This is where most agendas quietly fall apart — two topics assigned to the same person back-to-back, or a 45-minute agenda for a 30-minute slot.

Example: "Confirm each topic has exactly one owner. Add up the minutes and show the total at the end. If anything overflows, suggest what to cut."

You'll see a revised agenda where totals match and names sit beside each item. If the assistant doesn't add a "Total: X minutes" footer, ask for one — it's the simplest way to verify the math yourself before you copy anything out.

You'll know it worked when the total at the bottom equals your available time, and no item is missing a name.

Step 5 — Add a closing line with next steps

Ask the assistant to add a short "Next steps" block at the bottom: who sends the calendar invite, who brings what (slides, doc, numbers), and when the follow-up happens. This is the part most agendas skip, and it's the part that decides whether the meeting actually moves anything.

💬 Example: "Add a 'Next steps' block at the end: Maria sends the calendar invite by tomorrow morning; John brings the latest design mockup; follow-up sync the following Monday."

You'll see a new block appear below the agenda. If your team uses a specific tool — Notion, Google Docs, Slack — name it in your prompt and the assistant will adjust the closing line to fit that workflow.

You'll know it worked when each person at the meeting knows exactly what to do before it starts and after it ends.

Step 6 — Copy the result into your real-world tool

Select the final agenda in the chat and paste it into the place your team will actually see it — an email draft, a shared doc, a calendar description, or your meeting software. This step is small but it matters: the moment the agenda lives outside the chat, you can edit it, share it, and update it without re-asking the AI.

If your assistant offers a "copy" or "export" button, use it; otherwise highlight the text and use your system copy shortcut. The text lands in the new tool with light formatting shifts depending on the destination — that's normal.

You'll know it worked when the agenda sits in the same place your team already works, and you didn't need to retype a single bullet.

Step 7 — Save your prompt as a reusable template

Highlight the full chat — your original messy input, the structure request, the tone instructions, the owners-and-timing check, the closing block — and save it as a reusable template. A plain text file in your notes app works. If your assistant supports saved prompts or custom instructions, use that instead so the template loads automatically.

💬 Example: a doc titled "Agenda from messy notes" with three labeled blocks: [PASTE INPUT], [STRUCTURE REQUEST], [TONE NOTES].

Next time you have a chaotic thread, you paste the template, swap in the new content, and skip the prompt-writing step. The first time costs you ten minutes; every meeting after that costs you one.

You'll know it worked when the next agenda takes you under a minute to draft.

Common mistakes

- **Treating the first draft as final.** The first reply is a sketch. Always run one tone-and-owners pass — usually two — before you copy anything out.
- **Owners stay vague ("the team", "marketing").** Vague owners mean nobody follows up. One named person per topic, every time, no exceptions.
- **Forgetting the AI can't send the invite.** Once you copy the agenda out, the chat ends and the tool stops helping. You still need to send the calendar invite, share the doc, and chase the owners yourself.

What this means for you

- **In everyday life:** The same pattern tidies up a 10-minute prep before a school meeting, a volunteer group call, or a family planning session — paste the chaos, ask for structure, lock in who does what.
- **For your business or work:** Every recurring meeting — weekly stand-up, monthly review, project sync — gets a polished, owner-clear agenda in minutes instead of 30 minutes of wrestling with bullet points.
- **If you want to go further:** Try the same input-to-structure pattern on other meeting artefacts: a recap email after the meeting, a project kickoff brief from a Slack thread, or a status report from scattered updates. The skill is the same — stage the messy input, request the shape you want, iterate on tone and owners.

Try it now

Open any AI assistant, paste your messiest recent email thread, and send this single line: "Turn this into a 30-minute agenda with one named owner per topic and a polite opener." Watch what comes back — that's your starting point for the iteration loop above.

How to Clean Up a Flooded Email Inbox Using AI in One Afternoon

7 steps · Read online

> **In short:** You can use any AI assistant (ChatGPT, Claude, Gemini, or Copilot) to group your unread email into folders and write a one-paragraph summary of what needs attention today. Most people finish in under three hours.

By the end of this guide, you'll have a tidy inbox with folders that make sense, plus a single short paragraph at the top telling you what actually matters right now. The AI does the thinking; you do the clicking.

Before you start

- An account with any AI assistant you can already use — ChatGPT, Claude, Gemini, Copilot, or Le Chat all work. The free version is enough.
- Your email account (Gmail, Outlook, Yahoo, Apple Mail, anything) open in a web browser.
- About 2–3 hours of quiet time — most of it is waiting, not typing.
- A notes app or blank document where you can paste text (Notepad, Google Docs, Notes — anything that saves text).

Step 1 — Pick which AI you'll use, and open it in a new tab

Open the AI assistant of your choice in a browser tab and sign in. You'll see a chat box at the bottom of the screen with a blank text field waiting for your first message. The conversation is private to you — nothing you paste is sent back to your email provider.

If you've never used one before, ChatGPT, Claude, or Gemini are all safe picks. The free version is plenty for this task; paid plans just let you paste longer messages. If you can't find your bookmark, search "[tool name] login" in Google — the first result is almost always the real site.

Example: open chatgpt.com, claude.ai, gemini.google.com, or copilot.microsoft.com.

You'll know it worked when you see a typing box with a "Send" arrow or paper-plane icon, and the AI greets you when you say hello.

Step 2 — Make a copy of your unread emails the AI can read

In your email account, search for `is:unread` (Gmail), `unread:true` (Outlook web), or use the "Unread" filter in the left sidebar. The screen fills with every message you've ignored — that's your starting pile.

Pick the most recent 50–100 unread emails. Copy the subject lines and senders into a plain text document with `Ctrl+C` (Windows) or `Cmd+C` (Mac). If your email client offers "Export to .csv" for the unread folder, that's even faster — but only export unread mail, never your whole account.

💬 Example: in Gmail, tick the checkbox at the top of the inbox. If a "Select all conversations" link appears, click it. Open the ⋯ menu and choose "Export."

You'll know it worked when you have a file with rows like "Sender — Subject — Date — first 2 lines."

Step 3 — Ask the AI to group the emails into 5–7 folders

Switch back to your AI tab and paste the list in. Then ask it to suggest a folder structure — copy the prompt below as-is.

The AI will respond with a small list of folder names (5–7 is the sweet spot — too many and you stop using them) and which emails belong where. It will not move anything itself; it only suggests. That part is still on you.

💬 Example prompt: *"Here is a list of my unread emails. Group them into 5–7 folder names that make sense, and tell me which emails go in which folder. Use plain folder names a normal person would understand. Don't add explanations yet — just the list."*

You'll know it worked when you see something like: "Personal & Family — 12 emails. Bills & Receipts — 8 emails. Work — 27 emails. Newsletters — 41 emails. Calendar & RSVPs — 6 emails."

Step 4 — Create those folders in your email client

In your email account, create the folders the AI suggested. Most email apps have a "+" or "Create new label/folder" button in the left sidebar — if you don't see it, right-click an existing folder name and look for "New label," "New folder," or "Create."

Name them exactly as the AI wrote them, with the same spelling and capitalization. Lowercase, plural, no emoji — keeping it boring makes it stick. If the AI suggested something like "WorkStuff" and you'd prefer "Work," rename it now. The point is a structure you'll actually use.

💬 Example: in Gmail, look in the left sidebar, scroll to "Labels," click the "+" next to it, and type "Bills & Receipts." In Outlook web, right-click "Folders" in the sidebar and choose "New folder."

You'll know it worked when every new folder name shows up in the sidebar, ready to receive mail.

Step 5 — Move messages into folders, starting with the small groups

Go back to your inbox and tackle the smallest group first — usually "Calendar & RSVPs" or "Personal." Open each email, click "Move to," and pick the matching folder. Do not read anything. The goal is an empty inbox, not total comprehension.

For the biggest group (often "Newsletters"), speed this up: search `category:promotions` (Gmail), `category:updates`, or use the bulk-select checkbox at the top of the list, then choose "Move to" → "Newsletters." This is the only step that feels slow — the rest is fast.

💬 Example: type `label:unread from:newsletter` in the search box, select all results, move them to your "Newsletters" folder in one click.

You'll know it worked when your main inbox counter shows a much smaller number than before, and the moved messages no longer appear in the inbox list.

Step 6 — Ask the AI to write your "what matters now" summary

Now paste the remaining urgent emails (the ones still in your inbox, or the Work and Personal folders) back into the AI and ask it to write a short paragraph telling you what actually needs attention today.

The AI is good at this — it sorts through noise, spots deadlines and asks, and skips pleasantries. Your job is to give it enough context: who is waiting on you, what is due this week, what can wait a month.

💬 Example prompt: *"Here are the remaining unread emails. Write me one short paragraph (under 100 words) telling me what I actually need to act on today, what's due this week, and what I can ignore. Start with the most urgent."*

You'll know it worked when the AI returns a single paragraph beginning with the deadline or action that's most urgent, like: "Reply to the school about Thursday pickup by tomorrow morning — they need a yes or no. Confirm the dentist appointment on Friday. Ignore the three shipping notifications…"

Step 7 — Pin that summary somewhere you'll see it

Take the paragraph the AI wrote and put it where your eyes land first in the morning. The easiest spot: the very top of your inbox (most email apps let you "Pin" or "Star" a message), or a sticky note on your desktop, or the first line of today's to-do list.

This is the whole point of the exercise. A clean inbox feels good for ten minutes — a clear "what matters now" keeps you focused all week. The folders are just a parking lot; the paragraph is the compass.

💬 Example: copy the paragraph into a new email titled "Today" addressed to yourself, and pin it at the top of your inbox.

You'll know it worked when tomorrow morning, before you open any other email, you can read that one paragraph and know what to do first.

Common mistakes

- **Exporting your whole email account instead of just unread.** The AI does not need ten years of receipts to group your unread mail — and pasting your entire inbox into a chatbot is risky. Stick to unread, and only the last few hundred.
- **Letting the AI invent clever folder names like "CriticalPath_Q3_Items."** If a folder name is not something you'd write on a sticky note, change it. Boring folder names get used; clever ones get ignored.
- **Reading every email as you move it.** That's not cleanup, that's procrastination wearing a costume. Move first, read later — the summary tells you what actually matters.

What this means for you

- **In everyday life:** you'll stop dreading the inbox icon. The folder structure is a parking lot for things you don't need to act on; the summary is your morning compass.
- **For your work or business:** if you get a lot of client email, do this once a month. The folder names might shift slightly (a new project, a new client), but the pattern stays the same.
- **If you're just getting started:** pick one AI you already trust, do Steps 1–4 today, and save 5–7 for tomorrow. Two shorter sessions beat one exhausting one.

Try it now

Open your email account, search for unread messages, and copy the first 30 subject lines into your AI chat. Ask it for 5 folder names. That's it — ten minutes of work, and you'll already see the shape of the cleanup.

Check your resume for ATS-friendly keywords using Hugging Face's Inference API

7 steps · Read online

You'll learn how to scan your resume for keywords that applicant tracking systems (ATS) love, using free AI tools. This helps your resume get past the first filter and into a human's hands.

Before you start

- A **free Hugging Face account** (sign up here)
- A **resume in PDF or Word format** (saved on your device)
- A **computer or laptop** (phone works, but a larger screen is easier)
- **No coding skills needed** — this guide is for everyone
- **Total time:** about 10 minutes

💡 *What is ATS?* Think of it as a robot that scans your resume before a human does. It looks for keywords from the job ad. If your resume doesn't have enough of those words, it might get rejected before you even know.

Step 1 — Create a free Hugging Face account

Go to huggingface.co and click **Sign Up** in the top-right corner. Fill in your email, choose a username and password, then click **Create account**. Check your email and confirm your account.

You'll see your Hugging Face profile page. This is your dashboard — your home base for using AI models.

💬 Example:

```
Email: yourname@example.com
Username: resumehelper
Password: MyPassword123!
```

You'll know it worked when you see your profile picture and username in the top-right corner.

Step 2 — Open the Inference API page

On the Hugging Face homepage, look for **Products** in the top menu. Hover over it and click **Inference API**. This is a free tool that lets you run AI models without installing anything.

You'll see a page with a text box and a list of AI models. This is where you'll paste your job ad and resume.

💡 *What is an API?* Think of it like a waiter in a restaurant. You tell the waiter what you want (your job ad and resume), and the waiter brings it to the kitchen (the AI model) and brings back the answer — without you needing to cook.

You'll know it worked when you see a text box that says **Enter your prompt here** and a button that says **Submit**.

Step 3 — Choose an ATS-friendly keyword finder model

In the text box, type or paste this model name exactly:

```
microsoft/Phi-3-mini-4k-instruct
```

Then press **Enter** or click **Submit**.

This model is a smart AI assistant that can read your job ad and resume and tell you which keywords are missing.

💡 *What is a model?* Think of it like a super-smart student who has read thousands of job ads and resumes. It knows what words to look for.

You'll know it worked when you see a chat window open with a welcome message from the AI.

Step 4 — Write your job ad and resume into the chat

In the chat box, type or paste this exact prompt (copy the whole block):

```
I need help finding ATS-friendly keywords for a job application.
Here is the job ad:
[PASTE JOB AD HERE]

Here is my resume:
[PASTE RESUME TEXT HERE]

List the top 10 keywords from the job ad that are missing from my resume.
```

Replace `[PASTE JOB AD HERE]` with the full text of the job ad you're applying for. Replace `[PASTE RESUME TEXT HERE]` with the text of your resume. To copy text from a PDF, open the PDF, select the text, and press **Ctrl+C** (Windows) or **Cmd+C** (Mac). Paste it into the chat with **Ctrl+V** or **Cmd+V**.

💬 Example:

```
I need help finding ATS-friendly keywords for a job application.
Here is the job ad:
We are looking for a digital marketing specialist with experience in SEO, Google
Analytics, and content creation.

Here is my resume:
I have 3 years of experience in digital marketing, including social media management
and email campaigns.
```

You'll know it worked when the AI replies with a list of keywords, like:

> The top 10 keywords missing from your resume are: SEO, Google Analytics, content creation, PPC, keyword research, on-page optimization, conversion rate optimization, Google Ads, Ahrefs, backlink building.

Step 5 — Ask the AI to rewrite a sentence with the missing keywords

Copy one of the missing keywords from the AI's reply. Then, in the same chat, type:

```
Rewrite this sentence from my resume to include the keyword [KEYWORD]:
"I have 3 years of experience in digital marketing, including social media management
and email campaigns."
```

Replace `[KEYWORD]` with the keyword you copied, like **SEO** or **Google Analytics**.

💬 Example:

```
Rewrite this sentence from my resume to include the keyword SEO:
"I have 3 years of experience in digital marketing, including social media management
and email campaigns."
```

You'll know it worked when the AI replies with a rewritten sentence that includes the keyword naturally, like:

> I have 3 years of experience in digital marketing, including SEO strategy, social media management, and email campaigns.

Step 6 — Save your updated resume

Copy the rewritten sentence from the AI's reply. Open your resume file (Word or Google Docs), find the sentence you want to update, and paste the new version over it. Save the file with a new name like `Resume_ATS_Updated.pdf`.

💡 *Why save a new file?* This way, you always have your original resume to go back to if you need to.

You'll know it worked when you have a new file saved with your updated resume text.

Step 7 — Run the process again to double-check

Repeat Step 4 with your updated resume. Paste the same job ad and your new resume text into the chat. Ask the AI:

```
List the top 10 keywords from the job ad that are still missing from my resume.
```

You'll know it worked when the AI replies with a shorter list of missing keywords, or says:

> All the top keywords from the job ad are now included in your resume.

Common mistakes

- **Mistake:** Pasting the job ad or resume as an image or screenshot. **Fix:** Copy the text from the PDF or Word file and paste it as plain text. Images can't be read by the AI.
- **Mistake:** Using a phone to copy and paste large blocks of text. **Fix:** On a phone, it's easy to lose your place. Use a computer or laptop if you can. If you must use a phone, paste smaller chunks of text at a time.
- **Mistake:** Expecting the AI to rewrite your whole resume perfectly in one go. **Fix:** The AI is great at spotting missing keywords, but it can't rewrite your whole resume in one sentence. Update one sentence at a time, then check again.

What this means for you

- **In everyday life:** You'll spend less time guessing which keywords to add to your resume. The AI does the hard work for you, so you can focus on the rest of your job application.
- **For your business or work:** If you're hiring, you can use this same method to check if your job ads include the right keywords. It helps attract candidates whose resumes will pass the ATS filter.
- **If you're just getting started:** This is the easiest first move to make your resume ATS-friendly. No coding, no expensive tools — just free AI and a few minutes of your time.

Try it now

Open your resume file and copy the first sentence. Go to huggingface.co, open the Inference API, and paste this prompt:

```
I need help finding ATS-friendly keywords for a job application.
Here is the job ad:
[PASTE A JOB AD YOU'RE APPLYING FOR]

Here is my resume:
[PASTE THE FIRST SENTENCE OF YOUR RESUME]

List the top 5 keywords from the job ad that are missing from my resume.
```

Press **Submit** and see what the AI suggests. You'll have your first ATS keyword list in under 2 minutes.

Write a Polite Complaint Email to a Company with DeepSeek

7 steps · Read online

> **In short:** DeepSeek can draft a polite, well-structured complaint email in under a minute — but you'll need to add the real details (order numbers, dates, names) and send it from your own account. It can't contact the company for you.

When something goes wrong with a product or service, the email you write in frustration usually gets ignored. A short, calm message that says what happened, what you want, and by when is far more likely to get a real reply. DeepSeek is great at shaping that kind of message — you describe the problem in your own words, and it turns your rambling into a clean, firm but polite email.

Before you start

- A phone, tablet, or computer with internet access.
- A web browser (Chrome, Safari, Edge, Firefox — any modern one). Or the DeepSeek app if you prefer.
- A rough idea of what went wrong: when it happened, what you bought, and what you'd like the company to do about it.
- Any order numbers, receipts, or screenshots handy (helpful, but not required).
- About 10 to 15 minutes total.

Step 1 — Open DeepSeek in your browser

Go to **chat.deepseek.com** in your web browser. The page loads a chat screen with a text box near the bottom labeled something like "Send a message" or "Message DeepSeek." No account is needed for a basic chat, so you can start typing right away. If you have the DeepSeek app installed, open it instead and you'll see the same chat layout. 💬 Example: type nothing yet — just confirm the box is there. You'll know it worked when the chat screen is empty and ready, with a blinking cursor in the message box.

Step 2 — Describe what happened in plain words

In the chat box, write a short, honest summary of the problem. Don't worry about grammar or politeness — just dump the facts the way you'd tell a friend. Include the company name, the product or service, when it happened, and what went wrong. DeepSeek reads everything in this message and uses it to draft the email later. 💬 Example: "I bought a coffee maker from BrewCo on March 12, order #44291. It arrived broken — the carafe was cracked and water leaks everywhere. I emailed them a week ago and nobody replied. I want a replacement or a full

refund." You'll know it worked when DeepSeek replies by asking a follow-up question or jumping straight into a draft. If it asks "What outcome do you want?", move to Step 3.

Step 3 — Ask for a polite complaint email

Type a clear follow-up: "Please turn this into a polite complaint email I can send to the company. Keep it firm but respectful, around 150 words, with a subject line." DeepSeek now produces a full email — greeting, body paragraphs, and a sign-off — usually in a couple of seconds. The subject line usually appears at the top. 💬 Example: "Please turn this into a polite complaint email I can send to BrewCo customer service. Keep it firm but respectful, around 150 words, and give me a subject line too." You'll know it worked when you see a complete email in the chat, with a "Subject:" line, a "Dear…" greeting, and a closing like "Sincerely."

Step 4 — Tell DeepSeek what outcome you want

The first draft is usually generic. Tell DeepSeek specifically what you want the company to do — a refund, a replacement, an apology, a callback, a coupon. The clearer you are, the more useful the next draft becomes. You can also mention a reasonable deadline, like "within 7 business days." 💬 Example: "Please rewrite it so I clearly ask for a full refund of $89 within 7 business days, and mention that if I don't hear back I'll file a chargeback with my bank. Keep the same polite tone." You'll know it worked when the new version includes your specific ask, in plain language, not a vague "I hope you can help."

Step 5 — Adjust the tone if needed

Read the email out loud. Does it sound like you? If it's too stiff, say "make it sound warmer and more conversational." If it's too soft and you want to sound firmer, say "make the tone more direct without being rude." DeepSeek rewrites the whole email in seconds, so don't be shy about asking two or three times. 💬 Example: "Make it a bit warmer — I'm not angry, just disappointed. Keep the firmness but add one sentence saying I have liked their products in the past." You'll know it worked when the new version feels like something you'd actually be comfortable sending under your own name.

Step 6 — Add your real details and read it carefully

Now do the human part. Open the email in your own mail app (Gmail, Outlook, Apple Mail — anything), and copy the draft over. Fill in the missing pieces: the actual recipient email (usually support@example.com or a contact form address), the real order number, dates, your name, and any return or tracking numbers. Read it once more to make sure it sounds like you and every fact is correct. 💬 Example: replace "[Your Name]" with your real name, "[Order #]" with the

real order number, and double-check the dollar amounts. You'll know it worked when the email reads as a complete, accurate message with no placeholder text like "[insert date here]."

Step 7 — Send the email from your own account

Hit send from your own email account. DeepSeek can't send emails for you — it's a chat tool, not a mail program. After sending, keep the message handy; if you don't hear back in the time you mentioned (often 5 to 7 business days), you can simply reply to your own thread with "Just following up on my email below." 💬 Example: write a short follow-up — "Hi, I'm following up on my email from March 20 about order #44291. I haven't received a reply yet. Please let me know how you'd like to resolve this." You'll know it worked when the email lands in your "Sent" folder and you have a copy you can quote later if needed.

Common mistakes

- **Being too angry in the prompt.** DeepSeek often mirrors the tone of what you give it. If you rant, you may get a rant back. Stay matter-of-fact and let the tool do the polishing.
- **Forgetting the specific ask.** A complaint that doesn't say "refund, replacement, or reply by Friday" often gets a generic "we're sorry" reply. Always name what you want.
- **Sending without reading.** AI drafts sometimes invent details — a wrong date, a fake policy, a tone that isn't yours. Always read it once, top to bottom, before sending.
- **No receipt of what you sent.** Screenshot or forward the final email to yourself so you have proof of what was said and when.

What this means for you

- **In everyday life:** A broken toaster, a missed delivery, a noisy hotel room, a wrong phone bill — any of these can become a 10-minute fix instead of a week of ignored messages.
- **For your business or work:** If you invoice or buy from other companies, the same trick works for late payments, missing deliveries, or poor service. A polite, specific email usually beats a phone call.
- **If you're just getting started:** Pick the smallest, most recent annoyance and try one email. You'll see how much easier it is when you don't have to start from a blank page.

Try it now

Think of one small frustration from the last month — a delayed package, a wrong charge, a rude reply from a store. Open DeepSeek, type two sentences about what happened, and ask it to turn that into a polite email. Read the draft, fix one detail, and save it in your drafts folder. You'll have a real, ready-to-send message in under five minutes.

Drafting Professional Presentations from Scratch Using Microsoft Copilot in PowerPoint

7 steps · Read online

Imagine turning a messy idea or a blank slide deck into a polished, ready-to-use presentation in under 10 minutes. With Microsoft Copilot in PowerPoint, you can draft outlines, generate slide titles, and fill in speaker notes—all from a simple prompt. This guide walks you through every step, from opening PowerPoint to exporting your finished slides.

Before you start

- A **Microsoft 365 account** (free or paid)
- The **PowerPoint desktop app** (Windows or Mac) or **PowerPoint on the web**
- A **device with internet access**
- Roughly **10–15 minutes** of your time
- Optional: a **Microsoft 365 Copilot license** (paid) for full features

> ⚠️ Note: Some Copilot features may require a paid Microsoft 365 Copilot license. If you don't have it, you'll still be able to use basic Copilot features like generating outlines and slide titles.

Step 1 — Open PowerPoint and start a new presentation

Open the **PowerPoint app** on your computer or go to PowerPoint on the web. Click **New blank presentation** to start with a blank slide deck.

You'll see a single slide titled "Click to add title." The ribbon at the top shows tabs like **Home**, **Insert**, and **Design**.

If you don't see the **Home** tab, look for a menu at the top with words like **File**, **Home**, or **View**.

You'll know it worked when you see a blank slide with a title and subtitle placeholder.

Step 2 — Open Copilot in the PowerPoint ribbon

On the **Home** tab in the ribbon, look for a button labeled **Copilot** (it looks like a sparkle or a chat bubble with stars). Click it to open the Copilot pane on the right side of your screen.

You'll see a chat box that says "Ask me anything about this document." This is where you'll type your instructions.

If you don’t see the Copilot button, make sure you’re using the latest version of PowerPoint. Look for **File > Account > Update Options > Update Now**.

You'll know it worked when the Copilot pane opens on the right with a blinking cursor ready for your prompt.

💬 Example: Type or paste:

```
Help me create a presentation about the benefits of remote work for a team of 20 software developers.
```

Step 3 — Give Copilot a clear prompt to generate an outline

In the Copilot chat box, type a clear, specific prompt. Include:

- Your **topic**
- Your **audience**
- Any **key points** you want covered

After you press Enter, Copilot will process your request and show a message like “Generating outline…” in the pane.

You’ll see a new section titled **Outline** appear in the Copilot pane, with 4–6 suggested slide titles and bullet points.

If the outline isn’t quite right, you can type:

```
Add a slide about work-life balance and another about team collaboration tools.
```

You'll know it worked when you see a structured outline with slide titles and bullet points in the Copilot pane.

💬 Example:

```
Create an outline for a 10-minute presentation on 'How Remote Work Boosts Developer Productivity'. Include slides on flexibility, tools, and team culture.
```

Step 4 — Insert the outline into your slides

Click the **Add to presentation** button under the outline in the Copilot pane. Copilot will automatically create new slides based on the outline.

You'll see your slide deck grow from one blank slide to 5–7 slides with titles and bullet points.

If the slides look too plain, don't worry—you can edit them later. If the Copilot pane disappears, click the **Copilot** button again to reopen it.

You'll know it worked when you see multiple slides with titles and bullet points in your deck.

Step 5 — Ask Copilot to generate speaker notes for each slide

Select the first slide. In the Copilot pane, type:

```
Add speaker notes to this slide explaining the benefits of remote work for developers.
```

Copilot will write 2–3 sentences of speaker notes under the slide. Repeat this for each slide by selecting it and asking Copilot to add notes.

You'll see a small note icon appear on each slide. Click the icon to read the notes.

If speaker notes don't appear, make sure you've selected the slide first. If Copilot doesn't respond, refresh the pane or check your internet connection.

You'll know it worked when you see speaker notes appear under each slide.

💬 Example: For a slide titled "Flexibility," Copilot might add:

```
Speaker notes: Remote work lets developers choose their most productive hours. Studies show 60% of developers report higher focus at home. Mention our recent productivity report from Q2.
```

Step 6 — Refine content with follow-up prompts

With the outline and notes in place, ask Copilot to improve or expand any slide. For example:

```
Rewrite the slide titled 'Team Collaboration Tools' to include specific examples like Microsoft Teams and GitHub.
```

Copilot will update the slide content in real time. You can also ask it to:

- Shorten the text
- Add a quote
- Include a data point

You'll see the slide content change immediately after you press Enter.

If Copilot doesn't update the slide, try clicking **Refresh** in the Copilot pane or re-selecting the slide.

You'll know it worked when the slide content updates to match your new prompt.

💬 Example:

```
Add a quote from Satya Nadella about remote work culture to the last slide.
```

Step 7 — Export or present your presentation

When you're happy with the slides, save your presentation (**File > Save As**). To present, click **Slide Show > From Current Slide** or press **F5**.

You'll see your slides full-screen, ready to present.

If you want to share it, go to **File > Share > Send a Copy** to export as PDF or PowerPoint file.

You'll know it worked when you can present or share a complete, professional-looking deck.

Common mistakes

- **Mistake:** Writing a vague prompt like "Make a presentation." **Fix:** Be specific: include topic, audience, and key points. Example: "Create a 5-slide presentation on AI tools for small business owners."
- **Mistake:** Expecting Copilot to write perfect content in one try. **Fix:** Use follow-up prompts to refine slides. Ask Copilot to shorten, add data, or include a quote.
- **Mistake:** Not checking speaker notes or slide content for accuracy. **Fix:** Always review Copilot's output. Edit or remove anything that doesn't fit your message.

What this means for you

- **In everyday life:** Need to present a project update at work? Use Copilot to draft slides in minutes instead of hours. Great for team meetings, client pitches, or internal reviews.

- **For your business or work:** Create consistent, professional presentations without hiring a designer. Use it for quarterly reviews, training decks, or onboarding materials.
- **If you're just getting started:** Start with a simple prompt and let Copilot do the heavy lifting. You'll save time and look polished—even if you're not a designer.

Try it now

Open PowerPoint, start a new presentation, and open Copilot. Type:

```
Help me create a 5-slide presentation about the benefits of remote work for a team of 20 software developers.
```

Press Enter. Review the outline, add it to your slides, and start presenting in under 10 minutes.

Digitize Handwritten Notes into Editable Text with AI

6 steps · Read online

Do you have notebooks full of brilliant ideas, meeting notes, or shopping lists that you can never quite find when you need them? This guide will show you how to quickly turn those physical scribbles into digital text using AI, so they're always searchable and editable, right on your phone or computer. It's for anyone who wants to bring order to their paper notes without retyping a single word. While AI is great at "reading" your handwriting, it can't understand *meaning* or summarize for you yet – it just converts the words.

Before you start

- **A smartphone or tablet with a working camera.** Ensure your device has enough battery life and storage for a few photos.
- **Access to an app that offers AI-powered OCR.** This means **Optical Character Recognition**, a technology that "reads" text from images. Think of it like a very clever digital assistant looking at your picture and identifying each letter and word, then turning it into text you can actually edit. The "AI-powered" part means it uses advanced artificial intelligence to understand different handwritings and layouts, making it much more accurate than older systems. Many free note-taking or scanning apps now include this feature, such as Google Keep (available on Android and iOS) or Microsoft OneNote (available on Android and iOS).
- **A free Google or Microsoft account.** You'll likely need one of these to use the apps suggested, though you probably already have one set up on your device. Ensure you're logged in before you start.
- **Roughly 5-10 minutes** to follow these steps for your first digitized note.

Step 1 — Prepare your notes for capture

Before you take a photo, you need to set up your handwritten notes in a way that helps the AI "read" them clearly. You'll see your physical notes lying on a surface. Your goal is to give the AI the best possible image to work with, as good, even lighting and legible handwriting make a big difference for the AI to understand your scribbles correctly. Think of it like a librarian trying to read tiny, smudged print – the clearer, the better!

Once prepared, your notes will be ready for you to take a photo. If your usual handwriting is quite messy, try to write a little slower and clearer for important notes you plan to digitize, as this significantly improves the AI's accuracy.

Example: Place your handwritten notes on a plain, uncluttered surface, like a clear desk. Position them under natural light from a window, if possible, or use a well-lit lamp to avoid

casting shadows over your writing. Make sure the paper is flat and smooth, with no creases. You'll know it worked when your notes are lying flat, look easy to read in the available light, and are clear of any distracting objects, ready for a photo.

Step 2 — Open your scanning app

Now it's time to open your chosen app. You'll see your device's home screen with various app icons. You need to locate and tap on the icon for your OCR app. For this guide, we'll use Google Keep as the primary example, which is free and widely available. Microsoft OneNote is another option that includes document scanning. Once opened, the app will load, often showing your existing notes or presenting a default view.

If you don't already have one of these apps, you'll need to download it from your device's app store (Google Play Store for Android or Apple App Store for iOS) and install it first. Look for the app icon, which usually looks like a colorful note for Google Keep or a blue notebook icon for Microsoft OneNote.

Example: Tap the "Google Keep" app icon on your home screen or app drawer (it looks like a yellow and white notepad). Alternatively, tap the "Microsoft OneNote" app icon (a blue notebook icon). You'll know it worked when the app opens and you see its main interface, such as a list of notes or a camera view, indicating it's ready for you to take action.

Step 3 — Take a clear photo

With your chosen app open, you'll typically see options to create a new note or use the camera. You need to find and tap the camera or scan icon to activate your device's camera within the app. The screen will then switch to a camera interface, often with specific modes like "Document" or "Scan" in apps like OneNote to help you capture text accurately.

Position your device directly above your notes, aiming for a straight-on shot without shadows, ensuring all your writing fits within the frame. Then, tap the shutter button. The app will capture the image, and some might even automatically detect the edges of your document and crop the image for you, making it look neater.

Example: In Google Keep, tap the plus (+) icon at the bottom of the screen, then choose the camera icon (it might be labeled 'Take photo' or 'Image'). In Microsoft OneNote, select 'Camera' from the options at the bottom, then hold your phone steady over your notes and tap the shutter button. You'll know it worked when you see a sharp, unblurred image of your notes presented within the app, often with the document edges automatically detected and cropped, ready for the next step.

Step 4 — Initiate text recognition with AI

With your clear photo captured in the app, it's time for the AI-powered OCR to do its magic. You'll typically see your captured image on the screen, possibly with a menu or a button related to text. You need to find and tap an option that tells the app to "read" the text from the image. This button might say something like "Scan text," "Recognize text," "Grab image text," or simply show a 'T' icon.

Once you activate this function, the AI will process the image you've taken, which usually takes just a few seconds. A new screen or panel will then appear, displaying your handwritten words now typed out as digital text. If you don't see an obvious "Grab image text" button, try tapping on the image itself, or look for a three-dot menu (⋯) or a gear icon (⚙️) that might contain the option you need.

💬 **Example:** In Google Keep, after your photo is captured, tap on the image to view it full screen. Then, look for a three-dot menu (⋯) in the top-right corner of the screen and tap the option labeled 'Grab image text'. In Microsoft OneNote, after taking a photo, tap the 'Text' icon that appears on the captured image to extract the text. You'll know it worked when a new text box or a separate screen appears, displaying your handwritten words now typed out as digital text, ready for you to review.

Step 5 — Review and make corrections

The AI is incredibly smart, but it's not perfect! Especially with unique handwriting, faint lines, or unusual symbols, it might make a few mistakes. You'll see the extracted text, usually in an editable field, displayed on your screen. Take a moment to read through this digitized text carefully and correct any errors.

As you make corrections, the text on your screen will update instantly, just like in any word processor. If the text isn't directly editable, look for an "Edit" button or a cursor you can tap to activate the editing mode. Remember, the AI is a great starting point, but it usually cannot perfectly capture every nuance of human handwriting, so your review is crucial for accuracy and ensures your digital notes are truly useful.

💬 **Example:** Carefully read through the digitized text. If you spot a word like "flow" incorrectly appearing as "flew," tap on "flew" and type "flow" to correct it. Also, remove any extra spaces, line breaks, or unwanted characters the AI might have picked up from smudges or background textures on the paper. You'll know it worked when you've corrected all the obvious mistakes, and the digital text accurately reflects your original handwritten notes, making them much clearer than your original scrawl.

Step 6 — Save or export your editable text

After reviewing and correcting your digitized notes, you can now save them or share them. You'll see options on your screen to save the note within the app, copy the text to your device's clipboard, or share it to other applications. This final step means your old scribbles are now **editable text** (you can change it whenever you like) and **searchable text** (your device can find it instantly when you search).

Once you take action, your app will usually confirm that the note has been saved, or the text will be transferred to your chosen destination. If you don't see an explicit "Save" button, look for an arrow pointing left (←) or a "Done" button, which will usually save your changes within the app you're using. Free versions of apps might have limits on how many notes you can save or what export options are available, but basic saving and copying are almost always free.

💬 **Example:** Tap the 'Save' or 'Done' button within your chosen app to store the note. Alternatively, look for a 'Share' icon (often a three-dot connection or an upward-pointing arrow) to copy the text to your phone's clipboard or send it directly to an email, a document app like Google Docs or Microsoft Word, or cloud storage like Google Drive or OneDrive. You'll know it worked when your app confirms the note is saved (e.g., you see it in your main notes list), or when you successfully paste the text into another application, ready for editing, searching, or sharing.

Common mistakes

- **Poor lighting or blurry photos:** Shadows or dim light make it incredibly hard for the AI to "read" your notes accurately, and blurry images are almost impossible for it to decipher. Always aim for bright, even light, and hold your phone steady when taking the picture. If the first attempt is blurry, simply delete it and retake the photo.
- **Not reviewing the extracted text:** It's easy to assume the AI got it all right, but even the best AI can make mistakes, especially with unique handwriting or complex layouts. A quick check and correction can save you from future confusion or embarrassment when you try to use the notes later. Always perform a quick read-through.
- **Using an app without a strong OCR feature:** Not all camera or note-taking apps have robust AI-powered OCR. If your results are consistently bad (lots of errors, missing text), try downloading a dedicated scanning app like Google Keep or Microsoft OneNote, which are known for good text recognition capabilities.

What this means for you

- **In everyday life:** Imagine quickly finding that recipe idea you scribbled down weeks ago, or having a digital copy of your kids' school notes that you can share instantly with a co-parent. No more rummaging through stacks of paper! This tool helps you keep

important information organized and accessible, whether it's a shopping list, a quick thought, or important reminders.

- **For your business or work:** Effortlessly convert meeting notes, brainstorming sessions, or client feedback from paper to a digital format. This makes them searchable, easy to share with colleagues, and simple to integrate into digital documents, saving you valuable time on transcription.
- **If you're just getting started:** Pick one short, clearly written note – like a shopping list or a phone number – and try to digitize it. It's a low-pressure way to see how well the AI works with your handwriting.

Try it now

Grab a piece of paper and quickly jot down three items on your grocery list. Now, open Google Keep (or Microsoft OneNote), take a clear photo of your list, and use the app's text recognition feature. In less than two minutes, you'll have your shopping list in an editable format, ready to paste into a messaging app or email!

How to turn a long PDF contract into a plain-English summary in 5 minutes

6 steps · Read online

> **In short:** Open any AI assistant that accepts PDF uploads, attach the file, and ask a few focused questions about deadlines, fees, and cancellation. You'll get a plain-English summary in minutes — but always read the original key clauses yourself, because AI summaries can miss important details.

Contracts hide the important parts deep in the small print: the auto-renewal clause, the 30-day cancellation window, the fee that "may apply" later. An AI assistant can pull all of that into a short, readable summary, but only if you ask the right questions. This guide shows the technique that works in any major AI assistant that accepts PDF uploads — ChatGPT, Claude, Gemini, Copilot, Le Chat, and others.

Before you start

- An AI assistant that can read PDF files you upload. Most major assistants do this on their free tier, but very long contracts (over roughly 100 pages) may need a paid plan. Check the upload limit on the assistant's pricing page if your document is unusually long.
- The contract itself, saved as a PDF on your computer or phone.
- About 10 minutes. The summary itself is fast; the careful read-through at the end takes a moment.
- A quiet moment to actually read what the assistant tells you.

Step 1 — Start a fresh chat and upload the PDF

Open your AI assistant and begin a brand-new conversation. A clean slate gives you the cleanest answer — older chats can confuse what the assistant reads in the new document. Once the chat is open, look for a **paperclip**, **attachment**, or **plus icon** near the message box. Tap that, then choose the PDF from your files. Some assistants also let you drag and drop the file straight onto the chat window.

If you don't see a paperclip, look for a **+ icon** or a button labeled something like **Add files** or **Upload**. On phone apps the icon often hides inside a small menu at the side of the text box.

💬 Example: "I'm uploading a contract. Please don't analyze it yet — I'll ask my questions one at a time."

You'll know it worked when the file name appears above the message box (or as a small chip beside it) and the assistant confirms it can see the document. If nothing happens after 30

seconds, the file may be too large — check the assistant's file-size limit before assuming the tool is broken.

Step 2 — Ask for a plain-English overview first

Before you drill into the details, ask the assistant to give you a one-paragraph summary of what the document is, who it's between, and what it covers. This grounds the rest of the conversation: once the assistant has the big picture in mind, its later answers tend to be more accurate. Type your question in the message box where the uploaded file is already attached, and press Enter or the send arrow.

If the answer is longer than a paragraph or feels like a sales pitch, ask it to "shorten to one paragraph and skip the legal jargon." If the assistant gives you only a vague intro, ask a follow-up: "What kind of contract is this — a service agreement, a lease, a sales contract, or something else?"

Example: "In plain English, what is this document, who are the parties, and what is the main thing being agreed to?"

You'll know it worked when the assistant gives you back a short paragraph (3 to 6 sentences) that names the type of contract, the people or companies involved, and the core purpose. If any of those three is missing, ask one more focused question before moving on.

Step 3 — Ask the three questions that actually matter

Now the real work. In the same chat, type three focused questions, one after another: **deadlines** (when does this start, when can I leave, when does it auto-renew), **fees** (what does it cost, what changes after the first period, what extra charges might apply), and **cancellation** (how do I end this, what's the notice period, what happens if I miss it). Ask all three in one message so the assistant can answer them together. Press Enter and wait — usually under a minute.

If the assistant answers only one of the three, send the missing questions in a fresh message. If the answer feels too short, ask "Please list every date, dollar amount, and cancellation step you can find in the document."

Example: "From this contract, list: 1) every deadline or time limit I need to know, 2) every fee, charge, or price that might apply, and 3) exactly how I cancel, including any notice period. Use bullet points."

You'll know it worked when the assistant gives you a structured, easy-to-scan answer with three clear sections. If any answer is "I can't find that in the document," ask it to point to the page or section it looked at — sometimes the answer is there and the assistant missed it.

Step 4 — Push for the things people forget to ask

Once you have the big three, push one more layer. Contracts hide risk in the small print: what happens if you miss a payment, who pays legal fees in a dispute, whether the company can change terms mid-contract, and which country's laws apply. Ask these as a follow-up message in the same chat. The assistant now has the document **and** your earlier answers in its context window (the conversation history it can refer to), so it can build on what it already told you.

If the assistant's answer feels generic, paste in the actual clause: "I'm looking at section 8.2. What does this mean in plain English, and what's the worst-case scenario for me?"

💬 Example: "What happens if I miss a payment, who covers legal fees in a dispute, can the company change the terms mid-contract, and which state's laws apply? Point me to the section for each."

You'll know it worked when the assistant points to specific sections or page numbers and explains each risk in a sentence a non-lawyer can follow. If it just says "this is a standard clause" without explaining the actual risk, ask "what's the worst that could happen to me, in concrete terms?"

Step 5 — Ask for a one-page checklist you can save

Now ask the assistant to compress everything into a short, printable checklist you can keep next to the contract. A checklist is easier to skim than a long summary, and it's the format most people actually re-read later. Ask in the same chat — the assistant still has the document and your earlier conversation. The result should fit on one screen: action items, key dates, and money amounts, plain and short.

If the checklist comes back too long, ask it to "cut this to the 10 most important items I'd forget in six months."

💬 Example: "Summarize everything we just found into a one-page checklist I can save: dates I must remember, amounts I'll be charged, and the steps to cancel. Keep it under 15 lines."

You'll know it worked when you can copy the answer into a notes app and read it in under a minute. If the checklist still feels like a wall of text, ask the assistant to put each item on its own line with the date or amount in bold.

Step 6 — Spot-check three sections in the original PDF

AI summaries are good, but they're not perfect. The assistant can misread a clause, skip a footnote, or paraphrase something in a way that changes the meaning. Before you rely on the summary, open the original PDF and read three things yourself: the **cancellation clause**

(usually labeled "Termination"), the **fees section** (often "Fees," "Charges," or "Payment Terms"), and any **date or deadline** the assistant flagged. If those three match what the assistant told you, the rest of the summary is very likely accurate.

If you can't find the section the assistant mentioned, use the PDF's search function (Ctrl+F on a computer, or the magnifying glass icon in most phone PDF apps) and type a keyword from the clause. If something doesn't match, send a follow-up question: "Section X says something different — can you re-read and reconcile?"

💬 Example: "Cancellation clause says 30 days, fees say a charge after year one, and auto-renewal is 60 days before expiration — can you re-read sections 7, 9, and 12 and tell me if your summary matches?"

You'll know it worked when the three sections you read match the summary, and you feel comfortable either signing the contract or deciding not to. If anything is off, send the assistant the exact text and ask it to explain the difference — that's how you catch the small print that changes everything.

Common mistakes

- **Asking one giant question.** A single "summarize this contract" prompt usually gives a vague, hard-to-scan answer. Splitting it into overview, the three big questions, the hidden risks, and a checklist gets you a result you can actually use.
- **Trusting the summary without checking.** AI assistants can misread clauses, especially around fees and dates. Always read the cancellation, fees, and deadline sections yourself before you sign. The summary is a map, not a replacement for the territory.
- **Letting a long chat get messy.** After 20 or 30 messages, the assistant can lose track of what it told you earlier. If the conversation gets long, copy the checklist into a new chat, re-upload the PDF, and start fresh. Cleaner chats give cleaner answers.
- **Skipping the free-tier limits.** Some assistants let you upload small files on the free plan but cap the size on paid plans. If your upload fails, check the assistant's file-size limit before assuming the tool is broken.

What this means for you

- **In everyday life:** Leasing a new apartment, renewing a phone plan, signing up for a gym membership — the same five-step question set works for any contract. Once you've done it twice, the second one takes about three minutes.
- **For your business or work:** Vendor agreements, client contracts, partnership letters, freelancer terms — all of these can be summarized the same way. The checklist step is

especially useful when you need to brief a partner or a manager who hasn't read the document.

- **If you want to go further:** Try a second AI assistant on the same PDF and compare the two summaries. Two independent reads catch more than one — and they don't always agree on which clauses matter most. Just don't paste a confidential contract into a tool your company hasn't approved for outside use.

Try it now

Pick the last contract you signed — apartment lease, phone plan, gym membership, anything with more than five pages. Upload it to your AI assistant and ask the three big questions from Step 3. Save the answer somewhere you can find it in six months, when the renewal notice arrives. You'll never read a long contract the same way again.

Unlock Document Insights with Google NotebookLM

6 steps · Read online

Imagine having a smart assistant that can read all your scattered notes, long reports, or important documents and then answer your questions using *only* the information you've provided. That's the power of Google NotebookLM. By the end of this guide, you'll be able to upload your own documents, ask intelligent questions, and quickly pull out key information or summaries without sifting through pages yourself. It’s perfect for students, researchers, or anyone dealing with lots of text.

Before you start

- **A Google Account:** You'll need a standard Google account (like the one you use for Gmail or Google Drive) to access NotebookLM.
- **A Web Browser:** Access NotebookLM through a web browser on your computer. While it works on tablets, a computer provides the best experience for this guide.
- **An Internet Connection:** A stable connection is needed to upload documents and interact with the AI.
- **A few documents:** Have 1-3 documents (e.g., a PDF report, a text file of notes, or a Google Doc) ready to upload. These will be your "sources".
- **Cost:** NotebookLM is currently available for free.
- **Time:** This guide will take about 10-15 minutes to follow.

Step 1 — Access Google NotebookLM

First, you'll need to open Google NotebookLM in your web browser. This is where you'll create a dedicated space for your documents and questions. You'll see a welcome screen or an overview of your existing Notebooks, which are like individual project folders. If you're logged into your Google account, it should open directly. If it asks you to sign in, use your Google account details.

If it looks different: If you see a welcome tour or introductory pop-ups, simply click through them to get to the main interface. Look for a "Create new Notebook" or a large "+" button to start. **You'll know it worked when:** You see the NotebookLM interface, likely with a main content area and a sidebar, and perhaps a prompt to create your first Notebook.

Step 2 — Create a New Notebook

A "Notebook" in NotebookLM is like a digital binder where you keep all the related documents for a specific topic or project. It helps keep your information organized and ensures the AI only uses the sources relevant to that project. To create one, look for a prominent button or option to start a new project.

What you see: On the main screen, you'll typically find a large "New Notebook" button or a similar option in a sidebar. Clicking it usually prompts you to give your new Notebook a name. **What happens next:** A small pop-up window will appear, asking you to name your new workspace. Give it a descriptive title that reflects the documents you'll be adding, so you can easily find it later. **If it looks different:** If you don't see a "New Notebook" button, look for an icon like a "+", a folder, or text that says "Create". The exact placement might vary, but it's usually at the top or in a left-hand navigation pane. 💬 Example: `My Research Notes on Renewable Energy` **You'll know it worked when:** A new, empty Notebook with your chosen name appears on your screen, ready for you to add "sources".

Step 3 — Add Your Source Documents

"Sources" are the documents you want NotebookLM to read and analyze. These can be PDFs, Google Docs, text files, or even web pages saved as PDFs. The AI will then use *only* these sources to answer your questions, acting like a super-smart indexer for your own content. You can usually add multiple files to a single Notebook.

What you see: Inside your new Notebook, you'll see an area labeled "Sources" or "Add sources", often with an "Upload" button or an option to connect to Google Drive. Click the "Upload" button. **What happens next:** A file explorer window (or Google Drive selector) will open, allowing you to choose documents from your computer. Select one or more relevant files. NotebookLM will then process these files, which might take a few moments depending on their size and number. **If it looks different:** Look for an icon representing a document, a paperclip, or a cloud upload symbol if "Upload" isn't immediately visible. If you're using a Google Doc, there might be a direct integration option for Google Drive. **You'll know it worked when:** Your uploaded documents appear listed under the "Sources" section in your Notebook. You might see a "Processing" status initially, which will change to "Ready" when the AI has finished reading them.

Step 4 — Ask a Question Based on Your Sources

Now that NotebookLM has your documents, you can start asking it questions. Remember, the key here is that the AI (which is powered by a large language model, or LLM – a sophisticated computer program that understands and generates human-like text) will *only* refer to the information within the sources you provided. This prevents "hallucinations" – when an AI confidently makes up facts – and ensures its answers are grounded in your specific content.

What you see: In the main area of your Notebook, there's usually a chat box or a "Ask your sources" field, similar to what you'd find in a regular chat AI. Type your question here. **What happens next:** The AI will process your question, consult *only* the uploaded sources, and then generate an answer directly below your prompt. It will often include small citations (like "Source

1, Page 3") pointing back to where it found the information in your documents. **If it looks different:** The chat box might be labeled differently, perhaps "Query sources" or "Enter prompt". Just ensure you're typing your question into the area designated for interacting with your documents. 💬 Example: `What are the key challenges mentioned in the "Renewable Energy Report.pdf"?` **You'll know it worked when:** You receive a detailed answer from NotebookLM, specifically referencing your uploaded documents, with citations next to the relevant parts of its response.

Step 5 — Explore and Refine with Follow-up Questions

The conversation doesn't have to end with the first answer. NotebookLM is designed for back-and-forth interaction. If an answer is too broad, too short, or if you want to dive deeper into a specific point, you can ask follow-up questions. This helps you refine your understanding and extract exactly what you need.

What you see: After receiving an answer, the chat box typically remains open, or a "New prompt" button appears. You can type another question directly into this field, building on the previous interaction. **What happens next:** NotebookLM will take your new question and the context of your previous conversation into account, again pulling information solely from your sources. The new response will appear below the previous one, continuing the thread. **If it looks different:** If the chat box clears, simply re-type your follow-up question. The AI remembers the context of the current conversation within the Notebook. 💬 Example: `Can you summarise the solutions proposed for reducing these challenges?` **You'll know it worked when:** You get a new answer that directly addresses your follow-up question, further refining the information based on your sources, and continuing the conversational flow.

Step 6 — Generate an Overview or Summary

One of NotebookLM's most powerful features for a beginner is its ability to quickly generate an "Overview" or summary of your entire Notebook or specific documents. This is incredibly useful for getting the gist of large amounts of text without reading every word, or for quickly preparing for a meeting or study session.

What you see: Look for a button or option in your Notebook, often labeled "Overview", "Summarize Notebook", or "Generate Summary". It might be in the sidebar or at the top of the main conversation area. Click it. **What happens next:** NotebookLM will quickly process all the documents within that specific Notebook and generate a concise summary or a list of key topics and insights. This often appears as a new section or card in your Notebook interface. **If it looks different:** The feature might be represented by an icon, such as a bulleted list or a document with an arrow, or be found under a "Tools" or "More actions" menu. **You'll know it worked**

when: A new, structured summary or overview appears in your Notebook, highlighting the main points and themes from all your uploaded sources.

Common mistakes

- **Expecting general knowledge answers:** NotebookLM *only* answers from your uploaded sources. If you ask a question not covered in your documents, it won't be able to help, and might even tell you it can't find the information. **Fix:** Always ensure your question can be answered by the content you've provided.
- **Uploading unsupported file types:** While NotebookLM supports common formats like PDFs and Google Docs, it can't read every type of file (e.g., image files, spreadsheets, or proprietary formats). **Fix:** Convert your documents to a supported format like PDF or plain text before uploading.
- **Asking overly broad or vague questions:** A very general question might lead to a generic or incomplete answer, even with relevant sources. **Fix:** Be specific with your prompts. If you want details on a particular section, mention it. If you want a specific type of information, ask for it directly (e.g., "List the pros and cons of X").

What this means for you

- **In everyday life:** If you're planning a trip, you could upload travel guides and articles into a Notebook and ask it to highlight the best restaurants or activities in your destination, saving hours of reading.
- **For your business or work:** You can drop in client reports, internal policies, or research papers and quickly get summaries, compare different sections, or pull out specific data points for presentations, making your research much faster and more accurate.
- **If you're just getting started:** Try uploading a favorite recipe or a hobby-related article. Then, ask NotebookLM simple questions like "What are the main ingredients?" or "What's the first step?" to get a feel for how it extracts information.

Try it now

Open Google NotebookLM, create a new Notebook, and upload two or three of your own PDF documents, like a household bill, a product manual, or an article you've saved. Then, ask it: "What are the key points in these documents?"

Summarize Any Long Document or Article with AI

6 steps · Read online

Tired of scrolling through a 20-page policy, a lengthy news article, or a long email thread? This guide will show you how to use a free AI tool to quickly get a short, reliable summary. In just a few minutes, you'll be able to cut through the noise and focus only on what truly matters.

Before you start

- **A device with internet access:** A computer, tablet, or smartphone is perfect.
- **A free AI chat tool:** We'll use general instructions that work for popular options like ChatGPT (find it at `chat.openai.com`) or Gemini (at `gemini.google.com`). You'll need to create a free account and log in.
- **A document or article to shorten:** This could be a webpage, a PDF where you can select text, an email thread, or any plain text.
- **Rough total time:** About 6 minutes to follow all the steps.

Step 1 — Copy your document text

First, open the document or article that you want the AI to summarize. Once it's open, you need to select all the relevant text. On a computer, you can often do this by clicking your mouse at the start of the text and dragging to the end, or by using a keyboard shortcut like `Ctrl+A` (Windows) or `Cmd+A` (Mac) to select everything. With the text highlighted, press `Ctrl+C` (Windows) or `Cmd+C` (Mac) to copy it to your **clipboard**. The clipboard is a temporary storage area on your device, like a hidden notepad, where copied items wait to be pasted.

You'll know it worked when the text is highlighted (often in blue or gray), and you've used the copy command. Although nothing visibly changes on your screen after copying, the text is now safely on your clipboard, ready for the next step. If you can't select text, try opening the document in a different program (e.g., a web browser for PDFs) or look for a "Select All" option within the application's menu.

💬 Example: Open a long email thread. Press `Ctrl+A` (Windows) or `Cmd+A` (Mac) to select the entire email body, then press `Ctrl+C` (Windows) or `Cmd+C` (Mac).

Step 2 — Open your AI chat tool

Next, open your preferred free AI chat tool in your web browser or dedicated app. For instance, you might type `chat.openai.com` into your browser's address bar for ChatGPT, or `gemini.google.com` for Gemini. If you're using a phone or tablet, simply launch the app. If you're prompted, make sure you're logged into your free account.

You'll typically see the AI tool's main chat interface, which usually features a large, empty text box at the bottom of the screen. This is where you'll type your messages to the AI. The tool is now waiting for your instructions. If your screen looks different, or you see old conversations, look for a "New Chat" button, a plus sign icon (`+`), or a similar option to start a fresh, blank conversation.

💬 Example: Open your web browser and navigate to `gemini.google.com`. Log in to your account if prompted, then ensure you have a blank conversation ready.

You'll know it worked when you see an empty chat box, ready for you to type your first message, usually labeled "Message [AI name]..." or "Type your prompt here...".

Step 3 — Craft your summary request (Prompt)

Now, in the AI's chat box, you'll type your instruction for the summary. This instruction is called a **prompt** – it's like giving a clear job description to the AI. Be specific about the desired length, format (e.g., bullet points), and tone for your summary. A good prompt helps the AI understand exactly what you need.

You'll see the words you type appearing in the chat box. At this stage, the AI is just receiving your input, and you haven't sent anything yet. If the chat box shows pre-written suggestions, simply ignore them and type your own request. Being clear upfront saves time later.

💬 Example: Type this into the chat box (do not press Enter or Send yet):

> "Please summarize the following document in three concise bullet points. Keep the language neutral and factual, avoiding any personal opinions or interpretations."

You'll know it worked when your complete summary request is typed into the box, clearly visible, but you haven't sent it to the AI yet.

Step 4 — Paste your text and send

After typing your summary request, move your cursor to the end of your prompt, then immediately paste the long document text that you copied in Step 1. You can do this by pressing `Ctrl+V` (Windows) or `Cmd+V` (Mac). The pasted text will appear right after your instruction in the same chat box. Once both your prompt and the entire document text are in the box, press Enter or click the "Send" button (often a paper plane icon) to send everything to the AI.

You'll see your prompt and the long document text appear in the chat history. After a moment, the AI will start generating its response. This processing is called **inference**, which is when the AI "thinks" about your request and creates an output. If your document is extremely long (e.g., hundreds of pages), the AI might tell you it's too much text for its **context window**. The

context window is like the AI's short-term memory – the amount of information it can "remember" and process at once. If this happens, you might need to break your document into smaller chunks and summarize them separately, or consider a paid plan, as some free versions may have usage restrictions.

Example:

> "Please summarize the following document in three concise bullet points. Keep the language neutral and factual, avoiding any personal opinions or interpretations.
>
> [Now, paste your copied document text here, directly after the prompt.]"

You'll know it worked when you see your message disappear into the chat history, and the AI begins to type out its summary, usually appearing as new text below your input.

Step 5 — Check and refine the summary

Once the AI has generated its summary, carefully read through the bullet points it provided. Compare them against the original document to ensure they are accurate, don't miss any crucial details, and match the tone you requested. Sometimes, AI can "hallucinate" – meaning it confidently makes up information that isn't present in the original text, or misinterprets something. If anything feels off, needs clarification, or isn't quite what you wanted, simply type a follow-up question or instruction in the chat box.

You'll see the AI's initial summary in the chat. When you ask a follow-up question, your new prompt will appear, and then the AI will try to adjust its response based on your feedback. If the AI completely misunderstands or gives a poor response, you might need to try rephrasing your initial prompt more clearly or start a new chat conversation altogether.

Example:

> "The third bullet point mentions 'Section 5'. Please confirm the exact wording of Section 5 from the original text and update that bullet point to be more precise. Also, can you ensure the overall tone is strictly professional?"

You'll know it worked when the AI responds with an updated summary that addresses your feedback, making it more accurate and aligned with your specific needs.

Step 6 — Finalise and use your summary

Once you are completely happy with the AI's summary, it's time to make it your own. You can copy the final bullet points from the chat interface to your clipboard. Most AI tools allow you to highlight the text with your mouse (click and drag) and then use the copy command (`Ctrl+C` or

`Cmd+C`). Some tools may also have a specific "Copy" button or icon next to the AI's response to make this easier.

You'll see the polished, accurate summary on your screen. After copying, the text is on your clipboard, ready to be pasted into another application, like an email, a report, or your personal notes. If there isn't an obvious "copy" button, simply highlight the AI's text and use your device's standard copy function.

💬 Example: Paste the finalised bullet points into a new email, a team chat application like Microsoft Teams, or a note-taking app on your phone to share or save.

You'll know it worked when the summary text is successfully copied (you can test this by pasting it into a simple text editor like Notepad or Notes) and you can then paste it into your desired application.

Common mistakes

- **Too vague a request:** Asking for "a summary" without specifying length or format can result in a lengthy, unstructured response. **The fix:** Always be specific. Say "in three bullet points," "one short paragraph," or "a summary suitable for an executive" in your prompt.
- **AI "hallucinations" or opinions:** AI tools can sometimes confidently present incorrect information or add their own interpretations that aren't in the original text. **The fix:** Always include phrases like "keep the language neutral and factual; do not add your own opinions" in your initial prompt, and critically review the summary against the original document. Don't be afraid to ask for corrections.
- **Text too long for the AI:** Trying to summarize extremely long documents (e.g., an entire book or a very dense legal document) might hit the AI's **context window** limits. **The fix:** Modern AI tools have much larger context windows than before, but extremely long documents may still require breaking them into smaller sections, or using a paid plan if available.

What this means for you

- **In everyday life:** The next time you receive a long work email, a detailed news article, a community newsletter, or even a lengthy recipe blog post, you can quickly get the main points without reading every single word. This saves you valuable time and helps you stay informed efficiently.
- **For your business or work:** Condense lengthy reports, meeting minutes, customer feedback surveys, or policy documents into concise, actionable bullet points. This can help you prepare for presentations faster, communicate key information efficiently to colleagues, or quickly grasp the core of complex documents. Remember, while AI is

incredibly helpful, it cannot provide legal or medical advice, and human review is crucial for critical documents.

- **If you're just getting started:** Pick one long email thread you've received today, or a lengthy article you saved to read later. It's an easy first step that will quickly show you the power of AI summaries and build your confidence in using these tools.

Try it now

Open any long email thread you've received in the last week

How to Prepare for a Job Interview Using a Pasted Job Post

6 steps · Read online

> **In short:** Paste a job listing and a short note about your background into your AI assistant, and you'll come away with ten likely interview questions, two talking points for each, and five thoughtful questions to ask the interviewer — typically in under ten minutes of back-and-forth.

A good interview answer sounds rehearsed without sounding rehearsed. The trick is having specific stories and numbers ready before you sit down — and that's exactly what an AI assistant is good at when you give it the right inputs. In fifteen minutes, you can walk into a room with three lists ready: questions they'll probably ask, talking points for each, and questions you'll ask them. It can draft the prep, but it can't read the room or know which story from your past will actually land — that's still on you.

Before you start

- **An AI assistant you already use** — ChatGPT, Claude, Gemini, Copilot, or any other chat-style assistant. The steps are the same.
- **The full job description** copied to your clipboard. Don't paraphrase yet; the more detail the assistant sees, the sharper the output.
- **A short paragraph about you** — your current role, years of experience, and one or two wins you'd actually mention in an interview.
- **About 15 minutes** — enough time to do the back-and-forth and skim the result.

Step 1 — Open a fresh chat and name it after the role

Click **New chat** (or **New conversation** — look for a pencil-and-paper icon or a + sign) in your assistant. Before you type anything, give the chat a name you'll recognize later, like "Acme Senior Designer interview" or just the company plus the role.

Why a new chat? Older messages from unrelated work can leak in and confuse the assistant. A fresh, named thread keeps everything you build here tied to this one interview — easy to find the night before, easy to revisit if a second round is scheduled.

💬 Example: name it "Acme – Marketing Lead interview prep."

You'll know it worked when a new empty chat opens with the name you chose (or with an editable title field at the top of the screen).

Step 2 — Paste the job description and your background in one message

In your new chat, send a single message that contains two clearly labeled blocks: the job description, and a paragraph about you. Use clear labels like **"JOB POSTING:"** and **"ABOUT ME:"** so the assistant knows where one ends and the other begins. Very long descriptions (over a few thousand words) may need to be split into two messages if your app refuses the paste.

What happens next: the assistant usually replies with a short acknowledgement ("Got it — ready when you are") or jumps straight into a summary. Either is fine.

If it looks different: if pasting a long posting fails, try uploading the posting as a file — most assistants accept PDFs and text files in the same chat where you usually send text.

💬 Example:

```
JOB POSTING:
[Paste the full job description here — responsibilities, requirements, and any "nice to haves"]

ABOUT ME:
[Paste 4–6 sentences: current role, years of experience, two recent achievements with numbers if you have them]
```

You'll know it worked when the assistant confirms it has the context — usually with a brief recap, a clarifying question, or simply by being ready to answer your next prompt.

Step 3 — Ask for ten likely interview questions

Send a single, specific prompt asking for ten questions the hiring manager is likely to ask, based on what you just pasted. Ask for a mix of categories — behavioral, technical, and situational — so the list isn't one-note.

Why ten? Fewer than that and you may miss the ones that come up. More than ten and you'll feel overloaded. Ten is the sweet spot: enough to practice, small enough to actually read.

💬 Example: "Based on the job posting and my background above, give me ten interview questions a hiring manager is most likely to ask me. Mix behavioral, technical, and situational questions. For each, add one sentence on what they're really trying to find out."

If it looks different: if the list comes back generic, follow up with "make these specific to [the company or team from the posting]" and ask for a fresh list. A vague list is usually a sign the original prompt needed more context, not a sign the assistant is broken.

You'll know it worked when you have a numbered list of ten specific, role-tied questions — not stock ones like "tell me about yourself."

Step 4 — Get two concrete talking points for each question

Without starting a new chat, follow up with a second prompt asking for two concrete talking points under each of the ten questions. Ask for stories, examples, or numbers drawn from your own background — and tell the assistant to flag any question where your pasted background looks thin and you'll need to dig deeper yourself.

What happens next: the assistant writes a structured answer, often with a short headline per question and a couple of ideas underneath. This is the moment the prep actually becomes useful.

If it looks different: if the response is too long to scan, add "keep each talking point to two sentences max" in your next prompt, or ask it to rewrite the same list more concisely.

💬 Example: "Now for each of those ten questions, give me two concrete talking points I could use, drawing from my background. Keep each point to one or two sentences. Flag any question where my pasted background seems thin."

You'll know it worked when each of the ten questions now has two short, specific points you could picture yourself saying out loud.

Step 5 — Ask for five smart questions to ask back

In the same chat, ask for five questions you, the candidate, should ask the interviewer. Ask for ones that show you did your homework — about the team, the goals, or how the role is measured. The point is to leave the conversation with real signals, not rehearsed fluff.

💬 Example: "Give me five strong questions I, the candidate, can ask back at the end of the interview. Make them specific to this role and company. Skip anything generic like 'what's the culture like'."

If it looks different: if any of the questions come back feeling templated, follow up with "replace any you can't tailor to this exact role" — most assistants will rewrite them on the spot.

You'll know it worked when you have five questions you wouldn't have thought of on your own, and each one is something you'd genuinely want answered.

Step 6 — Pin or save the chat and skim the highlights

Take thirty seconds to skim the chat and star, pin, or copy the talking points that felt most true to you. Different apps handle this differently — look for a **star icon** next to an answer, a **pin**

option in a menu, or just copy the best points into a notes app where you'll actually see them again.

Why bother? Two reasons. First, you'll want to reread the chat the morning of the interview. Second, picking the strongest two or three talking points now means you practice those — not all twenty.

If it looks different: if your assistant has no pin or star option, the chat's title alone usually makes it easy to find later. Or paste the talking points into your phone's notes app under the company name.

💬 Example: pick the talking points that felt truest to your real work — maybe one about leadership, one about a measurable result, one about a hard problem you solved — and re-read them once before you walk into the interview.

You'll know it worked when you can open the chat or your notes in under ten seconds and find the three talking points you actually plan to use.

Common mistakes

- **Pasting only the company blurb.** If the description is six paragraphs long and you paste just the first one, the assistant's questions will be vague. Paste the whole thing — responsibilities, requirements, and any "nice to haves" — so the assistant has the full picture.
- **Asking for "interview questions" without context.** A generic request gets generic answers. Always include both the job description and your background in the same chat; that's what makes the list specific to *this* role and *your* level.
- **Stopping after the first list.** The ten questions are the start, not the end. The follow-up that asks for talking points is where the real prep happens. Skip it and you'll have a list of questions with nothing ready to say.
- **Treating the output as a script.** The assistant may suggest a story you don't actually have. Treat the talking points as a draft, not a script — swap in your own real examples whenever you can.
- **Reusing the same chat for unrelated work.** Once you start mixing topics, the assistant's answers drift. Keep this chat dedicated to one interview until you're done with it.

What this means for you

- **In everyday life:** the same recipe works for any "I have a document and I want to be ready" moment — a school presentation, a doctor's appointment where you want a list of questions, or a tough conversation with a landlord. Paste the source, share your context, ask for what you need.

- **For your business or work:** if you ever sit on the other side of the table, paste the job description and a candidate's resume into the same flow and you'll get a quick read on likely strengths, gaps, and good first-round questions — useful when screening applications.
- **If you want to go further:** the next level is saving this whole structure as a **Custom Instruction** or a saved template in your assistant, so the next time you prep, you only paste the job posting and your background — the rest of the structure is already loaded. Many assistants also let you create a dedicated "Project" or named workspace for one company, so multiple chats (prep, follow-up, salary research) live together.

Try it now

Open a new chat in the assistant you already use, name it after a real interview you have coming up (or one you're imagining), and paste one real job description plus four sentences about your background. Then send the prompt from Step 3. Ten minutes from now, you'll have a list you didn't have ten minutes ago.

Write a Resume From Scratch (Even With No Work Experience)

6 steps · Read online

> **In short:** A resume is a one-page story of what you've done, what you're good at, and what you want next — and you can build one from school, volunteering, caregiving, side jobs or hobbies. You don't need formal work experience; you need to describe your real life in the language employers recognize.

Writing a resume from zero feels harder than it is. Most "first resumes" stay blank because people think they have nothing to put on them. In reality, you've already handled deadlines, solved problems, dealt with people, and learned new things — you just haven't framed that work as work yet. This guide walks you through turning everyday experience into a clean, honest resume a hiring manager can read in under a minute.

Before you start

- **A computer or phone** with a notes app, Google Docs, Microsoft Word, or a free resume builder (look for one by searching "free resume template").
- **About 60–90 minutes** for your first draft. After that, updates take 10 minutes.
- **A list of what you've actually done** — school, classes, projects, babysitting, family jobs, club roles, volunteering, online courses, certifications. Don't filter yet; just dump everything.
- **A job (or two) you'd like to apply to.** Even a saved job posting helps you match words later.

Step 1 — Pick a clean, one-page layout

Open Google Docs and click **Template Gallery** in the top-right, or search for "resume" in Microsoft Word's template search. Pick something simple: name at the top, contact line below it, then clear section headings (think "section headers" — the bold labels that divide your resume into parts like Experience and Education). Avoid heavy designs, photos, or color blocks; hiring managers scan, they don't admire.

If templates feel overwhelming, just open a blank document and type your name big at the top. You'll add structure next. If you don't have Word or Docs, free builders like the ones from job search sites work too — search "free resume builder" and pick one with no signup wall.

Example: section headers to copy — `Summary`, `Experience`, `Education`, `Skills`, `Projects`.

You'll know it worked when you have a blank one-page document with your name at the top and five empty sections waiting for content.

Step 2 — Write a 2-sentence summary at the top

Under your name and contact info, add a **summary** — two short sentences that say who you are and what you're aiming for. Think of it as the cover of a book: it tells the reader what to expect inside. Use plain words, not job titles you haven't earned.

If you've never had a job, your summary can lean on what you're studying, what kind of work you want, and one strength. Keep it specific — "hardworking" doesn't say much, but "organized, comfortable with spreadsheets, and quick to learn new tools" paints a picture. If you don't know exactly what you want yet, write a **target** instead: "Looking for an entry-level role in retail, admin, or customer service."

💬 Example: "High school senior finishing in May 2026. Strong with customers, cash handling, and staying calm under pressure. Looking for part-time retail or front-desk work."

You'll know it worked when someone reading just those two sentences could guess what kind of job you'd be good at.

Step 3 — Turn life experience into an "Experience" section

Create an **Experience** section heading and list anything where you showed up, did something, and someone (even a family member or teacher) noticed. Jobs count, but so do babysitting siblings, helping at a family store, school projects, volunteering at a food bank, mowing lawns, Discord servers you moderate, or a class group you led. For each entry, use this format: **Title — Where — Dates** (even if dates are approximate), then 2–4 short **bullet points** underneath describing what you did.

Each bullet should follow a simple recipe: *did what, using what skill, with what result*. You don't need fancy words; you need honest ones. If you watched three younger cousins every afternoon while your parent worked, that's "supervised and cared for three children daily, planned meals, handled minor conflicts, and maintained a consistent routine." That's a real bullet.

💬 Example (babysitting): `Cared for three children (ages 4–9) after school each day. Prepared snacks, helped with homework, and organized weekly craft activities.`

💬 Example (school project): `Led a 4-person team on a science fair project. Researched topics, delegated tasks, and presented findings to a panel of 5 judges.`

You'll know it worked when your experience list has 3–6 entries and every bullet describes something you actually did.

Step 4 — Build an "Education" section and a "Skills" list

Add an **Education** section right below Experience. Include your school (or "Self-taught" / "Online coursework"), graduation year (or expected year), and one line of relevant classes or achievements — anything that hints at what you know. No degree yet? No problem: list what you ARE studying, plus any courses, workshops, or certificates, even free ones from sites like Coursera, Khan Academy, or YouTube tutorials.

Then add a **Skills** section. This is just a comma-separated list of tools and abilities: typing speed ("60+ WPM" if true — WPM means words per minute), languages, software (Google Sheets, Canva, TikTok editing, whatever you actually use), and people-skills ("customer service," "calm under pressure"). Skip anything you can't honestly do.

💬 Example education entry: `Lincoln High School — Expected graduation May 2027. Coursework: Business Math, Intro to Computer Science, Spanish II.`

💬 Example skills line: `Skills: Google Docs, Sheets, Canva, social media scheduling (Meta Business Suite), conversational Spanish, POS cash handling.`

You'll know it worked when a stranger could scan these two sections and quickly see what you know and where you learned it.

Step 5 — Add optional sections that strengthen your story

Two more sections often help first-time resumes: **Projects** and **Volunteering**. A project can be anything you built or organized — a class group, a club, a YouTube channel, a small Etsy shop, a Discord community, a coding exercise. Volunteering covers anything you did for free that helped people — church groups, school events, neighborhood cleanups, food drives, coaching a kid's team.

You can also add **Languages**, **Certifications** (CPR, food handler's permit, even a free online certificate you actually finished), or **Awards** (perfect attendance, honor roll, employee-of-the-month at a family business — yes, that counts). If you have nothing extra, skip these — a tight two-page-equivalent on one page beats a sparse two pages.

💬 Example project entry: `Project: Neighborhood plant-swap app — designed a simple scheduling tool in Google Sheets, promoted via local Facebook group, coordinated 3 successful swaps with 40+ attendees.`

You'll know it worked when your resume fills the page naturally and every line answers "what did this person actually do?"

Step 6 — Proofread, shorten, and save as a PDF

Read your resume out loud once. Fix anything that sounds awkward, too long, or that you can't defend in an interview. Most first drafts are too long; aim for one page. Cut filler like "hardworking team player" unless you back it up with an example. Replace vague verbs ("helped," "worked on") with specific ones ("organized," "built," "answered," "tracked," "taught").

Save the final file as a **PDF** (in Google Docs: **File → Download → PDF**; in Word: **File → Save As → PDF**). PDFs keep your formatting intact no matter who opens them. Name the file clearly: `FirstName-LastName-Resume.pdf`.

💬 Example: cut this — "Responsible for various tasks including helping customers." Replace with — "Greeted 80+ customers daily, restocked shelves, and handled returns at a family-run grocery store."

You'll know it worked when you can email the file, the formatting holds, and you feel ready to explain any line on it out loud.

Common mistakes

- **Mistake:** Leaving the page half-empty because you think you "don't have experience." **Fix:** List school projects, caregiving, volunteering, hobbies done seriously, or small paid gigs — anything with responsibility counts. If the page is still sparse, add a one-line "Interests" section (photography, gaming communities, language learning) and tie one to a skill.
- **Mistake:** Using only generic words like "hardworking" and "team player." **Fix:** Replace each with a concrete example — "trained two new volunteers on check-in procedures" beats "good trainer" every time.
- **Mistake:** Lying or stretching the truth. **Fix:** Stick to what really happened, then frame it precisely. A hiring manager will ask follow-ups; honest, well-described experience always beats padded claims.

What this means for you

- **In everyday life:** A resume is just a written summary of your life so far — once you've made one, updates for the next 10 years take minutes, not hours.
- **For your business or work:** If you hire people, a resume built this way tells you whether someone can actually describe what they did — a skill that predicts on-the-job clarity.
- **If you're just getting started:** Your first resume won't be perfect, and that's fine. Send it to one job you actually want, see what feedback comes back, and revise. Resume-writing is a skill, and it improves with each round.

Try it now

Open a blank document right now and type your name at the top. Underneath, write one sentence describing what you're good at, and one sentence describing the kind of work you'd like. That's your summary. You just started a real resume — the rest is filling it in, one honest line at a time.

How to Organize a Messy Email Inbox Into Simple Folders in One Sitting

6 steps · Read online

> **In short:** Sort by sender, archive or delete anything that doesn't need a reply, then create just three broad folders (To act on, To read later, To keep) so the system still works a month from now. The whole thing takes about 30 minutes.

A messy inbox isn't a personal failure - it's what happens when there's no system. Here's a 30-minute plan that creates one simple system, designed to keep working long after you've finished. It works in any major email app because the idea is the same everywhere: stop the noise at the door, then give the messages that matter a home.

Before you start

- A computer or phone with your email app or website open. The steps apply generally to Gmail, Outlook, Apple Mail, Yahoo, and similar services - exact button names vary, but the ideas don't.
- About 30-45 minutes of quiet time, with no other tabs open competing for attention.
- Your password nearby in case you get logged out, and a notepad (or a note on your phone) so you can jot down anything that needs a real follow-up.

Step 1 — Pick your window and open the inbox the right way

Choose a 30-45 minute block when you won't be interrupted - a Sunday morning, a lunch break, after the kids are asleep. Open your email app or website and sign in. The key move: switch your inbox view from "Unread" or "Priority" to something that shows everything - often labeled "All Mail" or "All Messages." This matters because you need to see the full mess before you can fix it. Think of it like clearing the table before you start sorting the dishes.

If your app doesn't have a clear "All Mail" option, look for a folder on the left sidebar labeled something like "All," "Everything," or "Archive." The goal is to see every message in one long list, not just the unread ones. On phone apps, the default view usually already shows everything in your inbox, which works fine.

💬 Example: tell yourself out loud, "I'm spending the next 35 minutes on this. No answering anything that needs more than 60 seconds."

You'll know it worked when your full inbox is visible in one list - usually hundreds or thousands of messages, not just the unread ones.

Step 2 — Sort by sender so the pattern jumps out

Find the option to sort or group your email by sender. In many webmail apps this lives under a sort menu (look for a downward arrow, a "Sort by" option, or three vertical dots near the top of the message list). On a phone app, look in the top-right corner for a filter or sort icon. When you group by sender, the same address suddenly appears 20, 30, or 50 times - that's your noise made visible.

If you can't find a "sort by sender" option, search for one big sender's name instead - the one whose name you recognize the most in the list. The point is the same: figure out who's filling your inbox before you start moving anything.

💬 Example: search "Amazon" and count how many emails you find. If the number is in the double digits, that's a sender you'll handle in Step 3.

You'll know it worked when you can see, at a glance, which senders dominate your inbox - usually a handful of marketing lists, a few retailers, and some auto-notifications.

Step 3 — Clear the noise: archive or delete

Start with your biggest offenders - the senders who appear most often. For each one, you have three honest choices: read and reply (rare), archive the whole batch (the default for anything that doesn't need a response), or delete it. Look for an "unsubscribe" link near the bottom of marketing emails and click it - that legally stops future mail from the same sender. "Archive" just means moving a message out of your inbox while keeping it findable, like putting a paper into a filing cabinet instead of leaving it on your desk.

If you're not sure which messages to keep, ask one question: "Will I need this in six months?" A receipt for a big purchase, yes. A flash sale from a store you visited once, no. Anything you can't answer becomes "archive" - still searchable, just not in your way.

💬 Example: for the Amazon thread - open one message, click "unsubscribe" near the bottom, wait for the confirmation page, then select all Amazon emails (look for a checkbox at the top of the list) and click "Archive" or "Delete."

You'll know it worked when unsubscribe confirmations start arriving in your inbox saying "you've been removed" - that's a sign you've cut the supply line.

Step 4 — Create your three broad folders

Now make just three folders. Pick simple, broad names that won't go stale next month: "To act on" (anything requiring a real reply or decision), "To read later" (interesting but not urgent - articles, newsletters you'll actually open), and "To keep" (receipts, confirmations, contracts, anything you might need to find again). A "folder" here is just a named container for related emails - like a labeled drawer in that same filing cabinet.

Most email apps have a "+" or "Create new folder/label" button - look for it in the left sidebar, often just above the existing folder list. Right-clicking the sidebar often brings up the same option. Keep the names short - one or two words each.

If your app uses "labels" instead of folders (some do), the idea is the same: three broad categories that you tag messages with. You can apply more than one label to the same message later if you need to.

💬 Example: name them "To act on," "To read later," "To keep" - exactly those three, no more. Resist the urge to add a fourth on the same day.

You'll know it worked when all three folders appear in your left sidebar and you can click into an empty one.

Step 5 — Sweep the remaining emails into the right folder

Go back to your inbox. Now that the noise is gone, what remains are messages that probably need your attention or your future self's attention. Work through them quickly, one at a time: open, decide, file. Most belong in one of your three folders. A few might need a quick reply - if it takes less than 60 seconds, reply now; if it takes longer, file it under "To act on" and keep moving.

If a message doesn't fit any of the three folders, archive it instead of creating a new one. The point of three folders is that they stay broad - resist the urge to add a fourth until you've lived with these for at least a month.

💬 Example: a newsletter you want to read on the weekend → drag to "To read later." A package tracking update → drag to "To keep." A friend's question that needs a real answer → drag to "To act on."

You'll know it worked when your actual inbox (the view that doesn't include your new folders) is mostly empty - ideally under 20 messages.

Step 6 — Add one simple rule so it keeps working

Find your email app's "filter" or "rule" feature - usually under Settings, then a tab labeled "Filters," "Rules," or "Inbox rules." A filter is just an automatic instruction: "when an email arrives from this address, move it to that folder." Create one rule for your biggest remaining sender (often a newsletter you actually like, or a frequent work contact).

This is the step that keeps the system working months from now. Without a rule, you'll have to re-sort every week. With even one rule, the inbox stays close to empty on its own.

If your email app doesn't support automatic rules, do the next best thing: make a tiny daily habit of glancing at your inbox once a day and dragging anything obvious into the right folder - two minutes is usually enough.

💬 Example: in Gmail settings, choose "Create new filter" → From: newsletter@example.com → "Apply the label: To read later." Every future issue lands in the right folder automatically, no work from you.

You'll know it worked when the next email from that sender shows up already filed - and you didn't have to do anything.

Common mistakes

- **Creating too many folders.** The temptation is to make one for "Bills," one for "Work," one for "Family" - and suddenly you have fifteen folders, half empty, and you can't find anything. Three broad folders beats ten narrow ones. Add a fourth only after a real, repeated need shows up.
- **Filing things you'll never look at again.** "To read later" can quietly become a graveyard of saved articles. If a newsletter has been sitting unread for two weeks, archive it unopened. The point of saving is to read, not to feel productive.
- **Skipping the unsubscribe step.** Just archiving the noise doesn't stop it from arriving tomorrow. The unsubscribe link is the actual fix - archive is just the bandage. Spend the extra 10 seconds per sender; it pays back forever.

What this means for you

- **In everyday life:** A clean inbox means you stop dreading the little red badge on your phone. You see the few messages that actually matter when you open the app, instead of scanning past 400 to find them.
- **For your business or work:** When you search for an old receipt or a past conversation, you find it in seconds because everything has a home. Tax season, warranty claims, and "what did they say last March?" all get easier.
- **If you're just getting started:** Tomorrow morning, spend 10 minutes only on Step 3 - clearing one big sender. Don't try to do all six steps at once yet. Get a feel for the rhythm before committing the full 30 minutes.

Try it now

Open your email app right now and find the single sender with the most emails in your inbox. Click "unsubscribe" at the bottom of one of their messages. That's it - one minute, one sender, one less source of noise tomorrow.

Practice job interview questions with Qwen before the real interview

6 steps · Read online

> **In short:** Qwen can play the role of an interviewer, ask you common questions, and give you honest feedback on how your answers come across — clarity, completeness, and confidence. You'll need about 20–30 minutes for one full practice run, and a free account is enough to start. It cannot read your body language or tone, so for video interviews, do at least one round on camera separately.

By the end of this guide, you'll have run a real mock interview with Qwen, gotten concrete feedback on your answers, and saved your best versions for the night before the real thing. This is for anyone with an interview coming up — first job, career change, or just wanting to sound more polished.

Before you start

- A Qwen account (free sign-up with email or a supported login).
- A device with a browser — a computer is easier for long answers, a phone works too.
- A job title or two in mind that you actually want to practice for.
- About 20–30 minutes for a full session, or 5 minutes for a quick drill.

Step 1 — Open Qwen and start a fresh chat

Go to Qwen's chat page in your browser, or open the Qwen mobile app if you have it installed. Look for a **"New Chat"** button or a small **"+"** or pencil icon — usually at the top of the sidebar on a computer, or right above the message box on a phone. Click or tap it. A fresh, empty conversation opens with a blank message box at the bottom. Starting fresh matters because leftover context from an earlier chat could push Qwen's questions in the wrong direction. If you don't see a New Chat button, try refreshing the page — sometimes the sidebar hides it on smaller screens.

💬 Example: just type "Hi" first to confirm the chat is responsive, then move to Step 2 to set up the real practice.

You'll know it worked when you see an empty message box with no previous messages above it.

Step 2 — Tell Qwen what role you're practicing for

In the message box, type a short paragraph that gives Qwen the context it needs to ask realistic questions. The more specific you are about the job, the better the questions will be. Mention the job title, the industry, the seniority level, and one or two skills you want to highlight. Press Enter

or tap send. Qwen usually replies with a short confirmation like, "Got it — let's practice for a marketing coordinator role." If the reply feels off or skips your details, scroll up, click the **regenerate** button (usually a circular arrow), or retype your setup more simply.

💬 Example: "I have an interview next week for a marketing coordinator role at a mid-size SaaS company. I have 2 years of experience in social media and email campaigns. Please act as the hiring manager and interview me one question at a time. After each answer, give me feedback on clarity, completeness, and confidence, then move to the next question."

You'll know it worked when Qwen replies acknowledging the role and confirms it will go one question at a time.

Step 3 — Answer the first question like it's real

Read Qwen's question carefully, then answer as if you were sitting in the actual interview. Type your answer in the same box — don't overthink it, don't use perfect essay language, just answer the way you naturally would. The point is to practice, not to write something polished. If you blank out, write "I'm not sure" the way you might say it out loud, then keep going. The more honestly you answer, the more useful Qwen's feedback will be. If the question feels too easy, ask for a harder version before answering — this is normal interview prep, not a sign you're doing it wrong.

💬 Example: if Qwen asks "Tell me about a time you handled a difficult customer," you might type something like: "At my old retail job, a customer wanted to return a laptop past the 30-day window. I listened to her frustration, explained the policy, then asked my manager if we could do a store credit as a one-time exception. She agreed, and the customer left happy. I learned that sometimes bending the rules a little keeps a customer for life."

You'll know it worked when you've typed a full, honest answer — anything between 3 and 12 sentences is a typical interview length.

Step 4 — Ask for specific feedback on your answer

After you finish typing, don't wait for Qwen to move on. Type a follow-up asking it to evaluate what you just said, and be specific about what kind of feedback you want — clarity, completeness, confidence, filler words, structure. Qwen usually gives you a paragraph or two of concrete suggestions, like "your answer is clear but you ramble in the second sentence" or "you named the result but didn't explain your specific role." If you want a number, ask for a score out of 10. If the feedback feels too soft ("great answer, well done!"), push back: ask Qwen to be tougher, since you'd rather hear it now than in a real interview room.

💬 Example: "Before the next question, can you give me feedback on my last answer? Specifically: (1) Was it clear and easy to follow? (2) Did I actually answer what was asked? (3) Did I sound confident or unsure? (4) What is the one thing I should improve? Please be honest — I'd rather hear it from you than from a real interviewer."

You'll know it worked when Qwen responds with specific, actionable feedback — not just a generic compliment, but actual comments about what was strong and what to improve.

Step 5 — Push yourself with follow-up and pressure questions

Real interviewers rarely stop at one question on a topic — they ask follow-ups like "and what would you do differently?" or "how did your manager react?" So once you've answered and gotten feedback, ask Qwen to dig deeper. Type a message asking for a follow-up, or simply say "ask me a harder version of that one." If you really want to stress-test yourself, ask Qwen to challenge you: "Now push back on my answer as if you don't believe me." This is where most of the real interview prep happens — the first answer is usually fine, but it's the second and third round that separates nervous candidates from confident ones.

💬 Example: "That feedback helps. Now please ask me a follow-up question about the same story — something a real interviewer would press me on. I'll answer, you give me feedback, and then we move on."

You'll know it worked when Qwen asks a sharper, more specific question that builds on what you said — not a brand-new topic.

Step 6 — Save your best answers for the day before

Once you've practiced for 20–30 minutes and feel you've landed some strong answers, ask Qwen to summarize your session. Then copy your best answers into a notes app, a Google Doc, or even a simple text file on your phone. The night before the interview, re-read your saved answers out loud — this refreshes the words in your mouth so they sound natural, not memorized. If Qwen gave you specific suggestions ("trim the first sentence, end with a number"), apply those edits now and save the cleaner version. Treat the saved file as your personal answer bank — you can reuse and refine it for future interviews.

💬 Example: "Can you give me a summary of the three strongest answers from this session, plus one thing I should remember to improve next time?"

You'll know it worked when you have a short document with your polished answers that you can re-read on the morning of the real interview.

Common mistakes

- **Asking Qwen to "just give me good answers."** That turns it into a cheat sheet, not practice. You learn by answering yourself and getting feedback, not by copying what Qwen writes.
- **Pasting ten questions into one prompt and answering them all at once.** That's a quiz, not an interview. Real interviews are conversational. Ask Qwen to go one question at a time so you can answer, get feedback, then continue.
- **Stopping after one round of feedback.** The first answer is rarely the best. Ask Qwen to ask the same question again with a follow-up, or to push back, or to make the question harder. Two or three rounds of pressure is where confidence actually builds.
- **Forgetting that Qwen can't read body language or tone.** Qwen evaluates what you *type*, not how you *sound*. If you're prepping for a video interview, do at least one practice round on camera or out loud, separately, to catch the things text can't show.

What this means for you

- **In everyday life:** Even outside job hunting, this same trick works for any high-stakes conversation — a parent-teacher meeting, a price negotiation, a tough email you're about to send. Practice the words before you have to say them.
- **For your business or work:** If you hire people or run interviews yourself, you can flip the script — ask Qwen to generate the five most common weak answers to a question you always ask, so you know what to listen for when you're the one in the chair.
- **If you're just getting started:** This is one of the most useful first things to try with any AI assistant. No technical setup, no cost to begin, and you'll immediately see why "AI as a practice partner" beats "AI as a cheat sheet."

Try it now

Open Qwen right now and run a single 60-second drill: ask it, "Give me one common interview question for a [your job title] role." Read the question, close the laptop, and answer it out loud in your room. That's it — your first mock interview is done.

How to write a polite but firm complaint email that actually gets a reply

6 steps · Read online

> **In short:** You can ask any AI assistant — ChatGPT, Gemini, Claude, Copilot, Le Chat, or another — to draft a complaint email that stays calm, lists the facts, and ends with one clear ask. The trick is giving it the right facts and tone instructions; it can write the draft in seconds, but you still need to read it aloud, cut it to one ask, and send it yourself.

You've had a bad experience with a company. Maybe a product arrived broken, a bill is wrong, or support stopped replying. You want to complain, but you don't want to sound rude — and you really want a reply this time. An AI assistant can help you write a message that's hard to ignore: short, factual, polite, and impossible to misread. Here's how to do it in about 15 minutes.

Before you start

- **An AI assistant you can type into.** Any major one works (ChatGPT, Gemini, Claude, Copilot, Le Chat, DeepSeek). A free account is enough for this task.
- **A real complaint in mind.** Pick one you actually need to send — vague practice complaints don't teach you much.
- **About 15 minutes.** Most of that time is reading the draft out loud and editing, not waiting for the AI.

One honest note before we start: the assistant can draft the email, but it can't send it for you, and it can't guarantee a reply. What it *can* do is remove the hardest part — finding the right words when you're frustrated.

Step 1 — Write down the five facts a rep will need

Before you open any AI tool, get the facts out of your head and onto paper (or a notes app). The biggest reason complaint emails get ignored is that they're vague. A sentence like "Your service is terrible" gives a customer-service rep nothing to act on.

Open a blank note and answer five questions in one or two short sentences each: (1) What did you buy or sign up for, and when? (2) What went wrong? (3) What have you already tried to fix it? (4) What do you want them to do now? (5) What's your deadline for a response?

💬 Example: "Bought a blue wool coat from Example Outfitters on March 4 (order #1234). It arrived March 10 with a torn lining. Emailed support on March 11, no reply. I'd like a replacement shipped by March 25, or a full refund."

You'll know it worked when you can read the note out loud in under 30 seconds and someone else would understand what happened without asking follow-up questions.

Step 2 — Open your AI assistant and start a fresh chat

Open the AI assistant you normally use, on whatever device is handy — phone, tablet, or computer. When the chat window opens, start a new conversation rather than continuing an old one. A fresh chat gives the assistant a clean slate, so it won't accidentally mix in tone or facts from a previous topic.

Look for a button like a **"New chat"**, **"+"**, or **pencil icon** near the top or side of the screen. On some apps it lives inside a menu (often three dots or three lines). If you can't find it, just refresh the page or close and reopen the app — most assistants start you on a blank conversation by default.

💬 Example: type "I want to draft a complaint email" as your first line, just to confirm the chat is responding.

You'll know it worked when you see a blank message box at the bottom of the screen and the assistant is ready for your prompt.

Step 3 — Paste a prompt that names the facts, the tone, and the ask

Now copy the five facts from Step 1 and paste them into the chat, followed by a short set of instructions. The clearer your instructions, the closer the first draft will be to something you'd actually send.

Type something like the prompt below — adjust the wording to your situation, but keep the same four pieces: facts, tone, length, and ask.

💬 Example: "Please draft a polite but firm complaint email. Facts: [paste your five facts here]. Tone: calm, professional, no sarcasm, no insults. Length: under 150 words. End with one specific ask and a reasonable deadline of 5 business days. Do not invent any order numbers, dates, or names I didn't give you."

You'll know it worked when the assistant produces a complete email — greeting, facts, what you want, deadline — in one block of text that reads like a real human wrote it. If the first draft comes out too long, too stiff, or too emotional, that's normal; you'll fix that in the next step.

Step 4 — Read the draft out loud and tame the tone

Don't send the first draft yet. Read it out loud, slowly. Your ear catches things your eyes miss — awkward phrases, sentences that sound angrier than you feel, or phrases that sound so polite they sound sarcastic.

If a sentence makes you wince, copy that one sentence back into the chat and ask the assistant to rewrite it in a calmer voice. Don't throw the whole draft away; most of it is probably fine. Common rewrites you'll want: "I'm extremely disappointed" becomes "I'd like this resolved"; "This is unacceptable" becomes "I'd like to understand what happened"; "You need to fix this immediately" becomes "I'd appreciate a response by [date]."

If the draft came out too apologetic — sounding like you're asking permission to complain — say so directly. Example: "Make this sound more confident. I'm being reasonable, but I expect a real answer."

You'll know it worked when you can read the email aloud without rolling your eyes, and a friend would describe the tone as "firm but polite" — not angry, not groveling.

Step 5 — Cut the email down to one clear ask

A complaint email that asks for three different things often gets three different "we'll look into it" replies and nothing actually happens. Read the draft and circle the single most important thing you want.

If the email mentions a refund, a replacement, an apology, and a discount, pick the one that matters most and cut the others. You can always bring up the others in a later email if the first ask is refused. Then paste the draft back into the chat and ask for the trim.

Example: "Shorten this to one ask: I want a replacement coat shipped by March 25. Remove the part about the discount."

You'll know it worked when the last paragraph of the email contains exactly one sentence that begins with "I'd like" or "Please" and ends with a specific date. Everything else in the email is supporting context, not a wish list.

Step 6 — Send it, save the thread, and set a follow-up reminder

Copy the final draft into your real email app, fill in the real order number, your name, and any attachments (photos of the damage, receipts, screenshots of the original chat). Read it one last time to make sure no detail was lost in the copy-paste — small errors like a wrong date can undo an otherwise strong email.

Before you hit send, take a screenshot of the email and save the AI chat somewhere you can find it. If you have to follow up later, you'll have a record of exactly what you asked for. Then set a

calendar reminder for two business days after your stated deadline.

💬 Example follow-up: "Following up on my email below — I haven't received a response. I'd still like a replacement coat shipped by March 25. Please let me know either way."

You'll know it worked when the email is sitting in your sent folder, you have a date on your calendar for the follow-up, and you feel — for the first time in this whole mess — that you've said exactly what you needed to say.

Common mistakes

- **Letting the AI invent facts.** If you say "I ordered something last month," the assistant may invent an order number or date that sounds real but isn't. Always paste your *own* facts, and add the line "Do not invent any details I didn't give you."
- **Sending the first draft without reading it.** AI drafts are usually 80% right and 20% wrong in tone. The 20% is what makes an email sound robotic or slightly off. Always read it out loud.
- **Asking for too many things at once.** "I want a refund, a replacement, an apology, and a discount" usually gets you nothing. Pick one ask, and bring up the others later if needed.
- **Copying a tone that doesn't sound like you.** If the draft sounds too stiff or too formal, paste in a sentence you've actually written before — a text to a friend, an older email — and ask the assistant to match that voice.

What this means for you

- **In everyday life:** The next time a delivery is wrong, a bill is off, or a service lets you down, you'll have a way to complain that takes 15 minutes instead of an evening of staring at a blank screen.
- **For your business or work:** The same method works for vendor disputes, late invoices, or client complaints. A calm, one-ask email is also the right tone when you're escalating an issue to a manager — you sound professional, not emotional.
- **If you're just getting started:** The very first email to draft this way is the hardest. After that, you'll have a prompt template you can reuse with new facts pasted in.

Try it now

Open your AI assistant, paste this prompt, and fill in the five facts from Step 1: *"Draft a polite but firm complaint email. Facts: [your facts]. Tone: calm and professional. Length: under 150 words. End with one specific ask and a 5-business-day deadline. Do not invent any details."* Read the result out loud, trim it to one ask, and you'll have a sendable email in under 10 minutes.

Translate and Explain a Confusing Email or Letter with Perplexity

6 steps · Read online

> **In short:** Open Perplexity, paste the confusing message, and ask it to translate, summarize, and draft a reply. Free to use on the web; Pro adds longer messages and file uploads. It can draft a reply, but it cannot send the email for you.

That long email from your landlord, the tax letter in German, the polite-but-passive-aggressive message from a client — they all sit in your inbox looking scary. Perplexity can turn any of them into two things you actually need: a plain-English summary of what they want, and a calm suggested reply you can edit before sending.

Before you start

- A free Perplexity account (sign up with email, Google, or Apple) — works in any modern browser on phone or computer.
- The confusing message itself: either copied as text, or as a clear photo/screenshot (uploading may require Pro).
- About **5–8 minutes** for your first message, faster once you get the hang of it.
- Honest note: Perplexity is great at language and summarizing, but it's not a lawyer, accountant, or doctor. For official notices (taxes, contracts, legal threats), treat its summary as a starting point and confirm with a professional before you act.

Step 1 — Open Perplexity and start a new thread

Go to **perplexity.ai** in your browser (or open the Perplexity app if you have it installed). On the home page you'll see a big text box in the middle that says something like **"Ask anything..."** — that's where the conversation happens. Click or tap inside the box once. A blinking cursor appears, and the page is ready for your question. If you're already signed in, you'll also see a small **"New Thread"** or + button at the top-left or top-right — tap it so your message doesn't mix with an older chat. If the layout looks different (Perplexity updates its design often), look for a clearly empty text area with a **Send** button (paper-plane icon ✈️ or arrow) at the end. 💬 Example: don't type yet — just leave the cursor blinking in the box. You'll know it worked when the text box is empty, your cursor is blinking inside it, and no earlier conversation is showing above.

Step 2 — Paste the confusing message into the chat

Open your email or letter in another tab, window, or app. Select **all the text** (Ctrl+A on Windows, Cmd+A on Mac, or long-press → Select All on a phone). Copy it (Ctrl+C / Cmd+C /

Copy in the menu). Then tap back into Perplexity's text box and **paste** (Ctrl+V / Cmd+V / Paste). The full message lands inside the box as one block of text — you should see the words from the letter sitting there. If the message is very long, Perplexity may show a small character count or a warning; trim the greeting and signature if needed. If your screen looks different, you may also see a **paperclip icon** 📎 or an **upload button** — that's for attaching files or photos of letters instead of pasting text. 💬 Example: paste a paragraph like "Sehr geehrte Mieter, wir möchten Sie daran erinnern, dass die Miete für November bis spätestens zum dritten Werktag des Monats zu überweisen ist…" — the exact wording matters, so paste the real thing. You'll know it worked when the entire message shows up in the box, no strange symbols, and you can still see the send button (✈️ or arrow) at the end.

Step 3 — Add a clear instruction below the pasted text

Don't hit send yet. Press **Enter** to make a new line inside the same box, then type what you want Perplexity to do with the message. Be specific — a vague "what does this mean?" gets a vague answer. The clearer the instruction, the more useful the reply. If you only see a single-line text box (some mobile views hide multi-line), look for a small **expand** icon or a + that opens a fuller composer. On the free plan you can usually ask anything; Pro users get longer pasted messages and file uploads — check the current plan details on Perplexity's pricing page if you're unsure which you have. 💬 Example: "Please translate this email into plain English, list the three things the sender wants me to do, and tell me how urgent it sounds." You'll know it worked when the box contains both the pasted message **and** your instruction underneath, and you can read both clearly before sending.

Step 4 — Send and read the summary

Tap the **Send** button (✈️ paper-plane icon, or an arrow ↵ at the end of the box). A short "thinking" animation appears — Perplexity is reading your message and writing its answer. After a few seconds, a reply appears under your message, with the translation at the top, then the bullet list of what the sender wants, then the urgency note. Skim it slowly: any word you don't understand, highlight and copy, then ask a follow-up like "What does [word] mean in everyday language?" If the answer doesn't quite hit the mark, click the small **Regenerate** button (often a circular arrow ↻ near the reply) to get a fresh attempt — sometimes the second version is clearer. If the screen looks different on your device, look for a **Rewrite**, **Try again**, or **New answer** button near the bottom of the reply. 💬 Example follow-up: "Can you explain 'bis spätestens zum dritten Werktag' like I'm 10 years old?" You'll know it worked when you see a clear English version of the email, a short list of what the sender wants, and a plain-language sense of how urgent it really is.

Step 5 — Ask for a suggested reply in the same chat

Stay in the **same conversation** (don't start a new thread) so Perplexity remembers the original message and its summary. In the text box at the bottom, type a request for a draft reply. Be polite about tone — Perplexity matches the energy you ask for. The reply will appear below your message, ready for you to copy. If you want a different vibe (firmer, more apologetic, more formal), just ask in a follow-up — it's faster than starting over. Free users can do this as many times as they want; Pro only matters if you need to paste very long messages or upload PDFs. 💬 Example: "Now write a short, polite reply in English confirming I'll pay by Friday and asking for a receipt." You'll know it worked when a reply draft appears that you could realistically send after light editing — and you can copy it with a single tap or click.

Step 6 — Copy, edit, and send the reply yourself

Highlight Perplexity's suggested reply and **copy** it (Ctrl+C / Cmd+C / long-press → Copy). Open your email app, start a new message, **paste** the draft into the body, and then **edit it in your own voice** — add your name, your real deadline, any details only you know. Perplexity can draft, but it cannot send email or access your inbox, so this last step is always yours. Read the whole thing once out loud before clicking send; if any sentence sounds weird in your voice, change it. If your interface looks different, look for a **Copy** icon (two overlapping squares 📋) at the corner of the reply, or simply long-press the reply on a phone to bring up the copy menu. 💬 Example edit: change "I will remit payment by Friday" to "I'll send the payment by Friday — thanks for the reminder, Maria." You'll know it worked when you've sent a clear, calm reply that you actually wrote (with Perplexity's help), and the original confusing email no longer feels scary.

Common mistakes

- **Pasting without an instruction.** If you only paste the email and hit send, Perplexity may summarize it but won't translate or draft a reply. Always add a short instruction line under the paste — it changes the result completely.
- **Trusting the summary on legal or medical letters.** Perplexity can miss legal nuance. For tax notices, court letters, or anything health-related, treat its plain-English version as a heads-up, not a final answer — confirm with the right professional before paying or signing anything.
- **Sending the draft without editing.** AI-written text sometimes has odd phrasing, wrong names, or invented details. Always read the whole reply once, fix anything that doesn't sound like you, and remove anything that isn't true.

What this means for you

- **In everyday life:** the next time a long email from your bank, landlord, school, or insurance company lands in your inbox, you'll have a two-minute way to understand it

and reply without stress — even if it arrives in another language.

- **For your business or work:** if clients or colleagues send you formal letters, contract updates, or foreign-language requests, you can draft a professional reply in the right tone in seconds, freeing you to focus on the actual decision rather than the wording.
- **If you're just getting started:** the easiest first move is to pick one email you've been ignoring because it felt too long or too formal, paste it into Perplexity, and ask for a three-bullet summary in plain English — that's it.

Try it now

Open perplexity.ai, find one email you've been avoiding, paste it into the box, and type underneath: *"Translate this into plain English and give me three bullet points of what the sender actually wants from me."* Send it, read the result, and notice how much lighter the email feels five minutes later.

Draft a Polite but Firm Refund Email with Le Chat

6 steps · Read online

> **In short:** Le Chat, Mistral's AI assistant, can turn a messy complaint into a calm, clear refund request in a few minutes. You describe what went wrong and what you want back; the assistant shapes it into a respectful email. You'll still review, personalize and send the message yourself from your own email account.

If a product arrived broken, a service didn't deliver what was promised, or a charge looks wrong, the hardest part is often just writing the email. You want to be firm enough to be taken seriously, but polite enough that the person reading it on the other end actually wants to help. That's exactly the kind of writing Le Chat is good at shaping. This guide walks you through it, start to finish, in roughly ten minutes.

Before you start

- A free or paid Le Chat account. Mistral says Le Chat has a free tier with a daily message limit and paid plans for heavier use; check the current tiers on Mistral's site before you start.
- Access to Le Chat in your browser (go to chat.mistral.ai) or through the official app.
- The key facts about your case: order or invoice number, what you bought, when, what went wrong, and what you'd like the company to do (refund, replacement, repair, credit).
- Roughly 10 minutes.

Step 1 — Open Le Chat and start a fresh conversation

Go to **chat.mistral.ai** in your browser, or open the Le Chat app on your phone. If you're not already signed in, log in with the Mistral account you created. Once you're on the main chat screen, look for a **"New chat"** button, a **pencil icon**, or a **plus (+) sign** — these are common ways to start a clean conversation. Click or tap it so you have an empty chat box ready.

💬 Example: you don't need to type anything yet. Just click the button that starts a new chat.

You'll know it worked when the screen shows an empty message box with no previous conversation in the list, and any old chats are out of the way. Starting fresh matters because Le Chat only uses what you put in *this* conversation, so an empty box means nothing will confuse the reply.

Step 2 — Tell Le Chat what happened, in plain words

In the empty message box, write a short, honest description of what went wrong. Don't try to make it sound polite yet — that's Le Chat's job. Just give the facts in a few sentences: what you

bought, when, what was supposed to happen, what actually happened, and how you found out. Think of it like writing a quick note to a friend, not a letter to a lawyer.

Example: "I bought a noise-canceling headphone set from AudioBrand on October 12, order #44291, for $189. It was supposed to arrive by October 19. It arrived on October 24 in a damaged box, and the right ear cup doesn't work. AudioBrand support replied once and then stopped responding."

You'll know it worked when Le Chat confirms it understands the situation, usually by restating the key facts back to you in one or two sentences. If the assistant seems to have misunderstood something, just reply with a short correction ("no, it was the headphones, not the charger") before moving on.

Step 3 — Tell Le Chat the tone you want

Right after describing the problem, in the same message or in a follow-up, tell Le Chat how you want the email to sound. The trick is to name the tone in everyday words: "polite but firm" works well, but so do phrases like "respectful and clear," "calm but serious," or "friendly but I want a real answer, not a brush-off." Also tell Le Chat who you are to the company — a regular customer, a long-time subscriber, whoever fits — so the tone matches.

Example: "Now please turn this into a refund request email. Keep the tone polite but firm — I'm a regular customer who's been patient, but I want a real resolution, not a generic reply. Don't threaten legal action. Keep it under 200 words."

You'll know it worked when Le Chat produces a short, readable email draft with a greeting, a clear statement of the problem, and a specific ask at the end. If the draft feels too soft or too aggressive, you can simply say "a bit firmer, please" or "softer please" and Le Chat will adjust in the next message.

Step 4 — Add the missing details Le Chat should include

Most refund requests get ignored or delayed because the reader on the company's side doesn't have enough information to act. Look at Le Chat's draft and check whether it includes: your name, the order or invoice number, the date of purchase, what the product or service was, what went wrong, what you want (full refund, partial credit, replacement), and a reasonable deadline for a reply. If anything is missing, ask Le Chat to add it.

Example: "Please add my order number (#44291), the date I bought it (October 12), and a line asking for a full refund or a replacement within 7 business days. Also add a short closing sentence thanking them for looking into it."

You'll know it worked when a fresh version of the email appears with those details woven naturally into the text, not pasted in as a random list. If Le Chat puts the details in a weird spot — like the middle of a sentence — tell it where you'd rather see them ("put the order number right after my name").

Step 5 — Ask Le Chat to give you two versions

Before you settle on a draft, ask Le Chat for one more option that's slightly different. This is a small step that pays off: it gives you something to choose from, and it also lets you see how Le Chat handles the same facts in a different style. You can ask for a "shorter and more direct" version, or a "warmer and more conversational" one.

💬 Example: "Can you give me one more version that's shorter and more direct — about 100 words? Keep the same tone and the same ask."

You'll know it worked when two clearly different drafts are sitting in your chat, and you can spot the small differences in tone and length. Pick the one that feels closest to how you'd actually write, and use the other as a source of small phrases you can borrow (an opening line, a polite but firm closing sentence, a good way to ask for a deadline).

Step 6 — Personalize the email and send it yourself

Copy your chosen version out of Le Chat into your own email app. Read it once out loud — does it sound like something *you* would say? If anything feels off, change the words. Add a real sign-off with your actual name. Attach any proof you have (a photo of the damaged item, the receipt, the original confirmation email). Then send it from your own email account, addressed to the company's support or billing address — not to Le Chat, since Le Chat can write the email but it can't send it for you.

💬 Example: your finished email might start, "Hi AudioBrand Support team, I'm writing about order #44291, a noise-canceling headphone set I purchased on October 12..."

You'll know it worked when the email in your outbox is calm, specific, and signed with your real name, and the company has everything it needs to act on it. If you don't hear back within the deadline you set, Le Chat can also help you draft a polite follow-up using the same conversation as a reference.

Common mistakes

- **Sending the AI's draft without reading it.** Le Chat usually writes well, but it doesn't know your real voice. Skim it once, fix anything that doesn't sound like you, and add your name before sending.

- **Leaving out the order number or date.** Customer support can't help with a complaint they can't trace. If your facts aren't in the email, the reply will almost always be "could you provide more details?" — which slows everything down.
- **Asking for a refund without saying what kind.** "I want my money back" is vague. Be specific: full refund, replacement, store credit, or repair. The clearer the ask, the faster the answer.
- **Being rude "to sound firm."** Strong wording is fine; insults or sarcasm aren't. A polite, factual email gets better results than an angry one, even when you're frustrated.

What this means for you

- **In everyday life:** the next time a package arrives broken, a subscription overcharges you, or a service quietly stops working, you'll have a ready way to write a calm, effective email in a few minutes instead of staring at a blank screen.
- **For your business or work:** the same prompt style works for vendor disputes, invoice corrections, and client complaints. You describe the situation, name the tone, ask for a clear action, and let the assistant shape the wording.
- **If you're just getting started:** the easiest first move is to paste one short paragraph about a real (or made-up) situation into Le Chat and ask for a polite version. Seeing the result once is the fastest way to learn what the assistant is good at.

Try it now

Open chat.mistral.ai, start a new chat, and paste a real complaint you actually have (or a short made-up one) followed by: "Turn this into a polite but firm refund request email, under 200 words, that asks for a full refund within 7 business days." Read what comes back — that's the whole skill, in one message.

How to Use Gemini Deep Research for Thorough Reports

6 steps · Read online

Imagine having a tireless personal research assistant who can scour dozens of websites, cross-reference facts, and write a polished report in minutes. This guide will show you how to use Gemini's powerful capabilities to gather deep insights for work, study, or personal projects, just like a professional researcher.

Before you start

- A Google account: You'll need a standard Google account to access Gemini.
- Access to Gemini: Ensure you can open Gemini in your web browser (gemini.google.com) or use the Gemini app on your mobile device. Some advanced features might perform better or offer more depth with a Gemini Advanced subscription, but the core research principles apply to all versions.
- A device: A computer, laptop, tablet, or smartphone with internet access.
- A clear topic or question: Have a specific subject in mind that you want Gemini to research.
- Rough total time: This process usually takes 5-10 minutes, depending on the complexity of your research topic and how many follow-up questions you ask.

Step 1 — Start a new chat and enable web search

First, let's open Gemini and make sure it has permission to search the internet. An AI (Artificial Intelligence) tool like Gemini is powered by an LLM (Large Language Model) – think of it as the super-smart brain that helps the AI understand and generate human-like text. To perform deep research, this brain needs access to the wider internet.

When you open Gemini, you'll see a clear text box at the bottom of the screen where you type your messages. Look towards the top-right or top-center of the screen for a toggle switch or button that usually says "Extensions" or shows a puzzle piece icon. Clicking this often reveals a list of tools, including "Google Workspace" and "Google Search." Ensure the "Google Search" extension is switched 'on'.

If it looks different: On some devices or versions, the "Extensions" might already be active or there might not be a visible toggle. If so, don't worry, Gemini often suggests using Google Search automatically when you ask a research question. Just ensure you're starting a fresh conversation, usually by clicking a "New Chat" button (often a '+' icon) in the side menu.

You'll know it worked when: The toggle for "Google Search" appears switched on, usually indicated by a blue or green color, or Gemini doesn't prompt you to enable search when you type

your first question.

Step 2 — Craft a detailed research prompt

A prompt is simply the instruction you give to an AI. To get a high-quality, in-depth report, you need to be very clear with your prompt. Don't just ask a one-sentence question. Instead, tell Gemini who the report is for, what specific areas to cover, and how you want the final document organized. Think of it like giving instructions to a new research assistant.

Type your detailed instructions into the text box. The more detail you provide, the better Gemini can understand your needs and focus its search.

Example:

> "Act as an expert researcher for a homeowner. Please conduct a deep, multi-step search of the web to investigate current trends in home solar battery storage for residential properties in the US. Focus on average cost ranges, typical lifespan, available government rebates or incentives, and environmentally friendly recycling options. Organize the report with clear headings, bullet points, and provide specific examples where possible. Include links to your sources."

You'll know it worked when: Gemini acknowledges your request (often by saying "Okay, I can help with that...") and begins to indicate it's searching the web or preparing to ask clarifying questions.

Step 3 — Answer Gemini's clarifying questions

One of the cleverest parts of Gemini's deep research process is that it doesn't just guess what you want. It will often reply with two or three clarifying questions to narrow down the scope of its search. This is Gemini using its "brain" (the LLM) to ensure it understands your request perfectly, saving you time later.

Read Gemini's questions carefully. These might be about geographical focus, specific types of technology, or the target audience for the report. Type your concise answers in the text box below.

If it looks different: Gemini might sometimes jump straight into generating the report if your initial prompt was extremely clear. If it doesn't ask questions, that's fine; simply move to the next step.

Example: If Gemini asks, "Are you interested in global statistics or a specific country for the solar battery research?" you might reply:

> "Focus primarily on American residential properties, especially those in suburban and regional areas, not just cities."

You'll know it worked when: Gemini confirms it has understood your answers, perhaps by saying "Got it," and then transitions to the next stage of gathering information.

Step 4 — Observe Gemini gathering information

Once you've clarified the details, Gemini will begin its thorough search. You'll often see a visual progress bar or a series of messages indicating what the AI is doing in real-time. These messages might pop up within the chat window, showing stages like "Searching for battery lifespans" or "Analyzing recycling policies." This is Gemini actively using its Google Search extension to find and process information from various websites, much like a human researcher would.

This step takes longer than a standard AI reply—often up to a minute or two—because the assistant is reading actual web pages rather than relying purely on its pre-existing training memory.

If it looks different: Sometimes, the progress indicators might be subtle, or Gemini might just show a spinning wheel. As long as the chat isn't frozen and you see some activity, it's working.

You'll know it worked when: You see clear messages indicating Gemini is searching the web ("Using Google Search...", "Searching for..."), followed by text beginning to appear that resembles your requested report structure.

Step 5 — Review and verify the generated report

When the report is ready, Gemini will present a structured document, usually complete with clear headings, bullet points, and, crucially, citations or links to the websites it used. This is your chance to put on your detective hat.

Before you copy the text for use, perform a quick spot-check. Click on a few of the provided links to make sure they are active and support the claims made in the report. This protects you against "hallucination," which is the term for when an AI confidently invents facts or links that do not actually exist. While rare with Gemini's search capabilities, it's always wise to verify important details.

If it looks different: Ensure you scroll to the end of the report. The citations are often small numbers or footnotes. If you don't see any, you can politely ask Gemini: "Can you provide the source links for this report?"

You'll know it worked when: You've clicked on a few links and confirmed that the information in the report aligns with the content on the source websites.

Step 6 — Refine or ask follow-up questions

The first report Gemini provides is a great starting point, but you might want to dive deeper or change the focus slightly. This is where the interactive nature of AI truly shines. You can continue the conversation, asking for more details on a specific section, requesting a different format, or even asking Gemini to summarize complex points.

Type your follow-up questions or requests into the text box. This allows you to iteratively improve the research until it perfectly meets your needs.

💬 Example:

> "Can you expand on the environmental impact of lithium-ion battery production? Also, please summarize the key points of this report in three bullet points, suitable for a quick presentation slide."

You'll know it worked when: Gemini provides new information, rephrases existing content, or summarizes the report based on your latest instructions, showing its ability to build on previous responses.

Common mistakes

- **Being too broad:** Asking for a report on "the history of technology" is too vast for a single deep dive. *The fix:* Narrow your focus to a specific decade, tool, or industry, or break it into smaller prompts.
- **Skipping the fact-check:** Never copy and paste a report straight into an important document or project without clicking the source links to verify the numbers and claims. *The fix:* Always dedicate a few minutes to spot-checking sources and ensuring accuracy.
- **Ignoring clarifying questions:** If Gemini asks for clarification, don't skip it or give a vague answer. Your response is crucial for guiding the AI to find the exact data you need, rather than general information. *The fix:* Take a moment to provide a clear, concise answer to each of Gemini's questions.

What this means for you

- **In everyday life:** You can quickly compare complex options

How to Draft a Polite Follow-up Email After a Job Interview

6 steps · Read online

> **In short:** Send a short, specific thank-you email within 24 hours of the interview — three or four sentences, one fresh detail from the chat, and a clear next step. Keep it warm, not needy; this guide shows you exactly what to write and what to avoid.

A good follow-up email is one of the few small things still fully in your control after an interview. It takes about ten minutes to write, and it quietly reminds the hiring manager who you are. Below is a simple, friendly process you can follow the same day, even if you've just walked out of the meeting room still buzzing.

Before you start

- **Time needed:** roughly 10–15 minutes, ideally within 24 hours of the interview.
- **What you need:** the interviewer's name, their email address, and the job title you applied for. A notes app (your phone is fine) where you scribbled two or three things that came up in the conversation.
- **Tool:** any email app or webmail — Gmail, Outlook, Apple Mail, whatever you already use. No special software required.
- **Optional:** an AI assistant (such as ChatGPT, Gemini or Claude) can help you polish your draft, but it cannot send the email for you, and it doesn't know what was actually said in the interview — that's your job to feed in.

Step 1 — Jot down three things you actually talked about

Open your notes app the moment you can after the interview — even two minutes in the parking lot is better than waiting. Write down three specific things that came up: a question they asked, a story you told, a detail about the role or the team. These are the bricks you'll build the email from. **What you see:** your phone's notes app, the date at the top, and probably a few messy half-sentences. **What happens next:** you'll have raw material to pull from, so the email can sound like *you* and not like a template you copied off the internet. **If it looks different:** if you didn't take notes, try to recall one question you were asked and one thing you said that landed well. Even one believable detail is enough. **You'll know it worked when** you have at least two concrete points you can reference in the email — something the interviewer would recognize as "ours, not just any interview." 💬 Example: *Asked about handling an angry customer → I told the café story. Mentioned the new Nashville office opens in March.*

Step 2 — Open a new email and set the subject line

Open your email app and click **Compose** (or **New Message**). For the subject line, keep it short and easy to scan — the interviewer gets a lot of mail. A clear subject line is like a polite wave from across a crowded room: it says "this is the person you met today, not a stranger." **What you see:** a blank email with a To, Subject and message area. **What happens next:** the moment you type their name correctly, the email already feels more personal than half the inbox they'll skim today. **If it looks different:** on some apps the button says **New email**, **Write**, or shows a pencil icon — any of these will do. **You'll know it worked when** the subject clearly says thank-you, names the role, and is short enough to read fully on a phone screen. 💬 Example: `Subject: Thank you – Marketing Coordinator interview`

Step 3 — Open with a warm, specific thank-you

In the first one or two sentences, thank the interviewer by name and mention something specific from your conversation. This is the part that turns a generic "thanks for the interview" into a real, human note. **What you see:** the cursor blinking at the top of the message body. **What happens next:** the reader instantly knows this email is about them and the meeting they actually had — not a copy-paste sent to ten companies. **If it looks different:** if you interviewed with a panel, address the lead person by name and use "and the team" for the rest. **You'll know it worked when** the first paragraph would only make sense to someone who was in that room. 💬 Example: *Hi Sarah, thank you so much for taking the time to chat with me this morning about the Marketing Coordinator role. I really enjoyed hearing how the team is planning the March launch.*

Step 4 — Briefly reinforce why you're a fit

Now add two or three sentences that gently remind them of your strengths, tied to something you actually discussed. This is where your notes from Step 1 come in. Think of it like adding one fresh photo to a friend's birthday card — small, but it shows you remembered. **What you see:** the email is taking shape — it's getting longer than a text but shorter than a letter. **What happens next:** the interviewer gets a quick, evidence-based recap of why you belong in the role, without you having to repeat your whole CV. **If it looks different:** if the conversation was mostly about the company or the role itself, you can skip this part and go straight to Step 5 — a thoughtful question can carry the same weight. **You'll know it worked when** the reader can see, in plain words, that you understood the role and can do it. 💬 Example: *Our chat about managing the café's social media during the weekend rush reminded me how much I enjoy fast-paced, customer-facing work — that's exactly the energy I'd love to bring to your team.*

Step 5 — Close with a clear, low-pressure next step

End with one short sentence that says you're looking forward to hearing from them, and sign off politely. Avoid phrases like "I really need this job" or "Please let me know ASAP" — they add

pressure, and pressure is the opposite of what you want. **What you see:** the email might be around 80–120 words in total. That's the right length. **What happens next:** the email lands in their inbox as a calm, professional note — easy to reply to, easy to forward, easy to remember. **If it looks different:** if you don't know the next steps (some processes are slow), just say "I'm happy to provide anything else you need" — it's honest and useful. **You'll know it worked when** the email closes without begging, without bragging, and without leaving the interviewer unsure about what you want. 💬 Example: *I'm excited about the role and would love to hear how the next steps look. Thanks again, Sarah — have a lovely rest of your week.*

Step 6 — Do a quick proofread and hit send

Read the whole email out loud once, slowly. Check the interviewer's name is spelled right, the role title matches the job ad, and there are no obvious typos. Then send. **What you see:** the final email, maybe 100 words on screen. **What happens next:** the email is on its way. Most of the work is done. **If it looks different:** if you spot a mistake after sending, don't panic — a small typo is not worth a follow-up apology. Save that energy for the next step in the process. **You'll know it worked when** you can read it back and think, "yes, that sounds like me on a good day." 💬 Example (final email): *Hi Sarah, thank you so much for taking the time to chat with me this morning about the Marketing Coordinator role. I especially enjoyed hearing how the team is planning the March launch. Our conversation about handling busy weekend rushes reminded me how much I love fast-paced, customer-facing work, and I'd be thrilled to bring that energy to Wynnum Co. I'm excited about the role and would love to hear how the next steps look. Warm regards, Jordan.*

Common mistakes

- **Sending one identical email to everyone.** If you can't tell which interviewer it was for, they can tell you didn't write it. Always swap at least one specific detail.
- **Writing five paragraphs.** A follow-up is a thank-you card, not a second cover letter. If it scrolls on a phone, it's too long.
- **Sounding desperate or pushy.** Lines like "I really need this opportunity" or "Any update would be appreciated" add pressure. Keep the tone calm and confident.
- **Forgetting to proofread the name.** Getting the interviewer's name wrong undoes the whole "personal" effect. Double-check before you send.

What this means for you

- **In everyday life:** The same habit — a short, specific, timely message — works after a coffee chat, a school meeting, a meeting with a tradie, or a volunteering shift. It makes you the person people remember kindly.

- **For your work or job search:** A clean follow-up email slightly tilts the odds in your favor because most candidates don't bother, and many who do bother send something generic. Yours doesn't have to be brilliant — it just has to be real.
- **If you're just getting started:** Don't try to write the perfect email. Write a rough one, read it out loud, and send it within 24 hours. Done beats perfect every time.

Try it now

Within the next five minutes, open your notes app and write down three things you remember from your last interview (or your last important meeting). That's the hard part done — the email writes itself from there.

How to practice for a job interview with an AI mock interviewer tonight

6 steps · Read online

> **In short:** You can run a realistic mock interview with any major AI assistant tonight by giving it the job description, asking it to play interviewer, and stress-testing one tough answer at a time. It won't replace a real human, but it will sharpen what you actually say.

Job interviews don't reward memorization — they reward clarity under pressure. Tonight you can build that clarity with a practice partner that already lives in your phone or browser: an AI assistant. The trick isn't to ask random questions. It's to run a short, repeatable session that stress-tests your real answers before a human ever hears them. Below is a method you can use tonight, and again the night before any interview.

Before you start

- An account with any major AI assistant — **ChatGPT, Claude, Gemini, Microsoft Copilot, and Le Chat** all work for this. Use whichever you already have; don't sign up for a new one tonight.
- The job description for the role. Copy-paste the whole posting, even if it's long — the AI uses every line.
- A quiet 30–45 minutes, ideally with headphones and a microphone so you can hear yourself.
- Optional: a notes app open on your phone so you can save weak spots at the end.

Step 1 — Open a fresh chat with your AI assistant

Open the AI assistant of your choice in a browser or app and start a brand-new conversation. What you'll see is a blank chat box — no prior context, no leftover questions from yesterday. This matters because the assistant draws on the whole chat thread when it replies, so a clean slate means a clean practice session. In ChatGPT, look for a top-left button that says **New chat**; in Claude, Gemini, and Copilot, look for **New chat** or a + icon near the conversation list. On mobile, you may need to swipe to your conversation history first before the button appears. 💬 Example: just open the chat — no typing yet. You'll know it worked when you see an empty message box with a placeholder like "Message…" or "Ask anything."

Step 2 — Tell the AI exactly what role you're practicing for

Paste the job description into the chat, then add two or three lines about your background. What you'll see is a long message appear, followed by the assistant's reply — usually a short

acknowledgment like "Got it, here's how I'll approach this." The assistant now has the context it needs to ask realistic questions instead of generic ones. If your tool has a **Custom instructions** or **Memory** setting, you can paste the role there so it persists across chats; otherwise just paste it into the first message of this session. On Claude and ChatGPT those settings live under your profile picture; on Gemini they're under **Settings**. 💬 Example: "I'm preparing for a Senior Customer Success interview at a B2B SaaS company. Here is the job description: [paste]. My background: 4 years in support, promoted once, led a small team. Please ask me one question at a time, like a real interviewer would." You'll know it worked when the assistant references your actual role or company in its next reply — not a generic "tell me about yourself."

Step 3 — Start with the one question you find hardest

Ask the assistant to ask you the single question you're most worried about — usually a behavioral or scenario question, not the opener. What you'll see is the assistant posing it the way a real interviewer would: *"Tell me about a time you dealt with a difficult customer."* Read the question, then **answer out loud** — even if you're typing, speak the answer aloud so you catch filler words, awkward pauses, and rambling. This is the most underused part of practice, and the part that actually changes how you sound in the room. If your assistant has a voice mode (ChatGPT and Gemini both do on desktop and mobile), switch to it now; otherwise read your typed answer out loud to yourself once you're done. 💬 Example: "Ask me one question at a time. Start with: 'Tell me about a time you had to win back an unhappy customer.'" You'll know it worked when you've actually spoken a full answer of roughly 60–90 seconds — not just thought one.

Step 4 — Ask for specific, structured feedback

Once you've answered, prompt the assistant to score your reply and point to one improvement. What you'll see is a short rubric — usually a 1-to-5 rating on clarity, specificity, or impact, plus one thing to tighten. This is where generic praise is useless, so ask for criticism directly; the assistant now becomes a coach instead of a yes-machine. If it forgets and just says "great answer!", tell it again to be tougher — most assistants default to polite, especially in voice mode. On the free plan of all the assistants listed above, text feedback is unlimited; voice mode may have a small daily cap. 💬 Example: "Rate my answer from 1 to 5 on clarity, specificity, and impact. Give me one strength and one specific thing to improve. Be honest — I'd rather hear it from you than from an interviewer." You'll know it worked when you can name one concrete sentence or word you'd change before tomorrow.

Step 5 — Run a five-question dress rehearsal

Once your warm-up answer feels solid, ask the assistant to run a full short interview. What you'll see is it moving through five questions in sequence — one at a time, waiting for your reply

before continuing. Treat this like the real thing: no peeking at notes, no redoing an answer, and try to keep each answer between 60 and 90 seconds. If your assistant has voice mode, use it for this round; speaking back-to-back is closer to the actual interview than typing. Set a phone timer for 10 minutes so you don't run over. If the screen shows a **Stop** or **Regenerate** button, ignore both and just answer naturally. 💬 Example: "Now run a 5-question mock interview, one question at a time. Wait for my answer before moving on. Mix behavioral and role-specific questions." You'll know it worked when you've answered five questions back-to-back out loud and feel a bit tired — that's the real interview muscle kicking in.

Step 6 — Capture your three weakest spots before you close the chat

End the session by asking the assistant to summarize the patterns it noticed — vague answers, missing numbers, weak openings. What you'll see is a short list of two to four patterns, plus the specific questions that exposed them. Screenshot it or copy it into your notes app before you close the chat, because tomorrow you'll want to practice those exact spots. This single summary is what turns a random practice session into a repeatable method: each night you run the same drill, you knock down one more weak pattern. Be aware — the assistant may flatter you here too, so push for the harshest pattern if its list feels too soft. 💬 Example: "Based on this whole session, list the three patterns that hurt my answers most, and the one question I should drill tomorrow." You'll know it worked when you have a 3-item checklist saved somewhere outside the chat.

Common mistakes

- **Reading your answer instead of saying it.** Typing an answer and reading it back to yourself feels safe but doesn't catch the rambling, filler words, or long pauses a real interviewer hears. Speak every answer out loud, even if it feels awkward.
- **Letting the assistant be too polite.** Most AI assistants default to "great answer!" — which is useless for practice. If your feedback feels like a hug, tell the assistant directly to score you from 1 to 5 and give one specific improvement.
- **Practicing only "tell me about yourself."** That opener is rarely the question that decides the interview. Spend at least half your session on the two or three harder behavioral and scenario questions for the actual role — that's where the real anxiety hides.

What this means for you

- **In everyday life:** If you're between jobs or thinking about a change, this same loop works for promotion conversations, salary reviews, or even asking a landlord to fix something — anywhere a clear, structured answer matters more than a clever one.

- **For your business or work:** Use this method to rehearse client pitches, performance reviews, or tough conversations with a colleague. The "AI asks, you answer out loud, get one specific critique" pattern generalizes to almost any high-stakes talk.
- **If you're just getting started:** Don't try to do all six steps tonight. Do Step 1, paste the job description, and answer one tough question out loud. That's already more practice than most candidates do.

Try it now

Open your AI assistant right now, paste the job description into a fresh chat, and type: *"Ask me one question at a time. Start with the hardest behavioral question for a [role] interview."* Then answer it out loud. That's the whole loop in two minutes — and the easiest way to find out where you actually stand tonight.